
From Noise to Diversity: Random Embedding Injection in LLM Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

Recent *soft prompt* research has tried to improve reasoning by inserting trained vectors into LLM inputs, yet whether the gain comes from the *learned content* or from the *act of injection itself* has not been carefully separated. We study Random Soft Prompts (RSPs), which drop the training step entirely and append a freshly drawn sequence of random embedding vectors to the input. Each RSP vector is sampled from an isotropic Gaussian fitted to the entrywise mean and variance of the pretrained embedding table; the sequence carries no learned content, and yet reaches accuracy comparable to optimized *soft prompts* on math reasoning benchmarks in several settings. The mechanism unfolds in two stages: because attention has to absorb a never-seen-before random position, the distribution over the first few generated tokens flattens and reasoning trajectories *branch*, and as generation continues this influence dilutes naturally so the response commits to a single completion. We show that during inference RSPs lift early-stage token diversity and, combined with temperature sampling, widen Pass@ N , the probability that at least one out of N attempts is correct. Beyond inference, we carry the same effect into DAPO training and demonstrate practical gains. Our contributions are: (i) RSP isolates the simplest form of *soft prompt* — training-free, freshly resampled — providing a unified lens for the structural effect of injection that variants otherwise differing in training and form all share; (ii) a theoretical and empirical validation of the underlying mechanism; and (iii) an extension from inference to training.

1 Introduction

This work began from a surprising observation: even *untrained*, randomly drawn embeddings appended to LLM inputs lift reasoning accuracy and reshape early decoding in several settings — with no optimization at all. As LLMs scale, full fine-tuning is prohibitive, motivating parameter-efficient methods [Hu et al., 2022, Houlby et al., 2019]. Among these, *soft prompts* insert learnable continuous embeddings into the input and steer behavior through attention [Li and Liang, 2021, Lester et al., 2021]; the paradigm is actively adopted for mathematical reasoning [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Hao et al., 2025, Zhang et al., 2026]. These methods diverge in injection position, training objective, and prompt form, and prior work attributes their gains to the *learned content* of the optimized vectors. The act of injecting embeddings into the input, common to all variants, has remained outside the analysis. Even though theory shows trained prefixes only bias attention in a fixed direction [Petrov et al., 2024], the source of the effect itself remains unsettled.

To separate *learned content* from the *act of injection*, we use Random Soft Prompts (RSPs) as a control — continuous vectors drawn from an isotropic Gaussian fitted to the embedding table’s mean and variance, with no learning. The control turns out to be non-neutral. Applying a chat template to a base model not fine-tuned for it degrades reasoning, yet appending RSPs recovers up to +29




Figure 1: Conceptual overview of RSP-induced trajectory diversity. The hidden states shown are the final-layer outputs that drive the next-token distribution. (a) Early decoding: RSP perturbation reaches a *branching state* (red; formalized in §3.1), whose top- K differs from the no-RSP baseline and induces a diverse output distribution. (b) Later decoding: as the KV cache grows, the bound on RSP attention mass narrows (Theorem 1); branching events become rare and the hidden state (blue) commits to the branch already opened. (c) Across rollouts, the trajectories opened by early RSP perturbation accumulate into independent paths.

pp on Qwen2.5-Math-1.5B; if simple noise were merely shaking the model, accuracy should have worsened instead. The injection itself, independent of any learned content, alters model behavior.

We analyze training-free RSPs on LLM reasoning. With each rollout receiving an *independent* draw, a single RSP injection *branches* hidden states into diverse trajectories early on, which then *stabilize* as decoding proceeds. Empirically, RSPs share the early-decoding entropy signature of optimized methods (LTPO, TTSV, SoftCoT) despite carrying no learned content; Pass@ N accumulates only under *independent* resampling; and the same input-side diversity composes with DAPO [Yu et al., 2025] training. Our contributions are: (i) RSP isolates the simplest form of *soft prompt* — training-free and freshly resampled per rollout — providing a unified lens for the structural effect of injection across variants that otherwise differ in training objective and prompt form, while still reaching accuracy comparable to optimized methods in several settings; (ii) a theoretical and empirical account of the mechanism: attention-mediated branch opening, isotropic coverage across all directions, and the automatic attenuation of RSP attention as the KV cache grows; and (iii) DAPO extension showing independent RSP resampling composes with rollout-diversity training.

2 Background

2.1 Soft Prompt Methods for LLM Reasoning

Prior *soft prompt* methods share a common structure: they inject continuous vectors that do not correspond to discrete vocabulary tokens into the model’s representation stream and *optimize* them under task-specific objectives. As parameter-efficient alternatives to full fine-tuning, Prefix-Tuning [Li and Liang, 2021], Prompt Tuning [Lester et al., 2021], and P-tuning [Liu et al., 2024] reached near-fine-tuning performance using learnable continuous vectors, after which LLM-reasoning variants followed. TTSV [Kang et al., 2026] optimizes *prefix* embeddings at test time via trajectory entropy minimization; SoftCoT [Xu et al., 2025] replaces explicit chain-of-thought tokens with soft tokens generated by an assistant model; LTPO [Ye et al., 2026] optimizes the *soft prompt* per problem to maximize its own confidence on that problem; COCONUT [Hao et al., 2025] feeds its hidden states back as input embeddings; MemGen [Zhang et al., 2026] uses embeddings as experience memory; and Pause tokens [Goyal et al., 2024] insert learnable tokens to provide additional computation steps.

These methods differ in position, training stage, and training target, yet share a common assumption: that the *learned representations* are the core driver of the effect. This paradigm faces several limitations. (i) Training is domain/benchmark-specific, requiring re-optimization for new settings and sensitive to seed/hyperparameter choices. (ii) Several have been observed to degrade on out-of-distribution challenging tasks [Ye et al., 2026]. (iii) Evidence for the effect is largely empirical, with limited theoretical explanation beyond the training objective. Furthermore, soft prompting and prefix-tuning are theoretically restricted to biasing attention outputs in a fixed direction, without

Table 1: Accuracy (%) across model configurations and benchmarks. Values in parentheses indicate the difference from the baseline. Full per-position breakdown is in Appendix A.

Model	MATH-500			GSM8K			AIME24		
	Baseline	RSP(prefix)	RSP(suffix)	Baseline	RSP(prefix)	RSP(suffix)	Baseline	RSP(prefix)	RSP(suffix)
<i>Instruct Models</i>									
LLaMA-3.1-8B-Inst	52.20	45.00 (−7.2)	49.40 (−2.8)	85.75	76.35 (−9.4)	86.73 (+1.0)	6.67	3.33 (−3.3)	10.00 (+3.3)
Qwen2.5-Math-7B-Inst	83.20	81.40 (−1.8)	83.40 (+0.2)	95.45	94.92 (−0.5)	95.68 (+0.2)	13.33	13.33 (+0.0)	16.67 (+3.3)
Qwen2.5-Math-1.5B-Inst	73.00	74.80 (+1.8)	74.20 (+1.2)	85.29	85.29 (+0.0)	84.69 (−0.6)	10.00	13.33 (+3.3)	20.00 (+10.0)
<i>Base Models (Raw Text)</i>									
Qwen2.5-Math-7B	72.20	71.60 (−0.6)	55.60 (−16.6)	84.15	84.31 (+0.2)	54.13 (−30.0)	6.67	16.67 (+10.0)	13.33 (+6.7)
Qwen2.5-Math-1.5B	61.40	64.40 (+3.0)	54.60 (−6.8)	77.86	74.30 (−3.6)	58.61 (−19.3)	16.67	10.00 (−6.7)	3.33 (−13.3)
<i>Base Models + ChatML (Format Mismatch)</i>									
Qwen2.5-Math-7B	52.20	59.20 (+7.0)	70.40 (+18.2)	58.30	56.41 (−1.9)	75.97 (+17.7)	23.33	3.33 (−20.0)	23.33 (+0.0)
Qwen2.5-Math-1.5B	34.40	63.40 (+29.0)	51.00 (+16.6)	37.45	67.02 (+29.6)	50.64 (+13.2)	13.33	20.00 (+6.7)	13.33 (+0.0)

changing the relative attention pattern among content tokens [Petrov et al., 2024]. Together, these findings cast doubt on whether the learned prompt vectors function as intended: they are suitable for *eliciting* or *combining* skills in the pretrained model, but they struggle to learn new capabilities.

Moreover, prior works have not separated the contribution of *injection itself* from that of *optimization*. In light of limitations (i)–(iii), this paper focuses on *injection itself*: we introduce *Random Soft Prompts* (RSPs) — vectors drawn from an isotropic Gaussian fitted to the pretrained embedding statistics, with the training stage entirely removed — and show that they produce output-distribution shifts and accuracy comparable to those of learned soft prompts, thereby disentangling the two contributions. Section 3 analyzes the mechanism theoretically; Section 4 validates it empirically.

2.2 Noise Injection in Neural Networks

Neural networks have incorporated noise at several stages and locations. During training, dropout [Srivastava et al., 2014] randomly deactivates *hidden activations*, and NEFTune [Jain et al., 2024] adds Gaussian noise to *input embeddings* during instruction fine-tuning to improve instruction following. At inference time, randomized smoothing [Cohen et al., 2019] injects Gaussian noise into *raw inputs* to obtain certified robustness, and SmoothLLM [Robey et al., 2025] perturbs prompts at the *discrete character* level to defend against jailbreaking. In RL, NoisyNet [Fortunato et al., 2018] adds learnable, trained noise to *network weights* to aid exploration. Rather than *adding* noise to embeddings, RSP *appends* it, and differs in two respects: (i) unlike NoisyNet, it is entirely training-free, with no model fine-tuning and no learned noise parameters; (ii) it appends within a single forward pass, eliminating the multi-pass aggregation that robustness methods require.

3 Random Soft Prompts and Why They Work

3.1 Random Soft Prompt

This section starts from an empirical observation. Random embedding vectors appended to LLM inputs lift reasoning accuracy across several models and benchmarks (Table 1) with no training; the gain holds across injection positions (*prefix*, *suffix*) (§4.2), traveling with the act of injection rather than a particular configuration. To explain why, we track one quantity — the *attention mass* w that the model places on RSP tokens. Early in decoding w is large enough to open a different decoding branch (§3.2); as the KV cache grows, the upper bound on w narrows automatically and the model commits to the branch already opened (§3.3) — an implicit explore-then-exploit. Independent resampling per rollout opens a fresh branch each time (proofs in Appendix E).

To isolate the act of injection from learned content while keeping the magnitude comparable to real tokens and not privileging any direction in embedding space, we draw RSP from an isotropic Gaussian fitted to the entrywise statistics of the pretrained embedding table. Let $\mathbf{W}_E \in \mathbb{R}^{V \times d}$ denote the pretrained token-embedding matrix, with V the vocabulary size and d the hidden dimension. Write its entrywise statistics as $\mu_E := \frac{1}{V \cdot d} \sum_{v,k} [\mathbf{W}_E]_{v,k}$ (the mean over all $V \cdot d$ entries) and $\sigma_E := \text{std}(\mathbf{W}_E)$. An RSP of length L is a sequence of continuous vectors $\mathbf{H} = [h_1, \dots, h_L] \in \mathbb{R}^{L \times d}$

drawn independently from

$$h_j \sim \mathcal{N}(\mu_E \mathbf{1}_d, \sigma_E^2 \mathbf{I}_d), \quad j = 1, \dots, L. \quad (1)$$

The centered form $\bar{h}_j := h_j - \mu_E \mathbf{1}_d \sim \mathcal{N}(0, \sigma_E^2 \mathbf{I}_d)$ is a zero-mean isotropic Gaussian. Letting $\mathbf{X} \in \mathbb{R}^{T \times d}$ denote the input token-embedding sequence (the T rows of $\mathbf{W}_E$ corresponding to the prompt tokens), the main text uses the *suffix* form $[\mathbf{X}; \mathbf{H}] \in \mathbb{R}^{(T+L) \times d}$ as the default; other positions (*prefix* $[\mathbf{H}; \mathbf{X}]$, *infix* ($\mathbf{H}$ inserted within $\mathbf{X}$) [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Zhang et al., 2026]) are reported as ablations in §4.2 and Appendix A. RSP is not a learnable parameter, so it receives no gradient and is freshly drawn for each rollout.

Call decoding step t a *branching event* and the last-layer hidden state $\mathbf{s}^{(t)}$ a *branching state* when the LM-head top- K differs from that of the no-RSP baseline $\bar{\mathbf{s}}^{(t)}$ (Figure 1, red).¹

3.2 Exploration: RSP strengthens early-stage exploration

Inserting a random prompt concentrates attention on it, and that concentration amplifies exploration. The effect of a single injection on the hidden state raises two questions — *how much* and *where* — decided respectively by the scalar w inside one attention head and by the distribution of $\mathbf{H}$. The isotropic Gaussian addresses both at once. Both questions trace back to the branching condition of §3.1: for a single RSP draw to shift top- K , its perturbation of the head output must be large enough to propagate through the layers (*how much*) and along directions that affect logit ordering (*where*).

Consider one attention head in self-attention layer ℓ . At query position i , the query attends $n \geq 1$ unmasked real tokens together with $L \geq 1$ RSP tokens, all attention logits finite. Write the attention weights as $\alpha_{ij}^{(\ell)}$ and the value vectors on the real and RSP sides as $\mathbf{v}_j^{(\ell)}, \mathbf{v}_{r_j}^{(\ell)}$. Defining the total attention mass on random tokens $w_{r,i}^{(\ell)} := \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)}$, the real-token mass $1 - w_{r,i}^{(\ell)}$ is positive and the head output $\mathbf{o}_i^{(\ell)}$ splits *exactly* into a renormalized real-token term $\tilde{\mathbf{o}}_i^{(\ell)}$ and an RSP-induced contribution $\boldsymbol{\eta}_i^{(\ell)}$:

$$\mathbf{o}_i^{(\ell)} = (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{o}}_i^{(\ell)} + \boldsymbol{\eta}_i^{(\ell)}, \quad \tilde{\mathbf{o}}_i^{(\ell)} := \sum_{j=1}^n \frac{\alpha_{ij}^{(\ell)}}{1 - w_{r,i}^{(\ell)}} \mathbf{v}_j^{(\ell)}, \quad \boldsymbol{\eta}_i^{(\ell)} := \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)}. \quad (2)$$

Derivation of Eq. (2). Write the head output as $\sum_{j=1}^n \alpha_{ij}^{(\ell)} \mathbf{v}_j^{(\ell)} + \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)}$, then factor $1 - w_{r,i}^{(\ell)} = \sum_{j=1}^n \alpha_{ij}^{(\ell)} > 0$ out of the real-token sum. No approximation beyond the softmax-normalization identity is used. $\square$

This single quantity $w_{r,i}^{(\ell)}$ controls both the attenuation ratio $1 - w_{r,i}^{(\ell)}$ on the real signal and the upper bound on the random contribution $\|\boldsymbol{\eta}_i^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_j \|\mathbf{v}_{r_j}^{(\ell)}\|$. The *magnitude* of RSP-induced variation thus reduces to one number in $[0, 1]$. Early in decoding the KV cache is short and this value is non-negligible, so a single RSP draw can perturb next-token decisions.

The scalar w controls *how much*; *where* is decided by the distribution of $\mathbf{H}$. For the local logit argument, isolate one RSP position and let $\bar{h} \in \mathbb{R}^d$ denote the i.i.d. marginal of any centered vector $\bar{h}_j$ in Eq. (1), with the other centered RSP positions fixed at zero. Let $z_a(\bar{h})$ be the under-RSP output logit at vocab token a and write the vocab-logit gap $\Delta_{ab}(\bar{h}) := z_a(\bar{h}) - z_b(\bar{h})$ ($\bar{h} = 0$ is the centered mean, not the no-RSP state; Appendix D). The transformer’s logit map is non-linear in $\bar{h}$, but a first-order Taylor expansion around $\bar{h} = 0$ gives the local surrogate $\Delta_{ab}(\bar{h}) \approx \Delta_{ab}(0) + b_{ab}^\top \bar{h}$, with $b_{ab} := \nabla_{\bar{h}}(z_a - z_b)|_{\bar{h}=0} \in \mathbb{R}^d$ the gradient direction along which vocab tokens a, b swap rank (not unit-norm in general). Because RSP is training-free, we cannot know in advance which b_{ab} opens a useful branch, so the distribution must avoid systematically under-perturbing any direction. A deterministic injection pins to one direction; vocabulary sampling inherits the embedding table’s anisotropy. The remaining design — maximize the worst-case directional variance at fixed budget — is uniquely solved by isotropy.

¹Italic h_j : RSP input space; bold $\mathbf{s}^{(t)}$: hidden state. K denotes the cardinality in top- K analyses; nucleus sampling uses a separate threshold p . t for decoding step, ℓ for self-attention layer.

149 **Proposition 1** (Maximin directional coverage). *Let $\mathcal{D}_\rho$ denote the family of zero-mean distributions*
 150 *on $\mathbb{R}^d$ whose covariance satisfies $\text{tr}(\Sigma_D) \leq \rho^2$. For each $D \in \mathcal{D}_\rho$, $\min_{\|u\|=1} \text{Var}_{h \sim D}(u^\top h) =$*
 151 *$\lambda_{\min}(\Sigma_D)$, and $\sup_{D \in \mathcal{D}_\rho} \lambda_{\min}(\Sigma_D) = \rho^2/d$, attained iff $\Sigma_D = (\rho^2/d)\mathbf{I}_d$.*

152 This is a *design criterion* for the prompt law, not a guarantee of correctness; correctness enters
 153 through the task-side $p_{\min}(x)$ assumption of §3.3 and independent resampling.²

154 Gaussian RSP attains the equality with $\Sigma = \sigma_E^2 \mathbf{I}_d$ ($\rho^2 = d\sigma_E^2$) and adds *full support* on $\mathbb{R}^d$
 155 (Appendix D). Isotropy excludes no direction; full support gives positive probability to the open
 156 branching-event region (§3.1) whenever it is non-empty (Appendix E.3.3). Monotone temperature
 157 preserves ranking and cannot reach this region, so RSP’s *rank-changing* branching follows from the
 158 two properties together, splitting the hidden state into multiple branching states early in reasoning.

159 3.3 Annealing: KV cache growth dilutes RSP attention

160 The early branching stabilizes as decoding proceeds: the upper bound on $w_{r,i}^{(\ell)}$ narrows along the
 161 sequence axis with autoregressive cache growth. Each step adds one real-token term while RSP terms
 162 stay fixed; at frozen attention-logit gap, this asymmetry yields the following bound.

163 **Theorem 1** (Attention-mass decay under KV cache growth). *Consider a single attention head with*
 164 *per-head key dimension d_{head} , where query $\mathbf{q}_i$ at position i attends $n \geq 1$ real keys $\mathbf{k}_j$ and $L \geq 1$*
 165 *RSP keys $\mathbf{k}_{r,j'}$ under finite logits. Define the pre-scale attention-logit gap $\Delta_i := \min_{j \in [n]} \mathbf{q}_i^\top \mathbf{k}_j -$*
 166 *$\max_{j' \in [L]} \mathbf{q}_i^\top \mathbf{k}_{r,j'}$ (distinct from the vocab-logit gap Δ_{ab} of §3.2). In scaled dot-product attention,*

$$w_{r,i} \leq \frac{L}{L + n \exp(\Delta_i / \sqrt{d_{\text{head}}})},$$

167 *and the right-hand side is strictly decreasing in n and tends to 0 at fixed L and Δ_i .*

168 *Proof.* Let $s_j := \mathbf{q}_i^\top \mathbf{k}_j / \sqrt{d_{\text{head}}}$, $s_{r,j'} := \mathbf{q}_i^\top \mathbf{k}_{r,j'} / \sqrt{d_{\text{head}}}$, $s_{r,\max} := \max_{j'} s_{r,j'}$, and $R :=$
 169 $\sum_{j'=1}^L \exp(s_{r,j'})$. The gap definition gives $s_j \geq \Delta_i / \sqrt{d_{\text{head}}} + s_{r,\max}$ for every real j , and
 170 $\exp(s_{r,\max}) \geq R/L$ holds because the max dominates the average; combining the two and sum-
 171 ming over the n real keys yields $\sum_{j=1}^n \exp(s_j) \geq (n/L) \exp(\Delta_i / \sqrt{d_{\text{head}}}) R$. Substituting into
 172 $w_{r,i} = R / (\sum_j \exp(s_j) + R)$ gives

$$w_{r,i} \leq \frac{R}{(n/L) \exp(\Delta_i / \sqrt{d_{\text{head}}}) R + R} = \frac{L}{L + n \exp(\Delta_i / \sqrt{d_{\text{head}}})}.$$

173 The derivative of the right-hand side in n is $-L \exp(\Delta_i / \sqrt{d_{\text{head}}}) / (L + n \exp(\Delta_i / \sqrt{d_{\text{head}}}))^2 < 0$, so
 174 the bound is strictly decreasing and tends to 0 as $n \rightarrow \infty$. $\square$

175 Applied at each layer ℓ to $w_{r,i}^{(\ell)}$ with gap $\Delta_i^{(\ell)}$, L caps the maximum influence while n grows each
 176 step. The theorem guarantees only sequence-axis attenuation at fixed (layer, query, gap); the gap may
 177 sharpen at deeper layers, as suggested empirically by Figure 2, but that effect is outside the theorem’s
 178 scope. The accurate reading is “the attainable upper bound at a fixed gap decreases monotonically.”
 179 Because Eq. (2) ties the random contribution to the same scalar, branching-event reachability inherits
 180 this envelope — as the cache grows, events become rare and the response commits to the branch
 181 already opened, an implicit explore-then-exploit. Figure 2 visualizes this pattern along both axes.

182 Lifting single-shot branching to task accuracy across N rollouts requires (M1) a positive lower bound
 183 $p_{\min}(x) > 0$ on the per-rollout probability that one execution reaches a correct trajectory and (M2)
 184 the N executions are mutually independent. Under both, $\Pr(\text{Pass@}N \text{ on } x) \geq 1 - (1 - p_{\min}(x))^N$
 185 (Appendix E.3.4). §3.2’s isotropy and full support make local branch opening reachable when
 186 the open region is nonempty, but they do not by themselves guarantee the task-side correctness
 187 condition (M1); *independent resampling* supplies (M2). Sharing one RSP breaks (M2) and removes
 188 the accumulation — a signature §4 verifies empirically.

²The maximin is stated in input space; isotropy is not claimed to survive transformer layers. Any logit direction pulls back
 via the model Jacobian to an input-space b , and isotropy guarantees only that no such pulled-back direction (in the role of u
 above) has zero projected variance.




Figure 2: Per-token attention mass on Qwen2.5-Math-7B (*suffix*, 500 MATH-500 problems, independent RSP per problem). (a) RSP-token attention decreases along the reasoning-position axis, matching Theorem 1’s KV cache-growth bound; the additional layer-axis decrease is an empirical phenomenon (deeper layers sharpen the real-vs-random gap) outside Theorem 1’s scope. (b) Question-token attention stays comparatively uniform. Reported quantity: per-token mass of Appendix F, equal to $w_{r,i}^{(\ell)} / L$ averaged over heads and samples.

4 Empirical Validation

The mechanism of §3 runs from early *branching* through KV cache-growth narrowing to $\text{Pass}@N$ accumulation under independent resampling. We verify this flow across three scenes. First, §4.2 tests whether RSPs preserve baseline accuracy and even recover it in some settings. Next, §4.3 examines whether the predicted early diversification and later stabilization appear in entropy and attention signals. §4.4 asks whether the $\text{Pass}@N$ gain depends on *independence* of the RSP draw. Finally, §4.5 extends the same perturbation from inference to RL training, where rollout diversity is critical.

4.1 Experimental Setup

We evaluate on three mathematical reasoning benchmarks, MATH-500 [Lightman et al., 2024], GSM8K [Cobbe et al., 2021], and AIME24, using LLaMA-3.1-8B-Instruct [Grattafiori et al., 2024] and the Qwen2.5-Math model family [Yang et al., 2024] across instruct models, base models, and base models with mismatched chat templates. Tables 1 and 2 use greedy decoding to isolate the effect of RSPs from sampling stochasticity; the diversity and $\text{Pass}@N$ experiments in §4.4 use the sampling settings specified there. RSPs follow the definition of §3.1, using the default *suffix* position with a freshly drawn $\mathbf{H}$ per rollout. Since Theorem 1 identifies the RSP length L as the design parameter capping the maximum perturbation influence, we ablate $L \in \{10, 15, 20\}$ per model to identify an appropriate perturbation strength. Tables 1 and 2 report results at the per-model selected length. The mechanism analyses in §4.3–§4.4 use a fixed setting ($L=20$, *suffix*), which insulates them from this selection effect. Full experimental details and per-position results are in Appendix A.

4.2 How Does Untrained RSP Compare to Baselines & Optimized Soft Prompts?

We unpack the per-configuration patterns of Table 1 (presented in §3). On instruct models RSP stays within a few percentage points of the baseline. The most striking pattern is format-mismatch recovery: when a ChatML template is applied to a base model not trained for that format, *Suffix* recovers +18.2/ +17.7 pp on Qwen2.5-Math-7B (MATH-500/GSM8K) and *prefix* recovers up to +29.0/ +29.6 pp on Qwen2.5-Math-1.5B — if simple noise were responsible, accuracy should have dropped further. In the base raw-text setting the picture inverts: *suffix* incurs sharp losses (−16.6/ −30.0 pp on Qwen2.5-Math-7B) while *prefix* changes accuracy less, suggesting base models treat EOS-adjacent signals as a derailment trigger that *prefix* avoids.

RSP’s *direct accuracy gains* concentrate in misaligned-input recovery and stay within $\pm$ a few percentage points on well-aligned instruct models. The effects this paper centers on are the token-, trajectory-, and outcome-level diversity in §4.4, with accuracy recovery as a secondary outcome.

Compared with three optimized methods — TTSV, SoftCoT, LTPO — under a unified pipeline (Appendix A.2), RSP lands in the same accuracy band without any training, optimization, or task-specific tuning. It wins outright on several model–benchmark cells and shows meaningful gains even on the challenging AIME24 setting. What we want to highlight is the converse: a fixed, training-free

Table 2: Comparison with prior *soft prompt* methods (TTSV, SoftCoT, LTPO) under unified evaluation. Baseline and RSP values are from Table 1. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Benchmark	Baseline	RSP	TTSV	SoftCoT	LTPO
Qwen2.5-Math-7B-Instruct	MATH-500	83.20	84.20	<u>83.80</u>	82.80	83.00
	GSM8K	<u>95.45</u>	95.68	95.30	95.00	94.84
	AIME24	13.33	<u>16.67</u>	23.30	<u>16.67</u>	13.33
Qwen2.5-Math-1.5B-Instruct	MATH-500	73.00	<u>75.20</u>	<u>75.20</u>	76.00	72.40
	GSM8K	85.29	85.75	<u>85.70</u>	84.46	84.53
	AIME24	<u>10.00</u>	20.00	6.70	6.67	<u>10.00</u>

distribution reaching this range is the paper’s central evidence that part of the *soft-prompt* effect comes less from learned content than from the random directional perturbation that injection induces — maximin coverage (§3.2) at a magnitude bounded by the KV cache (§3.3).

4.3 How Does RSP Affect the Output Distribution?

We next examine how RSP injection reshapes the output distribution. The setting is Qwen2.5-Math-1.5B-Instruct on MATH-500, and we measure per-token entropy, top-1 probability, and varentropy across generation steps. Figure 3 compares these metrics for the first 5% of generation steps across RSP variants and prior *soft prompt* methods (LTPO, SoftCoT, TTSV); full curves are in Appendix A.4.




Figure 3: Mean entropy, top-1 probability, and varentropy during the first 5% of generation steps (Qwen2.5-Math-1.5B-Instruct, MATH-500). Shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: prior *soft prompt* methods (LTPO, SoftCoT, TTSV).

Latent prompt methods including RSP share the same signature: *a rise in early-generation entropy followed by convergence to the baseline later in generation*. A natural mechanistic explanation is the out-of-vocabulary (OOV) effect — when continuous embeddings outside the vocabulary enter the input, the model has to attend to a never-before-seen key position on top of its familiar token distribution, redistributing probability mass across multiple candidates within the early next-token distribution. All three RSP injection positions show early-stage entropy $\uparrow$, top-1 probability $\downarrow$, and varentropy $\uparrow$, jointly consistent with multimodal distributions [Ahmed et al., 2026]. The change is confined to the early stage and all three metrics return to baseline in the middle and final portions, consistent with the *early influence with automatic KV cache-growth decay* pattern predicted in §3.3.

Prior methods trained under different objectives (LTPO, SoftCoT, TTSV) share this signature. Even TTSV, whose reward explicitly *lowers* mean reasoning-trajectory entropy, shows early-stage entropy comparable to the baseline. The entropy value itself is not a quality score — higher entropy is not always better; rather, the signature reveals a *shared mechanism* in which OOV embedding injection produces the same fingerprint independently of the optimization objective, supporting the hypothesis that the key driver of latent reasoning is the *act of injection*, not the learned content.

4.4 Does RSP Induce Trajectory Diversity?

Section 4.3 empirically established that RSP reshapes the early-stage output distribution. We now examine whether this shift translates into reasoning diversity at three complementary levels of analysis: token rankings, semantic content of trajectories, and task-level outcomes. The three analyses converge

on a directional picture: RSP’s first-token perturbation propagates into semantically more diverse reasoning trajectories, which in turn yields greater output diversity at the task level. All three analyses share the setting: Qwen2.5-Math-1.5B-Instruct, MATH-500, *suffix* injection, and $L = 20$.

Token-level: distribution beyond temperature. We quantify how much RSP shifts the first-token distribution — the critical forking point for reasoning trajectories [Wang et al., 2025]. Using the baseline at $\tau=1.0$ as reference, we compare baselines at $\tau=2.0, 3.0$ (16 samples per problem) against RSP at $\tau=1.0$ (averaged over 16 seeds) on three metrics: Spearman rank correlation ρ to detect rank reorganization, the probability mass placed outside the baseline’s top-10 to measure support expansion, and Jensen–Shannon divergence to quantify the overall magnitude of distributional change.

Table 3 shows the contrast: RSP’s distributional shift falls between $\tau=2.0$ and $\tau=3.0$ in magnitude, but the mechanism differs — temperature is a monotone transform that preserves ranking ($\rho = 1$ at any τ), while RSP partially preserves but reranks ($\rho = 0.709$). The Pass@1 row exposes the cost: matching RSP’s magnitude with temperature toward $\tau=3.0$ collapses accuracy to 1.42%, whereas RSP preserves 73.30%. Thus, RSP produces a fundamentally different perturbation than temperature scaling.

Moreover, RSP’s accuracy variability across 16 seeds at $\tau=1.0$ is ± 1.07 , comparable to the ± 1.47 that temperature sampling produces across 16 samples at $\tau=2.0$ (detailed definitions in Appendix B).

Table 3: First-token distribution metrics measured against the baseline at $\tau=1.0$, comparing baselines at higher temperatures ($\tau=2.0, 3.0$) with RSP at $\tau=1.0$. RSP metrics are averaged over 16 random seeds; Acc reports mean $\pm$ std across 16 samples per problem.

Metric	$\tau=2.0$	$\tau=3.0$	RSP
Spearman ρ	1.000	1.000	0.709
Mass outside top-10	5.18%	29.11%	21.86%
JS divergence	0.060	0.218	0.131
Acc (%)	20.77 ± 1.47	1.42 ± 0.35	73.30 ± 1.07

Trajectory-level: semantic diversity. Token-level reranking raises a second question: do first-token perturbations diverge into distinct trajectories, or converge to similar ones? We sample 64 problems from MATH-500 and generate 300 independent trajectories per problem under Baseline and RSP at $\tau = 1.0$. Each trajectory is encoded by BGE-M3 [Chen et al., 2024] into a dense semantic vector, and we compute three metrics: *pairwise cosine distance* (overall semantic spread), *inter-cluster distance* between centroids of DBSCAN [Ester et al., 1996], a density-based clustering algorithm (separation between distinct trajectory groups), and *intra-cluster distance* (coherence within groups). Table 4 reveals three concurrent patterns: pairwise distance rises by $1.4\times$, inter-cluster distance jumps by $5.4\times$, and intra-cluster distance is essentially unchanged — the trajectory set becomes more diverse and splits into more separated groups while preserving within-group coherence. RSP also yields larger pairwise distance than baseline on 56 of 64 problems.

Table 4: Semantic diversity metrics (64 problems, 300 trajectories per condition).

Metric	Baseline	RSP
Pairwise cos. dist.	0.0612	0.0879
Inter-cluster dist.	0.0183	0.0982
Intra-cluster dist.	0.0526	0.0528

Outcome-level: Pass@ N scaling. Does this token- and trajectory-level diversity translate into task-level outcome diversity, measured by Pass@ N [Chen et al., 2021] (the probability that at least one of N sampled solutions is correct)? We generate 16 samples per problem on MATH-500 and AIME24 and compare four conditions: (i) *Baseline* at $\tau = 1.0$; (ii) *RSP (single seed)* at $\tau = 1.0$, with one RSP shared across all samples; (iii) *RSP (indep. seed)* with greedy decoding, a fresh RSP per sample; and (iv) *RSP (indep. seed)* at $\tau = 1.0$, a fresh RSP per sample.

Condition (iv) achieves the highest Pass@ N on both benchmarks, reaching 92.80% on MATH-500 and 36.67% on AIME24 at $N = 16$, while (ii) shows no improvement over the baseline — the diversity gain depends on varying the embeddings across samples rather than fixing a single perturbation. Pass@1 stays comparable to the baseline on MATH-500, and on the harder AIME24 (iv) even surpasses it, suggesting that independent RSPs combined with temperature sampling expand trajectory diversity without sacrificing single-sample accuracy. Qualitative analysis is in Appendix C.

4.5 Application: DAPO Training with RSP

Beyond inference (§4), does the same effect transfer to training? DAPO [Yu et al., 2025] is a recent GRPO [Shao et al., 2024] variant that preserves rollout diversity at the loss level through asymmetric clipping. We use it as a testbed to verify whether input-side branch opening composes with loss-side




Figure 4: Pass@ N scaling on (a) MATH-500 and (b) AIME24 with Qwen2.5-Math-1.5B-Instruct, 16 samples per problem. *Baseline*: temperature sampling only. *RSP (single seed)*: single RSP shared across samples combined with temperature. *RSP (indep. seed)*: a different RSP per sample, with or without temperature.

304 preservation to yield further gains. In DAPO+RSP, we incorporate the input-side perturbation of §4.4
 305 into the RL rollout stage: each rollout is paired with an independently sampled suffix RSP. The full
 306 objective and implementation details are in Appendix G.

Table 5: Per-benchmark accuracy (%) at each method’s peak step on Qwen2.5-Math-7B. **Bold** denotes the better of the two RL methods on each benchmark. Avg is the unweighted five-benchmark mean.

Method	GSM8K	MATH-500	College	Minerva	AIME24	Avg
Base	57.2	52.6	18.3	13.6	8.6	30.06
DAPO	86.3	78.2	41.4	37.5	22.6	53.20
DAPO + RSP	87.7	78.6	42.2	39.7	23.6	54.36

307 We implement DAPO on top of SimpleRL-Zoo [Zeng et al., 2025] and train Qwen2.5-Math-7B on a
 308 MATH Level-3–5 subset (~8.5K prompts): each step uses $G = 8$ rollouts per prompt at $\tau=1.0$, batch
 309 1,024, 4 PPO mini-batches of 256, for 100 steps. *DAPO* is compared against *DAPO + RSP*, where
 310 *suffix* RSP ($L=20$) is applied during rollouts only and evaluation uses no RSP. We evaluate on five
 311 math benchmarks (GSM8K [Cobbe et al., 2021], MATH-500 [Lightman et al., 2024], College Math,
 312 Minerva Math [Lewkowycz et al., 2022], AIME24) and report the unweighted mean; AIME24 uses
 313 Avg@32 at $\tau=1.0$, the others greedy.

314 Figure 5 and Table 5 show two patterns: a late-
 315 stage advantage and a delayed early reward growth.
 316 DAPO + RSP reaches 54.36% at step 90 (vs DAPO’s
 317 53.20% peak, +1.16 pp) and widens the gap to +3.10 pp
 318 at step 100 with DAPO’s post-peak decline delayed; ear-
 319 lier the curves cross around step 20. The patterns follow
 320 from the methods acting at separate stages of the roll-
 321 out cycle: DAPO preserves diversity at the loss level
 322 over already-generated rollouts while RSP perturbs the
 323 hidden states that produce them (§3.2), exposing the pol-
 324 icy to broader learning signals. Empirically, this is the
 325 exploration–exploitation dynamics of RL [Sutton and
 326 Barto, 2018] induced from the input side.




Figure 5: Five-benchmark average accuracy on Qwen2.5-Math-7B. DAPO + RSP reaches a higher peak (step 90) and stays stable through step 100.

327 5 Conclusion

328 We studied Random Soft Prompts (RSPs), training-free isotropic Gaussian vectors whose per-layer
 329 contribution is bounded by an attention-mass scalar that shrinks with KV cache growth and whose
 330 isotropic resampling reaches top- K entry regions inaccessible to rank-preserving temperature (§3).
 331 Empirically, RSP shares the early-decoding entropy signature of optimized *soft prompts*, requires
 332 independent resampling for Pass@ N gains, and composes with DAPO training (§4–§4.5), suggesting
 333 that part of what trained *soft prompts* deliver is a *structural by-product of injection itself*.

334 Open directions: (i) generalization beyond RoPE decoders and math reasoning; (ii) deep-layer
 335 attention-gap sharpening outside Theorem 1’s envelope; (iii) multi-seed evaluation, especially on
 336 small competition sets; (iv) held-out validation for RSP length L and injection position.

References

- Farhan Ahmed, Yuya Jeremy Ong, and Chad DeLuca. LogitScope: A Framework for Analyzing LLM Uncertainty through Information Metrics. In *ICLR Workshop*, 2026.
- Nora Belrose, Igor Ostrovsky, Lev McKinney, Zach Furman, Logan Smith, Danny Halawi, Stella Biderman, and Jacob Steinhardt. Eliciting Latent Predictions from Transformers with the Tuned Lens. *arXiv:2303.08112*, 2023.
- Jianlyu Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. M3-Embedding: Multi-Linguality, Multi-Functionality, Multi-Granularity Text Embeddings Through Self-Knowledge Distillation. In *Findings of ACL*, 2024.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgen Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr, Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating Large Language Models Trained on Code. *arXiv:2107.03374*, 2021.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training Verifiers to Solve Math Word Problems. *arXiv:2110.14168*, 2021.
- Jeremy M. Cohen, Elan Rosenfeld, and J. Zico Kolter. Certified Adversarial Robustness via Randomized Smoothing. In *ICML*, 2019.
- Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu. A Density-Based Algorithm for Discovering Clusters in Large Spatial Databases with Noise. In *KDD*, 1996.
- Meire Fortunato, Mohammad Gheshlaghi Azar, Bilal Piot, Jacob Menick, Ian Osband, Alex Graves, Vlad Mnih, Remi Munos, Demis Hassabis, Olivier Pietquin, Charles Blundell, and Shane Legg. Noisy Networks for Exploration. In *ICLR*, 2018.
- Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. Transformer Feed-Forward Layers Are Key-Value Memories. In *EMNLP*, 2021.
- Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh Nagarajan. Think before You Speak: Training Language Models with Pause Tokens. In *ICLR*, 2024.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurelien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Roziere, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song, Danielle Pintz, Danny Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina Lobanova, Emily Dinan, Eric Michael Smith, Filip Radenovic, Francisco Guzmán, Frank Zhang, Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, Graeme Nail, Gregoire Mialon, Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Iliyan Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang, Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelfer van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang, Jiawen Liu, Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua

388 Johnstun, Joshua Saxe, Junteng Jia, Kalyan Vasuden Alwala, Karthik Prasad, Kartikeya Upasani,
 389 Kate Plawiak, Ke Li, Kenneth Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz
 390 Malik, Kuenley Chiu, Kunal Bhalla, Kushal Lakhotia, Lauren Rantala-Yearly, Laurens van der
 391 Maaten, Lawrence Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo,
 392 Lukas Blecher, Lukas Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat
 393 Singh, Manohar Paluri, Marcin Kardas, Maria Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya
 394 Pavlova, Melanie Kambadur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman
 395 Goyal, Narjes Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri Chatterji, Ning Zhang,
 396 Olivier Duchenne, Onur Celebi, Patrick Alrassy, Pengchuan Zhang, Pengwei Li, Petar Vasic,
 397 Peter Weng, Prajjwal Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Puxin Xu,
 398 Qing He, Qingxiao Dong, Ragavan Srinivasan, Raj Ganapathy, Ramon Calderer, Ricardo Silveira
 399 Cabral, Robert Stojnic, Roberta Raileanu, Rohan Maheswari, Rohit Girdhar, Rohit Patel, Romain
 400 Sauvestre, Ronnie Polidoro, Roshan Sumbaly, Ross Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar
 401 Hosseini, Sahana Chennabasappa, Sanjay Singh, Sean Bell, Seohyun Sonia Kim, Sergey Edunov,
 402 Shao-liang Nie, Sharan Narang, Sharath Raparthy, Sheng Shen, Shengye Wan, Shruti Bhosale,
 403 Shun Zhang, Simon Vandenhende, Soumya Batra, Spencer Whitman, Sten Sootla, Stephane
 404 Collot, Suchin Gururangan, Sydney Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha,
 405 Thomas Georgiou, Thomas Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal
 406 Karn, Vedanuj Goswami, Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Gouget,
 407 Virginie Do, Vish Vogeti, Vitor Albiero, Vladan Petrovic, Weiwei Chu, Wenhan Xiong, Wenyin
 408 Fu, Whitney Meers, Xavier Martinet, Xiaodong Wang, Xiaofang Wang, Xiaoqing Ellen Tan, Xide
 409 Xia, Xinfeng Xie, Xuchao Jia, Xuwei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei,
 410 Yi Wen, Yiwen Song, Yuchen Zhang, Yue Li, Yuning Mao, Zacharie Delpierre Coudert, Zheng Yan,
 411 Zhengxing Chen, Zoe Papakipos, Aaditya Singh, Aayushi Srivastava, Abha Jain, Adam Kelsey,
 412 Adam Shajnfeld, Adithya Gangidi, Adolfo Victoria, Ahuva Goldstand, Ajay Menon, Ajay Sharma,
 413 Alex Boesenberg, Alexei Baevski, Allie Feinstein, Amanda Kallet, Amit Sangani, Amos Teo,
 414 Anam Yunus, Andrei Lupu, Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew
 415 Poulton, Andrew Ryan, Ankit Ramchandani, Annie Dong, Annie Franco, Anuj Goyal, Aparajita
 416 Saraf, Arkabandhu Chowdhury, Ashley Gabriel, Ashwin Bharambe, Assaf Eisenman, Azadeh
 417 Yazdan, Beau James, Ben Maurer, Benjamin Leonhardi, Bernie Huang, Beth Loyd, Beto De Paola,
 418 Bhargavi Paranjape, Bing Liu, Bo Wu, Boyu Ni, Braden Hancock, Bram Wasti, Brandon Spence,
 419 Brani Stojkovic, Brian Gamido, Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu,
 420 Changan Wang, Changkyu Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris
 421 Tindal, Christoph Feichtenhofer, Cynthia Gao, Damon Civin, Dana Beaty, Daniel Kreymer, Daniel
 422 Li, David Adkins, David Xu, Davide Testuggine, Delia David, Devi Parikh, Diana Liskovich,
 423 Didem Foss, Dingkan Wang, Duc Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine
 424 Montgomery, Eleonora Presani, Emily Hahn, Emily Wood, Eric-Tuan Le, Erik Brinkman, Esteban
 425 Arcaute, Evan Dunbar, Evan Smothers, Fei Sun, Felix Kreuk, Feng Tian, Filippos Kokkinos, Firat
 426 Ozgenel, Francesco Caggioni, Frank Kanayet, Frank Seide, Gabriela Medina Florez, Gabriella
 427 Schwarz, Gada Badeer, Georgia Swee, Gil Halpern, Grant Herman, Grigory Sizov, Guangyi Zhang,
 428 Guna Lakshminarayanan, Hakan Inan, Hamid Shojanazeri, Han Zou, Hannah Wang, Hanwen Zha,
 429 Haroun Habeeb, Harrison Rudolph, Helen Suk, Henry Aspegren, Hunter Goldman, Hongyuan
 430 Zhan, Ibrahim Damlaj, Igor Molybog, Igor Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai
 431 Gat, Jake Weissman, James Geboski, James Kohli, Janice Lam, Japhet Asher, Jean-Baptiste Gaya,
 432 Jeff Marcus, Jeff Tang, Jennifer Chan, Jenny Zhen, Jeremy Reizenstein, Jeremy Teboul, Jessica
 433 Zhong, Jian Jin, Jingyi Yang, Joe Cummings, Jon Carvill, Jon Shepard, Jonathan McPhie, Jonathan
 434 Torres, Josh Ginsburg, Junjie Wang, Kai Wu, Kam Hou U, Karan Saxena, Kartikay Khandelwal,
 435 Katayoun Zand, Kathy Matosich, Kaushik Veeraraghavan, Kelly Michelena, Keqian Li, Kiran
 436 Jagadeesh, Kun Huang, Kunal Chawla, Kyle Huang, Lailin Chen, Lakshya Garg, Lavender A,
 437 Leandro Silva, Lee Bell, Lei Zhang, Liangpeng Guo, Licheng Yu, Liron Moshkovich, Luca
 438 Wehrstedt, Madian Khabza, Manav Avalani, Manish Bhatt, Martynas Mankus, Matan Hasson,
 439 Matthew Lennie, Matthias Reso, Maxim Groshev, Maxim Naumov, Maya Lathi, Meghan Keneally,
 440 Miao Liu, Michael L. Seltzer, Michal Valko, Michelle Restrepo, Mihir Patel, Mik Vyatskov,
 441 Mikayel Samvelyan, Mike Clark, Mike Macey, Mike Wang, Miquel Jubert Hermoso, Mo Metanat,
 442 Mohammad Rastegari, Munish Bansal, Nandhini Santhanam, Natascha Parks, Natasha White,
 443 Navyata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier, Nikhil Mehta, Nikolay Pavlovich
 444 Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, Olivia Hart, Omkar Salpekar, Ozlem
 445 Kalinli, Parkin Kent, Parth Parekh, Paul Saab, Pavan Balaji, Pedro Rittner, Philip Bontrager,
 446 Pierre Roux, Piotr Dollar, Polina Zvyagina, Prashant Ratanchandani, Pritish Yuvraj, Qian Liang,

447 Rachad Alao, Rachel Rodriguez, Rafi Ayub, Raghotham Murthy, Raghu Nayani, Rahul Mitra,
448 Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin Battey, Rocky Wang, Russ
449 Howes, Ruty Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu, Samyak Datta, Sara Chugh,
450 Sara Hunt, Sargun Dhillon, Sasha Sidorov, Satadru Pan, Saurabh Mahajan, Saurabh Verma, Seiji
451 Yamamoto, Sharadh Ramaswamy, Shaun Lindsay, Sheng Feng, Shenghao Lin,
452 Shengxin Cindy Zha, Shishir Patil, Shiva Shankar, Shuqiang Zhang, Shuqiang Zhang, Sinong Wang,
453 Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala, Stephanie Max, Stephen Chen, Steve Kehoe,
454 Steve Satterfield, Sudarshan Govindaprasad, Sumit Gupta, Summer Deng, Sungmin Cho, Sunny
455 Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman, Tal Remez, Tamar Glaser, Tamara
456 Best, Thilo Koehler, Thomas Robinson, Tianhe Li, Tianjun Zhang, Tim Matthews, Timothy Chou,
457 Tzook Shaked, Varun Vontimitta, Victoria Ajayi, Victoria Montanez, Vijai Mohan, Vinay Satish
458 Kumar, Vishal Mangla, Vlad Ionescu, Vlad Poenaru, Vlad Tiberiu Mihailescu, Vladimir Ivanov,
459 Wei Li, Wenchen Wang, Wenwen Jiang, Wes Bouaziz, Will Constable, Xiaocheng Tang, Xiaojian
460 Wu, Xiaolan Wang, Xilun Wu, Xinbo Gao, Yaniv Kleinman, Yanjun Chen, Ye Hu, Ye Jia, Ye Qi,
461 Yenda Li, Yilin Zhang, Ying Zhang, Yossi Adi, Youngjin Nam, Yu Wang, Yu Zhao, Yuchen Hao,
462 Yundi Qian, Yunlu Li, Yuzi He, Zach Rait, Zachary DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhenyu
463 Yang, Zhiwei Zhao, and Zhiyu Ma. The Llama 3 Herd of Models. *arXiv:2407.21783*, 2024.

464 Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian.
465 Training Large Language Models to Reason in a Continuous Latent Space. In *ICLR*, 2025.

466 Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Thai, Junhao Shen, Jinyi Hu, Xu Han,
467 Yujie Huang, Yuxiang Zhang, Jie Liu, Lei Qi, Zhiyuan Liu, and Maosong Sun. OlympiadBench: A
468 Challenging Benchmark for Promoting AGI with Olympiad-Level Bilingual Multimodal Scientific
469 Problems. In *ACL*, 2024.

470 Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin de Laroussilhe, Andrea
471 Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-Efficient Transfer Learning for NLP.
472 In *ICML*, 2019.

473 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang,
474 and Weizhu Chen. LoRA: Low-Rank Adaptation of Large Language Models. In *ICLR*, 2022.

475 Neel Jain, Ping-yeh Chiang, Yuxin Wen, John Kirchenbauer, Hong-Min Chu, Gowthami Somepalli,
476 Brian R. Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Aniruddha Saha, Micah Goldblum,
477 Jonas Geiping, and Tom Goldstein. NEFTune: Noisy Embeddings Improve Instruction Finetuning.
478 In *ICLR*, 2024.

479 Xinyue Kang, Diwei Shi, and Li Chen. Model Whisper: Steering Vectors Unlock Large Language
480 Models’ Potential in Test-time. In *AAAI*, 2026.

481 Brian Lester, Rami Al-Rfou, and Noah Constant. The Power of Scale for Parameter-Efficient Prompt
482 Tuning. In *EMNLP*, 2021.

483 Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay
484 Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, Yuhuai Wu, Behnam
485 Neyshabur, Guy Gur-Ari, and Vedant Misra. Solving Quantitative Reasoning Problems with
486 Language Models. In *NeurIPS*, 2022.

487 Xiang Lisa Li and Percy Liang. Prefix-Tuning: Optimizing Continuous Prompts for Generation. In
488 *ACL-IJCNLP*, 2021.

489 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike,
490 John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s Verify Step by Step. In *ICLR*, 2024.

491 Xiao Liu, Yanan Zheng, Zhengxiao Du, Ming Ding, Yujie Qian, Zhilin Yang, and Jie Tang. GPT
492 Understands, Too. *AI Open*, 5:208–215, 2024.

493 Aman Madaan and Amir Yazdanbakhsh. Text and Patterns: For Effective Chain of Thought, It Takes
494 Two to Tango. *arXiv:2209.07686*, 2022.

495 Aleksandar Petrov, Philip H. S. Torr, and Adel Bibi. When Do Prompting and Prefix-Tuning Work?
496 A Theory of Capabilities and Limitations. In *ICLR*, 2024.

497 Alexander Robey, Eric Wong, Hamed Hassani, and George J. Pappas. SmoothLLM: Defending Large
498 Language Models Against Jailbreaking Attacks. *Transactions on Machine Learning Research*,
499 2025.

500 Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang,
501 Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the Limits of
502 Mathematical Reasoning in Open Language Models. *arXiv:2402.03300*, 2024.

503 Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng,
504 Haibin Lin, and Chuan Wu. HybridFlow: A Flexible and Efficient RLHF Framework. In *EuroSys*,
505 2025.

506 Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov.
507 Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *Journal of Machine*
508 *Learning Research*, 15:1929–1958, 2014.

509 Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2
510 edition, 2018.

511 Shenzi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen,
512 Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen
513 Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 Rule: High-Entropy Minority Tokens Drive
514 Effective Reinforcement Learning for LLM Reasoning. In *NeurIPS*, 2025.

515 Yige Xu, Xu Guo, Zhiwei Zeng, and Chunyan Miao. SoftCoT: Soft Chain-of-Thought for Efficient
516 Reasoning with LLMs. In *ACL*, 2025.

517 An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu,
518 Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu,
519 Xingzhang Ren, and Zhenru Zhang. Qwen2.5-Math Technical Report: Toward Mathematical
520 Expert Model via Self-Improvement. *arXiv:2409.12122*, 2024.

521 Wengao Ye, Yan Liang, and Lianlei Shan. Thinking on the Fly: Test-Time Reasoning Enhancement
522 via Latent Thought Policy Optimization. In *ICLR*, 2026.

523 Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian
524 Fan, Gaohong Liu, Juncai Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming
525 Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Ru Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu,
526 Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Yuxuan Song, Xiangpeng Wei, Hao Zhou,
527 Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Yonghui Wu, and Mingxuan Wang. DAPO:
528 An Open-Source LLM Reinforcement Learning System at Scale. In *NeurIPS*, 2025.

529 Weihao Zeng, Yuzhen Huang, Qian Liu, Wei Liu, Keqing He, Zejun Ma, and Junxian He. SimpleRL-
530 Zoo: Investigating and Taming Zero Reinforcement Learning for Open Base Models in the Wild.
531 In *COLM*, 2025.

532 Guibin Zhang, Muxin Fu, and Shuicheng Yan. MemGen: Weaving Generative Latent Memory for
533 Self-Evolving Agents. In *ICLR*, 2026.

534 A Experimental Setup and Full Accuracy Breakdown

535 All experiments are conducted on a single NVIDIA A6000 40GB GPU with greedy decoding, a
536 maximum generation length of 3,072 tokens for MATH-500 and GSM8K, and 4,096 tokens for
537 AIME24. Batch size is fixed per model (16 for LLaMA-3.1-8B-Instruct and Qwen2.5-Math-7B
538 variants, 32 for Qwen2.5-Math-1.5B variants). The RSP token count L is selected per (model,
539 benchmark) pair from $\{10, 15, 20\}$ (Table 7); the per-position breakdown across *prefix* ($[\mathbf{H}; \mathbf{X}]$), *infix*
540 ($\mathbf{H}$ inserted within $\mathbf{X}$), and *suffix* ($[\mathbf{X}; \mathbf{H}]$) is in Table 6. The RSP injection harness used for these
541 experiments is included as anonymized supplementary material, with run commands documented in
542 the accompanying README.

Table 6: Accuracy (%) across model configurations, injection positions, and benchmarks. Values in parentheses indicate the difference from the baseline. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Mode	MATH-500	GSM8K	AIME24
<i>Instruct Models</i>				
LLaMA-3.1-8B-Instruct	Baseline	52.20	85.75	6.67
	<i>Prefix</i>	45.00 (−7.2)	76.35 (−9.4)	3.33 (−3.3)
	<i>Infix</i>	49.40 (−2.8)	85.97 (+0.2)	10.00 (+3.3)
	<i>Suffix</i>	49.40 (−2.8)	86.73 (+1.0)	10.00 (+3.3)
Qwen2.5-Math-7B-Instruct	Baseline	83.20	95.45	<u>13.33</u>
	<i>Prefix</i>	81.40 (−1.8)	94.92 (−0.5)	13.33 (+0.0)
	<i>Infix</i>	84.20 (+1.0)	95.60 (+0.2)	16.67 (+3.3)
	<i>Suffix</i>	<u>83.40</u> (+0.2)	95.68 (+0.2)	16.67 (+3.3)
Qwen2.5-Math-1.5B-Instruct	Baseline	73.00	85.29	10.00
	<i>Prefix</i>	74.80 (+1.8)	<u>85.29</u> (+0.0)	<u>13.33</u> (+3.3)
	<i>Infix</i>	75.20 (+2.2)	85.75 (+0.5)	6.67 (−3.3)
	<i>Suffix</i>	74.20 (+1.2)	84.69 (−0.6)	20.00 (+10.0)
<i>Base Models (Raw Text)</i>				
Qwen2.5-Math-7B	Baseline	72.20	84.15	6.67
	<i>Prefix</i>	<u>71.60</u> (−0.6)	84.31 (+0.2)	16.67 (+10.0)
	<i>Infix</i>	63.20 (−9.0)	64.29 (−19.9)	16.67 (+10.0)
	<i>Suffix</i>	55.60 (−16.6)	54.13 (−30.0)	<u>13.33</u> (+6.7)
Qwen2.5-Math-1.5B	Baseline	<u>61.40</u>	77.86	16.67
	<i>Prefix</i>	64.40 (+3.0)	74.30 (−3.6)	<u>10.00</u> (−6.7)
	<i>Infix</i>	63.20 (+1.8)	73.92 (−3.9)	<u>10.00</u> (−6.7)
	<i>Suffix</i>	54.60 (−6.8)	58.61 (−19.3)	3.33 (−13.3)
<i>Base Models + ChatML (Format Mismatch)</i>				
Qwen2.5-Math-7B	Baseline	52.20	58.30	23.33
	<i>Prefix</i>	59.20 (+7.0)	56.41 (−1.9)	<u>3.33</u> (−20.0)
	<i>Infix</i>	<u>62.00</u> (+9.8)	69.07 (+10.8)	23.33 (+0.0)
	<i>Suffix</i>	70.40 (+18.2)	75.97 (+17.7)	23.33 (+0.0)
Qwen2.5-Math-1.5B	Baseline	34.40	37.45	<u>13.33</u>
	<i>Prefix</i>	63.40 (+29.0)	67.02 (+29.6)	20.00 (+6.7)
	<i>Infix</i>	39.20 (+4.8)	50.42 (+13.0)	6.67 (−6.7)
	<i>Suffix</i>	<u>51.00</u> (+16.6)	<u>50.64</u> (+13.2)	<u>13.33</u> (+0.0)

Table 7: Number of RSP tokens (L) used for each model and benchmark.

Model	MATH-500	GSM8K	AIME24
LLaMA-3.1-8B-Instruct	15	20	15
Qwen2.5-Math-7B-Instruct	20	20	20
Qwen2.5-Math-1.5B-Instruct	20	10	20
Qwen2.5-Math-7B	10	10	20
Qwen2.5-Math-1.5B	10	10	20
Qwen2.5-Math-1.5B (ChatML)	20	20	10
Qwen2.5-Math-7B (ChatML)	20	20	20

543 A.1 RSP Injection Positions and Prompt Templates

544 Below we show the prompt structure for each injection position across all model types. Blue text
545 indicates RSP embeddings, and purple text indicates special tokens.

546 A.1.1 Prefix

547 RSP embeddings are concatenated before the prompt embeddings. No text modification is applied.

548 ChatML (Qwen Instruct / Base + ChatML).

```
[Random Embeddings (L tokens)]
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{>
<|im_start|>user
{question}<|im_end|>
<|im_start|>assistant
```

550 LLaMA Chat Template.

```
[Random Embeddings (L tokens)]
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{>
<|start_header_id|>user<|end_header_id|>
{question}<|eot_id|>
<|start_header_id|>assistant<|end_header_id|>
```

552 Raw Text (Qwen Base).

```
[Random Embeddings (L tokens)]
{question}
Please reason step by step, and put your final answer within \boxed{>
```

554 A.1.2 Infix

555 A description text and special tokens are inserted into the prompt. The special token positions are
556 then replaced with RSP embeddings at the embedding level.

557 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{>
<|im_start|>user
{question}

There are {L} special tokens that contain compressed latent reasoning information that
might be useful for your reasoning. If these tokens are useful for your case, you can
use them as reference. If these tokens are not useful for your case, you can ignore
them and focus back to solving the problem.

Here are the {L} special tokens: <special_token_1>...<special_token_L>
<|im_end|>
<|im_start|>assistant
```

559 LLaMA Chat Template.

```
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{>
<|start_header_id|>user<|end_header_id|>
```

{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens:

```
<reserved_special_token_0>...  
<reserved_special_token_L>  
<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

561

562 Raw Text (Qwen Base).

{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens: $\langle \text{special_token_1} \rangle \dots \langle \text{special_token_L} \rangle$

Please reason step by step, and put your final answer within `\boxed{\}`.

563

564 A.1.3 Suffix

565 For chat-templated models, RSP embeddings are inserted immediately before the end-of-turn token.

566 For raw text models, RSP embeddings are concatenated at the end of the prompt.

567 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system  
Please reason step by step, and put your final answer within \boxed{ }.<|im_end|>  
<|im_start|>user  
{question}[Random Embeddings ( $L$  tokens)]<|im_end|>  
<|im_start|>assistant
```

568

569 LLaMA Chat Template.

```
<|begin_of_text|>  
<|start_header_id|>system<|end_header_id|>  
Cutting Knowledge Date: December 2023  
Today Date: 26 Jul 2024  
Please reason step by step, and put your final answer within \boxed{ }.<|eot_id|>  
<|start_header_id|>user<|end_header_id|>  
{question}[Random Embeddings ( $L$  tokens)]<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

570

571 Raw Text (Qwen Base).

```
{question}  
Please reason step by step, and put your final answer within \boxed{ }. [Random  
Embeddings ( $L$  tokens)]
```

572

573 A.2 Answer Extraction and Grading

574 The final answer is extracted from the model output through a fallback chain. Each step is attempted

575 in order; if extraction fails, the next step is tried.

Answer Extraction Fallback Chain

1. "final answer is \$...\$. I hope" pattern (Minerva format)
2. `\boxed{...}`: last non-empty box is used; empty `\boxed{}` are skipped
3. Text following "the answer is"
4. Text following "final answer is"
5. Last number in the output (regex fallback)

Extracted answers are normalized by removing \$, \left, \right, and unit strings. Grading is performed via symbolic equivalence checking: exact string match, numerical comparison (`math.isclose, rel_tol=1e-4`), and SymPy symbolic simplification.

A.3 Reproduced Prior Work Hyperparameters

All prior methods are reproduced on Qwen2.5-Math-1.5B-Instruct and Qwen2.5-Math-7B-Instruct, evaluated with greedy decoding (temperature = 0) and our unified answer extraction pipeline.

TTSV. *Prefix* length 20, trained for 20 epochs with AdamW optimizer, batch size 2, gradient accumulation 8 steps, and linear learning rate schedule. Learning rate is 1e-3 for Qwen-1.5B and 1e-5 for Qwen-7B.

LTPO. Table 8 reports the per-benchmark hyperparameters.

Table 8: LTPO hyperparameters.

Parameter	GSM8K	MATH-500	AIME24
tokens	8	8	8
steps	20	20	20
top_k	10	10	10
sigma	20	20	5
sigma_decay	0.95	0.95	0.95
lr	1e-2	5e-3	5e-2

SoftCoT. Projection module trained for 10 epochs. Thought tokens: 32 during training, 4 during evaluation. Batch size 4 for GSM8K, 1 for MATH with gradient accumulation 4. The small (assistant) model is Qwen2.5-Math-1.5B-Instruct in both configurations, while the large model is Qwen2.5-Math-1.5B-Instruct (learning rate 2e-5) or Qwen2.5-Math-7B-Instruct (learning rate 5e-6). For MATH-500 and AIME24 evaluation, we use the projection module trained on GSM8K, as it yields higher accuracy than the one trained on MATH.

A.4 Full Entropy Curves

Figure 6 shows the full normalized (0–100%) generation trajectories for entropy, top-1 probability, and varentropy, and Table 9 reports the corresponding scalar statistics including overall means, first-5% means, and the average number of generated tokens per problem. All methods converge to similar levels after the initial generation stage, confirming that the distributional changes induced by RSP are localized to the early reasoning phase.

Table 9: Per-step statistics for Qwen2.5-Math-1.5B-Instruct on MATH-500 across the full generation and the first 5% of generation steps, together with the average number of generated tokens per problem. Top-1 Prob (overall) is reported as a weighted average over the first-10%/middle-last-10% segment means with weights 0.1/0.8/0.1. Extension of Figure 3 (main text) and Figure 6.

Metric	Baseline	Prefix	Infix	Suffix	LTPO	SoftCoT	TTSV
Tokens (mean)	575.7	558.4	549.9	573.3	592.3	594.4	577.4
Entropy (first 5%)	0.1332	0.1366	0.1402	0.1941	0.1662	0.4218	0.1359
Entropy (overall)	0.0871	0.0881	0.0899	0.1049	0.0909	0.1059	0.0864
Top-1 Prob (first 5%)	0.9535	0.9523	0.9525	0.9373	0.9437	0.9156	0.9534
Top-1 Prob (overall)	0.9684	0.9681	0.9678	0.9647	0.9683	0.9651	0.9685
Varentropy (first 5%)	0.2181	0.2207	0.2463	0.4010	0.2953	2.2234	0.2174
Varentropy (overall)	0.1353	0.1375	0.1422	0.2038	0.1452	0.2404	0.1361




Figure 6: Mean entropy (top), top-1 probability (middle), and varentropy (bottom) over the full normalized generation trajectory (0–100%). Averaged over MATH-500, shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: prior *soft prompt* methods. All methods converge after the initial stage.

B Token-Level Distribution Metrics

This appendix provides formal definitions for the metrics used in Section 4.4. Numerical results are reported in Table 3 of the main text. We extract first-token logits via a single forward pass (no autoregressive generation) and store the top-100 token ids and logit values per problem. Temperature conditions reuse the baseline logits scaled by $1/\tau$, requiring no separate inference. RSP draws 16 independent seeds per problem (MATH-500), and reported metrics are the mean across the 16 seeds. Pass@1 accuracy is computed over 16 samples per problem, with the baseline at $\tau=1.0$ reaching $73.92 \pm 0.78\%$. Let P_{base} denote the baseline next-token distribution at $\tau=1.0$ and P_{target} the candidate distribution being compared (e.g., baselines at $\tau=2.0, 3.0$ or RSP at $\tau=1.0$). All metrics are computed analytically from saved logits at the first generation step.

Spearman rank correlation. Spearman’s ρ is the Pearson correlation coefficient between the token ranks of two distributions:

$$\rho = \text{Pearson}(\text{rank}(P_{\text{target}}), \text{rank}(P_{\text{base}})). \quad (3)$$

Because temperature scaling $P^{(\tau)} \propto P^{1/\tau}$ is a strictly monotone transform, the rank order of tokens is preserved exactly, so $\rho = 1$ for any $\tau > 0$ as a mathematical identity. Any value $\rho < 1$ therefore indicates rank reorganization beyond what temperature scaling can produce.

Mass outside top- K . Let $\text{TopK}(P_{\text{base}})$ be the set of K tokens with the highest probability under P_{base} . The probability mass that the candidate distribution places *outside* this set is

$$\text{MassOutside}_K = 1 - \sum_{t \in \text{TopK}(P_{\text{base}})} P_{\text{target}}(t). \quad (4)$$

This directly measures support expansion: how much probability the candidate assigns to tokens that are not in the baseline’s preferred set. Reported values use $K = 10$.

Jensen–Shannon divergence. The symmetrized, bounded version of Kullback–Leibler divergence:

$$\text{JS}(P, Q) = \frac{1}{2} \text{KL}(P \parallel M) + \frac{1}{2} \text{KL}(Q \parallel M), \quad M = \frac{1}{2}(P + Q). \quad (5)$$

We use base-2 logarithms so that $\text{JS} \in [0, 1]$. Tokens absent from the stored top-100 are assigned a sentinel logit (-10^9) before softmax, which is negligible for our setting.

C Qualitative Analysis: How RSP Reshapes Reasoning Trajectories

This appendix complements the outcome-level Pass@ N results of §4.4 with a qualitative case study on AIME24, using the same $N=16$ sampling budget. We compare two of the four conditions defined in §4.4: *Baseline* (condition (i), $\tau=1.0$) and *RSP (indep. seed)* (condition (iv), $\tau=1.0$, *suffix*, $L=20$). Setting: Qwen2.5-Math-1.5B-Instruct, 30 AIME24 problems, 16 samples per problem per condition.

C.1 Aggregate Pattern Frequencies

Table 10 tabulates how RSP shifts coverage relative to Baseline across the 30 AIME24 problems.

Table 10: Coverage patterns across 30 AIME24 problems, $N=16$ samples per condition.

Pattern	Frequency	Problems
RSP-only correct (Baseline=0/16, RSP $\geq$ 1/16)	5/30 (16.7%)	#6, #8, #14, #22, #25
Baseline-only correct (Baseline $\geq$ 1/16, RSP=0/16)	0/30 (0%)	—
Both correct	6/30 (20.0%)	—
Both wrong	19/30 (63.3%)	—

Within the $N=16$ budget, RSP produces strict gains on five problems and never causes a strict regression. The two case studies below illustrate two faces of the same phenomenon: divergent solution framing (#6, §C.2) and divergent structural assumption (#22, §C.3). In both cases, RSP and Baseline diverge already in the opening sentences of the trajectory, and the divergence determines whether the trajectory reaches the correct answer.

633 **C.2 Case Study 1: AIME24 Problem 6 (Constrained Optimization)**

634 **Problem.**

Let $\mathcal{B}$ be the set of rectangular boxes with surface area 54 and volume 23. Let r be the radius of the smallest sphere that can contain each of the rectangular boxes that are elements of $\mathcal{B}$. The value of r^2 can be written as $\frac{p}{q}$, where p and q are relatively prime positive integers. Find $p + q$.

635
636 Ground truth: $\boxed{721}$, with the optimal box at $a = b = 2$, $c = 23/4$ giving $r^2 = 657/64$.

637 **Aggregate result.** Baseline produces 0/16 correct; RSP produces 1/16 correct. Baseline’s 16 wrong
638 final answers span 10 distinct values, with the modal answer 31 appearing in 4/16 samples; the correct
639 value 721 does not appear. RSP’s 16 samples produce 11 distinct values and include 721 exactly
640 once.

641 **Failure mode (Baseline).** Baseline’s representative trajectory (final answer 49) opens with an
642 *algebraic identity* frame and never solves for the actual box dimensions.

Baseline trajectory (final answer 49) — excerpts

Opening: To solve the problem, we start by defining the dimensions of a rectangular box as a , b , and c . The surface area and volume of the box are given by $2(ab + bc + ca) = 54$ and $abc = 23$. We need to find the radius r of the smallest sphere that can contain the box, given by $r = \frac{1}{2}\sqrt{a^2 + b^2 + c^2}$. First, we need to find the expression for $a^2 + b^2 + c^2$ in terms of the given equations.

Identity manipulation: Squaring the surface area equation yields $(ab + bc + ca)^2 = 2916$, which expands to $a^2b^2 + b^2c^2 + c^2a^2 + 2abc(a + b + c) = 2916$. Substituting $2abc = 46$ and $(a + b + c)^2 = x + 54$ where $x := a^2 + b^2 + c^2$, then attempting to express $a^4 + b^4 + c^4$ in terms of x , the manipulation collapses to the tautology $x^2 = x^2 - 2y + 2y$.

Termination by guess: “This equation is always true, so we need another way to find x After solving, we find $x = 45$.” Thus $r^2 = 45/4$, and $p + q = 45 + 4 = \boxed{49}$.

643
644 **Diagnosis.** The trajectory frames the task as *find the closed form of $a^2 + b^2 + c^2$ from the given con-*
645 *straints* and never attempts to solve for the actual dimensions. Two consecutive identity manipulations
646 collapse to a tautology that does not yield x , after which the trajectory commits to $x = 45$ without
647 justification (“After solving, we find”). This pattern — identity manipulation, tautology, terminal
648 guess — recurs across Baseline’s 16 samples; the constraint set under-determines $a^2 + b^2 + c^2$ in
649 closed form, and the algebraic-identity framing has no path to the correct value.

650 **Success mode (RSP).** The correct RSP trajectory (final answer 721) opens with an *optimization*
651 frame and a symmetry assumption that reduces the system to a single cubic.

RSP trajectory (final answer 721) — excerpts

Opening: To solve the problem, we need to find the dimensions of the rectangular box that *minimizes the radius of the* smallest sphere that can contain it, given the surface area and volume constraints. The radius is $r = \frac{1}{2}\sqrt{a^2 + b^2 + c^2}$. To minimize r , we minimize $a^2 + b^2 + c^2$.

Symmetry reduction: A more straightforward approach is to use symmetry: $a^2 + b^2 + c^2$ is minimized when a , b , c are as close to each other as possible. Assume $a = b$ and solve for c : $2a^2 + 2ac = 27$ and $a^2c = 23$. From $a^2c = 23$, $c = 23/a^2$. Substituting: $2a^2 + 46/a = 27$, multiplying through by a gives $2a^3 - 27a + 46 = 0$.

Rational root and termination: By the Rational Root Theorem, $a = 2$ is a root. Then $c = 23/4$, $a^2 + b^2 + c^2 = 4 + 4 + 529/16 = 657/16$, $r^2 = 657/64$, and $p + q = \boxed{721}$.

652
653 **Diagnosis.** Three properties of the opening drive the success: (i) the task is framed as *minimizing r*
654 (and thus $a^2 + b^2 + c^2$) rather than as a closed-form expression hunt; (ii) the path is committed to
655 solving for a, b, c instead of for $a^2 + b^2 + c^2$ alone; (iii) a symmetry assumption $a = b$ reduces the
656 system to one cubic in one unknown. The cubic $2a^3 - 27a + 46 = 0$ has $a = 2$ as a rational root,
657 leading to the explicit configuration $(2, 2, 23/4)$ and the correct $r^2 = 657/64$.

658 **Interpretation.** Baseline’s failure is rooted in the *opening sentence*: framing the task as “find the
659 expression for $a^2 + b^2 + c^2$ from the given equations” commits the trajectory to a path that does
660 not have a closed-form solution under the given constraints. RSP’s perturbation produces, in 1 of

661 16 attempts, a trajectory whose opening frame is *constrained optimization* with a symmetry-based
 662 variable reduction. The two trajectories share the same model weights, the same problem constraints,
 663 and the same sampling temperature; the input-side RSP perturbation is the only intervention, and its
 664 visible effect is the different framing the model commits to in the first few sentences.

665 C.3 Case Study 2: AIME24 Problem 22 (Constrained List Construction)

666 Problem.

A list of positive integers has the following properties:

- The sum of the items in the list is 30.
- The unique mode of the list is 9.
- The median of the list is a positive integer that does not appear in the list itself.

Find the sum of the squares of all the items in the list.

667

668 Ground truth: $\boxed{236}$, corresponding to $\{5, 7, 9, 9\}$ with median $\frac{7+9}{2} = 8$.

669 **Aggregate result.** Baseline produces 0/16 correct; RSP produces 3/16 correct. Baseline’s 16
 670 samples produce 16 distinct final answers (no two agree); the correct value 236 does not appear.
 671 RSP’s 16 samples concentrate slightly more, with 236 (3 times) and 248 (3 times) tied as modal
 672 answers and the remaining mass spread across 10 distinct values.

673 **Structural-assumption split.** The constraint “the median is a positive integer that does not appear
 674 in the list” *requires* an even-length list: for an odd-length list the median is by definition an element
 675 of the list. Across Baseline’s 16 samples, 6 commit to an odd-length list (typically 5 elements) at the
 676 opening and only 3 consider an even-length list. Across RSP’s 16 samples, 8 commit to odd and 6
 677 consider even. The shift in the even-length consideration rate ($3/16 \rightarrow 6/16$) is the structural correlate
 678 of RSP’s accuracy gain.

679 **Failure mode (Baseline).** Baseline’s representative trajectory (final answer 224) commits to an
 680 odd-length list *at the opening* and never reconsiders.

Baseline trajectory (final answer 224) — excerpts

Opening commitment: “Let’s start by assuming the list has an odd number of elements. If the list has $2k + 1$ elements, the median is the $(k + 1)$ -th element. Let’s try with the smallest possible odd number. . .”

Construction: 5 elements, $a_1 + a_2 + 9 + a_4 + 9 = 30$ so $a_1 + a_2 + a_4 = 12$. Try $a_1 = 2, a_2 = 3, a_4 = 7$: list $\{2, 3, 7, 9, 9\}$, median 7, “which does not appear in the list, so this is a valid candidate.”

Termination: Sum of squares = $4 + 9 + 49 + 81 + 81 = \boxed{224}$.

681

682 **Diagnosis.** Two errors compound: (i) the trajectory *defaults to odd length* at the opening without
 683 justification, and the constraint “median not in the list” forces even length, which the trajectory never
 684 derives; (ii) the median value 7 is verified as a positive integer but *not* verified as absent from the list
 685 (7 is in fact an element). The same structural assumption recurs across most of Baseline’s 16 samples;
 686 the elements vary, but the odd-length commitment and the incomplete constraint check are constant.

687 **Success mode (RSP).** The correct RSP trajectory (final answer 236) considers $n = 4$ from the
 688 outset and treats the median formula by parity.

RSP trajectory (final answer 236) — excerpts

Opening (Step 2: median by parity). “If the list has an odd number of elements, the median is the middle element. If the list has an even number of elements, the median is the average of the two middle elements. Since the median is a positive integer that does not appear in the list, it must be an integer and not equal to any of the other numbers in the list.”

Step 3: starting at $n = 4$. “ $n \geq 4$ (at least two 9’s). Case 1: $n = 4$. Sum: $9 + 9 + x + y = 30$, so $x + y = 12$.”

Construction: Try $x = 5, y = 7$: list $\{5, 7, 9, 9\}$, median $\frac{7+9}{2} = 8$ (integer, not in list). All conditions satisfied; sum of squares = $25 + 49 + 81 + 81 = \boxed{236}$.

689

Diagnosis. Three properties distinguish this trajectory: **(i)** *even length is considered first* ($n = 4$ enumerated as the initial case); **(ii)** the median is computed as an arithmetic mean for even length, producing $8 \notin \{5, 7, 9, 9\}$; **(iii)** the constraint check is performed against the candidate *value* (8), not against the structural position. The other two correct RSP samples reach 236 through partially incorrect intermediate reasoning, suggesting that some RSP samples that begin with the same odd-length assumption as Baseline can still arrive at the correct numerical answer.

Interpretation. Baseline’s failure is rooted in the *first structural assumption* after the conditions are listed: the trajectory commits to an odd-length list within its first few hundred tokens and never reconsiders. RSP’s perturbation increases the fraction of samples that consider $n = 4$ from 3/16 to 6/16, and one of these reaches the correct constraint check. The split between odd-length and even-length openings is the empirical fingerprint of the template-selection shift; on a small fraction of samples, the alternative template happens to be the structurally consistent one.

C.4 Summary of Case Studies

Both case studies exhibit the same shape. Baseline’s 16 samples cluster around a single failure template (algebraic-identity framing for #6; odd-length-list commitment for #22) and never break out of it within the sample budget. RSP does not change the model’s underlying knowledge: most RSP samples fall into the same templates as Baseline. What RSP changes is *which template the model commits to in the opening sentences*. On a small fraction of trajectories, the alternative template happens to admit a solvable subproblem (the optimization framing for #6, the even-length list for #22), and the trajectory reaches the correct answer. The accuracy lift ($0/16 \rightarrow 1/16$ on #6, $0/16 \rightarrow 3/16$ on #22) is the outcome-level signature of this opening-sentence divergence, consistent with the input-side branch-opening mechanism of §3.2: the perturbation is largest while the KV cache is short and the trajectory has not yet committed to a structural template.

D Prompt-Law Properties Used by the Theory

This appendix isolates the only properties of the RSP law used in the main text: centeredness, direction-agnostic covariance under a variance budget, and positive probability on every nonempty open set. We avoid stronger semantic claims because Section 3.2 needs only these geometric and measure-theoretic facts.

D.1 Centered Perturbations Do Not Impose Systematic First-Order Drift

Let z_a^{Base} denote the no-RSP output logit at vocab token a and $z_a^{\text{RSP}}(\bar{h})$ the corresponding logit when the centered RSP perturbation takes value $\bar{h}$. The (true) vocab-logit gap under RSP is $\Delta_{ab}(\bar{h}) := z_a^{\text{RSP}}(\bar{h}) - z_b^{\text{RSP}}(\bar{h})$, distinct from the no-RSP baseline gap $\Delta_{ab}^{\text{Base}} := z_a^{\text{Base}} - z_b^{\text{Base}}$. Note that $\bar{h} = 0$ corresponds to the centered RSP at its entrywise mean rather than to the no-RSP forward pass, so in general $\Delta_{ab}(0) \neq \Delta_{ab}^{\text{Base}}$. The transformer’s logit map is non-linear in $\bar{h}$, but a first-order Taylor expansion around $\bar{h} = 0$ yields the local surrogate

$$\Delta_{ab}^{\text{lin}}(\bar{h}) := \Delta_{ab} + b_{ab}^\top \bar{h},$$

with $\Delta_{ab} := \Delta_{ab}(0)$ and $b_{ab} := \nabla_{\bar{h}}(z_a^{\text{RSP}} - z_b^{\text{RSP}})|_{\bar{h}=0}$. The proposition below is stated for the surrogate Δ_{ab}^{lin} . For the true non-linear gap $\Delta_{ab}(\bar{h})$, the first-order term has zero mean under centered perturbations, but the Taylor remainder $r_{ab}(\bar{h}) = O(\|\bar{h}\|^2)$ may contribute a second-order mean shift; we make no claim about its sign or magnitude.

Proposition 2 (No systematic first-order drift in the surrogate). *If $\mathbb{E}[\bar{h}] = 0$, then*

$$\mathbb{E}[\Delta_{ab}^{\text{lin}}(\bar{h})] = \Delta_{ab}$$

for every token pair (a, b) . If a deterministic centered prompt $\bar{h}_0$ is reused across runs, then

$$\Delta_{ab}^{\text{lin}}(\bar{h}_0) - \Delta_{ab} = b_{ab}^\top \bar{h}_0$$

is a fixed directional bias.

732 *Proof.* The centered case is immediate from linearity of expectation:

$$\mathbb{E}[\Delta_{ab}^{\text{lin}}(\bar{h})] = \Delta_{ab} + b_{ab}^\top \mathbb{E}[\bar{h}] = \Delta_{ab}.$$

733 The deterministic case follows by substitution. $\square$

734 This proposition rules out a run-averaged first-order bias. It does not say that individual prompt draws
735 leave the logits unchanged.

736 D.2 Proof of Proposition 1

737 For a unit vector u , the first-order perturbation energy available along u is $\text{Var}_{h \sim D}(u^\top h) = u^\top \Sigma_D u$.
738 The worst-case directional energy is therefore the minimum Rayleigh quotient of Σ_D .

739 *Proof.* For any symmetric covariance matrix Σ_D ,

$$\min_{\|u\|=1} u^\top \Sigma_D u = \lambda_{\min}(\Sigma_D).$$

740 Let the eigenvalues of Σ_D be $\lambda_1, \dots, \lambda_d$. Then

$$\lambda_{\min}(\Sigma_D) \leq \frac{1}{d} \sum_{m=1}^d \lambda_m = \frac{\text{tr}(\Sigma_D)}{d} \leq \frac{\rho^2}{d}.$$

741 Equality holds if and only if all eigenvalues are equal, i.e., $\Sigma_D = (\rho^2/d)\mathbf{I}_d$. $\square$

742 The proof uses only covariance. We choose a Gaussian law for RSP because it adds the open-set
743 reachability property proved next.

744 D.3 Open-Set Reachability of Isotropic Gaussian Prompts

745 **Proposition 3** (Open-set reachability). *Let $\bar{h} \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_d)$ with $\sigma > 0$. Then every nonempty open*
746 *set $U \subseteq \mathbb{R}^d$ satisfies*

$$\Pr(\bar{h} \in U) > 0.$$

747 *Proof.* The Gaussian density

$$(2\pi\sigma^2)^{-d/2} \exp\left(-\frac{\|\bar{h}\|^2}{2\sigma^2}\right)$$

748 is strictly positive at every point of $\mathbb{R}^d$. Every nonempty open set contains a ball of positive Lebesgue
749 measure, and integrating a strictly positive density over that ball gives positive probability. $\square$

750 This is the only support property used in the informal branching argument of §3.2 and in Ap-
751 pendix E.3.3.

752 E Proofs for the Transformer-Level Mechanism

753 This appendix follows the same order as the main theory section: exploration, annealing, and
754 performance gain. The prompt-law facts used before these proofs are collected in Appendix D.

755 E.1 Exploration: attention decomposition and perturbation size

756 **Proposition 4** (Attention decomposition). *For one attention head at query position i in layer ℓ , let*
757 *$\alpha_{ij}^{(\ell)}$ be the attention weights over $n \geq 1$ unmasked real tokens and $L \geq 1$ random tokens, with finite*
758 *attention logits. Define*

$$w_{r,i}^{(\ell)} := \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)}$$

759 and, for $j \in [n]$,

$$\tilde{\alpha}_{ij}^{(\ell)} := \frac{\alpha_{ij}^{(\ell)}}{1 - w_{r,i}^{(\ell)}}.$$

760 Then

$$\mathbf{o}_i^{(\ell)} = (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{o}}_i^{(\ell)} + \boldsymbol{\eta}_i^{(\ell)},$$

761 where

$$\tilde{\mathbf{o}}_i^{(\ell)} := \sum_{j=1}^n \tilde{\alpha}_{ij}^{(\ell)} \mathbf{v}_j^{(\ell)} \quad \text{and} \quad \boldsymbol{\eta}_i^{(\ell)} := \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)}.$$

762 *Proof.* The head output is

$$\mathbf{o}_i^{(\ell)} = \sum_{j=1}^n \alpha_{ij}^{(\ell)} \mathbf{v}_j^{(\ell)} + \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)}.$$

763 By the softmax positivity and the existence of at least one unmasked real token, $\sum_{j=1}^n \alpha_{ij}^{(\ell)} =$
 764 $1 - w_{r,i}^{(\ell)} > 0$, so

$$\sum_{j=1}^n \alpha_{ij}^{(\ell)} \mathbf{v}_j^{(\ell)} = (1 - w_{r,i}^{(\ell)}) \sum_{j=1}^n \frac{\alpha_{ij}^{(\ell)}}{1 - w_{r,i}^{(\ell)}} \mathbf{v}_j^{(\ell)} = (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{o}}_i^{(\ell)}.$$

765 Substituting this identity gives the decomposition. □

766 **Corollary 1** (Magnitude scales with random-token attention). *For every realization of the random*
 767 *values,*

$$\|\boldsymbol{\eta}_i^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_{j \in [L]} \|\mathbf{v}_{r_j}^{(\ell)}\|.$$

768 *Proof.* By the triangle inequality,

$$\|\boldsymbol{\eta}_i^{(\ell)}\| = \left\| \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)} \right\| \leq \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \|\mathbf{v}_{r_j}^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_{j \in [L]} \|\mathbf{v}_{r_j}^{(\ell)}\|.$$

769 □

770 Corollary 1 is the only quantitative fact needed in the main text. No independence assumption
 771 between attention weights and random keys is required. Once $w_{r,i}^{(\ell)}$ is small, the random contribution
 772 must be small as well.

773 E.2 Corollaries of Theorem 1 on KV cache growth

774 The fixed-gap decay bound yields the corollaries below.

775 **Corollary 2** (Sequence-wise annealing of the fixed-gap envelope). *For fixed L and $\Delta_i^{(\ell)}$, the upper*
 776 *bound*

$$f(n) = \frac{L}{L + n \exp(\Delta_i^{(\ell)} / \sqrt{d_{\text{head}}})}$$

777 *is strictly decreasing in n and tends to 0 as $n \rightarrow \infty$. KV cache growth alone therefore narrows the*
 778 *largest RSP attention mass compatible with that gap.*

779 During autoregressive decoding $\Delta_i^{(\ell)}$ also moves because queries, keys, and hidden states co-evolve.
 780 The accurate reading is not “the realized mass decreases at every step” but “the envelope compatible
 781 with a fixed gap shrinks as the KV cache grows.”

782 *Proof.* Differentiate $f(n)$ with respect to n :

$$f'(n) = -\frac{L \exp(\Delta_i^{(\ell)} / \sqrt{d_{\text{head}}})}{\left(L + n \exp(\Delta_i^{(\ell)} / \sqrt{d_{\text{head}}})\right)^2} < 0.$$

783 Thus $f(n)$ is strictly decreasing, and its denominator diverges linearly in n while its numerator is
784 fixed, so $f(n) \rightarrow 0$. $\square$

785 **Corollary 3** (Vanishing random contribution under annealing). *If $n \rightarrow \infty$ while L and $\Delta_i^{(\ell)}$ remain
786 fixed, $w_{r,i}^{(\ell)} \rightarrow 0$, and for every realization of the random values*

$$\|\boldsymbol{\eta}_i^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_{j \in [L]} \|\mathbf{v}_{r_j}^{(\ell)}\| \rightarrow 0.$$

787 *Proof.* $w_{r,i}^{(\ell)} \rightarrow 0$ by Corollary 2. Apply Corollary 1 to bound $\|\boldsymbol{\eta}_i^{(\ell)}\|$. The maximum norm is finite
788 for any realization of finitely many sampled values. $\square$

789 E.3 Performance gain: from local admission to Pass@N

790 From here we work at the vocab-logit and task levels, dropping the per-layer index ℓ . The main text
791 uses a first-order surrogate to state two claims: the §3.2 claim that a baseline-excluded vocab token
792 can enter the local top- K set, and the §3.3 closing that independent reruns turn such local openings
793 into Pass@N gains. The proofs below use only one support condition on the centered prompt law D :

$$(S1) \quad D(U) > 0 \quad \text{for every nonempty open set } U \subseteq \mathbb{R}^d.$$

794 Proposition 3 shows that the isotropic Gaussian law of §3.2 satisfies (S1).

795 E.3.1 Autoregressive Formulation

796 In autoregressive decoding, the joint distribution over a sequence $y_{1:T}$ factorizes as

$$\Pr(y_{1:T} \mid x) = \prod_{t=1}^T p(y_t \mid x, y_{<t}),$$

797 where each step is governed by a prefix-conditioned distribution. Any perturbation introduced by
798 RSP affects generation through a sequence of local changes to $p(y_t \mid x, y_{<t})$, rather than a single
799 global distribution.

800 E.3.2 Local Linearization of Logits under RSP

801 Throughout this subsection, $\bar{h} \in \mathbb{R}^d$ denotes *one* centered prompt vector from the RSP definition
802 (Eq. (1) in §3.1), treated as a per-token surrogate. The actual implementation uses a length- L prompt
803 $\mathbf{H} = (h_1, \dots, h_L)$ with L independent draws; the arguments below isolate how one local perturbation
804 can shift the logits. We model the perturbed logits as a function $z_a(\bar{h})$ at vocab token a , assuming
805 local smoothness in a neighborhood of $\bar{h} = 0$:

$$z_a(\bar{h}) = z_a(0) + \nabla_{\bar{h}} z_a(0)^\top \bar{h} + r_a(\bar{h}),$$

806 where $r_a(\bar{h}) = O(\|\bar{h}\|^2)$. Defining $c_a := \nabla_{\bar{h}} z_a(0)$ and $z_a := z_a(0)$, we obtain the local affine
807 surrogate

$$z_a^{\text{lin}}(\bar{h}) := z_a + c_a^\top \bar{h}.$$

808 This approximation is used as a first-order surrogate for branch opening, not as an exact global model
809 of the transformer logits.

810 E.3.3 Top- K entry region (Remark)

811 This complements the informal branching argument in §3.2. Define the entry region

$$\mathcal{G}_{a,K} := \{\bar{h} \in \mathbb{R}^d : \#\{b \neq a : z_b^{\text{lin}}(\bar{h}) > z_a^{\text{lin}}(\bar{h})\} \leq K - 1\}.$$

812 If $\mathcal{G}_{a,K}$ contains a nonempty open set U , then by (S1) $D(U) > 0$, and $U \subseteq \mathcal{G}_{a,K}$ gives $\Pr_{\bar{h}}(\bar{h} \in$
813 $\mathcal{G}_{a,K}) \geq D(U) > 0$. A useful sufficient condition is the existence of some $\bar{h}_0$ at which token a strictly
814 outranks all but at most $K - 1$ other tokens in the affine surrogate; by continuity, a neighborhood of
815 $\bar{h}_0$ is then contained in $\mathcal{G}_{a,K}$. This is a rank-changing event unreachable by the monotone temperature
816 transform, but it does not by itself imply (M1).

817 E.3.4 Pass@ N ensemble bound (Remark)

818 The bound $\Pr(\text{Pass}@N \text{ on } x) \geq 1 - (1 - p_{\min}(x))^N$ in the main text is a standard independent-
819 Bernoulli union argument. Let A_n be the event that the n -th run reaches a correct trajectory on x ;
820 under (M1) $\Pr(A_n) \geq p_{\min}(x)$, and under (M2)

$$\Pr\left(\bigcap_{n=1}^N A_n^c\right) \leq (1 - p_{\min}(x))^N, \quad \Pr\left(\bigcup_{n=1}^N A_n\right) \geq 1 - (1 - p_{\min}(x))^N \geq 1 - e^{-Np_{\min}(x)}.$$

821 This is a generic fact, not specific to RSP; RSP’s role is to make branch opening possible and
822 to support (M2) through independent resampling, while whether a branch satisfies the task-side
823 correctness condition (M1) remains task-dependent.

824 F Attention Mass Analysis

825 This appendix formalizes the per-token attention mass reported in Figure 2 and the exclusion criteria
826 applied to the reasoning span and the layer axis. For each of the 500 MATH-500 problems, a fresh
827 RSP $\mathbf{H} \in \mathbb{R}^{L \times d}$ with $L = 10$ is sampled independently, so the reported heatmaps average over both
828 problem variability and RSP-vector variability.

829 F.1 Per-Token Attention Mass

830 For each problem, we measure attention on the concatenated sequence [prompt; $\mathbf{H}$; generation], where
831 the generation is the trajectory produced by the same model under matching RSP injection. Let
832 $\alpha_{i,j}^{(\ell,m)}$ denote the post-softmax attention weight at layer ℓ , head index m (using m to avoid collision
833 with the vocab token a of §3.2 and the RSP matrix $\mathbf{H}$), from query position i to key position j ,
834 satisfying $\sum_j \alpha_{i,j}^{(\ell,m)} = 1$ under the causal mask. Let Q and R denote the index sets of question
835 and RSP tokens with sizes $|Q|$ and $|R| = L$, and let N_{head} be the number of heads. The per-token
836 attention mass that query i directs to each group at layer ℓ is

$$\text{PTM}_Q[\ell, i] = \frac{1}{|Q| N_{\text{head}}} \sum_{m=1}^{N_{\text{head}}} \sum_{j \in Q} \alpha_{i,j}^{(\ell,m)}, \quad \text{PTM}_R[\ell, i] = \frac{1}{|R| N_{\text{head}}} \sum_{m=1}^{N_{\text{head}}} \sum_{j \in R} \alpha_{i,j}^{(\ell,m)}. \quad (6)$$

837 Normalization by $|Q|$ and $|R|$ controls for group size, so both quantities express how much attention a
838 single question (resp. RSP) token receives on average under the same softmax denominator. Question
839 tokens thereby serve as a size-matched baseline that absorbs sequence-length effects common to both
840 groups.

841 F.2 Heatmap Aggregation

842 We partition the layer indices and the reasoning position indices each into five equal-size bins
843 $\{B_L^{(b)}\}_{b=1}^5$ and $\{B_P^{(b)}\}_{b=1}^5$. For sample s and target group $\bullet \in \{Q, R\}$, the value in cell (b_L, b_P) is

$$M_{\bullet}^{(s)}[b_L, b_P] = \frac{1}{|B_L^{(b_L)}| \cdot |B_P^{(b_P)}|} \sum_{\ell \in B_L^{(b_L)}} \sum_{i \in B_P^{(b_P)}} \text{PTM}_{\bullet}[\ell, i]. \quad (7)$$

844 Figure 2 reports the sample average of Eq. (7) over the 500 valid samples.

F.3 Excluding the `\boxed{...}` Span

The reasoning position range terminates strictly before the final `\boxed{...}` span. Chain-of-thought outputs decompose into a *text* (content) component and a *pattern* (template) component that contribute independently to downstream performance [Madaan and Yazdanbakhsh, 2022]. The `\boxed{...}` wrapper is a canonical pattern: a short, stereotyped span whose role is to complete the answer format rather than to extend reasoning. Including it would therefore inject a template-driven attention signature unrelated to reasoning dynamics.

F.4 Excluding the Final Two Layers

The layer axis further excludes the last two transformer layers. This choice is motivated by a standard late-layer effect rather than by any model-specific tuning decision.

Output-projection specialization. Late-layer hidden states lie in an approximately linear relationship with vocabulary logits and therefore function as a progressive linear readout rather than as representation-transforming blocks [Belrose et al., 2023]. Consistently, late-layer feed-forward modules have been characterized as output-distribution shaping mechanisms [Geva et al., 2021]. Attention in these layers therefore reflects output formation rather than reasoning computation.

Excluding the final two layers keeps the heatmap focused on reasoning-phase dynamics instead of readout-phase dynamics. This exclusion is not load-bearing for the main claim: the sequence-direction attenuation emphasized in the main text is also visible without it, and Theorem 1 concerns sequence-direction attenuation rather than layer-wise monotonicity.

F.5 Additional Suffix Heatmaps

For the late-layer reasons discussed above, the pattern in Figure 7 does not generalize uniformly across every cell; nevertheless, the overall trend is consistent across both models: question-token attention varies broadly across layers, whereas RSP attention shows sequence-wise decay consistent with Theorem 1 alongside an empirical layer-wise attenuation that the theorem itself does not predict.




Figure 7: Per-token RSP attention mass under *suffix* injection for the remaining two models (500 MATH-500 samples each). Axes and preprocessing match Figure 2: x-axis is reasoning position binned into five quantiles, y-axis is layer depth binned into five quantiles (top = shallow), color encodes the 500-sample mean per-token attention assigned to RSP tokens, and tokens inside `\boxed{...}` spans and the last two layers are excluded.

G DAPO Implementation

This appendix complements Section 4.5 with the DAPO loss modifications relative to GRPO, the DAPO + RSP implementation, and full per-step results. The training code is included in the anonymized supplementary material; pretrained DAPO and DAPO + RSP checkpoints across training steps 10–100 will be released at the camera-ready stage.

G.1 From GRPO to DAPO

The vanilla GRPO objective [Shao et al., 2024] for a group of G rollouts $\{y_i\}_{i=1}^G$ (each $y_i = (y_{i,1}, \dots, y_{i,|y_i|})$ a token sequence; we use y rather than o throughout this appendix to avoid collision with the head output $\mathbf{o}_i^{(\ell)}$ of §3.2) with rewards $\{R_i\}_{i=1}^G$ is

$$\mathcal{J}_{\text{GRPO}}(\theta) = \mathbb{E} \left[\frac{1}{G} \sum_{i=1}^G \frac{1}{|y_i|} \sum_{t=1}^{|y_i|} \left(\min(r_{i,t} \hat{A}_{i,t}, \text{clip}(r_{i,t}, 1-\varepsilon, 1+\varepsilon) \hat{A}_{i,t}) - \beta D_{\text{KL}}(\pi_\theta \| \pi_{\text{ref}}) \right) \right], \quad (8)$$

with importance ratio $r_{i,t} = \pi_\theta(y_{i,t} | q, y_{i,<t}) / \pi_{\theta_{\text{old}}}(y_{i,t} | q, y_{i,<t})$, group-relative advantage $\hat{A}_{i,t} = (R_i - \text{mean}(\{R_i\})) / \text{std}(\{R_i\})$, clipping range ε , and KL penalty βD_{KL} against the reference π_{ref} .

We build on the SimpleRL-Zoo [Zeng et al., 2025] training pipeline and implement DAPO on top of VeRL [Sheng et al., 2025]. DAPO modifies Eq. (8) in four ways: (i) token-level loss aggregation $1 / \sum_i |y_i|$ instead of $1 / |y_i|$, (ii) asymmetric clipping with $(\varepsilon_{\text{low}}, \varepsilon_{\text{high}}) = (0.2, 0.28)$, (iii) an implementation-side negative-advantage dual-clipping safeguard $\max(L_{\text{clipped}}, c\hat{A})$ with $c = 10$ (active only when $\hat{A} < 0$; omitted from Eq. (9) below for readability), and (iv) KL terms removed. The resulting DAPO objective is

$$\mathcal{J}_{\text{DAPO}}(\theta) = \mathbb{E} \left[\frac{1}{\sum_i |y_i|} \sum_{i,t} \min(r_{i,t} \hat{A}_{i,t}, \text{clip}(r_{i,t}, 1-\varepsilon_{\text{low}}, 1+\varepsilon_{\text{high}}) \hat{A}_{i,t}) \right]. \quad (9)$$

G.2 DAPO + RSP Implementation

For DAPO + RSP, we additionally inject a per-rollout RSP $\mathbf{H}_i \in \mathbb{R}^{L \times d}$ ($L = 20$) into each rollout. Each $\mathbf{H}_i$ is generated deterministically from a per-step rollout seed and injected via a forward hook; it is not a trainable parameter, so the learning signal is the policy adapting to arbitrary $\mathbf{H}_i$ rather than a learned prompt. Including $\mathbf{H}_i$ in the log-probabilities at both rollout and update time keeps the update on-policy with respect to the RSP-conditioned distribution (avoiding off-policy mismatch). With $G = 8$ rollouts per prompt, each rollout receives an independently drawn $\mathbf{H}_i$.

G.3 Evaluation Protocol

Each saved checkpoint is converted to HuggingFace format and evaluated using the SimpleRL-Zoo `eval_math_nodes.sh` pipeline, which selects sampling parameters per benchmark to reduce variance on the smaller competition sets:

Table 11: Per-benchmark evaluation protocol used by SimpleRL-Zoo `sh/eval.sh`.

Benchmark	# Problems	Temperature	n_{sampling}	Metric
AIME 2024	30	1.0	32	Avg@32
AMC 2023	40	1.0	32	Avg@32
MATH-500	500	0.0	1	accuracy (mean@1)
GSM8K	1,319	0.0	1	accuracy
OlympiadBench	675	0.0	1	accuracy
Minerva Math	272	0.0	1	accuracy
GaoKao 2023-EN	385	0.0	1	accuracy

The maximum number of generated tokens is 16,000 across all benchmarks. The reported average is the unweighted mean over the five benchmarks (GSM8K, MATH-500, College Math, Minerva Math, AIME 2024).

900 G.4 Data, Batch, and Hyperparameters

Table 12: DAPO / DAPO + RSP training setup on Qwen2.5-Math-7B.

Field	Value
Train data	MATH Level 3–5 subset (simple1r_math_35, ~8,500 prompts)
Prompt template	qwen-boxed
Train batch size	1024
PPO mini-batch size	256 (4 PPO updates per rollout)
Rollouts per prompt G	8
Max prompt / response length	2,028 / 2,048
Rollout sampling	temperature 1.0, top- p 1.0, top- k 50
$(\epsilon_{\text{low}}, \epsilon_{\text{high}})$	(0.2, 0.28)
Dual clip c	10.0
Loss aggregation	token-mean ($1 / \sum_i y_i $)
KL terms	disabled
Entropy coefficient	0
Optimizer	AdamW, learning rate 5×10^{-7} (constant), no warmup
Total training	~10 epochs, ≥ 80 optimization steps
Save / eval frequency	every 10 steps
RSP (DAPO + RSP only)	suffix injection, $L = 20$, freshly drawn per rollout
Hardware	4 $\times$ NVIDIA B200 GPUs
Wall-clock training time	~12 hours per run

901 G.5 Per-Step Average Accuracy

Table 13: Five-benchmark average accuracy (%) at every checkpoint, computed over GSM8K, MATH-500, College Math, Minerva Math, and AIME 2024. **Bold** marks each method’s peak.

Step	DAPO	DAPO + RSP
0	30.06	30.06
10	46.40	46.22
20	49.02	49.54
30	49.58	50.72
40	50.62	52.06
50	50.88	52.40
60	51.84	52.96
70	52.52	53.86
80	53.20	54.04
90	52.52	54.36
100	50.54	53.64

902 H Broader Impacts

903 This work provides empirical and theoretical analysis of how random embedding injection affects
 904 the reasoning behavior of large language models. The contributions are primarily foundational,
 905 with potential downstream implications for reinforcement learning-based post-training of reasoning
 906 models such as GRPO.

907 **Positive impacts.** RSP-induced trajectory diversity could improve the sample efficiency of RL-
 908 based LLM post-training by providing richer learning signals within group-relative advantage estima-
 909 tion. This has the potential to reduce compute costs associated with alignment and reasoning-oriented
 910 fine-tuning, making such methods more accessible to research groups with limited resources.

911 **Potential concerns.** Methods that expand the effective output support of LLMs may increase the
 912 reachability of low-probability tokens, which includes both desirable alternative reasoning paths
 913 and potentially unsafe or off-distribution outputs. Practitioners applying RSP in production systems
 914 should combine it with existing safety filtering and alignment mechanisms rather than treating it as a
 915 standalone diversity source.

916 **Limitations on impact scope.** Since this work focuses on analysis rather than deploying a new
917 model or dataset, the direct societal impact is limited. The broader implications described above are
918 contingent on future work integrating RSP into applied training pipelines.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and §1 advance four scoped claims, each backed by a specific theorem, table, or figure in the paper. The unifying message is that the *simplest possible soft prompt*—training-free, freshly resampled per rollout, carrying no learned content—already reproduces the entropy-shaping signature and Pass@ N benefit of optimized methods, while admitting a clean theoretical account.

(i) The per-layer RSP contribution decomposes exactly (Eq. (2)) into a real-token term and a contribution governed by a single attention-mass scalar $w_{r,i}^{(\ell)}$, whose fixed-gap upper bound is shown in Theorem 1 to shrink monotonically with KV cache growth, yielding an implicit explore-then-exploit dynamic without any scheduler (§3.3).

(ii) Proposition 1 shows that isotropic Gaussians uniquely maximize worst-case directional variance at a fixed second-moment budget; combined with full support, this lets RSP reach top- K entry regions inaccessible to any rank-preserving temperature transform. Empirically, RSP at $\tau = 1.0$ produces Spearman $\rho = 0.709$ against the $\tau = 1.0$ baseline first-token distribution (vs $\rho = 1$ for any temperature) while preserving 73.30% Pass@1 (vs 1.42% at the $\tau = 3.0$ that would be needed to match RSP’s distributional spread; Table 3).

(iii) RSP shares the early-decoding entropy↑/varentropy↑ signature of LTPO, SoftCoT, and TTSV—three methods trained under *different* objectives—supporting the interpretation that part of what trained *soft prompts* deliver is a structural by-product of injection rather than of learned content (Figure 3). Pass@ N accumulates only under *independent* per-rollout resampling: at $N=16$, independent-seed RSP at $\tau=1.0$ reaches 92.80% on MATH-500 and 36.67% on AIME 2024, while a shared single-seed RSP shows no improvement over the temperature-only baseline (Figure 4).

(iv) The same input-side perturbation composes with DAPO’s loss-side diversity preservation: DAPO + RSP attains a five-benchmark average of 54.36% (vs DAPO’s 53.20% peak; +1.16 pp at peak, +3.10 pp at step 100), with the post-peak DAPO decline delayed (Table 5, Figure 5).

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Limitations are noted across the paper rather than in a single dedicated section, with future-work directions consolidated in §5. (a) The Pass@ N accumulation argument requires the task-side assumption $p_{\min}(x) > 0$, explicitly flagged in §3.3; the isotropy and full-support arguments in §3.2 cover only the local branch-opening step. (b) The +18–29 pp recovery on the format-mismatch rows of Table 1 is reported as an empirical observation, with the EOS-derailment hypothesis stated informally in §4.2 rather than derived from any theorem in §3. (c) Evaluation rigor: RSP length $L \in \{10, 15, 20\}$ and injection position are selected per model on the evaluation set without held-out validation, and Table 1 together with the DAPO + RSP run uses a single seed; multi-seed evaluation and held-out validation are listed as future work in §5 (iii)(iv) and discussed in item 7. (d) Formal-result scope: Theorem 1 guarantees only sequence-axis attenuation at fixed (layer, query, gap), with the layer-axis decay in Figure 2 outside its envelope (§5 (ii)); Proposition 1 is a design criterion stated in input space and does not claim that isotropy survives transformer layers (§3.2 footnote). (e) The analysis is confined to RoPE-based decoders and mathematical reasoning; whether the same mechanism transfers to other position encodings (ALiBi, NoPE) and domains (code, open-ended reasoning) is unverified and listed as future work (§5 (i)).

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The paper contains four formal theoretical results, each with explicit assumptions and a proof. (i) Equation (2) (head-output decomposition into a renormalized real-token term and an RSP-induced contribution) carries a one-line in-text derivation in §3.2, using only the softmax-normalization identity under finite attention logits with $n \geq 1$ real and $L \geq 1$ RSP tokens. (ii) Theorem 1 (KV cache-growth attenuation) is proved inline in §3.3, assuming scaled dot-product attention with $n \geq 1$ real keys, $L \geq 1$ RSP keys, finite logits, and the pre-scale attention-logit gap Δ_i as defined in the statement. (iii) Proposition 1 (maximin directional coverage) is stated in §3.2 with its full proof in Appendix D, under the second-moment budget $\text{tr}(\Sigma_D) \leq \rho^2$ on zero-mean distributions over $\mathbb{R}^d$. (iv) The Pass@ N ensemble bound $\Pr(\text{Pass}@N \text{ on } x) \geq 1 - (1 - p_{\min}(x))^N$ is stated in §3.3 and proved formally in Appendix E.3.4, requiring (M1) the task-side condition $p_{\min}(x) > 0$ and (M2) mutual independence of the N rollouts. The supporting top- K entry argument in §3.2, which combines full support of the RSP law with a locally affine surrogate $\Delta_{ab}(\bar{h}) \approx \Delta_{ab}(0) + b_{ab}^\top \bar{h}$ around $\bar{h} = 0$, is presented informally in the main text and written out formally in Appendix E.3.3.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: All hyperparameters required to reproduce the main experiments are disclosed in §4.1, §4.4, §4.5, and Appendices A–G. For the inference experiments, we report the decoding strategy (greedy for Tables 1–2; $\tau = 1.0$ sampling for the diversity and Pass@ N experiments in §4.4), the RSP length $L \in \{10, 15, 20\}$ together with its per-model selection protocol on the evaluation set (Tables 1–2) and the fixed $L = 10$, *suffix* configuration used in the mechanism analyses (§4.3–§4.4), the injection position grid $\{\text{prefix}, \text{suffix}\}$ (with per-position breakdown in Appendix A, and the exact prompt templates for all three positions $\{\text{prefix}, \text{infix}, \text{suffix}\}$ across the ChatML, LLaMA, and raw-text formats in Appendix A.1), the number of samples per problem (16 for the token-level and Pass@ N analyses), the number of RSP seeds (16 for the token metrics in Table 3), and the trajectory-level setting (64 problems with 300 independent trajectories per condition encoded by BGE-M3 [Chen et al., 2024] and clustered by DBSCAN [Ester et al., 1996]). The unified answer-extraction pipeline (Appendix A.2) and per-benchmark generation-length budgets are also reported. For the DAPO + RSP training experiments, Appendix G discloses the optimizer, the prompt batch (1,024 prompts), the per-prompt rollout count ($G = 8$), the rollout temperature ($\tau = 1.0$), the four PPO mini-batch updates of 256 prompts each, the total number of training steps (100), the asymmetric clipping range $(\varepsilon_{\text{low}}, \varepsilon_{\text{high}}) = (0.2, 0.28)$, the dual-clipping safeguard, the suffix RSP length used during rollouts ($L = 20$), the rollout-only application protocol (no RSP at evaluation), and the dataset (a Level-3–5 subset of MATH with $\sim 8.5\text{K}$ prompts). The five-benchmark evaluation protocol (GSM8K, MATH-500, College Math, Minerva Math, AIME 2024) specifies Avg@32 at $\tau = 1.0$ for AIME 2024 and greedy decoding elsewhere. The inference batch sizes are also disclosed (16 for LLaMA-3.1-8B-Instruct, 32 for the Qwen2.5-Math 1.5B variants).

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All evaluation data and base models used in the paper are publicly available and cited at first use. Beyond the description in the main text (Eq. (1), §4.1, Appendix G), an anonymized implementation of the RSP injection harness and the DAPO + RSP training pipeline is included as supplementary material, with a README detailing the run commands for the inference experiments in §4.2–§4.4 and the DAPO + RSP training in §4.5. Pretrained DAPO and DAPO + RSP checkpoints across training steps 10–100 will be released at the camera-ready stage to allow reviewers and downstream users to reproduce the main DAPO + RSP results in §4.5 without re-running the training stage.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Beyond the hyperparameter values disclosed for reproducibility (item 4), this item details the data splits, hyperparameter selection protocol, optimizer choices, and prior-work hyperparameters. **Data splits.** The inference experiments (§4.2–§4.4) involve no training and use the standard public test sets of MATH-500, GSM8K, and AIME 2024. The DAPO + RSP training (§4.5) uses a Level-3–5 subset of MATH (`simplelr_math_35`, ~8.5K prompts) for training and evaluates on five held-out benchmarks (GSM8K, MATH-500, College Math, Minerva Math, AIME 2024) under the per-benchmark sampling protocol of Appendix G (Table 11). **Selection protocol.** The RSP token count $L \in \{10, 15, 20\}$ is selected per (model, benchmark) on the evaluation set, with the per-row choice in Appendix A (Table 7); the injection position is reported as a grid (*prefix*, *infix*, *suffix*) in Appendix A (Table 6) rather than tuned, and the mechanism analyses in §4.3–§4.4 fix $L = 10$ and *suffix* to insulate them from this selection effect. The absence of held-out validation is acknowledged in item 2 and §5 (iv). **Optimizer.** DAPO + RSP training uses AdamW with a constant learning rate of 5×10^{-7} and no warmup; total training is approximately 10 epochs with checkpoints saved every 10 steps, as listed in Appendix G (Table 12). **Prior work hyperparameters.** The reproduced TTSV, SoftCoT, and LTPO baselines used in Table 2 are documented with their full training and evaluation hyperparameters in Appendix A.3: TTSV (prefix length, AdamW, batch size, gradient accumulation, learning-rate schedule, per-model learning rates), LTPO (per-benchmark token count, optimization steps, top- k , sigma and decay, learning rate; Table 8), and SoftCoT (projection-module training schedule, thought-token counts during training and evaluation, batch size, gradient accumulation, and per-model learning rates).

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: Variability is reported for the experiments most directly affected by sampling stochasticity. Table 3 reports Pass@1 as mean $\pm$ standard deviation over 16 samples per problem, with the token-level metrics additionally averaged over 16 independent RSP seeds. Figure 3 draws shading at ± 1 SEM across MATH-500 problems. The Pass@ N analysis in Figure 4 uses $N = 16$ samples per problem on MATH-500 and AIME 2024. The trajectory-level analysis (Table 4) is computed over 64 problems with 300 independent trajectories per condition and additionally reports the per-problem majority (RSP yields larger pairwise distance than baseline on 56 of 64 problems). On the DAPO + RSP evaluation side, AIME 2024 uses Avg@32 at $\tau = 1.0$ to reduce competition-set variance. By contrast, Table 1 and Table 2 report single-run greedy-decoded accuracies without bootstrap confidence intervals or paired tests, and the DAPO + RSP training (Table 5, Figure 5) is run with a single training seed. Multi-seed evaluation across both inference and training, particularly on the small competition sets (AIME 2024, AMC 2023) where the per-problem variance is highest, is identified as future work in §5 (iii).

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All inference experiments (Tables 1–4, Figures 3–4) run on a single NVIDIA A6000 40GB GPU with greedy decoding (Appendix A). DAPO + RSP training (§4.5) runs on $4 \times$ NVIDIA B200 GPUs for approximately 12 hours per run (Appendix G, Table 12, hardware and wall-clock rows).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work uses publicly available math reasoning benchmarks and pretrained language models, involves no human subjects or personal data, and does not introduce a release that bypasses standard model-safety considerations. We have reviewed the NeurIPS Code of Ethics and identify no conflicts.

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Appendix H discusses both directions: positive impact via more sample-efficient RL post-training of reasoning LLMs, and a potential downside that broadening the effective output support also increases reachability of undesirable trajectories, which we argue should keep RSP coupled with existing safety filtering rather than replacing it.

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The DAPO and DAPO + RSP checkpoints (to be released at the camera-ready stage, item 5) are math-domain fine-tunes of the publicly available Qwen2.5-Math-7B [Yang et al., 2024] on a Level-3–5 subset of MATH and do not extend the base model’s capabilities beyond mathematical reasoning, so they inherit the upstream safety profile without introducing new misuse vectors. RSP itself is a training-free generation-time perturbation applied to existing pretrained models and adds no separate artifact requiring safeguards. The camera-ready checkpoint release will be accompanied by a model card stating the intended use (mathematical reasoning research) and the upstream license under which the base model is distributed.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: All datasets (MATH-500 [Lightman et al., 2024], GSM8K [Cobbe et al., 2021], AIME 2024, AMC 2023, OlympiadBench [He et al., 2024], Minerva Math [Lewkowycz et al., 2022], GaoKao 2023-EN, College Math) and pretrained models (Qwen2.5-Math [Yang et al., 2024], LLaMA-3.1 [Grattafiori et al., 2024]) are publicly available and cited at first use. The training pipeline builds on SimpleRL-Zoo [Zeng et al., 2025] and VeRL [Sheng et al., 2025], also cited. Each asset is used under its original publicly stated terms.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: Two new assets are introduced to support reproducibility (item 5). (i) An anonymized implementation of the RSP injection harness and the DAPO + RSP training pipeline is included as supplementary material, accompanied by a README that documents the environment, run commands for the inference experiments in §4.2–§4.4, and the DAPO + RSP training entry point used in §4.5. (ii) Pretrained DAPO and DAPO + RSP checkpoints across training steps 10–100 will be released at the camera-ready stage, accompanied by a model card stating the intended use (mathematical reasoning research), a reference to the training data used (the Level-3–5 subset of MATH from SimpleRL-Zoo’s publicly available `simplelr_math_35`, ~8.5K prompts; the dataset itself is not

redistributed), and the upstream Qwen2.5-Math-7B [Yang et al., 2024] license that both checkpoints inherit. No new dataset is introduced and no human subjects are involved, so no consent procedure applies. Limitations of the released artifacts (single-seed training, RoPE/math-only scope) are documented in item 2 and §5.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper involves no crowdsourcing and no research with human subjects. All evaluations use static, publicly available math reasoning benchmarks.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: As noted in item 14, the paper involves no human subjects, so IRB review is not applicable.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: Pretrained LLMs (Qwen2.5-Math, LLaMA-3.1) appear in the paper as the experimental subject of analysis, not as a non-standard component of the methods themselves. The methods (RSP definition, attention decomposition, DAPO + RSP training) do not rely on auxiliary LLMs in any non-standard or originality-affecting role.