
From Noise to Diversity: Random Embedding Injection in LLM Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

Recent *soft prompt* methods, which prepend or insert trainable continuous embeddings into LLM inputs as virtual tokens, optimize these embeddings to enhance mathematical reasoning. However, the mechanism by which embedding injection itself affects model behavior, independent of optimization, remains unexplored. We investigate Random Soft Prompts (RSPs), continuous vectors sampled from an isotropic Gaussian fitted to the entrywise mean and variance of the pretrained embedding table and injected without any training. The randomness in RSP is not a model of the embedding distribution but the property that lets the same prompt produce independent perturbations across rollouts, which is precisely what exploration requires. Empirically, RSPs increase early-stage token entropy while generation stabilizes thereafter, and combining RSPs with temperature sampling yields higher trajectory diversity as measured by Pass@ n . Theoretically, we show that the RSP contribution at each attention layer is exactly captured by a scalar random-token attention mass that is structurally diluted as autoregressive cache growth accumulates real tokens, and that under a local affine analysis of the perturbed logits isotropic resampling admits previously excluded tokens into the effective top- K candidate set in a way temperature sampling provably cannot. The connection to a Hopfield-style energy landscape is recorded in the appendix as an interpretive lens but is not used in any formal result. We further discuss the implications for GRPO training, where trajectory diversity is critical for learning signal quality.

1 Introduction

As large language models (LLMs) scale to billions of parameters, full fine-tuning becomes increasingly expensive. This has motivated parameter-efficient methods that adapt frozen models by introducing a small number of trainable parameters [Hu et al., 2022, Housby et al., 2019]. Among these, *soft prompts* prepend or insert learnable continuous embeddings into the input sequence, steering model behavior through the attention mechanism [Li and Liang, 2021, Lester et al., 2021]. This paradigm has been actively adopted for enhancing mathematical reasoning in LLMs [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Hao et al., 2025, Zhang et al., 2026]. Despite their diverse designs, these methods share a common structure: they inject continuous embeddings that lie outside the pretrained vocabulary and optimize them to improve downstream performance.

However, we find that optimization is not a necessary condition for these effects. We observe that injecting randomly initialized, untrained continuous embeddings can also shift the output distribution. For instance, applying a chat template to a base model not fine-tuned with such format causes a prompt-format mismatch that degrades reasoning capability, but appending random embeddings to such mismatched inputs mitigates this degradation. This suggests that injecting out-of-vocabulary continuous embeddings itself, independent of any learned content, plays a non-trivial role in shaping model behavior. Existing studies primarily evaluate *soft prompt* methods through end-task

38 accuracy, and the mechanisms by which embedding injection influences the output distribution remain
39 insufficiently explored.

40 In this work, we investigate the role of Random Soft Prompts (RSPs), continuous vectors sampled
41 from an isotropic Gaussian fitted to the entrywise mean and variance of the pretrained embedding
42 table and injected without any training. The randomness here is not a model of the embedding
43 distribution: it is what allows the same prompt to produce *independent* perturbations across rollouts,
44 and exploration requires nothing less. Empirically, RSPs increase early-stage token entropy while the
45 output stabilizes in later generation stages, and combining RSPs with temperature sampling yields
46 more diverse reasoning trajectories as measured by Pass@ n scaling. Theoretically, we identify an
47 attention-level mechanism in which the RSP contribution is exactly captured by a scalar random-token
48 attention mass that is structurally diluted by autoregressive cache growth, and we show that, under a
49 local affine analysis of the perturbed logits, isotropic resampling admits previously excluded tokens
50 into the effective top- K candidate set in a way temperature sampling cannot. We further discuss the
51 implications of these findings for Group Relative Policy Optimization (GRPO) [Shao et al., 2024],
52 where trajectory diversity is essential for effective learning signals.

53 Our contributions are as follows:

- 54 • We empirically analyze how RSPs affect LLM generation, showing that RSPs increase
55 early token entropy and produce more diverse reasoning trajectories when combined with
56 temperature sampling.
- 57 • We give a transformer-level mechanism in which the RSP contribution is controlled by a
58 scalar attention mass with structural cache-driven dilution, and a distributional analysis under
59 a local affine surrogate showing that the isotropic, resampling-based randomness—rather
60 than any fixed perturbation—is what opens previously excluded tokens into the effective
61 top- K candidate set in a way temperature sampling cannot.
- 62 • We discuss the implications of RSP-induced diversity for GRPO training and argue that, be-
63 cause RSP is rank-changing while temperature is rank-preserving, RSP provides exploration
64 gains complementary to existing temperature-based approaches.

65 2 Background

66 2.1 Soft Prompt Methods for LLM Reasoning

67 *Soft prompts* emerged as a parameter-efficient alternative to full fine-tuning. Prefix-Tuning [Li and
68 Liang, 2021] introduced the paradigm by prepending trainable continuous vectors to every transformer
69 layer, Prompt Tuning [Lester et al., 2021] simplified this to input-layer-only tuning and showed that
70 it matches full fine-tuning at sufficient scale, and P-tuning [Liu et al., 2024] extended the approach
71 with an LSTM-based prompt encoder to model inter-position dependencies.

72 More recently, *soft prompts* have been actively applied to LLM reasoning. TTSV [Kang et al.,
73 2026] optimizes prefix embeddings at test time via entropy minimization, steering the model toward
74 high-certainty states to activate latent reasoning capabilities. SoftCoT [Xu et al., 2025] replaces
75 explicit chain-of-thought tokens with soft tokens generated by an assistant model, leveraging the
76 stronger representational capacity of continuous representations over discrete language. LTPO [Ye
77 et al., 2026] optimizes latent thought embeddings per test instance through policy gradient with a
78 confidence-based intrinsic reward. COCONUT [Hao et al., 2025] feeds the model’s own hidden
79 states back as input embeddings, enabling reasoning in a continuous latent space. MemGen [Zhang
80 et al., 2026] demonstrates that embedding vectors can serve as experience memory by constructing
81 latent token sequences that enrich the model’s reasoning. Pause tokens [Goyal et al., 2024] insert
82 learnable tokens into the input to provide additional computation steps.

83 These approaches all inject out-of-vocabulary continuous embeddings optimized for task-specific
84 objectives. Because evaluation centers on end-task accuracy, the contribution of the injection itself
85 and that of the optimization are not separated, and the mechanisms by which such out-of-vocabulary
86 continuous embeddings shape model behavior remain a separate and underexplored question.

87 2.2 Noise Injection in Neural Networks

88 Neural networks have incorporated noise at various stages. During training, dropout [Srivastava
89 et al., 2014] randomly deactivates hidden units to prevent overfitting, and NEFTune [Jain et al., 2024]
90 adds noise to token embeddings during instruction fine-tuning, reporting improvements in instruction
91 following. These methods use noise as a regularizer and do not apply perturbations at inference time.

92 At inference time, perturbation-based approaches primarily target robustness and safety. Randomized
93 smoothing [Cohen et al., 2019] injects Gaussian noise into inputs to obtain certified robustness
94 guarantees. In the LLM setting, SmoothLLM [Robey et al., 2025] perturbs input prompts at the
95 character level to defend against adversarial jailbreaking. These methods require multiple noisy
96 forward passes with output aggregation, and target prediction stability rather than output diversity.

97 In reinforcement learning, NoisyNet [Fortunato et al., 2018] injects learnable noise into network
98 weights to improve exploration, demonstrating that the location and structure of noise injection affect
99 exploration quality. This motivation is closest to RSPs, which similarly target diversity, but differ in
100 injection target (input embeddings vs. network weights) and require no learned noise parameters.

101 RSPs differ from these approaches in the following respects. (1) RSPs operate solely at inference
102 without any training, whereas dropout and NEFTune inject noise during training, modifying the
103 resulting model parameters accordingly. (2) RSPs preserve the original input and append fresh
104 embedding vectors in a single forward pass, whereas robustness methods corrupt the input and
105 aggregate over multiple noisy passes. (3) RSPs use purely random vectors sampled from the
106 pretrained embedding statistics, whereas NoisyNet learns noise parameters through optimization.

107 3 Experimental Analysis

108 3.1 Random Soft Prompts: Definition and Experimental Setup

109 Let $\mathbf{W}_E \in \mathbb{R}^{V \times d}$ denote the token embedding matrix of a pretrained LLM, where V is the vocabulary
110 size and d is the hidden dimension. We define an RSP $\mathbf{H} = [h_1, \dots, h_L] \in \mathbb{R}^{L \times d}$ as a sequence
111 of L continuous vectors, where each vector is independently sampled from an isotropic Gaussian
112 distribution whose mean and variance are matched to those of the pretrained token embedding matrix:

$$h_j \sim \mathcal{N}(\mu_E, \sigma_E^2 \mathbf{I}), \quad j = 1, \dots, L, \quad (1)$$

113 with $\mu_E = \frac{1}{Vd} \sum_{i,k} [\mathbf{W}_E]_{i,k}$ and $\sigma_E = \text{std}(\mathbf{W}_E)$. Let $\mathbf{X} = [x_1, \dots, x_T] \in \mathbb{R}^{T \times d}$ denote the em-
114 bedding sequence of the original input. The RSP-augmented input $\tilde{\mathbf{X}}$ is constructed by concatenating
115 $\mathbf{H}$ with $\mathbf{X}$ at one of three positions without modifying the original prompt text: *prefix* ($[\mathbf{H}; \mathbf{X}]$), *infix*
116 ($\mathbf{H}$ inserted within $\mathbf{X}$), or *suffix* ($[\mathbf{X}; \mathbf{H}]$), following the injection strategies adopted in prior work
117 [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Zhang et al., 2026]. Sampling a new $\mathbf{H}$ for each
118 batch ensures that the results are not dependent on any realization of the random embeddings.

119 We evaluate on mathematical reasoning benchmarks: MATH-500 [Lightman et al., 2024], GSM8K
120 [Cobbe et al., 2021], and AIME24, using LLaMA-3.1-8B-Instruct [Grattafiori et al., 2024] and the
121 Qwen2.5-Math model family [Yang et al., 2024] across instruct models, base models, and base
122 models with mismatched chat templates. All experiments use greedy decoding to isolate the effect of
123 RSPs from sampling stochasticity. Full experimental details are provided in Appendix A.

124 3.2 How Does Untrained RSP Compare to Baselines & Optimized Soft Prompts?

125 We first examine how injecting random embeddings affects the model’s reasoning performance.
126 Table 1 reports the accuracy for each model configuration, using the best RSP result across injection
127 positions (Prefix, Infix, Suffix) and $L \in \{10, 15, 20\}$ for each benchmark.

128 For instruct models and base models with their native prompt formats, RSP injection maintains or
129 slightly shifts baseline performance within a few percentage points. Most strikingly, under prompt-
130 format mismatch — where base models receive ChatML templates they were not trained on — RSP
131 yields pronounced improvements of up to +29 pp. These gains are difficult to attribute to simple
132 perturbation effects, as noise that merely raises model uncertainty would further degrade performance.

133 We compare RSPs with prior soft prompt methods (TTSV, SoftCoT, LTPO) that optimize their
134 embeddings for distinct objectives, under a unified evaluation protocol reported in Table 2. All

Table 1: Accuracy (%) across model configurations and benchmarks. RSP reports the best result across injection positions (Prefix, Infix, Suffix). Full per-position breakdown is in Appendix A. Values in parentheses indicate the difference from the baseline.

Model	MATH-500		GSM8K		AIME24	
	Baseline	RSP	Baseline	RSP	Baseline	RSP
<i>Instruct Models</i>						
LLaMA-3.1-8B-Inst	52.20	49.40 (−2.8)	85.75	86.73 (+1.0)	6.67	10.00 (+3.3)
Qwen2.5-Math-7B-Inst	83.20	84.20 (+1.0)	95.45	95.68 (+0.2)	13.33	16.67 (+3.3)
Qwen2.5-Math-1.5B-Inst	73.00	75.20 (+2.2)	85.29	85.75 (+0.5)	10.00	20.00 (+10.0)
<i>Base Models (Raw Text)</i>						
Qwen2.5-Math-7B	72.20	71.60 (−0.6)	84.15	84.31 (+0.2)	6.67	16.67 (+10.0)
Qwen2.5-Math-1.5B	61.40	64.40 (+3.0)	77.86	74.30 (−3.6)	16.67	10.00 (−6.7)
<i>Base Models + ChatML (Format Mismatch)</i>						
Qwen2.5-Math-7B	52.20	70.40 (+18.2)	58.30	75.97 (+17.7)	23.33	23.33 (+0.0)
Qwen2.5-Math-1.5B	34.40	63.40 (+29.0)	37.45	67.02 (+29.6)	13.33	20.00 (+6.7)

Table 2: Comparison with existing soft prompt methods (TTSV, SoftCoT, LTPO) under unified evaluation. Baseline and RSP values are from Table 1. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Benchmark	Baseline	RSP	TTSV	SoftCoT	LTPO
Qwen2.5-Math-7B-Instruct	MATH-500	83.20	84.20	<u>83.80</u>	82.80	83.00
	GSM8K	<u>95.45</u>	95.68	<u>95.30</u>	95.00	94.84
	AIME24	13.33	<u>16.67</u>	23.30	<u>16.67</u>	13.33
Qwen2.5-Math-1.5B-Instruct	MATH-500	73.00	<u>75.20</u>	<u>75.20</u>	76.00	72.40
	GSM8K	85.29	85.75	<u>85.70</u>	84.46	84.53
	AIME24	<u>10.00</u>	20.00	6.70	6.67	<u>10.00</u>

methods are re-evaluated with the same unified extraction pipeline (Appendix A.2), enabling a format-agnostic comparison despite each method’s distinct prompt format and grading scheme. The results place RSPs on par with the optimized methods without requiring any training.

3.3 How Does RSP Affect the Output Distribution?

We analyze how RSP injection changes the model’s output distribution on Qwen2.5-Math-1.5B-Instruct with MATH-500, measuring per-token entropy, top-1 probability, and varentropy across the generation process. Figure 1 compares these metrics for the first 5% of generation steps across RSP variants and prior soft prompt methods (LTPO, SoftCoT, TTSV). Full statistics are in Appendix A.4.




Figure 1: Mean entropy, top-1 probability, and varentropy during the first 5% of generation steps (Qwen2.5-Math-1.5B-Instruct, MATH-500). Shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: prior soft prompt methods (LTPO, SoftCoT, TTSV).

These results show that RSP modifies the early-stage output distribution across all injection positions: entropy increases (the token distribution flattens), top-1 probability decreases (commitment to a dominant token weakens), and varentropy increases (indicating multimodal distributions [Ahmed

et al., 2026]). High-entropy positions are known to act as critical forking points that determine reasoning trajectories [Wang et al., 2025], so this shift moves the model toward a state where multiple reasoning paths can coexist. The change is localized to the initial reasoning stage. In the middle and final portions of generation, all three metrics converge to baseline levels.

Although existing soft prompt methods are optimized for different objectives, they share a common pattern of increasing early-stage entropy. Notably, TTSV, which optimizes soft prompts using a reward that reduces the mean entropy of reasoning trajectories, still exhibits early entropy comparable to the baseline. This suggests that elevating early-stage entropy is an inherent effect of injecting non-pretrained continuous embeddings, independent of the optimization objective.

3.4 Does RSP Induce Trajectory Diversity?

Section 3.3 empirically established that RSP reshapes the early-stage output distribution. We now examine whether this shift translates into reasoning diversity at three complementary levels of analysis: token rankings, semantic content of trajectories, and task-level outcomes. The three analyses converge on a directional picture: RSP’s first-token perturbation propagates into semantically more diverse reasoning trajectories, which in turn yields greater output diversity at the task level. All three analyses share the setting: Qwen2.5-Math-1.5B-Instruct, MATH-500, suffix injection, and $L = 10$.

Token-level: distribution beyond temperature. We test whether RSP modifies the first-token distribution in a qualitatively different way from temperature scaling, which preserves token rankings while only flattening probabilities. Under autoregressive decoding, any trajectory-level divergence between RSP and the baseline must originate from the first generation step, a critical forking point for reasoning trajectories [Wang et al., 2025]. Using the baseline at $\tau=1.0$ as reference, we compare RSP at $\tau=1.0$ against the baseline at $\tau=2.0$ (the high-temperature ceiling of what flattening alone can achieve) on three complementary metrics: Spearman rank correlation ρ to detect rank reorganization, the probability mass placed outside the baseline’s top-10 to measure support expansion, and Jensen–Shannon divergence to quantify the overall magnitude of distributional change (formal definitions in Appendix B).

Table 3 reports the contrast. Even at $\tau=2.0$, only 5.15% of probability mass falls outside the baseline’s top-10. RSP at $\tau=1.0$ reaches 27.64%. The Spearman rank correlation tells the story: because temperature scaling $p_i^{(\tau)} \propto p_i^{1/\tau}$ is a monotone transform that preserves token ranking ($p_i > p_j \iff p_i^{(\tau)} > p_j^{(\tau)}$) for any $\tau > 0$, ρ relative to the baseline is exactly 1. RSP, by contrast, reranks the token distribution itself, reducing ρ to 0.651—an effect temperature cannot produce at any τ . The Jensen–Shannon divergence (0.231) quantifies the overall magnitude of this distributional change. Since $\rho < 1$ rules out rank-preserving flattening as the sole cause, the high divergence reflects a structural reshape of the distribution rather than a uniform flattening of the same ranking.

Table 3: First-token distribution metrics measured against the baseline at $\tau=1.0$, comparing the baseline at $\tau=2.0$ with RSP at $\tau=1.0$.

Metric	$\tau=2.0$	RSP
Spearman ρ	1.0	0.651
Mass outside top-10	5.15%	27.64%
JS divergence	0.086	0.231

Trajectory-level: semantic diversity. Token-level reranking raises a second question: do first-token perturbations produce semantically diverse reasoning trajectories, or do independently perturbed trajectories converge to semantically similar ones? We randomly sample 64 problems from MATH-500 and generate 300 independent trajectories per problem under Baseline and RSP at temperature $\tau = 1.0$. Each trajectory is encoded into a dense semantic vector by BGE-M3 [Chen et al., 2024], a sentence-level multilingual embedding model, and we compute three complementary metrics over these embeddings: *pairwise cosine distance* between trajectories captures the overall semantic spread of the trajectory set; *inter-cluster distance* between centroids identified by DBSCAN [Ester et al., 1996], a density-based clustering algorithm that groups nearby trajectories without prespecifying the number of clusters, captures the separation between distinct trajectory groups; and *intra-cluster distance* captures semantic coherence within each group.

Table 4: Semantic diversity metrics (64 problems, 300 trajectories per condition).

Metric	Baseline	RSP
Pairwise cos. dist.	0.0612	0.0879
Inter-cluster dist.	0.0183	0.0982
Intra-cluster dist.	0.0526	0.0528

Table 4 reveals three concurrent patterns. Pairwise cosine distance rises by about $1.4\times$, indicating that the trajectory set as a whole becomes more semantically diverse. Inter-cluster distance jumps by about $5.4\times$, showing that the trajectories split into substantially more separated groups. Intra-cluster distance is essentially unchanged, indicating that semantic coherence within each group remains comparable to the baseline. RSP also yields larger pairwise distance than baseline on 56 of 64 problems (87.5%).

Outcome-level: Pass@ n scaling. Finally, we evaluate whether the token- and trajectory-level diversity translates into task-level outcome diversity through Pass@ n [Chen et al., 2021], the probability that at least one of n sampled solutions is correct. We compare four conditions with $n = 16$ samples per problem on MATH-500 and AIME24: (i) *Baseline*, $\tau = 1.0$: temperature sampling only; (ii) *RSP (fixed seed)*, $\tau = 1.0$: a single RSP shared across all samples; (iii) *RSP (indep. seed), greedy*: a different RSP per sample without temperature; and (iv) *RSP (indep. seed)*, $\tau = 1.0$: a different RSP per sample combined with temperature sampling.




Figure 2: Pass@ n scaling on (a) MATH-500 and (b) AIME24 with Qwen2.5-Math-1.5B-Instruct, $n = 16$ samples per problem. Baseline: temperature sampling only. Fixed seed: single RSP shared across samples combined with temperature. Indep. seed: a different RSP per sample, with or without temperature.

Condition (iv) achieves the highest Pass@ n on both benchmarks, reaching 92.80% on MATH-500 and 36.67% on AIME24 at $n = 16$. Condition (ii) shows no improvement over the baseline, suggesting that the diversity gain likely depends on varying the random embeddings across samples rather than a single fixed perturbation. Pass@1 remains comparable to the baseline on MATH-500, and on the more challenging AIME24, (iv) achieves higher Pass@1 than the baseline. Together these observations suggest that combining independent RSPs with temperature sampling may yield greater trajectory diversity while preserving single-sample accuracy.

4 Theoretical Analysis: RSP as Attention-Mediated Exploration

The empirical regularities of Section 3—raised early-stage entropy, later stabilization, and Pass@ N gains under repeated decoding—admit a single transformer-level explanation. RSP perturbs the hidden state before the output softmax through one scalar quantity, the random-token attention mass, which autoregressive cache growth structurally dilutes. The randomness in RSP is not a model of the embedding table: it is what lets the same prompt produce *independent* perturbations across rollouts, and exploration requires nothing less than independent perturbations. The argument that follows uses only standard self-attention identities, and in particular assumes neither $\mathbf{K} = \mathbf{V}$, weight tying, nor a global energy function.

4.1 Scale-Matched Isotropic Perturbation

Definition 1 (Random Soft Prompt). Let $\mathbf{W}_E \in \mathbb{R}^{V \times d}$ be the pretrained input-embedding table with entrywise scalar mean μ_E and standard deviation σ_E . An RSP of length k is a sequence $(h_1, \dots, h_k)$ with $h_j = \mu_E \mathbf{1}_d + \sigma_E \xi_j$ and $\xi_j \stackrel{iid}{\sim} \mathcal{N}(0, \mathbf{I}_d)$. The centered prompt $\bar{h}_j := h_j - \mu_E \mathbf{1}_d$ then satisfies $\bar{h}_j \sim \mathcal{N}(0, \sigma_E^2 \mathbf{I}_d)$. For repeated-sampling protocols (Pass@ N), prompt draws and any decoding randomness across compared trials are assumed independent.

The empirical embedding distribution is generally anisotropic [Mu and Viswanath, 2018, Ethayarajh, 2019], and RSP departs from it on purpose. Borrowing σ_E keeps the perturbation in the magnitude

regime the model expects, while the identity covariance discards directional structure deliberately. The branch-sensitive direction at any given decoding state depends on prompt, layer, position, and the specific token pair compared, so a generic exploration perturbation should not commit in advance to a small set of embedding directions.

4.2 Why Randomness Is Necessary

Exploration requires variability across trials, and that variability cannot be supplied by a fixed perturbation. A trained continuous prompt can steer a model, but reusing the same vector across rollouts produces the same perturbation direction every time and so cannot generate independent alternatives. Discrete random tokens fail differently: each vocabulary token carries a lexical and semantic direction inherited from pretraining, so token-level sampling injects uncontrolled semantic priors. RSP avoids both failure modes by being continuous, centered, and isotropic, and the next two results pin down what each of those properties buys.

Lemma 1 (Semantic neutrality). *Let $\bar{h} \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_d)$ and $\hat{h} := \bar{h} / \|\bar{h}\|$. For any fixed unit vectors $u_1, \dots, u_M \in \mathbb{S}^{d-1}$ and $\delta \in (0, 1)$,*

$$\Pr \left[\max_{r \in [M]} |\hat{h}^\top u_r| \leq \sqrt{\frac{2 \log(2M/\delta)}{d-1}} \right] \geq 1 - \delta.$$

A token embedding fails this property: any fixed direction $u = e_s / \|e_s\|$ gives $|\hat{e}_s^\top u| = 1$. Lemma 1 formalizes only the weaker, sufficient claim that an isotropic random direction is nearly orthogonal to any fixed finite collection of semantic anchors with high probability; it does not assert that a single random vector is free of every incidental correlation.

Theorem 1 (Direction-agnostic coverage). *Let D be any zero-mean perturbation on $\mathbb{R}^d$ with covariance Σ_D and trace budget $\text{tr}(\Sigma_D) \leq \rho^2$. Then $\min_{\|b\|=1} \text{Var}_{h \sim D}(b^\top h) = \lambda_{\min}(\Sigma_D) \leq \rho^2/d$, with equality if and only if $\Sigma_D = (\rho^2/d) \mathbf{I}_d$.*

Theorem 1 is a minimax statement: among zero-mean perturbations sharing a total-variance budget, isotropic covariance uniquely maximizes the least-explored direction. Together with Lemma 1 this rules out the natural alternatives at once. A fixed prompt has $\Sigma_D = 0$ and provides no exploration; PCA-aligned noise is anisotropic and underexplores the orthogonal complement; vocabulary-token sampling is anisotropic and semantically committing. The same theorem also rules out the apparently natural alternative of matching the empirical embedding covariance: writing Σ_E for the row-covariance of $\mathbf{W}_E$ around the same scalar center $\mu_E \mathbf{1}_d$ with $\text{tr}(\Sigma_E) = d\sigma_E^2$, the embedding-matched law has $\lambda_{\min}(\Sigma_E) < \sigma_E^2$ whenever Σ_E is non-isotropic, so it leaves at least one direction strictly less explored than the isotropic comparator. Discarding directional geometry while keeping scale is therefore a deliberate optimization, not a simplification.

4.3 Exact Transformer-Level Mechanism

The effect of RSP in a standard transformer is controlled by the attention mass assigned to the random tokens. Consider self-attention layer ℓ with n real and k RSP tokens, attention weights $\alpha_{ij}^{(\ell)}$, and value vectors $\mathbf{v}_j^{(\ell)}, \mathbf{v}_{r,j}^{(\ell)}$ for real and RSP tokens. Let $w_{r,i}^{(\ell)} = \sum_{j=1}^k \alpha_{i,n+j}^{(\ell)}$ denote the random-token attention mass. The attention output then decomposes *exactly* as $\mathbf{x}_i^{(\ell+1)} = (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{x}}_i^{(\ell+1)} + \boldsymbol{\eta}_i^{(\ell)}$, where $\tilde{\mathbf{x}}_i^{(\ell+1)} = \sum_{j=1}^n \frac{\alpha_{ij}^{(\ell)}}{1 - w_{r,i}^{(\ell)}} \mathbf{v}_j^{(\ell)}$ and the random contribution $\boldsymbol{\eta}_i^{(\ell)} = \sum_{j=1}^k \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r,j}^{(\ell)}$ obeys $\|\boldsymbol{\eta}_i^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_j \|\mathbf{v}_{r,j}^{(\ell)}\|$. The same scalar $w_{r,i}^{(\ell)}$ thus controls both how much real-token signal is displaced and how large the random contribution can be; this decomposition is an algebraic identity rather than an approximation, and the deterministic norm bound is recorded in Appendix E.1. (A complementary interpretation as iterative retrieval on a Hopfield-style energy landscape [Ramsauer et al., 2021] is offered in Appendix D; the formal results in this section do not depend on it.)

The dilution structure follows directly. With $\Delta_i^{(\ell)} := \min_{j \in [n]} \mathbf{q}_i^\top \mathbf{k}_j - \max_{j' \in [k]} \mathbf{q}_i^\top \mathbf{k}_{r,j'}$, the real-vs-random pre-scaling logit gap, scaled dot-product attention gives

$$w_{r,i}^{(\ell)} \leq \frac{k}{k + n \exp(\Delta_i^{(\ell)} / \sqrt{d})}, \quad (2)$$

279 which is strictly decreasing in n at fixed k and $\Delta_i^{(\ell)}$. Autoregressive cache growth alone therefore
 280 imposes a monotone decrease on the largest random-token attention compatible with the gap; this
 281 is structural dilution of an upper bound, not monotone decay of every realized $w_{r,i}^{(\ell)}$, since $\Delta_i^{(\ell)}$
 282 itself co-evolves with the state. Layer-wise sharpening of $\Delta_i^{(\ell)}$ is a model-side hypothesis verified
 283 empirically by Figure 3. A full block adds residual connections, normalization, and an FFN; under a
 284 common local Lipschitz constant on a ball containing both perturbed and unperturbed trajectories
 285 [Xiong et al., 2020, Kim et al., 2021], Appendix E.3 shows that the per-block perturbation amplitude
 286 scales linearly with $w_{r,i}^{(\ell)}$, so iterating across finite depth yields a finite cumulative bound. Combined
 287 with Eq. (2), the perturbation enters strongly at early blocks and shrinks as evidence accumulates,
 288 without requiring persistence throughout the chain.




Figure 3: Per-token attention mass on Qwen2.5-Math-7B (suffix, 500 MATH-500 problems, each with an independently sampled RSP). Layer and reasoning-position axes are partitioned into five quantile bins. Attention to RSP tokens (a) decays along both axes; attention to question tokens (b) remains comparatively uniform.

289 4.4 Output Support Expansion Beyond Temperature

290 Temperature sampling is rank-preserving: $p_i^{(\tau)} \propto p_i^{1/\tau}$ implies $p_i > p_j \iff p_i^{(\tau)} > p_j^{(\tau)}$ for
 291 any $\tau > 0$, so it cannot change token rankings at a fixed decoding state. RSP, by acting before the
 292 output softmax, can. To isolate the conditions under which it does, we work with the local affine
 293 surrogate $z_i^{\text{lin}}(\bar{h}) := z_i + a_i^\top \bar{h}$, $a_i := \nabla_{\bar{h}} z_i(0)$, valid in a neighborhood of $\bar{h} = 0$ (Appendix E).
 294 Throughout this subsection $\bar{h}$ is treated as a per-token surrogate; the multi-token compounding
 295 through softmax-weighted attention is reported empirically in Section 3.4.

296 **Proposition 1** (Stochastic ranking). *If $a_i \neq a_j$ for some token pair (i, j) and the centered prompt
 297 law assigns positive probability to every nonempty open subset of $\mathbb{R}^d$, then $0 < \Pr_{\bar{h}}(z_i^{\text{lin}}(\bar{h}) >
 298 z_j^{\text{lin}}(\bar{h})) < 1$.*

299 **Proposition 2** (Top- K entry). *Let $J_K(s_t)$ be the baseline top- K token set at decoding state s_t and
 300 let $i \notin J_K(s_t)$. Define $\mathcal{G}_{i,K} := \{\bar{h} \in \mathbb{R}^d : \#\{j \neq i : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h})\} \leq K - 1\}$. A sufficient
 301 condition for $\mathcal{G}_{i,K}$ to contain a nonempty open set is the existence of $\bar{h}_0$ at which token i outranks all
 302 but at most $K - 1$ other tokens with a strict margin; the affine surrogate is continuous, so an open
 303 ball around $\bar{h}_0$ lies in $\mathcal{G}_{i,K}$ (Appendix E). If this condition holds and the centered prompt law has full
 304 open support, then $\Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0$.*

305 Proposition 2 is intentionally conditional: it makes no claim for tokens dominated by the baseline
 306 top- K in the surrogate. Temperature scaling reaches no $\mathcal{G}_{i,K}$ at all, because it never changes rank;
 307 the surrogate-level mechanism here is branch opening, not probability flattening.

308 4.5 From Independent Perturbations to Pass@ N

309 Set $\mu_i := \Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0$. Across N independent fresh RSP draws, the probability that at least
 310 one admits i into top- K is $1 - (1 - \mu_i)^N \geq 1 - \exp(-N\mu_i)$, so independent resampling converts a
 311 small per-draw branch-opening probability into an exponentially saturating coverage rate. Lifting
 312 this to task accuracy requires one task-side hypothesis on whether opened branches are productive.

313 **Proposition 3** (Exploration meets multi-path). *Fix a problem x . Run the deterministic baseline once
 314 and draw N independent fresh RSPs, decoding once per sample. Let $\mathcal{T}^*(x)$ be the nonempty set of
 315 correct trajectories and $p_{\text{base}}(x) \in \{0, 1\}$ the baseline correctness indicator. Assume (M1) for each n ,*

316 the marginal probability that the n -th RSP run produces a trajectory in $\mathcal{T}^*(x)$ is at least $p_{\min}(x) > 0$,
 317 and (M2) the N RSP runs are mutually independent. Then $\Pr(\text{Pass}@ (N+1) \text{ correct on } x) \geq$
 318 $\max\{p_{\text{base}}(x), 1 - (1 - p_{\min}(x))^N\} \geq 1 - e^{-Np_{\min}(x)}.$

319 Assumption (M1) is task-side and not derivable from the mechanism alone. Section 3.4 reports its
 320 empirical signature: RSP-induced trajectories include both improvements and regressions, with net
 321 gain accumulating under repeated sampling exactly as the proposition predicts.

322 5 Conclusion

323 We studied Random Soft Prompts (RSPs), isotropic Gaussian vectors fitted to the entrywise mean and
 324 variance of the pretrained embedding table and injected without training. The randomness in RSP is
 325 not a model of the embedding distribution; it is the property that supplies independent perturbations
 326 across rollouts, which is what exploration requires. Empirically RSPs raise early-stage entropy while
 327 later generation stabilizes (Section 3); theoretically the RSP contribution at each attention layer is
 328 captured by a scalar attention mass that cache growth structurally dilutes, and under a local affine
 329 analysis isotropic resampling admits baseline-excluded tokens into the effective top- K candidate set
 330 in a way temperature scaling cannot (Section 4). Because RSP is rank-changing while temperature
 331 is rank-preserving, the two are complementary, suggesting a sampling-level remedy for entropy
 332 collapse in GRPO [Shao et al., 2024] beyond DAPO [Yu et al., 2025] and CE-GPPO [Su et al., 2026].
 333 Training-time validation against these baselines, extension beyond mathematical reasoning, and a
 334 principled choice of RSP length and injection position remain open.

335 References

- 336 Farhan Ahmed, Yuya Jeremy Ong, and Chad DeLuca. LogitScope: A Framework for Analyzing
 337 LLM Uncertainty through Information Metrics. In *ICLR Workshop*, 2026.
- 338 Nora Belrose, Igor Ostrovsky, Lev McKinney, Zach Furman, Logan Smith, Danny Halawi, Stella
 339 Biderman, and Jacob Steinhardt. Eliciting Latent Predictions from Transformers with the Tuned
 340 Lens. *arXiv:2303.08112*, 2023.
- 341 Stéphane Boucheron, Gábor Lugosi, and Pascal Massart. *Concentration Inequalities: A Nonasymptotic*
 342 *Theory of Independence*. Oxford University Press, 2013.
- 343 Jianlyu Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. M3-
 344 Embedding: Multi-Linguality, Multi-Functionality, Multi-Granularity Text Embeddings Through
 345 Self-Knowledge Distillation. In *Findings of ACL*, 2024.
- 346 Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared
 347 Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating Large
 348 Language Models Trained on Code. *arXiv:2107.03374*, 2021.
- 349 Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser,
 350 Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John
 351 Schulman. Training Verifiers to Solve Math Word Problems. *arXiv:2110.14168*, 2021.
- 352 Jeremy M. Cohen, Elan Rosenfeld, and J. Zico Kolter. Certified Adversarial Robustness via Random-
 353 ized Smoothing. In *ICML*, 2019.
- 354 Timothée Darcet, Maxime Oquab, Julien Mairal, and Piotr Bojanowski. Vision Transformers Need
 355 Registers. In *ICLR*, 2024.
- 356 Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu. A Density-Based Algorithm for
 357 Discovering Clusters in Large Spatial Databases with Noise. In *KDD*, 1996.
- 358 Kawin Ethayarajh. How Contextual are Contextualized Word Representations? Comparing the
 359 Geometry of BERT, ELMo, and GPT-2 Embeddings. In *Proceedings of the 2019 Conference on*
 360 *Empirical Methods in Natural Language Processing (EMNLP)*, 2019.

361 Meire Fortunato, Mohammad Gheshlaghi Azar, Bilal Piot, Jacob Menick, Ian Osband, Alex Graves,
362 Vlad Mnih, Remi Munos, Demis Hassabis, Olivier Pietquin, Charles Blundell, and Shane Legg.
363 Noisy Networks for Exploration. In *ICLR*, 2018.

364 Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. Transformer Feed-Forward Layers Are
365 Key-Value Memories. In *EMNLP*, 2021.

366 Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh
367 Nagarajan. Think before You Speak: Training Language Models with Pause Tokens. In *ICLR*,
368 2024.

369 Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad
370 Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela
371 Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem
372 Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurelien Rodriguez, Austen Gregerson, Ava
373 Spataru, Baptiste Roziere, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya
374 Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang
375 Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song,
376 Danielle Pintz, Danny Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan,
377 Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina
378 Lobanova, Emily Dinan, Eric Michael Smith, Filip Radenovic, Francisco Guzmán, Frank Zhang,
379 Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, Graeme Nail, Gregoire
380 Mialon, Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron,
381 Iliyan Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang,
382 Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelmer
383 van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang,
384 Jiawen Liu, Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua
385 Johnstun, Joshua Saxe, Junteng Jia, Kalyan Vasuden Alwala, Karthik Prasad, Kartikeya Upasani,
386 Kate Plawiak, Ke Li, Kenneth Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz
387 Malik, Kuenley Chiu, Kunal Bhalla, Kushal Lakhotia, Lauren Rantala-Yeary, Laurens van der
388 Maaten, Lawrence Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo,
389 Lukas Blecher, Lukas Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat
390 Singh, Manohar Paluri, Marcin Kardas, Maria Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya
391 Pavlova, Melanie Kambadur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman
392 Goyal, Narjes Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri Chatterji, Ning Zhang,
393 Olivier Duchenne, Onur Celebi, Patrick Alrassy, Pengchuan Zhang, Pengwei Li, Petar Vasic,
394 Peter Weng, Prajjwal Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Puxin Xu,
395 Qing He, Qingxiao Dong, Ragavan Srinivasan, Raj Ganapathy, Ramon Calderer, Ricardo Silveira
396 Cabral, Robert Stojnic, Roberta Raileanu, Rohan Maheswari, Rohit Girdhar, Rohit Patel, Romain
397 Sauvestre, Ronnie Polidoro, Roshan Sumbaly, Ross Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar
398 Hosseini, Sahana Chennabasappa, Sanjay Singh, Sean Bell, Seohyun Sonia Kim, Sergey Edunov,
399 Shaoliang Nie, Sharan Narang, Sharath Raparthy, Sheng Shen, Shengye Wan, Shruti Bhosale,
400 Shun Zhang, Simon Vandenhende, Soumya Batra, Spencer Whitman, Sten Sootla, Stephane
401 Collot, Suchin Gururangan, Sydney Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha,
402 Thomas Georgiou, Thomas Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal
403 Karn, Vedanuj Goswami, Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Conguet,
404 Virginie Do, Vish Vogeti, Vitor Albiero, Vladan Petrovic, Weiwei Chu, Wenhan Xiong, Wenxin
405 Fu, Whitney Meers, Xavier Martinet, Xiaodong Wang, Xiaofang Wang, Xiaoqing Ellen Tan, Xide
406 Xia, Xinfeng Xie, Xuchao Jia, Xuwei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei,
407 Yi Wen, Yiwen Song, Yuchen Zhang, Yue Li, Yuning Mao, Zacharie Delpierre Coudert, Zheng Yan,
408 Zhengxing Chen, Zoe Papakipos, Aaditya Singh, Aayushi Srivastava, Abha Jain, Adam Kelsey,
409 Adam Shajnfeld, Adithya Gangidi, Adolfo Victoria, Ahuva Goldstand, Ajay Menon, Ajay Sharma,
410 Alex Boesenberg, Alexei Baevski, Allie Feinstein, Amanda Kallet, Amit Sangani, Amos Teo,
411 Anam Yunus, Andrei Lupu, Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew
412 Poulton, Andrew Ryan, Ankit Ramchandani, Annie Dong, Annie Franco, Anuj Goyal, Aparajita
413 Saraf, Arkabandhu Chowdhury, Ashley Gabriel, Ashwin Bharambe, Assaf Eisenman, Azadeh
414 Yazdan, Beau James, Ben Maurer, Benjamin Leonhardi, Bernie Huang, Beth Loyd, Beto De Paola,
415 Bhargavi Paranjape, Bing Liu, Bo Wu, Boyu Ni, Braden Hancock, Bram Wasti, Brandon Spence,
416 Brani Stojkovic, Brian Gamido, Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu,
417 Changhan Wang, Changkyu Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris

418 Tindal, Christoph Feichtenhofer, Cynthia Gao, Damon Civin, Dana Beaty, Daniel Kreymer, Daniel
 419 Li, David Adkins, David Xu, Davide Testuggine, Delia David, Devi Parikh, Diana Liskovich,
 420 Didem Foss, Dingkan Wang, Duc Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine
 421 Montgomery, Eleonora Presani, Emily Hahn, Emily Wood, Eric-Tuan Le, Erik Brinkman, Esteban
 422 Arcaute, Evan Dunbar, Evan Smothers, Fei Sun, Felix Kreuk, Feng Tian, Filippos Kokkinos, Firat
 423 Ozgenel, Francesco Caggioni, Frank Kanayet, Frank Seide, Gabriela Medina Florez, Gabriella
 424 Schwarz, Gada Badeer, Georgia Swee, Gil Halpern, Grant Herman, Grigory Sizov, Guangyi Zhang,
 425 Guna Lakshminarayanan, Hakan Inan, Hamid Shojanazeri, Han Zou, Hannah Wang, Hanwen Zha,
 426 Haroun Habeeb, Harrison Rudolph, Helen Suk, Henry Aspegren, Hunter Goldman, Hongyuan
 427 Zhan, Ibrahim Damlaj, Igor Molybog, Igor Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai
 428 Gat, Jake Weissman, James Geboski, James Kohli, Janice Lam, Japhet Asher, Jean-Baptiste Gaya,
 429 Jeff Marcus, Jeff Tang, Jennifer Chan, Jenny Zhen, Jeremy Reizenstein, Jeremy Teboul, Jessica
 430 Zhong, Jian Jin, Jingyi Yang, Joe Cummings, Jon Carvill, Jon Shepard, Jonathan McPhie, Jonathan
 431 Torres, Josh Ginsburg, Junjie Wang, Kai Wu, Kam Hou U, Karan Saxena, Kartikay Khandelwal,
 432 Katayoun Zand, Kathy Matosich, Kaushik Veeraraghavan, Kelly Michelena, Keqian Li, Kiran
 433 Jagadeesh, Kun Huang, Kunal Chawla, Kyle Huang, Lailin Chen, Lakshya Garg, Lavender A,
 434 Leandro Silva, Lee Bell, Lei Zhang, Liangpeng Guo, Licheng Yu, Liron Moshkovich, Luca
 435 Wehrstedt, Madian Khabza, Manav Avalani, Manish Bhatt, Martynas Mankus, Matan Hasson,
 436 Matthew Lennie, Matthias Reso, Maxim Groshev, Maxim Naumov, Maya Lathi, Meghan Keneally,
 437 Miao Liu, Michael L. Seltzer, Michal Valko, Michelle Restrepo, Mihir Patel, Mik Vyatskov,
 438 Mikayel Samvelyan, Mike Clark, Mike Macey, Mike Wang, Miquel Jubert Hermoso, Mo Metanat,
 439 Mohammad Rastegari, Munish Bansal, Nandhini Santhanam, Natascha Parks, Natasha White,
 440 Navyata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier, Nikhil Mehta, Nikolay Pavlovich
 441 Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, Olivia Hart, Omkar Salpekar, Ozlem
 442 Kalinli, Parkin Kent, Parth Parekh, Paul Saab, Pavan Balaji, Pedro Rittner, Philip Bontrager,
 443 Pierre Roux, Piotr Dollar, Polina Zvyagina, Prashant Ratanachandani, Pritish Yuvraj, Qian Liang,
 444 Rachad Alao, Rachel Rodriguez, Rafi Ayub, Raghotham Murthy, Raghu Nayani, Rahul Mitra,
 445 Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin Battey, Rocky Wang, Russ
 446 Howes, Ruty Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu, Samyak Datta, Sara Chugh,
 447 Sara Hunt, Sargun Dhillon, Sasha Sidorov, Satadru Pan, Saurabh Mahajan, Saurabh Verma, Seiji
 448 Yamamoto, Sharadh Ramaswamy, Shaun Lindsay, Shaun Lindsay, Sheng Feng, Shenghao Lin,
 449 Shengxin Cindy Zha, Shishir Patil, Shiva Shankar, Shuqiang Zhang, Shuqiang Zhang, Sinong Wang,
 450 Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala, Stephanie Max, Stephen Chen, Steve Kehoe,
 451 Steve Satterfield, Sudarshan Govindaprasad, Sumit Gupta, Summer Deng, Sungmin Cho, Sunny
 452 Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman, Tal Remez, Tamar Glaser, Tamara
 453 Best, Thilo Koehler, Thomas Robinson, Tianhe Li, Tianjun Zhang, Tim Matthews, Timothy Chou,
 454 Tzook Shaked, Varun Vontimitta, Victoria Ajayi, Victoria Montanez, Vijai Mohan, Vinay Satish
 455 Kumar, Vishal Mangla, Vlad Ionescu, Vlad Poenaru, Vlad Tiberiu Mihailescu, Vladimir Ivanov,
 456 Wei Li, Wenchen Wang, Wenwen Jiang, Wes Bouaziz, Will Constable, Xiaocheng Tang, Xiaojian
 457 Wu, Xiaolan Wang, Xilun Wu, Xinbo Gao, Yaniv Kleinman, Yanjun Chen, Ye Hu, Ye Jia, Ye Qi,
 458 Yenda Li, Yilin Zhang, Ying Zhang, Yossi Adi, Youngjin Nam, Yu Wang, Yu Zhao, Yuchen Hao,
 459 Yundi Qian, Yunlu Li, Yuzi He, Zach Rait, Zachary DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhenyu
 460 Yang, Zhiwei Zhao, and Zhiyu Ma. The Llama 3 Herd of Models. *arXiv:2407.21783*, 2024.

461 Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian.
 462 Training Large Language Models to Reason in a Continuous Latent Space. In *ICLR*, 2025.

463 Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin de Laroussilhe, Andrea
 464 Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-Efficient Transfer Learning for NLP.
 465 In *ICML*, 2019.

466 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang,
 467 and Weizhu Chen. LoRA: Low-Rank Adaptation of Large Language Models. In *ICLR*, 2022.

468 Neel Jain, Ping-yeh Chiang, Yuxin Wen, John Kirchenbauer, Hong-Min Chu, Gowthami Somepalli,
 469 Brian R. Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Aniruddha Saha, Micah Goldblum,
 470 Jonas Geiping, and Tom Goldstein. NEFTune: Noisy Embeddings Improve Instruction Finetuning.
 471 In *ICLR*, 2024.

472 Xinyue Kang, Diwei Shi, and Li Chen. Model Whisper: Steering Vectors Unlock Large Language
 473 Models’ Potential in Test-time. In *AAAI*, 2026.

474 Hyunjik Kim, George Papamakarios, and Andriy Mnih. The Lipschitz Constant of Self-Attention.
475 *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021.

476 Brian Lester, Rami Al-Rfou, and Noah Constant. The Power of Scale for Parameter-Efficient Prompt
477 Tuning. In *EMNLP*, 2021.

478 Xiang Lisa Li and Percy Liang. Prefix-Tuning: Optimizing Continuous Prompts for Generation. In
479 *ACL-IJCNLP*, 2021.

480 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike,
481 John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s Verify Step by Step. In *ICLR*, 2024.

482 Xiao Liu, Yanan Zheng, Zhengxiao Du, Ming Ding, Yujie Qian, Zhilin Yang, and Jie Tang. GPT
483 Understands, Too. *AI Open*, 5:208–215, 2024.

484 Aman Madaan and Amir Yazdanbakhsh. Text and Patterns: For Effective Chain of Thought, It Takes
485 Two to Tango. *arXiv:2209.07686*, 2022.

486 Jiaqi Mu and Pramod Viswanath. All-but-the-Top: Simple and effective postprocessing for word
487 representations. In *International Conference on Learning Representations (ICLR)*, 2018.

488 Yurii Nesterov. *Lectures on Convex Optimization*. Springer, 2 edition, 2018.

489 Hubert Ramsauer, Bernhard Schöfl, Johannes Lehner, Philipp Seidl, Michael Widrich, Lukas Gruber,
490 Markus Holzleitner, Thomas Adler, David Kreil, Michael K Kopp, Günter Klambauer, Johannes
491 Brandstetter, and Sepp Hochreiter. Hopfield Networks Is All You Need. In *ICLR*, 2021.

492 Alexander Robey, Eric Wong, Hamed Hassani, and George J. Pappas. SmoothLLM: Defending Large
493 Language Models Against Jailbreaking Attacks. *Transactions on Machine Learning Research*,
494 2025.

495 Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang,
496 Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the Limits of
497 Mathematical Reasoning in Open Language Models. *arXiv:2402.03300*, 2024.

498 Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov.
499 Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *Journal of Machine*
500 *Learning Research*, 15:1929–1958, 2014.

501 Zhenpeng Su, Leiyu Pan, Minxuan Lv, Yuntao Li, Wenping Hu, Fuzheng Zhang, Kun Gai, and Guorui
502 Zhou. CE-GPPO: Coordinating Entropy via Gradient-Preserving Clipping Policy Optimization in
503 Reinforcement Learning. In *ACL*, 2026.

504 Roman Vershynin. *High-Dimensional Probability: An Introduction with Applications in Data Science*.
505 Cambridge University Press, 2018.

506 Shenzi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen,
507 Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen
508 Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 Rule: High-Entropy Minority Tokens Drive
509 Effective Reinforcement Learning for LLM Reasoning. In *NeurIPS*, 2025.

510 Ruibin Xiong, Yunchang Yang, Di He, Kai Zheng, Shuxin Zheng, Chen Xing, Huishuai Zhang,
511 Yanyan Lan, Liwei Wang, and Tie-Yan Liu. On Layer Normalization in the Transformer Architec-
512 ture. In *International Conference on Machine Learning (ICML)*, 2020.

513 Yige Xu, Xu Guo, Zhiwei Zeng, and Chunyan Miao. SoftCoT: Soft Chain-of-Thought for Efficient
514 Reasoning with LLMs. In *ACL*, 2025.

515 An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu,
516 Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu,
517 Xingzhang Ren, and Zhenru Zhang. Qwen2.5-Math Technical Report: Toward Mathematical
518 Expert Model via Self-Improvement. *arXiv:2409.12122*, 2024.

519 Wengao Ye, Yan Liang, and Lianlei Shan. Thinking on the Fly: Test-Time Reasoning Enhancement
520 via Latent Thought Policy Optimization. In *ICLR*, 2026.

- 521 Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian
522 Fan, Gaohong Liu, Juncai Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming
523 Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Ru Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu,
524 Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Yuxuan Song, Xiangpeng Wei, Hao Zhou,
525 Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Yonghui Wu, and Mingxuan Wang. DAPO:
526 An Open-Source LLM Reinforcement Learning System at Scale. In *NeurIPS*, 2025.
- 527 Guibin Zhang, Muxin Fu, and Shuicheng Yan. MemGen: Weaving Generative Latent Memory for
528 Self-Evolving Agents. In *ICLR*, 2026.

529 A Experimental Setup and Full Accuracy Breakdown

530 Table 5 provides the full per-position accuracy breakdown (Prefix, Infix, Suffix) corresponding to
531 the best-across-positions results summarized in Table 1 in the main text. Table 6 lists the number of
532 RSP tokens L used for each model–benchmark pair, selected from $\{10, 15, 20\}$. All experiments are
533 conducted on a single NVIDIA A6000 40GB GPU with greedy decoding, a maximum generation
534 length of 3,072 tokens for MATH-500 and GSM8K, and 4,096 tokens for AIME24. Batch size is
535 fixed per model (16 for LLaMA-3.1-8B-Instruct and Qwen2.5-Math-7B variants, 32 for Qwen2.5-
536 Math-1.5B variants).

Table 5: Accuracy (%) across model configurations, injection positions, and benchmarks. Values in parentheses indicate the difference from the baseline. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Mode	MATH-500	GSM8K	AIME24
<i>Instruct Models</i>				
LLaMA-3.1-8B-Instruct	Baseline	52.20	85.75	<u>6.67</u>
	Prefix	45.00 (−7.2)	76.35 (−9.4)	3.33 (−3.3)
	Infix	49.40 (−2.8)	85.97 (+0.2)	10.00 (+3.3)
	Suffix	<u>49.40</u> (−2.8)	86.73 (+1.0)	10.00 (+3.3)
Qwen2.5-Math-7B-Instruct	Baseline	83.20	95.45	<u>13.33</u>
	Prefix	81.40 (−1.8)	94.92 (−0.5)	13.33 (+0.0)
	Infix	84.20 (+1.0)	95.60 (+0.2)	16.67 (+3.3)
	Suffix	<u>83.40</u> (+0.2)	95.68 (+0.2)	16.67 (+3.3)
Qwen2.5-Math-1.5B-Instruct	Baseline	73.00	<u>85.29</u>	10.00
	Prefix	74.80 (+1.8)	85.29 (+0.0)	<u>13.33</u> (+3.3)
	Infix	75.20 (+2.2)	85.75 (+0.5)	6.67 (−3.3)
	Suffix	74.20 (+1.2)	84.69 (−0.6)	20.00 (+10.0)
<i>Base Models (Raw Text)</i>				
Qwen2.5-Math-7B	Baseline	72.20	84.15	6.67
	Prefix	<u>71.60</u> (−0.6)	84.31 (+0.2)	16.67 (+10.0)
	Infix	63.20 (−9.0)	64.29 (−19.9)	16.67 (+10.0)
	Suffix	55.60 (−16.6)	54.13 (−30.0)	<u>13.33</u> (+6.7)
Qwen2.5-Math-1.5B	Baseline	<u>61.40</u>	77.86	16.67
	Prefix	64.40 (+3.0)	74.30 (−3.6)	10.00 (−6.7)
	Infix	63.20 (+1.8)	73.92 (−3.9)	<u>10.00</u> (−6.7)
	Suffix	54.60 (−6.8)	58.61 (−19.3)	3.33 (−13.3)
<i>Base Models + ChatML (Format Mismatch)</i>				
Qwen2.5-Math-7B	Baseline	52.20	58.30	23.33
	Prefix	59.20 (+7.0)	56.41 (−1.9)	3.33 (−20.0)
	Infix	<u>62.00</u> (+9.8)	69.07 (+10.8)	23.33 (+0.0)
	Suffix	70.40 (+18.2)	75.97 (+17.7)	23.33 (+0.0)
Qwen2.5-Math-1.5B	Baseline	34.40	37.45	<u>13.33</u>
	Prefix	63.40 (+29.0)	67.02 (+29.6)	20.00 (+6.7)
	Infix	39.20 (+4.8)	50.42 (+13.0)	6.67 (−6.7)
	Suffix	<u>51.00</u> (+16.6)	<u>50.64</u> (+13.2)	<u>13.33</u> (+0.0)

Table 6: Number of RSP tokens (L) used for each model and benchmark.

Model	MATH-500	GSM8K	AIME24
LLaMA-3.1-8B-Instruct	15	20	15
Qwen2.5-Math-7B-Instruct	20	20	20
Qwen2.5-Math-1.5B-Instruct	20	10	20
Qwen2.5-Math-7B	10	10	20
Qwen2.5-Math-1.5B	10	10	20
Qwen2.5-Math-1.5B (ChatML)	20	20	10
Qwen2.5-Math-7B (ChatML)	20	20	20

537 A.1 RSP Injection Positions and Prompt Templates

538 Below we show the prompt structure for each injection position across all model types. Blue text
539 indicates RSP embeddings, and purple text indicates special tokens.

540 A.1.1 Prefix

541 RSP embeddings are concatenated before the prompt embeddings. No text modification is applied.

542 ChatML (Qwen Instruct / Base + ChatML).

```
[Random Embeddings (L tokens)]
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{<|im_end|>
<|im_start|>user
{question}<|im_end|>
<|im_start|>assistant
```

544 LLaMA Chat Template.

```
[Random Embeddings (L tokens)]
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{<|eot_id|>
<|start_header_id|>user<|end_header_id|>
{question}<|eot_id|>
<|start_header_id|>assistant<|end_header_id|>
```

546 Raw Text (Qwen Base).

```
[Random Embeddings (L tokens)]
{question}
Please reason step by step, and put your final answer within \boxed{<|eot_id|>
```

548 A.1.2 Infix

549 A description text and special tokens are inserted into the prompt. The special token positions are
550 then replaced with RSP embeddings at the embedding level.

551 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{<|im_end|>
<|im_start|>user
{question}

There are {L} special tokens that contain compressed latent reasoning information that
might be useful for your reasoning. If these tokens are useful for your case, you can
use them as reference. If these tokens are not useful for your case, you can ignore
them and focus back to solving the problem.

Here are the {L} special tokens: <special_token_1>...<special_token_L>
<|im_end|>
<|im_start|>assistant
```

553 LLaMA Chat Template.

```
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{<|eot_id|>
<|start_header_id|>user<|end_header_id|>
```


{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens:

```
<reserved_special_token_0>...  
<reserved_special_token_L>  
<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

555

556 Raw Text (Qwen Base).

{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens: $\langle \text{special_token_1} \rangle \dots \langle \text{special_token_L} \rangle$

Please reason step by step, and put your final answer within `\boxed{\}`.

557

558 A.1.3 Suffix

559 For chat-templated models, RSP embeddings are inserted immediately before the end-of-turn token.

560 For raw text models, RSP embeddings are concatenated at the end of the prompt.

561 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system  
Please reason step by step, and put your final answer within \boxed{<|im_end|>  
<|im_start|>user  
{question}[Random Embeddings (L tokens)]<|im_end|>  
<|im_start|>assistant
```

562

563 LLaMA Chat Template.

```
<|begin_of_text|>  
<|start_header_id|>system<|end_header_id|>  
Cutting Knowledge Date: December 2023  
Today Date: 26 Jul 2024  
Please reason step by step, and put your final answer within \boxed{<|eot_id|>  
<|start_header_id|>user<|end_header_id|>  
{question}[Random Embeddings (L tokens)]<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

564

565 Raw Text (Qwen Base).

```
{question}  
Please reason step by step, and put your final answer within \boxed{<|eot_id|>[Random  
Embeddings (L tokens)]
```

566

567 A.2 Answer Extraction and Grading

568 The final answer is extracted from the model output through a fallback chain. Each step is attempted
569 in order; if extraction fails, the next step is tried.

Answer Extraction Fallback Chain

1. "final answer is \$...\$. I hope" pattern (Minerva format)
2. `\boxed{...}`: last non-empty box is used; empty `\boxed{}` are skipped
3. Text following "the answer is"
4. Text following "final answer is"
5. Last number in the output (regex fallback)

Extracted answers are normalized by removing \$, \left, \right, and unit strings. Grading is performed via symbolic equivalence checking: exact string match, numerical comparison (`math.isclose, rel_tol=1e-4`), and SymPy symbolic simplification.

A.3 Reproduced Prior Work Hyperparameters

All prior methods are reproduced on Qwen2.5-Math-1.5B-Instruct and Qwen2.5-Math-7B-Instruct, evaluated with greedy decoding (temperature = 0) and our unified answer extraction pipeline.

TTSV. Prefix length 20, trained for 20 epochs with AdamW optimizer, batch size 2, gradient accumulation 8 steps, and linear learning rate schedule. Learning rate is 1e-3 for Qwen-1.5B and 1e-5 for Qwen-7B.

LTPO. Table 7 reports the per-benchmark hyperparameters.

Table 7: LTPO hyperparameters.

Parameter	GSM8K	MATH-500	AIME24
tokens	8	8	8
steps	20	20	20
top_k	10	10	10
sigma	20	20	5
sigma_decay	0.95	0.95	0.95
lr	1e-2	5e-3	5e-2

SoftCoT. Projection module trained for 10 epochs. Thought tokens: 32 during training, 4 during evaluation. Batch size 4 for GSM8K, 1 for MATH with gradient accumulation 4. The small (assistant) model is Qwen2.5-Math-1.5B-Instruct in both configurations, while the large model is Qwen2.5-Math-1.5B-Instruct (learning rate 2e-5) or Qwen2.5-Math-7B-Instruct (learning rate 5e-6). For MATH-500 and AIME24 evaluation, we use the projection module trained on GSM8K, as it yields higher accuracy than the one trained on MATH.

A.4 Full Entropy Curves

Figure 4 shows the full normalized (0–100%) generation trajectories for entropy, top-1 probability, and varentropy, and Table 8 reports the corresponding scalar statistics including overall means, first-5% means, and the average number of generated tokens per problem. All methods converge to similar levels after the initial generation stage, confirming that the distributional changes induced by RSP are localized to the early reasoning phase.

Table 8: Per-step statistics for Qwen2.5-Math-1.5B-Instruct on MATH-500 across the full generation and the first 5% of generation steps, together with the average number of generated tokens per problem. Top-1 Prob (overall) is reported as a weighted average over the first-10%/middle-last-10% segment means with weights 0.1/0.8/0.1. Extension of Figure 1 (main text) and Figure 4.

Metric	Baseline	Prefix	Infix	Suffix	LTPO	SoftCoT	TTSV
Tokens (mean)	575.7	558.4	549.9	573.3	592.3	594.4	577.4
Entropy (first 5%)	0.1332	0.1366	0.1402	0.1941	0.1662	0.4218	0.1359
Entropy (overall)	0.0871	0.0881	0.0899	0.1049	0.0909	0.1059	0.0864
Top-1 Prob (first 5%)	0.9535	0.9523	0.9525	0.9373	0.9437	0.9156	0.9534
Top-1 Prob (overall)	0.9684	0.9681	0.9678	0.9647	0.9683	0.9651	0.9685
Varentropy (first 5%)	0.2181	0.2207	0.2463	0.4010	0.2953	2.2234	0.2174
Varentropy (overall)	0.1353	0.1375	0.1422	0.2038	0.1452	0.2404	0.1361




Figure 4: Mean entropy (top), top-1 probability (middle), and varentropy (bottom) over the full normalized generation trajectory (0–100%). Averaged over MATH-500, shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: existing soft prompt methods. All methods converge after the initial stage.

593 B Token-Level Distribution Metrics

594 This appendix provides formal definitions for the three metrics used in the Token-level analysis of
 595 Section 3.4. Let P_{base} denote the baseline next-token distribution at $\tau = 1.0$ and P_{target} the candidate
 596 distribution being compared (e.g., the baseline at $\tau = 2.0$ or RSP at $\tau = 1.0$). All metrics are
 597 computed analytically from saved logits at the first generation step.

598 **Spearman rank correlation.** Spearman’s ρ is the Pearson correlation coefficient between the token
 599 ranks of two distributions:

$$\rho = \text{Pearson}(\text{rank}(P_{\text{target}}), \text{rank}(P_{\text{base}})). \quad (3)$$

600 Because temperature scaling $P^{(\tau)} \propto P^{1/\tau}$ is a strictly monotone transform, the rank order of tokens
 601 is preserved exactly, so $\rho = 1$ for any $\tau > 0$ as a mathematical identity. Any value $\rho < 1$ therefore
 602 indicates rank reorganization beyond what temperature scaling can produce.

603 **Mass outside top- K .** Let $\text{TopK}(P_{\text{base}})$ be the set of K tokens with the highest probability under
 604 P_{base} . The probability mass that the candidate distribution places *outside* this set is

$$\text{MassOutside}_K = 1 - \sum_{t \in \text{TopK}(P_{\text{base}})} P_{\text{target}}(t). \quad (4)$$

605 This directly measures support expansion: how much probability the candidate assigns to tokens that
 606 are not in the baseline’s preferred set. Reported values use $K = 10$.

607 **Jensen–Shannon divergence.** The symmetrized, bounded version of Kullback–Leibler divergence:

$$\text{JS}(P, Q) = \frac{1}{2} \text{KL}(P \parallel M) + \frac{1}{2} \text{KL}(Q \parallel M), \quad M = \frac{1}{2}(P + Q). \quad (5)$$

608 We use base-2 logarithms so that $\text{JS} \in [0, 1]$. Tokens absent from the stored top-100 are assigned a
 609 sentinel logit (-10^9) before softmax, which is negligible for our setting since the top-5 covers more
 610 than 99.9% of the probability mass.

C Why Random Prompts Are Semantically Neutral and Direction-Agnostic

This appendix formalizes the “why random” claim used in Section 4.2. The point is not that random prompts contain hidden task information. The point is structural: if the perturbation is meant to act as a generic exploration device rather than as a learned prompt prior, then it should be resampled across runs, centered so that it does not induce a fixed first-order drift, and direction-agnostic so that it does not under-serve whichever local branch direction matters for the current prompt.

C.1 Semantic Neutrality

Let $\bar{h} \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_d)$ and $\hat{h} := \bar{h}/\|\bar{h}\|$ whenever $\bar{h} \neq 0$. Let $u_1, \dots, u_M \in \mathbb{S}^{d-1}$ be any fixed unit vectors representing semantic anchor directions.

Lemma 2 (Semantic neutrality). *For any $\delta \in (0, 1)$,*

$$\Pr \left[\max_{r \in [M]} |\hat{h}^\top u_r| \leq \sqrt{\frac{2 \log(2M/\delta)}{d-1}} \right] \geq 1 - \delta.$$

Hence the normalized direction of an isotropic Gaussian prompt is simultaneously nearly orthogonal to every fixed finite collection of semantic anchors with high probability.

Proof. Because $\bar{h}$ is isotropic Gaussian, its normalized direction $\hat{h}$ is uniformly distributed on the sphere $\mathbb{S}^{d-1}$. For a fixed unit vector $u \in \mathbb{S}^{d-1}$, the marginal $\hat{h}^\top u$ satisfies the standard spherical concentration bound [Vershynin, 2018, Theorem 3.4.6][Boucheron et al., 2013, Section 3.2],

$$\Pr(|\hat{h}^\top u| \geq t) \leq 2 \exp\left(-\frac{(d-1)t^2}{2}\right) \quad \text{for } t \in (0, 1),$$

which follows from the spherical-cap area bound $\sigma\{\hat{h} \in \mathbb{S}^{d-1} : \hat{h}^\top u \geq t\} \leq \frac{1}{2}(1-t^2)^{(d-1)/2}$ together with $(1-t^2)^{(d-1)/2} \leq \exp(-(d-1)t^2/2)$ derived from $\log(1-x) \leq -x$ for $x \in [0, 1)$. Applying this bound to each u_r and using a union bound,

$$\Pr \left[\max_{r \in [M]} |\hat{h}^\top u_r| \geq t \right] \leq 2M \exp\left(-\frac{(d-1)t^2}{2}\right).$$

Setting

$$t = \sqrt{\frac{2 \log(2M/\delta)}{d-1}}$$

gives the claim. If the displayed value of t exceeds 1, the statement is trivial because $|\hat{h}^\top u_r| \leq 1$ always. $\square$

The contrast with a token embedding is immediate. If $e_s \neq 0$ is a fixed token embedding and $\hat{e}_s := e_s/\|e_s\|$, then choosing the anchor direction $u = \hat{e}_s$ gives

$$|\hat{e}_s^\top u| = 1.$$

A token prompt is therefore maximally aligned with at least one semantic direction, whereas an isotropic random prompt is nearly orthogonal to every fixed finite set of such directions. This is the sense in which tokens inject semantic commitment while random prompts act as semantically uncommitted perturbations.

C.2 Centered Perturbations Do Not Impose Systematic First-Order Drift

Let Δ_{ij} be the baseline logit gap between tokens i and j , and write the first-order surrogate as

$$\Delta_{ij}(\bar{h}) = \Delta_{ij} + b_{ij}^\top \bar{h}.$$

Proposition 4 (No systematic first-order drift). *If $\mathbb{E}[\bar{h}] = 0$, then*

$$\mathbb{E}[\Delta_{ij}(\bar{h})] = \Delta_{ij}$$

for every token pair (i, j) . If a deterministic prompt h_0 is reused across runs, then

$$\Delta_{ij}(h_0) - \Delta_{ij} = b_{ij}^\top h_0$$

is a fixed directional bias.

643 *Proof.* The centered case is immediate from linearity of expectation:

$$\mathbb{E}[\Delta_{ij}(\bar{h})] = \Delta_{ij} + b_{ij}^\top \mathbb{E}[\bar{h}] = \Delta_{ij}.$$

644 The deterministic case follows by substitution. $\square$

645 The proposition does not say that a centered perturbation leaves every *realization* unchanged; it says
 646 there is no systematic first-order push toward one branch across runs. This is exactly the behavior
 647 desired from a generic exploration device. A fixed token or fixed soft prompt may still be useful, but
 648 it acts as a persistent prompt prior rather than as an unbiased exploration mechanism.

649 C.3 Direction-Agnostic Coverage

650 The second requirement is directional coverage. Let D be any zero-mean perturbation distribution
 651 in $\mathbb{R}^d$ with covariance Σ_D and variance budget $\text{tr}(\Sigma_D) \leq \rho^2$. For a unit vector b , the first-order
 652 perturbation energy available along b is $\text{Var}_{h \sim D}(b^\top h) = b^\top \Sigma_D b$. This is the quantity that appears
 653 in the local linear logit surrogate used in the main text.

654 **Theorem 2** (Direction-agnostic coverage). *Define*

$$\mathcal{E}(D) := \min_{\|b\|=1} \text{Var}_{h \sim D}(b^\top h).$$

655 *Then*

$$\mathcal{E}(D) = \lambda_{\min}(\Sigma_D) \leq \frac{\rho^2}{d},$$

656 *with equality if and only if $\Sigma_D = (\rho^2/d)\mathbf{I}_d$.*

657 *Proof.* For any symmetric covariance matrix Σ_D ,

$$\min_{\|b\|=1} b^\top \Sigma_D b = \lambda_{\min}(\Sigma_D),$$

658 so $\mathcal{E}(D) = \lambda_{\min}(\Sigma_D)$. Let the eigenvalues of Σ_D be $\lambda_1, \dots, \lambda_d$. Then

$$\lambda_{\min}(\Sigma_D) \leq \frac{1}{d} \sum_{m=1}^d \lambda_m = \frac{\text{tr}(\Sigma_D)}{d} \leq \frac{\rho^2}{d}.$$

659 Equality holds if and only if all eigenvalues are equal, i.e., $\Sigma_D = (\rho^2/d)\mathbf{I}_d$. $\square$

660 This theorem is a statement about second moments, not about Gaussianity by itself. Many distributions
 661 could realize isotropic covariance. We use a Gaussian in RSP because it is the simplest centered,
 662 rotation-symmetric family that matches the scale of pretrained embeddings. The theorem clarifies the
 663 structural advantage: a fixed prompt has $\Sigma_D = 0$ and therefore no across-run exploration, whereas
 664 any anisotropic perturbation leaves at least one branch-sensitive direction less explored than the
 665 isotropic optimum.

666 **Proposition 5** (Covariance pushforward under the value projection). *Let $\bar{h}$ be a centered random*
 667 *vector with covariance $\sigma_E^2 \mathbf{I}_d$, and let $v = W\bar{h}$ for a matrix $W \in \mathbb{R}^{d_v \times d}$. Then*

$$\text{Cov}(v) = \sigma_E^2 W W^\top.$$

668 *Consequently, for every unit vector $u \in \mathbb{S}^{d_v-1}$,*

$$\sigma_E^2 \lambda_{\min}(W W^\top) \leq \text{Var}(u^\top v) = \sigma_E^2 u^\top W W^\top u \leq \sigma_E^2 \lambda_{\max}(W W^\top).$$

669 *In particular, if W is full rank then the projected perturbation has strictly positive variance in every*
 670 *output-space direction.*

671 *Proof.* Since $v = W\bar{h}$ and $\mathbb{E}[\bar{h}] = 0$,

$$\text{Cov}(v) = \mathbb{E}[v v^\top] = W \mathbb{E}[\bar{h} \bar{h}^\top] W^\top = \sigma_E^2 W W^\top.$$

672 For any unit vector u ,

$$\text{Var}(u^\top v) = u^\top \text{Cov}(v) u = \sigma_E^2 u^\top W W^\top u,$$

673 and the Rayleigh quotient bounds give the stated inequalities. $\square$

674 D Interpretive Connection to Hopfield Retrieval (Optional)

675 *This appendix is included for completeness and is not load-bearing for any quantitative result in*
676 *the main text.* Section 4.3 of the main paper deliberately uses only standard self-attention identities
677 (the exact decomposition of Eq. (??) and the decay bound of Eq. (2)); the formal claims there do
678 not assume $\mathbf{K} = \mathbf{V}$, weight-tied layers, or the existence of a global energy. The present appendix
679 records the alternative reading—originating from the Modern Hopfield correspondence [Ramsauer
680 et al., 2021]—under which a single attention update with $\mathbf{K} = \mathbf{V}$ becomes one gradient step on a
681 Hopfield energy. This reading provides additional intuition for why a deterministic forward pass can
682 be exploration-poor and has been a useful lens in the related literature, but it is not used in any proof
683 of Section 4.

684 We restate the Modern Hopfield correspondence in the form used in this interpretive lens. For a key
685 matrix $\mathbf{K} = (\mathbf{k}_1, \dots, \mathbf{k}_n)$ at a fixed query position, the Hopfield energy is

$$E_{\mathbf{K}}(\sigma) := -\beta^{-1} \log \sum_{j=1}^n \exp(\beta \mathbf{k}_j^\top \sigma) + \frac{1}{2} \|\sigma\|^2.$$

686 A direct calculation gives the gradient

$$\nabla E_{\mathbf{K}}(\sigma) = \sigma - \mathbf{K}^\top \text{softmax}(\beta \mathbf{K} \sigma).$$

687 Under (A1) $\mathbf{K} = \mathbf{V}$, the attention update with this $\mathbf{K}$ is therefore $\sigma^{\text{new}} = \sigma - \nabla E_{\mathbf{K}}(\sigma) =$
688 $\mathbf{K}^\top \text{softmax}(\beta \mathbf{K} \sigma)$, which is one gradient step on $E_{\mathbf{K}}$ at unit step size [Ramsauer et al., 2021].

689 **Lemma 3** (Descent inequality). *If $E_{\mathbf{K}}$ is L_E -smooth on a neighborhood of σ , then the unit-step*
690 *gradient update $\sigma^{\text{new}} := \sigma - \nabla E_{\mathbf{K}}(\sigma)$ satisfies*

$$E_{\mathbf{K}}(\sigma^{\text{new}}) \leq E_{\mathbf{K}}(\sigma) - (1 - \frac{L_E}{2}) \|\nabla E_{\mathbf{K}}(\sigma)\|^2.$$

691 *Proof.* The standard L_E -smooth quadratic upper bound [Nesterov, 2018, Lemma 1.2.3] states $E(\sigma +$
692 $v) \leq E(\sigma) + \nabla E(\sigma)^\top v + (L_E/2) \|v\|^2$ for any v in the smoothness neighborhood. Applying
693 this to $v = -\nabla E(\sigma)$ gives $E(\sigma - \nabla E) \leq E(\sigma) - \|\nabla E\|^2 + (L_E/2) \|\nabla E\|^2 = E(\sigma) - (1 -$
694 $L_E/2) \|\nabla E\|^2$. $\square$

695 **Theorem 3** (Deterministic descent and sublevel-set invariance). *Assume (A1) $\mathbf{K} = \mathbf{V}$, (A2) the*
696 *attention component is weight-tied across blocks, and (A3) $E_{\mathbf{K}}$ is L_E -smooth on a neighborhood*
697 *of the unperturbed trajectory with $L_E < 2$. Then the unit-step iterate $\sigma^{(l+1)} = \sigma^{(l)} - \nabla E_{\mathbf{K}}(\sigma^{(l)})$*
698 *is strictly monotone-descending in $E_{\mathbf{K}}$ at every non-stationary point, and the sublevel set $\{\sigma :$
699 $E_{\mathbf{K}}(\sigma) \leq E_{\mathbf{K}}(\sigma^{(0)})\}$ *is forward-invariant.**

700 *Proof.* By Lemma 3, $E_{\mathbf{K}}(\sigma^{(l+1)}) \leq E_{\mathbf{K}}(\sigma^{(l)}) - (1 - L_E/2) \|\nabla E_{\mathbf{K}}(\sigma^{(l)})\|^2$. Under (A3), the
701 coefficient $1 - L_E/2 > 0$, so $E_{\mathbf{K}}(\sigma^{(l+1)}) \leq E_{\mathbf{K}}(\sigma^{(l)})$ with strict inequality unless $\nabla E_{\mathbf{K}}(\sigma^{(l)}) = 0$,
702 which proves forward invariance. $\square$

703 **Corollary 1** (Connected-component containment for the unit-step iterate). *Let $S := \{\sigma : E_{\mathbf{K}}(\sigma) \leq$*
704 *$E_{\mathbf{K}}(\sigma^{(0)})\}$ and let $C(\sigma^{(0)})$ denote the connected component of S containing $\sigma^{(0)}$. Define*

$$\text{sep}(C(\sigma^{(0)})) := \inf\{\|x - y\| : x \in C(\sigma^{(0)}), y \in S \setminus C(\sigma^{(0)})\},$$

705 *with the convention $\text{sep} = +\infty$ if S is connected. Assume $\text{sep}(C(\sigma^{(0)})) > 0$. Then the unit-step*
706 *iterate $\sigma^{(l+1)} = \sigma^{(l)} - \nabla E_{\mathbf{K}}(\sigma^{(l)})$ remains in $C(\sigma^{(0)})$ for all $l \geq 0$ whenever*

$$\|\nabla E_{\mathbf{K}}(\sigma^{(l)})\| < \text{sep}(C(\sigma^{(0)})) \quad \text{for every } l.$$

707 *Proof.* By Theorem 3, every iterate lies in S . Suppose by induction $\sigma^{(l)} \in C(\sigma^{(0)})$. The displacement
708 $\sigma^{(l+1)} - \sigma^{(l)} = -\nabla E_{\mathbf{K}}(\sigma^{(l)})$ has Euclidean length $\|\nabla E_{\mathbf{K}}(\sigma^{(l)})\|$. A jump to a distinct component
709 $C' \subseteq S$ would require this length to be at least $\inf\{\|x - y\| : x \in C(\sigma^{(0)}), y \in C'\} \geq \text{sep}(C(\sigma^{(0)}))$,
710 which the hypothesis excludes. Hence $\sigma^{(l+1)} \in C(\sigma^{(0)})$. $\square$

711 **Positivity of sep.** $\text{sep}(C) > 0$ holds whenever the components of S are spatially separated by a
712 region of strictly higher energy than $E_{\mathbf{K}}(\sigma^{(0)})$, which we treat as a problem-dependent assumption
713 rather than a structural property: it requires that the relevant components of the Hopfield energy do
714 not share closure boundary points at level $E_{\mathbf{K}}(\sigma^{(0)})$. If $\text{sep} = 0$ the corollary becomes vacuous; in
715 that case the trajectory is still in S by Theorem 3, but is not formally trapped in $C(\sigma^{(0)})$.

716 **Caveats and scope.** Theorem 3 depends on three assumptions that are individually mild but jointly
717 idealized for real transformers. (A1) $\mathbf{K} = \mathbf{V}$ collapses the key and value projections; this is the
718 standard Modern Hopfield assumption [Ramsauer et al., 2021] and isolates the energy-iteration aspect
719 of attention from the additional value-projection degree of freedom. (A2) weight-tying treats the
720 attention component as a single repeated map; without it, each block has its own energy and the
721 trajectory does not minimize a single fixed function across blocks. (A3) is a smoothness condition on
722 $E_{\mathbf{K}}$. For the Hopfield energy used here, $\nabla^2 E_{\mathbf{K}}(\sigma) = \mathbf{I} - \beta \mathbf{K}^\top (\text{diag}(p_\sigma) - p_\sigma p_\sigma^\top) \mathbf{K}$ where $p_\sigma =$
723 $\text{softmax}(\beta \mathbf{K} \sigma)$, so the operator-norm bound on $\nabla^2 E_{\mathbf{K}}$ depends on the inverse-temperature β and
724 on $\|\mathbf{K}\|$. (A3) holds whenever $\beta \|\mathbf{K}\|^2$ is sufficiently small, which is the standard regime in which the
725 Hopfield iterate converges in a small number of steps. The theorem therefore characterizes an idealized
726 tendency of attention-only iteration in this regime; it is not a literal property of every transformer
727 block. In particular, sublevel-set confinement follows from Theorem 3, whereas connected-component
728 confinement additionally requires the separation and per-step gradient assumptions of Corollary 1. Its
729 role in the main text is to supply an idealized source of exploration deficit that an injected stochastic
730 perturbation addresses; the empirical link is the cross-layer attention decay observed in Figure 3 and
731 the trajectory diversity gain reported in Section 3.4.

E Proofs for the Transformer-Level Mechanism

This appendix collects the formal proofs underlying Sections 4.3–4.5 of the main text.

E.1 Exact Attention Decomposition and Magnitude Control

Proposition 6 (Attention decomposition). *Let $\alpha_{ij}^{(l)}$ be the attention weights for query position i at layer l , with n real tokens and k random tokens. Define*

$$w_{r,i}^{(l)} := \sum_{j=1}^k \alpha_{i,n+j}^{(l)}$$

and, for $j \in [n]$,

$$\tilde{\alpha}_{ij}^{(l)} := \frac{\alpha_{ij}^{(l)}}{1 - w_{r,i}^{(l)}}.$$

Then

$$\mathbf{x}_i^{(l+1)} = (1 - w_{r,i}^{(l)}) \tilde{\mathbf{x}}_i^{(l+1)} + \boldsymbol{\eta}_i^{(l)},$$

where

$$\tilde{\mathbf{x}}_i^{(l+1)} := \sum_{j=1}^n \tilde{\alpha}_{ij}^{(l)} \mathbf{v}_j^{(l)} \quad \text{and} \quad \boldsymbol{\eta}_i^{(l)} := \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)}.$$

Proof. The attention output is

$$\mathbf{x}_i^{(l+1)} = \sum_{j=1}^n \alpha_{ij}^{(l)} \mathbf{v}_j^{(l)} + \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)}.$$

By definition, $\sum_{j=1}^n \alpha_{ij}^{(l)} = 1 - w_{r,i}^{(l)}$, so

$$\sum_{j=1}^n \alpha_{ij}^{(l)} \mathbf{v}_j^{(l)} = (1 - w_{r,i}^{(l)}) \sum_{j=1}^n \frac{\alpha_{ij}^{(l)}}{1 - w_{r,i}^{(l)}} \mathbf{v}_j^{(l)} = (1 - w_{r,i}^{(l)}) \tilde{\mathbf{x}}_i^{(l+1)}.$$

Substituting this identity gives the decomposition. □

Corollary 2 (Magnitude scales with random-token attention). *For every realization of the random values,*

$$\|\boldsymbol{\eta}_i^{(l)}\| \leq w_{r,i}^{(l)} \max_{j \in [k]} \|\mathbf{v}_{r_j}^{(l)}\|.$$

Proof. By the triangle inequality,

$$\|\boldsymbol{\eta}_i^{(l)}\| = \left\| \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)} \right\| \leq \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \|\mathbf{v}_{r_j}^{(l)}\| \leq w_{r,i}^{(l)} \max_{j \in [k]} \|\mathbf{v}_{r_j}^{(l)}\|.$$

□

Corollary 2 is the only quantitative fact needed in the main text. No independence assumption between attention weights and random keys is required. Once $w_{r,i}^{(l)}$ is small, the random contribution must be small as well.

750 E.2 Attention Decay Bound

751 **Theorem 4** (Attention decay). *The total attention on random tokens satisfies*

$$w_{r,i}^{(l)} \leq \frac{k}{k + n \cdot \exp(\Delta_i^{(l)} / \sqrt{d})}, \quad (6)$$

752 where $\Delta_i^{(l)} := \min_{j \in [n]} \mathbf{q}_i^\top \mathbf{k}_j - \max_{j' \in [k]} \mathbf{q}_i^\top \mathbf{k}_{r_{j'}}$, is the real-vs-random logit gap.

753 *Proof.* Let $s_j = \mathbf{q}_i^\top \mathbf{k}_j / \sqrt{d}$ for real keys and $s_{r_{j'}} = \mathbf{q}_i^\top \mathbf{k}_{r_{j'}} / \sqrt{d}$ for random keys. By definition
 754 of $\Delta_i^{(l)}$, each real key satisfies $s_j \geq \Delta_i^{(l)} / \sqrt{d} + s_{r,\max}$ where $s_{r,\max} = \max_{j'} s_{r_{j'}}$. Therefore
 755 $\exp(s_j) \geq \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot \exp(s_{r,\max})$ for every real key. Summing over n real keys gives

$$\sum_{j=1}^n \exp(s_j) \geq n \cdot \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot \exp(s_{r,\max}) \geq n \cdot \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot \frac{1}{k} \sum_{j'=1}^k \exp(s_{r_{j'}}).$$

756 Let $R := \sum_{j'=1}^k \exp(s_{r_{j'}})$. Then

$$w_{r,i}^{(l)} = \frac{R}{\sum_{j=1}^n \exp(s_j) + R} \leq \frac{R}{n \cdot \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot R/k + R} = \frac{k}{k + n \cdot \exp(\Delta_i^{(l)} / \sqrt{d})}.$$

757 □

758 **Corollary 3** (Sequence-wise annealing under a fixed logit gap). *For fixed k and $\Delta_i^{(l)}$, the upper*
 759 *bound*

$$f(n) = \frac{k}{k + n \exp(\Delta_i^{(l)} / \sqrt{d})}$$

760 *is strictly decreasing in n . Consequently, cache growth alone decreases the maximum perturbation*
 761 *mass compatible with that gap.*

762 *Proof.* Differentiate $f(n)$ with respect to n :

$$f'(n) = -\frac{k \exp(\Delta_i^{(l)} / \sqrt{d})}{\left(k + n \exp(\Delta_i^{(l)} / \sqrt{d})\right)^2} < 0.$$

763 Thus $f(n)$ is strictly decreasing. □

764 **Corollary 4** (Vanishing random contribution under annealing). *If $n \rightarrow \infty$ along decoding while k*
 765 *and $\Delta_i^{(l)}$ remain fixed, then $w_{r,i}^{(l)} \rightarrow 0$. Moreover, for any fixed realization of the k random values,*

$$\|\boldsymbol{\eta}_i^{(l)}\| \leq w_{r,i}^{(l)} \max_{j \in [k]} \|\mathbf{v}_{r_j}^{(l)}\| \rightarrow 0.$$

766 *Since the implementation samples finite-dimensional vectors and k is finite, the displayed maximum*
 767 *is almost surely finite.*

768 *Proof.* Corollary 3 implies $w_{r,i}^{(l)} \rightarrow 0$. The bound on $\|\boldsymbol{\eta}_i^{(l)}\|$ is Corollary 2. Since the multiplier
 769 $\max_{j \in [k]} \|\mathbf{v}_{r_j}^{(l)}\|$ is finite for any fixed realization of finitely many sampled values, the product
 770 converges to zero. □

771 E.3 Feed-Forward Propagation under Local Lipschitz Control

772 A transformer block applies an attention sublayer followed by a feed-forward sublayer with residual
 773 connections:

$$\mathbf{x}_i^{(l,\text{attn})} = \mathbf{x}_i^{(l)} + \text{Attn}^{(l)}(\mathbf{x}_i^{(l)}), \quad \mathbf{x}_i^{(l+1)} = \mathbf{x}_i^{(l,\text{attn})} + F^{(l)}(\mathbf{x}_i^{(l,\text{attn})}).$$

774 Let $\mathbf{x}_i^{(l,\text{attn},\text{RSP})}$ and $\mathbf{x}_i^{(l,\text{attn},\text{base})}$ denote the post-attention states under RSP and the unperturbed forward
 775 pass at the same input $\mathbf{x}_i^{(l)}$.

776 **Lemma 4** (Per-block FFN propagation). *The post-attention difference satisfies*

$$\|\mathbf{x}_i^{(l, \text{attn}, \text{RSP})} - \mathbf{x}_i^{(l, \text{attn}, \text{base})}\| \leq \|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|.$$

777 *If, in addition, $F^{(l)}$ is $L^{(l)}$ -Lipschitz on a bounded set containing both $\mathbf{x}_i^{(l, \text{attn}, \text{RSP})}$ and $\mathbf{x}_i^{(l, \text{attn}, \text{base})}$,*
 778 *then*

$$\|\mathbf{x}_i^{(l+1, \text{RSP})} - \mathbf{x}_i^{(l+1, \text{base})}\| \leq (1 + L^{(l)}) (\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|).$$

779 *Proof.* The unperturbed attention sublayer produces $\mathbf{x}_i^{(l, \text{attn}, \text{base})} = \mathbf{x}_i^{(l)} + \tilde{\mathbf{x}}_i^{(l+1)}$, while the RSP
 780 attention sublayer produces $\mathbf{x}_i^{(l, \text{attn}, \text{RSP})} = \mathbf{x}_i^{(l)} + (1 - w_{r,i}^{(l)}) \tilde{\mathbf{x}}_i^{(l+1)} + \boldsymbol{\eta}_i^{(l)}$ by Proposition 6. Subtracting,

$$\mathbf{x}_i^{(l, \text{attn}, \text{RSP})} - \mathbf{x}_i^{(l, \text{attn}, \text{base})} = \boldsymbol{\eta}_i^{(l)} - w_{r,i}^{(l)} \tilde{\mathbf{x}}_i^{(l+1)},$$

781 and the triangle inequality gives the first bound. For the second bound, write the post-FFN difference
 782 as

$$[\mathbf{x}_i^{(l, \text{attn}, \text{RSP})} + F^{(l)}(\mathbf{x}_i^{(l, \text{attn}, \text{RSP})})] - [\mathbf{x}_i^{(l, \text{attn}, \text{base})} + F^{(l)}(\mathbf{x}_i^{(l, \text{attn}, \text{base})})].$$

783 By Lipschitz continuity, $\|F^{(l)}(\mathbf{x}_i^{(l, \text{attn}, \text{RSP})}) - F^{(l)}(\mathbf{x}_i^{(l, \text{attn}, \text{base})})\| \leq L^{(l)} \|\mathbf{x}_i^{(l, \text{attn}, \text{RSP})} - \mathbf{x}_i^{(l, \text{attn}, \text{base})}\|$.
 784 The triangle inequality combined with the first bound gives the second. $\square$

785 **Corollary 5** (Cross-block propagation control). *Define $\Delta^{(l)} := \mathbf{x}_i^{(l, \text{RSP})} - \mathbf{x}_i^{(l, \text{base})}$ with $\Delta^{(0)} = \mathbf{0}$*
 786 *(matched inputs at layer 0), and let $L_{\text{block}}^{(l)}$ denote a common local Lipschitz constant valid for both the*
 787 *baseline block map $B_{\text{base}}^{(l)} : \mathbf{x} \mapsto \mathbf{x} + \text{Attn}^{(l)}(\mathbf{x}) + F^{(l)}(\mathbf{x} + \text{Attn}^{(l)}(\mathbf{x}))$ and the RSP-augmented*
 788 *block map $B_{\text{RSP}}^{(l)}$ (which replaces $\text{Attn}^{(l)}$ with the attention computation that includes the random*
 789 *tokens) on a ball containing both trajectories. Such a constant exists for standard transformer blocks*
 790 *because Attn , LayerNorm , and the FFN sublayer are each locally Lipschitz on any bounded set*
 791 *[Xiong et al., 2020, Kim et al., 2021], and the random-key augmentation only adds a finite number of*
 792 *additional softmax inputs at every layer. Then*

$$\|\Delta^{(l+1)}\| \leq (1 + L_{\text{block}}^{(l)}) \|\Delta^{(l)}\| + (1 + L^{(l)}) (\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|).$$

793 *Iterating across L blocks with $\Delta^{(0)} = \mathbf{0}$ yields*

$$\|\Delta^{(L+1)}\| \leq \sum_{l=0}^L \prod_{l'=l+1}^L (1 + L_{\text{block}}^{(l')}) \cdot (1 + L^{(l)}) (\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|).$$

794 *Proof.* Decompose the cross-block discrepancy at layer $l + 1$ into two contributions:

$$\Delta^{(l+1)} = \underbrace{B_{\text{RSP}}^{(l)}(\mathbf{x}^{(l, \text{RSP})}) - B_{\text{RSP}}^{(l)}(\mathbf{x}^{(l, \text{base})})}_{\text{(I) propagation of prior } \Delta^{(l)}} + \underbrace{B_{\text{RSP}}^{(l)}(\mathbf{x}^{(l, \text{base})}) - B_{\text{base}}^{(l)}(\mathbf{x}^{(l, \text{base})})}_{\text{(II) RSP-induced perturbation injected at block } l},$$

795 where $B_{\text{RSP}}^{(l)}$ replaces $\text{Attn}^{(l)}$ with the RSP-augmented attention. Term (I) is bounded via the block
 796 Lipschitz constant: $\|(I)\| \leq L_{\text{block}}^{(l)} \|\Delta^{(l)}\|$. With residuals included one gets the factor $(1 + L_{\text{block}}^{(l)})$
 797 when also propagating the residual term $\Delta^{(l)}$ itself. Term (II) is exactly the per-block injection from
 798 Lemma 4, bounded by $(1 + L^{(l)}) (\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|)$. Combining via the triangle inequality
 799 gives the recursion. The summed bound follows by unrolling the recursion: with $\Delta^{(0)} = \mathbf{0}$, induction
 800 on L shows that $\|\Delta^{(L+1)}\|$ is at most the displayed sum. $\square$

801 **Boundedness of $L^{(l)}$ for standard architectures.** Modern transformer blocks combine LayerNorm,
 802 GELU or SwiGLU activations, and linear projections. LayerNorm has Lipschitz constant on the
 803 order of $\|\gamma\| / \inf \|x\|$ on every set bounded away from the origin, where γ is the LayerNorm scale
 804 parameter [Xiong et al., 2020]. GELU and SwiGLU are smooth and therefore locally Lipschitz on
 805 bounded activation regions, and linear projections have Lipschitz constant equal to the spectral norm
 806 of the weight matrix. The composed self-attention sublayer has a Lipschitz constant that depends
 807 on the spectral norms of $\mathbf{W}_Q$, $\mathbf{W}_K$, $\mathbf{W}_V$ and the inverse-temperature β as analyzed by Kim et al.
 808 [2021]. On the bounded activation regions traversed during a single forward pass, the per-block
 809 Lipschitz constants in Corollary 5 are therefore finite. For a fixed model depth this yields a finite
 810 cumulative bound on $\|\Delta^{(L+1)}\|$; stronger depth-uniform non-expansion would require additional
 811 contraction assumptions not imposed here. This is the formal sense in which the injected perturbation
 812 stays controlled across a finite forward pass, while the size of each new injection is set by $w_{r,i}^{(l)}$.

813 E.4 Proofs for Output Distribution Shift

814 The main text makes two output-level claims. Proposition 1 establishes pairwise rank randomness,
815 and Proposition 2 upgrades that statement to top- K entry. Both are statements about a local first-order
816 surrogate of the perturbed logits.

817 E.4.1 Autoregressive Formulation

818 In autoregressive decoding, the joint distribution over a sequence $y_{1:T}$ factorizes as

$$\Pr(y_{1:T} \mid x) = \prod_{t=1}^T p(y_t \mid x, y_{<t}),$$

819 where each step is governed by a prefix-conditioned distribution. Any perturbation introduced by
820 RSP affects generation through a sequence of local changes to $p(y_t \mid x, y_{<t})$, rather than a single
821 global distribution.

822 E.4.2 Local Linearization of Logits under RSP

823 Throughout this subsection, $\bar{h} \in \mathbb{R}^d$ denotes *one* centered prompt vector $\bar{h}_l$ from Definition 1, treated
824 as a per-token surrogate. The actual implementation uses a length- L prompt $\mathbf{H} = (h_1, \dots, h_L)$ with
825 L independent draws; the propositions below identify how each individual $\bar{h}_l$ shifts the logits, and
826 the multi-token compounding through softmax-weighted attention is reported empirically in Figure 3
827 and Section 3.4. We model the perturbed logits as a function $z_i(\bar{h})$, assuming local smoothness in a
828 neighborhood of $\bar{h} = 0$:

$$z_i(\bar{h}) = z_i(0) + \nabla_h z_i(0)^\top \bar{h} + r_i(\bar{h}),$$

829 where $r_i(\bar{h}) = O(\|\bar{h}\|^2)$. Defining $a_i := \nabla_h z_i(0)$, we obtain the local affine surrogate

$$z_i^{\text{lin}}(\bar{h}) := z_i + a_i^\top \bar{h}.$$

830 This approximation is valid for sufficiently small $\|\bar{h}\|$ and is used as a first-order surrogate for the
831 direction of perturbation rather than as an exact global model of the transformer logits.

832 For the two propositions below we isolate the only support property that is actually used. Let D
833 denote the centered prompt law. We assume

$$(S1) \quad D(U) > 0 \quad \text{for every nonempty open set } U \subseteq \mathbb{R}^d.$$

834 The Gaussian instantiation in Section 4.2 satisfies (S1), but the proofs below require only this support
835 condition.

836 E.4.3 Proof of Proposition 1 (Stochastic Ranking)

837 *If there exists a token pair (i, j) such that $a_i \neq a_j$, and if the centered prompt law satisfies (S1), then*
838 $0 < \Pr_{\bar{h}}(z_i^{\text{lin}}(\bar{h}) > z_j^{\text{lin}}(\bar{h})) < 1$.

839 *Proof.* Under the local affine surrogate,

$$z_i^{\text{lin}}(\bar{h}) - z_j^{\text{lin}}(\bar{h}) = (z_i - z_j) + (a_i - a_j)^\top \bar{h}.$$

840 Let $b_{ij} := a_i - a_j \neq 0$. The perturbed logit difference

$$\Delta_{ij}(\bar{h}) := (z_i - z_j) + b_{ij}^\top \bar{h}$$

841 vanishes on the affine hyperplane

$$H_0 := \{\bar{h} \in \mathbb{R}^d : (z_i - z_j) + b_{ij}^\top \bar{h} = 0\}.$$

842 Because $b_{ij} \neq 0$, H_0 has codimension one and partitions $\mathbb{R}^d$ into the two nonempty open half-spaces

$$H_+ := \{\bar{h} : (z_i - z_j) + b_{ij}^\top \bar{h} > 0\}, \quad H_- := \{\bar{h} : (z_i - z_j) + b_{ij}^\top \bar{h} < 0\}.$$

843 By (S1), both H_+ and H_- have positive probability, which yields

$$0 < \Pr_{\bar{h}}(\Delta_{ij}(\bar{h}) > 0) < 1.$$

844

□

845 **E.4.4 Proof of Proposition 2 (Top-K Entry)**

846 Let $J_K(s_t)$ be the baseline top- K set at state s_t , and let $i \notin J_K(s_t)$. Define

$$\mathcal{G}_{i,K} := \{\bar{h} \in \mathbb{R}^d : \#\{j \neq i : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h})\} \leq K-1\}.$$

847 If $\mathcal{G}_{i,K}$ contains a nonempty open set and the centered prompt law satisfies (S1), then

$$\Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0.$$

848 *Proof.* Under the local affine surrogate, token i belongs to the surrogate top- K set if and only if at
 849 most $K-1$ tokens outrank it, i.e., if and only if $\bar{h} \in \mathcal{G}_{i,K}$. For each subset S of token indices with
 850 $i \notin S$ and $|S| \leq K-1$, define

$$\mathcal{O}_S := \{\bar{h} : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h}) \text{ for all } j \in S, z_i^{\text{lin}}(\bar{h}) > z_j^{\text{lin}}(\bar{h}) \text{ for all } j \notin S \cup \{i\}\}.$$

851 Each $\mathcal{O}_S$ is an intersection of open half-spaces and is therefore an open subset of $\mathbb{R}^d$. By construction,
 852 every $\bar{h} \in \mathcal{O}_S$ has token i ranked $|S|+1 \leq K$ in the affine surrogate, so $\mathcal{O}_S \subseteq \mathcal{G}_{i,K}$. Hence

$$\bigcup_{S: i \notin S, |S| \leq K-1} \mathcal{O}_S \subseteq \mathcal{G}_{i,K}.$$

853 The reverse inclusion fails only on the measure-zero set of $\bar{h}$ that produce ties, which we do not need:
 854 the assumption that $\mathcal{G}_{i,K}$ contains a nonempty open set is sufficient. Condition (S1) therefore implies

$$\Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0,$$

855 which is exactly the event that token i enters the surrogate top- K set. $\square$

856 **E.4.5 Proof of Corollary ?? (Sample Complexity)**

857 Let $\mu_i = \Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0$. Within N independent RSP samples $\bar{h}_1, \dots, \bar{h}_N$, the probability of at
 858 least one admission is at least $1 - (1 - \mu_i)^N \geq 1 - \exp(-N\mu_i)$.

859 *Proof.* The samples $\{\bar{h}_n\}_{n=1}^N$ are independent, so

$$\Pr(\text{no admission within } N \text{ samples}) = \prod_{n=1}^N \Pr(\bar{h}_n \notin \mathcal{G}_{i,K}) = (1 - \mu_i)^N.$$

860 The complement gives the first inequality. The second inequality is standard: $1 - \mu_i \leq \exp(-\mu_i)$,
 861 hence $(1 - \mu_i)^N \leq \exp(-N\mu_i)$. $\square$

862 **E.4.6 Proof of Proposition 3 (Exploration-meets-multi-path Pass@ N gain)**

863 Fix a problem x and let $p_n(x) \geq p_{\min}(x) > 0$ be the per-run success probability of the n -th RSP run
 864 as in (M1). Under (M1)+(M2),

$$\Pr(\text{Pass@}(N+1) \text{ correct on } x) \geq \max\{p_{\text{base}}(x), 1 - (1 - p_{\min}(x))^N\} \geq 1 - e^{-Np_{\min}(x)}.$$

865 *Proof.* Let A_0 be the event that the deterministic baseline run reaches a trajectory in $\mathcal{T}^*(x)$, and
 866 let A_n ($n = 1, \dots, N$) be the event that the n -th RSP run reaches such a trajectory. Pass@ $(N+1)$
 867 succeeds on x iff at least one of $A_0, A_1, \dots, A_N$ occurs.

868 By assumption $p_{\text{base}}(x) := \Pr(A_0) \in \{0, 1\}$ and $\Pr(A_n) = p_n(x) \geq p_{\min}(x)$ for each $n \geq 1$.
 869 Independence of the N RSP runs (M2) gives

$$\Pr\left(\bigcap_{n=1}^N A_n^c\right) = \prod_{n=1}^N (1 - p_n(x)) \leq \prod_{n=1}^N (1 - p_{\min}(x)) = (1 - p_{\min}(x))^N,$$

870 using $p_n(x) \geq p_{\min}(x)$ in the inequality. Therefore

$$\Pr\left(\bigcup_{n=1}^N A_n\right) = 1 - \Pr\left(\bigcap_{n=1}^N A_n^c\right) \geq 1 - (1 - p_{\min}(x))^N.$$

871 The Pass@ $(N+1)$ event $A_0 \cup A_1 \cup \dots \cup A_N$ has probability at least $\max\{\Pr(A_0), \Pr(\bigcup_{n \geq 1} A_n)\}$,
 872 giving the first claimed inequality. The second inequality follows from the standard bound $1-x \leq e^{-x}$
 873 for $x \in [0, 1]$ applied with $x = p_{\min}(x)$:

$$(1 - p_{\min}(x))^N \leq e^{-Np_{\min}(x)},$$

874 hence $1 - (1 - p_{\min}(x))^N \geq 1 - e^{-Np_{\min}(x)}$. □

875 **E.4.7 Trajectory-Level Compounding**

876 Since autoregressive generation composes conditional distributions across steps, perturbations at
 877 early steps can lead to distinct trajectories. Let $\mathcal{R}_{k,K}(x)$ denote the set of trajectories containing
 878 at least one token outside the baseline top- K set. If RSP increases the accessibility of such tokens
 879 at early steps, then the probability of sampling trajectories in $\mathcal{R}_{k,K}(x)$ increases under repeated
 880 sampling. The key point is temporal compounding: one early top- K entry changes the next decoding
 881 state, which changes the next candidate set, and so on. This is the heuristic bridge from local rank
 882 perturbations to global trajectory diversity that motivates assumption (M1) of Proposition 3.

F Attention Mass Analysis

This appendix formalizes the per-token attention mass reported in Figure 3 and the exclusion criteria applied to the reasoning span and the layer axis. For each of the 500 MATH-500 problems, a fresh RSP $\mathbf{H} \in \mathbb{R}^{L \times d}$ with $L = 10$ is sampled independently, so the reported heatmaps average over both problem variability and RSP-vector variability.

F.1 Per-Token Attention Mass

For each problem, we measure attention on the concatenated sequence [prompt; $\mathbf{H}$; generation], where the generation is the trajectory produced by the same model under matching RSP injection. Let $\alpha_{i,j}^{(\ell,h)}$ denote the post-softmax attention weight at layer ℓ , head h , from query position i to key position j , satisfying $\sum_j \alpha_{i,j}^{(\ell,h)} = 1$ under the causal mask. Let Q and R denote the index sets of question and RSP tokens with sizes $|Q|$ and $|R| = L$, and let H be the number of heads. The per-token attention mass that query i directs to each group at layer ℓ is

$$\text{PTM}_Q[\ell, i] = \frac{1}{|Q|H} \sum_{h=1}^H \sum_{j \in Q} \alpha_{i,j}^{(\ell,h)}, \quad \text{PTM}_R[\ell, i] = \frac{1}{|R|H} \sum_{h=1}^H \sum_{j \in R} \alpha_{i,j}^{(\ell,h)}. \quad (7)$$

Normalization by $|Q|$ and $|R|$ controls for group size, so both quantities express how much attention a single question (resp. RSP) token receives on average under the same softmax denominator. Question tokens thereby serve as a size-matched baseline that absorbs sequence-length effects common to both groups.

F.2 Heatmap Aggregation

We partition the layer indices and the reasoning position indices each into five equal-size bins $\{B_L^{(b)}\}_{b=1}^5$ and $\{B_P^{(b)}\}_{b=1}^5$. For sample s and target group $\bullet \in \{Q, R\}$, the value in cell (b_L, b_P) is

$$M_{\bullet}^{(s)}[b_L, b_P] = \frac{1}{|B_L^{(b_L)}| \cdot |B_P^{(b_P)}|} \sum_{\ell \in B_L^{(b_L)}} \sum_{i \in B_P^{(b_P)}} \text{PTM}_{\bullet}[\ell, i]. \quad (8)$$

Figure 3 reports the sample average of Eq. (8) over the 500 valid samples.

F.3 Excluding the `\boxed{\dots}` Span

The reasoning position range terminates strictly before the final `\boxed{\dots}` span. Chain-of-thought outputs decompose into a *text* (content) component and a *pattern* (template) component that contribute independently to downstream performance [Madaan and Yazdanbakhsh, 2022]. The `\boxed{\dots}` wrapper is a canonical pattern: a short, stereotyped span whose role is to complete the answer format rather than to extend reasoning. Including it would therefore inject a template-driven attention signature unrelated to reasoning dynamics.

F.4 Excluding the Final Two Layers

The layer axis further excludes the last two transformer layers. Two well-documented properties of late layers, none specific to our setup, motivate this choice.

Output-projection specialization. Late-layer hidden states lie in an approximately linear relationship with vocabulary logits and therefore function as a progressive linear readout rather than as representation-transforming blocks [Belrose et al., 2023]. Consistently, late-layer feed-forward modules have been characterized as output-distribution shaping mechanisms [Geva et al., 2021]. Attention in these layers therefore reflects output formation rather than reasoning computation.

Register-token hypothesis. Uninformative patch tokens in Vision Transformers have been shown to be repurposed as *registers*: storage slots that absorb attention in deeper layers [Darcet et al., 2024]. Because RSPs share a similar profile as continuous tokens without pretrained semantic content, a

similar repurposing in late transformer layers cannot be ruled out. Excluding the final two layers removes this potential register-induced confound from the reasoning-phase signal.

The two lines of evidence above identify known late-layer phenomena — linear readout and potential register repurposing — whose attention signatures are orthogonal to reasoning-phase dynamics. Removing the final two layers preempts their contamination of the reasoning-phase signal, yielding a principled, literature-grounded exclusion rather than a model-specific tuning decision.

F.5 Additional Suffix Heatmaps

Figure 5 reports the same per-token RSP attention analysis under suffix injection for Qwen2.5-Math-1.5B-Instruct and LLaMA-3.1-8B-Instruct. For the late-layer reasons discussed above, the pattern does not generalize uniformly across every cell; nevertheless, the overall trend is consistent across both models: question-token attention varies broadly across layers, whereas RSP attention shows the layer-wise and sequence-wise decay predicted by Eq. (2).




Figure 5: Per-token RSP attention mass under suffix injection for the remaining two models (500 MATH-500 samples each). Axes and preprocessing match Figure 3: x-axis is reasoning position binned into five quantiles, y-axis is layer depth binned into five quantiles (top = shallow), color encodes the 500-sample mean per-token attention assigned to RSP tokens, and tokens inside `\boxed{}` spans and the last two layers are excluded.

933 G GRPO Formulation

934 For a given prompt q , the behavior policy $\pi_{\theta_{\text{old}}}$ samples a group of G responses $\{o_i\}_{i=1}^G$ with
 935 corresponding rewards $\{R_i\}_{i=1}^G$. The advantage of each response is computed by normalizing the
 936 group-level rewards:

$$\hat{A}_{i,t} = \frac{R_i - \text{mean}(\{R_i\}_{i=1}^G)}{\text{std}(\{R_i\}_{i=1}^G)}. \quad (9)$$

937 The policy is updated via a clipped objective with a KL penalty:

$$\mathcal{J}_{\text{GRPO}}(\theta) = \mathbb{E} \left[\frac{1}{G} \sum_{i=1}^G \frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \left(\min \left(r_{i,t} \hat{A}_{i,t}, \text{clip}(r_{i,t}, 1-\varepsilon, 1+\varepsilon) \hat{A}_{i,t} \right) - \beta D_{\text{KL}}(\pi_{\theta} \parallel \pi_{\text{ref}}) \right) \right], \quad (10)$$

938 where $r_{i,t} = \pi_{\theta}(o_{i,t} \mid q, o_{i,<t}) / \pi_{\theta_{\text{old}}}(o_{i,t} \mid q, o_{i,<t})$ is the importance sampling ratio, ε is the clipping
 939 range, and β controls the strength of the KL penalty against the reference policy π_{ref} .

940 H Broader Impacts

941 This work provides empirical and theoretical analysis of how random embedding injection affects
942 the reasoning behavior of large language models. The contributions are primarily foundational,
943 with potential downstream implications for reinforcement learning-based post-training of reasoning
944 models such as GRPO.

945 **Positive impacts.** RSP-induced trajectory diversity could improve the sample efficiency of RL-
946 based LLM post-training by providing richer learning signals within group-relative advantage estima-
947 tion. This has the potential to reduce compute costs associated with alignment and reasoning-oriented
948 fine-tuning, making such methods more accessible to research groups with limited resources.

949 **Potential concerns.** Methods that expand the effective output support of LLMs may increase the
950 reachability of low-probability tokens, which includes both desirable alternative reasoning paths
951 and potentially unsafe or off-distribution outputs. Practitioners applying RSP in production systems
952 should combine it with existing safety filtering and alignment mechanisms rather than treating it as a
953 standalone diversity source.

954 **Limitations on impact scope.** Since this work focuses on analysis rather than deploying a new
955 model or dataset, the direct societal impact is limited. The broader implications described above are
956 contingent on future work integrating RSP into applied training pipelines.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.
- Keep the checklist subsection headings, questions/answers and guidelines below.
- Do not modify the questions and only use the provided macros for your answers.

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer **[Yes]** if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- 1161 • The paper should disclose whether the full research project required more compute
1162 than the experiments reported in the paper (e.g., preliminary or failed experiments that
1163 didn't make it into the paper).

1164 9. Code of ethics

1165 Question: Does the research conducted in the paper conform, in every respect, with the
1166 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

1167 Answer: **[TODO]**

1168 Justification: **[TODO]**

1169 Guidelines:

- 1170 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of
1171 Ethics.
1172 • If the authors answer **[No]**, they should explain the special circumstances that require a
1173 deviation from the Code of Ethics.
1174 • The authors should make sure to preserve anonymity (e.g., if there is a special consid-
1175 eration due to laws or regulations in their jurisdiction).

1176 10. Broader impacts

1177 Question: Does the paper discuss both potential positive societal impacts and negative
1178 societal impacts of the work performed?

1179 Answer: **[TODO]**

1180 Justification: **[TODO]**

1181 Guidelines:

- 1182 • The answer [N/A] means that there is no societal impact of the work performed.
1183 • If the authors answer [N/A] or **[No]**, they should explain why their work has no societal
1184 impact or why the paper does not address societal impact.
1185 • Examples of negative societal impacts include potential malicious or unintended uses
1186 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations
1187 (e.g., deployment of technologies that could make decisions that unfairly impact specific
1188 groups), privacy considerations, and security considerations.
1189 • The conference expects that many papers will be foundational research and not tied
1190 to particular applications, let alone deployments. However, if there is a direct path to
1191 any negative applications, the authors should point it out. For example, it is legitimate
1192 to point out that an improvement in the quality of generative models could be used to
1193 generate Deepfakes for disinformation. On the other hand, it is not needed to point out
1194 that a generic algorithm for optimizing neural networks could enable people to train
1195 models that generate Deepfakes faster.
1196 • The authors should consider possible harms that could arise when the technology is
1197 being used as intended and functioning correctly, harms that could arise when the
1198 technology is being used as intended but gives incorrect results, and harms following
1199 from (intentional or unintentional) misuse of the technology.
1200 • If there are negative societal impacts, the authors could also discuss possible mitigation
1201 strategies (e.g., gated release of models, providing defenses in addition to attacks,
1202 mechanisms for monitoring misuse, mechanisms to monitor how a system learns from
1203 feedback over time, improving the efficiency and accessibility of ML).

1204 11. Safeguards

1205 Question: Does the paper describe safeguards that have been put in place for responsible
1206 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
1207 image generators, or scraped datasets)?

1208 Answer: **[TODO]**

1209 Justification: **[TODO]**

1210 Guidelines:

- 1211 • The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

1263 Guidelines:

1264 • The answer [N/A] means that the paper does not involve crowdsourcing nor research

1265 with human subjects.

1266 • Including this information in the supplemental material is fine, but if the main contribu-

1267 tion of the paper involves human subjects, then as much detail as possible should be

1268 included in the main paper.

1269 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,

1270 or other labor should be paid at least the minimum wage in the country of the data

1271 collector.

1272 **15. Institutional review board (IRB) approvals or equivalent for research with human**

1273 **subjects**

1274 Question: Does the paper describe potential risks incurred by study participants, whether

1275 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)

1276 approvals (or an equivalent approval/review based on the requirements of your country or

1277 institution) were obtained?

1278 Answer: [TODO]

1279 Justification: [TODO]

1280 Guidelines:

1281 • The answer [N/A] means that the paper does not involve crowdsourcing nor research

1282 with human subjects.

1283 • Depending on the country in which research is conducted, IRB approval (or equivalent)

1284 may be required for any human subjects research. If you obtained IRB approval, you

1285 should clearly state this in the paper.

1286 • We recognize that the procedures for this may vary significantly between institutions

1287 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the

1288 guidelines for their institution.

1289 • For initial submissions, do not include any information that would break anonymity (if

1290 applicable), such as the institution conducting the review.

1291 **16. Declaration of LLM usage**

1292 Question: Does the paper describe the usage of LLMs if it is an important, original, or

1293 non-standard component of the core methods in this research? Note that if the LLM is used

1294 only for writing, editing, or formatting purposes and does *not* impact the core methodology,

1295 scientific rigor, or originality of the research, declaration is not required.

1296 Answer: [TODO]

1297 Justification: [TODO]

1298 Guidelines:

1299 • The answer [N/A] means that the core method development in this research does not

1300 involve LLMs as any important, original, or non-standard components.

1301 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not

1302 be described.