
From Noise to Diversity: Random Embedding Injection in LLM Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

Recent *soft prompt* methods, which prepend or insert trainable continuous embeddings into LLM inputs as virtual tokens, optimize these embeddings to enhance mathematical reasoning. However, the mechanism by which embedding injection itself affects model behavior, independent of optimization, remains unexplored. We investigate Random Soft Prompts (RSPs), continuous vectors sampled from the model’s token embedding distribution and injected without any training. Empirically, RSPs increase early-stage token entropy while generation stabilizes thereafter, and combining RSPs with temperature sampling yields higher trajectory diversity as measured by Pass@ n . Theoretically, we provide two complementary analyses: an energy-based framework showing that RSP introduces Langevin-like perturbations on the Hopfield energy landscape with natural dual decay across layers and generation steps, and a distributional analysis showing that RSP induces stochastic token rankings and expands the output support beyond temperature sampling. We further discuss the implications for GRPO training, where trajectory diversity is critical for learning signal quality.

1 Introduction

As large language models (LLMs) scale to billions of parameters, full fine-tuning becomes increasingly expensive. This has motivated parameter-efficient methods that adapt frozen models by introducing a small number of trainable parameters [Hu et al., 2022, Houlsby et al., 2019]. Among these, *soft prompts* prepend or insert learnable continuous embeddings into the input sequence, steering model behavior through the attention mechanism [Li and Liang, 2021, Lester et al., 2021]. This paradigm has been actively adopted for enhancing mathematical reasoning in LLMs [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Hao et al., 2025, Zhang et al., 2025]. Despite their diverse designs, these methods share a common structure: they inject continuous embeddings that lie outside the pretrained vocabulary and optimize them to improve downstream performance.

However, we find that optimization is not a necessary condition for these effects. We observe that injecting randomly initialized, untrained continuous embeddings can also shift the output distribution. For instance, applying a chat template to a base model not fine-tuned with such format causes a prompt-format mismatch that degrades reasoning capability; appending random embeddings to such mismatched inputs mitigates this degradation. This suggests that injecting out-of-vocabulary continuous embeddings itself, independent of any learned content, plays a non-trivial role in shaping model behavior. Existing studies primarily evaluate *soft prompt* methods through end-task accuracy, and the mechanisms by which embedding injection influences the output distribution remain insufficiently explored.

In this work, we investigate the role of Random Soft Prompts (RSPs), continuous vectors sampled from the Gaussian distribution of the model’s token embedding matrix and injected without any

training. Empirically, RSPs increase early-stage token entropy while the output stabilizes in later generation stages, and combining RSPs with temperature sampling yields more diverse reasoning trajectories as measured by Pass@ n scaling. Theoretically, we analyze RSPs from two perspectives: an energy-based framework showing that RSP introduces controlled stochastic perturbations with natural decay through the attention mechanism, and a distributional analysis showing that RSP induces stochastic token rankings and expands the output support beyond what temperature sampling can achieve. We further discuss the implications of these findings for Group Relative Policy Optimization (GRPO) [Shao et al., 2024], where trajectory diversity is essential for effective learning signals.

Our contributions are as follows:

- We empirically analyze how RSPs affect LLM generation, showing that RSPs increase early token entropy and produce more diverse reasoning trajectories when combined with temperature sampling.
- We provide theoretical analyses from two perspectives: an energy-based framework explaining why RSP can perturb the output distribution without degrading performance, and a distributional analysis showing that RSP induces stochastic token rankings and expands the output support beyond the reach of temperature sampling.
- We discuss the implications of RSP-induced diversity for GRPO training and argue that RSP provides exploration gains complementary to existing temperature-based approaches.

2 Background

2.1 Soft Prompt Methods for LLM Reasoning

Soft prompts emerged as a parameter-efficient alternative to full fine-tuning. Prefix-Tuning [Li and Liang, 2021] introduced the paradigm by prepending trainable continuous vectors to every transformer layer, Prompt Tuning [Lester et al., 2021] simplified this to input-layer-only tuning and showed that it matches full fine-tuning at sufficient scale, and P-tuning [Liu et al., 2023] extended the approach with an LSTM-based prompt encoder to model inter-position dependencies.

More recently, *soft prompts* have been actively applied to LLM reasoning. TTSV [Kang et al., 2026] optimizes prefix embeddings at test time via entropy minimization, steering the model toward high-certainty states to activate latent reasoning capabilities. SoftCoT [Xu et al., 2025] replaces explicit chain-of-thought tokens with soft tokens generated by an assistant model, leveraging the stronger representational capacity of continuous representations over discrete language. LTPO [Ye et al., 2026] optimizes latent thought embeddings per test instance through policy gradient with a confidence-based intrinsic reward. COCONUT [Hao et al., 2025] feeds the model’s own hidden states back as input embeddings, enabling reasoning in a continuous latent space. MemGen [Zhang et al., 2025] demonstrates that embedding vectors can serve as experience memory by constructing latent token sequences that enrich the model’s reasoning. Pause tokens [Goyal et al., 2024] insert learnable tokens into the input to provide additional computation steps.

These approaches all inject out-of-vocabulary continuous embeddings optimized for task-specific objectives. Because evaluation centers on end-task accuracy, the contribution of the injection itself and that of the optimization are not separated, and the mechanisms by which such out-of-vocabulary continuous embeddings shape model behavior remain a separate and underexplored question.

2.2 Noise Injection in Neural Networks

Neural networks have incorporated noise at various stages. During training, dropout [Srivastava et al., 2014] randomly deactivates hidden units to prevent overfitting, and NEFTune [Jain et al., 2024] adds noise to token embeddings during instruction fine-tuning, reporting improvements in instruction following. These methods use noise as a regularizer and do not apply perturbations at inference time.

At inference time, perturbation-based approaches primarily target robustness and safety. Randomized smoothing [Cohen et al., 2019] injects Gaussian noise into inputs to obtain certified robustness guarantees. In the LLM setting, SmoothLLM [Robey et al., 2023] perturbs input prompts at the character level to defend against adversarial jailbreaking. These methods require multiple noisy forward passes with output aggregation, and target prediction stability rather than output diversity.

Table 1: Accuracy (%) across model configurations and benchmarks. RSP reports the best result across injection positions (Prefix, Infix, Suffix); full per-position breakdown is in Appendix A. Values in parentheses indicate the difference from the baseline.

Model	MATH-500		GSM8K		AIME24	
	Baseline	RSP	Baseline	RSP	Baseline	RSP
<i>Instruct Models</i>						
LLaMA-3.1-8B-Inst	52.20	49.40 (−2.8)	85.75	86.73 (+1.0)	6.67	10.00 (+3.3)
Qwen2.5-Math-7B-Inst	83.20	84.20 (+1.0)	95.45	95.68 (+0.2)	13.33	16.67 (+3.3)
Qwen2.5-Math-1.5B-Inst	73.00	75.20 (+2.2)	85.29	85.75 (+0.5)	10.00	20.00 (+10.0)
<i>Base Models (Raw Text)</i>						
Qwen2.5-Math-7B	72.20	71.60 (−0.6)	84.15	84.31 (+0.2)	6.67	16.67 (+10.0)
Qwen2.5-Math-1.5B	61.40	64.40 (+3.0)	77.86	74.30 (−3.6)	16.67	10.00 (−6.7)
<i>Base Models + ChatML (Format Mismatch)</i>						
Qwen2.5-Math-7B	52.20	70.40 (+18.2)	58.30	75.97 (+17.7)	23.33	23.33 (+0.0)
Qwen2.5-Math-1.5B	34.40	63.40 (+29.0)	37.45	67.02 (+29.6)	13.33	20.00 (+6.7)

In reinforcement learning, NoisyNet [Fortunato et al., 2018] injects learnable noise into network weights to improve exploration, demonstrating that the location and structure of noise injection affect exploration quality. This motivation is closest to RSPs, which similarly target diversity, but differ in injection target (input embeddings vs. network weights) and require no learned noise parameters.

RSPs differ from these approaches in the following respects. (1) RSPs operate solely at inference without any training, whereas dropout and NEFTune inject noise during training, modifying the resulting model parameters accordingly. (2) RSPs preserve the original input and append fresh embedding vectors in a single forward pass, whereas robustness methods corrupt the input and aggregate over multiple noisy passes. (3) RSPs use purely random vectors sampled from the pretrained embedding statistics, whereas NoisyNet learns noise parameters through optimization.

3 Experimental Analysis

3.1 Random Soft Prompt: Definition and Setup

Let $\mathbf{W}_E \in \mathbb{R}^{V \times d}$ denote the token embedding matrix of a pretrained LLM, where V is the vocabulary size and d is the hidden dimension. We define an RSP $\mathbf{H} = [h_1, \dots, h_L] \in \mathbb{R}^{L \times d}$ as a sequence of L continuous vectors, where each vector is independently sampled from:

$$h_j \sim \mathcal{N}(\mu_E, \sigma_E^2 \mathbf{I}), \quad j = 1, \dots, L, \quad (1)$$

with $\mu_E = \frac{1}{Vd} \sum_{i,k} [\mathbf{W}_E]_{i,k}$ and $\sigma_E = \text{std}(\mathbf{W}_E)$. Let $\mathbf{X} = [x_1, \dots, x_T] \in \mathbb{R}^{T \times d}$ denote the embedding sequence of the original input. The RSP-augmented input $\tilde{\mathbf{X}}$ is constructed by concatenating $\mathbf{H}$ with $\mathbf{X}$ at one of three positions without modifying the original prompt text: *prefix* ($[\mathbf{H}; \mathbf{X}]$), *infix* ($\mathbf{H}$ inserted within $\mathbf{X}$), or *suffix* ($[\mathbf{X}; \mathbf{H}]$), following the injection strategies adopted in prior work [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Zhang et al., 2025]. Sampling a new $\mathbf{H}$ for each batch ensures that the results are not dependent on any realization of the random embeddings.

We evaluate on three mathematical reasoning benchmarks: MATH-500 [Lightman et al., 2024], GSM8K [Cobbe et al., 2021], and AIME24, using LLaMA-3.1-8B-Instruct [Grattafiori et al., 2024] and the Qwen2.5-Math model family [Yang et al., 2024] across seven configurations including instruct models, base models, and base models with mismatched chat templates. All experiments use greedy decoding to isolate the effect of RSPs from sampling stochasticity. Only the final answer is evaluated through a fallback-based extraction pipeline. Full experimental details are provided in Appendix A.

3.2 Does RSP Degrade Reasoning Performance?

We first examine how injecting random embeddings affects the model’s reasoning performance. Table 1 reports the accuracy for each model configuration, using the best RSP result across injection positions (Prefix, Infix, Suffix) and $L \in \{10, 15, 20\}$ for each benchmark.

Table 2: Comparison with existing soft prompt methods (TTSV, SoftCoT, LTPO) under unified evaluation. Baseline and RSP values are from Table 1. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Benchmark	Baseline	RSP	TTSV	SoftCoT	LTPO
Qwen2.5-Math-7B-Instruct	MATH-500	83.20	84.20	<u>83.80</u>	82.80	83.00
	GSM8K	<u>95.45</u>	95.68	95.30	95.00	94.84
	AIME24	13.33	<u>16.67</u>	23.30	<u>16.67</u>	13.33
Qwen2.5-Math-1.5B-Instruct	MATH-500	73.00	<u>75.20</u>	<u>75.20</u>	76.00	72.40
	GSM8K	85.29	85.75	<u>85.70</u>	84.46	84.53
	AIME24	<u>10.00</u>	20.00	6.70	6.67	<u>10.00</u>

For instruct models and base models with their native prompt formats, RSP injection maintains or slightly shifts baseline performance within a few percentage points. Most strikingly, under prompt-format mismatch — where base models receive ChatML templates they were not trained on — RSP yields pronounced improvements of up to +29%p. These gains are difficult to attribute to simple perturbation effects, as noise that merely raises model uncertainty would further degrade performance.

We compare RSPs with prior soft prompt methods (TTSV, SoftCoT, LTPO) that optimize their embeddings for distinct objectives, under a unified evaluation protocol reported in Table 2. All methods are re-evaluated with a shared answer extraction pipeline that applies fallback strategies, enabling a format-agnostic comparison despite each method’s distinct prompt format and grading scheme. The results place RSPs on par with the optimized methods without requiring any training.

3.3 How Does RSP Affect the Output Distribution?

We analyze how RSP injection changes the model’s output distribution during generation on Qwen2.5-Math-1.5B-Instruct with MATH-500, measuring per-token entropy, top-1 probability, and varentropy across the generation process. Figure 1 and Table 3 compare these metrics for the first 5% of generation steps across RSP variants and prior soft prompt methods.




Figure 1: Mean entropy, top-1 probability, and varentropy during the first 5% of generation steps (Qwen2.5-Math-1.5B-Instruct, MATH-500). Shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: existing soft prompt methods (LTPO, SoftCoT, TTSV).

Table 3: Mean entropy, top-1 probability, and varentropy for the first 5% of generation steps (Qwen2.5-Math-1.5B-Instruct, MATH-500).

Metric	Baseline	Prefix	Infix	Suffix	LTPO	SoftCoT	TTSV
Entropy	0.1332	0.1366	0.1402	0.1941	0.1662	0.4218	0.1359
Top-1 Prob	0.9535	0.9523	0.9525	0.9373	0.9437	0.9156	0.9534
Varentropy	0.2181	0.2207	0.2463	0.4010	0.2953	2.2234	0.2174

These results show that RSP modifies the early-stage output distribution across all injection positions: entropy increases (the token distribution flattens), top-1 probability decreases (commitment to a dominant token weakens), and varentropy increases (indicating multimodal distributions [Ahmed et al., 2026]). High-entropy positions are known to act as critical forking points that determine reasoning trajectories [Wang et al., 2025], so this shift moves the model toward a state where multiple

reasoning paths can coexist. The change is localized to the initial reasoning stage; in the middle and final portions of generation, all three metrics converge to baseline levels.

Although existing soft prompt methods are optimized for different objectives, they share a common pattern of increasing early-stage entropy. Notably, TTSV, which optimizes soft prompts using a reward that reduces the mean entropy of reasoning trajectories, still exhibits early entropy comparable to the baseline. This suggests that elevating early-stage entropy is an inherent effect of injecting non-pretrained continuous embeddings, independent of the optimization objective.

3.4 Does RSP Induce Trajectory Diversity?

Section 3.3 empirically established that RSP reshapes the early-stage output distribution. We now examine whether this shift translates into reasoning diversity at three complementary levels of analysis: token rankings, semantic content of trajectories, and task-level outcomes. All three analyses share the setting: Qwen2.5-Math-1.5B-Instruct, MATH-500, suffix injection, and $L = 10$.

Token-level: distribution beyond temperature. Under autoregressive decoding, any trajectory-level divergence between RSP and the baseline must originate from the first generation step, a critical forking point for reasoning trajectories [Wang et al., 2025]. Using the baseline at $\tau=1.0$ as reference, we compare RSP at $\tau=1.0$ against the baseline at $\tau=2.0$ (the high-temperature ceiling of what flattening alone can achieve) on three complementary metrics: Spearman rank correlation ρ to detect rank reorganization, the probability mass placed outside the baseline’s top-10 to measure support expansion, and Jensen–Shannon divergence to quantify the overall magnitude of distributional change (formal definitions in Appendix C).

Table 4: First-token distribution metrics measured against the baseline at $\tau=1.0$, comparing the baseline at $\tau=2.0$ with RSP at $\tau=1.0$.

Metric	$\tau=2.0$	RSP
Spearman ρ	1.0	0.651
Mass outside top-10	5.15%	27.64%
JS divergence	0.086	0.231

Table 4 reports the contrast. Even at $\tau=2.0$, only 5.15% of probability mass falls outside the baseline’s top-10; RSP at $\tau=1.0$ reaches 27.64%. The Spearman rank correlation tells the story more starkly: because temperature scaling $p_i^{(\tau)} \propto p_i^{1/\tau}$ is a monotone transform that preserves token ranking ($p_i > p_j \iff p_i^{(\tau)} > p_j^{(\tau)}$) for any $\tau > 0$, ρ relative to the baseline is exactly 1 as a mathematical identity. RSP, by contrast, reranks the token distribution itself, reducing ρ to 0.651—an effect temperature cannot produce at any τ . The Jensen–Shannon divergence (0.231) quantifies the overall magnitude of this distributional change. Since $\rho < 1$ rules out rank-preserving flattening as the sole cause, the high divergence reflects a structural reshape of the distribution rather than a uniform flattening of the same ranking.

Trajectory-level: semantic diversity. Token-level reranking raises a second question: do first-token perturbations produce semantically distinct reasoning paths, or do independently perturbed trajectories converge to the same solution strategy? We randomly sample 64 problems from MATH-500 and generate 300 independent trajectories per problem under two conditions, Baseline and RSP, with temperature $\tau = 1.0$. The trajectories are embedded via BGE-M3 [Chen et al., 2024] and compared on three metrics: mean pairwise cosine distance within each problem, inter-cluster distance between DBSCAN [Ester et al., 1996] centroids, and intra-cluster distance. Table 5 shows that RSP substantially increases the inter-cluster distance while leaving the intra-cluster distance unchanged: clusters become more separated while reasoning within each cluster remains consistent. RSP also wins on pairwise distance in 56 of 64 problems (87.5%).

Table 5: Semantic diversity metrics averaged across the 64 problems.

Metric	Baseline	RSP
Pairwise cos. dist.	0.0612	0.0879
Inter-cluster dist.	0.0183	0.0982
Intra-cluster dist.	0.0526	0.0528

Figure 2(a) illustrates the qualitative difference for the representative problem: the RSP samples (blue) cover the region occupied by baseline samples (red) while extending into clusters that the baseline does not reach; the expanded regions also contain correct solutions (filled markers), showing that RSP’s exploration is not limited to incorrect alternatives but reaches new reasoning paths that still yield the correct answer. Figures 2(b) and 2(c) show that this pattern holds across the 64 problems: on most problems RSP produces more clusters and larger pairwise cosine distance than baseline. The combination of large inter-cluster increase and unchanged intra-cluster distance mirrors the token-




Figure 2: Semantic diversity analysis on 64 randomly sampled MATH-500 problems with 300 trajectories per condition (Qwen2.5-Math-1.5B-Instruct, $\tau = 1.0$, BGE-M3 embeddings). (a) t-SNE projection for one representative problem (test/prealgebra/951): baseline (red) concentrates around one region while RSP suffix (blue) spreads across distinct clusters; filled markers denote correct answers and open markers denote incorrect. (b) Per-problem mean pairwise cosine distance; points above the diagonal indicate RSP produces more diverse trajectories than baseline. (c) Per-problem DBSCAN cluster count; RSP yields more distinct clusters than baseline on most problems.

level finding: RSP alters which cluster a trajectory enters while preserving the local continuation within each cluster. Additional per-problem t-SNE plots are provided in Figure 6 (Appendix B).

Outcome-level: Pass@ n scaling. Finally, we evaluate whether the token- and trajectory-level diversity translates into task-level outcome diversity through Pass@ n , the probability that at least one of n sampled solutions is correct. We compare four conditions with $n = 16$ samples per problem on MATH-500 and AIME24: (i) *Baseline*, $\tau = 1.0$: temperature sampling only; (ii) *RSP (fixed seed)*, $\tau = 1.0$: a single RSP shared across all samples; (iii) *RSP (indep. seed), greedy*: a different RSP per sample without temperature; and (iv) *RSP (indep. seed)*, $\tau = 1.0$: a different RSP per sample combined with temperature sampling.




Figure 3: Pass@ n scaling on MATH-500 (left) and AIME24 (right) with Qwen2.5-Math-1.5B-Instruct, $n = 16$ samples per problem. Baseline: temperature sampling only. Fixed seed: single RSP shared across samples combined with temperature. Indep. seed: a different RSP per sample, with or without temperature.

Condition (iv) achieves the highest Pass@ n on both benchmarks, reaching 92.80% on MATH-500 and 36.67% on AIME24 at $n = 16$. Condition (ii) shows no improvement over the baseline, confirming that the diversity gain originates from varying the random embeddings across samples, not from a single fixed perturbation. Pass@1 remains comparable to the baseline on MATH-500, and on the more challenging AIME24, (iv) achieves higher Pass@1 than the baseline. Independent RSPs combined with temperature sampling produce the most diverse reasoning trajectories while maintaining or improving single-sample accuracy.

Synthesis across the three levels. The three analyses converge on a consistent picture. At the token level, RSP induces support expansion concentrated at the first token and decaying monotonically across steps. At the trajectory level, this localized perturbation translates into distinct reasoning strategies (large inter-cluster separation) while preserving local coherence within each strategy

(unchanged intra-cluster distance). At the outcome level, the resulting trajectories cover more correct solutions without sacrificing single-sample accuracy.

4 Theoretical Analysis: Why RSP Works

4.1 RSP as Controlled Exploration on the Energy Landscape

Under the two simplifying assumptions introduced below, this section provides a theoretical account of two properties observed in our experiments: (1) RSP can change the model’s output distribution, and (2) RSP does not degrade reasoning performance. We show that RSP converts the transformer’s deterministic forward pass into a controlled stochastic process on the Hopfield energy landscape, where the noise introduced through attention naturally decays with both layer depth and generation step. We build on the equivalence between softmax attention and gradient descent on the Modern Hopfield energy [Ramsauer et al., 2021]. The analysis relies on two simplifying assumptions: (A1) $\mathbf{K} = \mathbf{V}$ in attention (exact in Hopfield models, approximate in standard transformers) and (A2) weight-tied layers (exact in deep equilibrium models [Bai et al., 2019], approximate otherwise). Full proofs are in Appendix D.

Transformer layers as energy minimization. The Modern Hopfield Energy for stored patterns $\Xi \in \mathbb{R}^{d \times m}$ and state $\sigma \in \mathbb{R}^d$ is $E(\sigma; \Xi) = -\beta^{-1} \log \sum_{j=1}^m \exp(\beta \xi_j^\top \sigma) + \frac{1}{2} \|\sigma\|^2$, where $\beta = 1/\sqrt{d}$. A full gradient step on this energy yields $\sigma^{\text{new}} = \Xi \cdot \text{softmax}(\beta \Xi^\top \sigma)$, which is identical to softmax attention [Ramsauer et al., 2021]. Each transformer layer thus performs one gradient descent step on this energy landscape: the hidden state moves toward a lower-energy configuration at every layer.

The trapping problem. This monotone descent implies that, without perturbation, the forward pass converges to the local minimum in the basin of attraction of $\sigma^{(0)}$. If a better minimum σ_g^* exists in a different basin, the deterministic forward pass cannot reach it because it cannot cross energy barriers.

Why temperature scaling and additive noise do not resolve this. Temperature scaling modifies only the final output distribution: $p^{(\tau)}(v) = \text{softmax}(\mathbf{z}^{(L)}/\tau)$, leaving the internal trajectory $\mathbf{x}^{(l+1)}|_\tau = \mathbf{x}^{(l+1)}|_{\tau=1}$ for all $l < L$ identical. It cannot help escape a local minimum during inference. Additive Gaussian noise $\mathbf{x}^{(0)} + \epsilon$ perturbs the initial state, but in subsequent layers the perturbation magnitude is bounded by $\prod_{l'} L_{f_{l'}} \cdot \sigma$ where $L_{f_{l'}}$ is the Lipschitz constant of layer l' . When $L_{f_{l'}} > 1$, the noise amplifies rather than decays, providing no controlled exploration schedule.

RSP as controlled exploration via attention. RSP resolves both limitations through two mechanisms.

First, concatenating k random vectors sampled from $\mathcal{N}(\mu_E, \sigma_E^2 \mathbf{I})$ causes each token to attend to both real and random keys, yielding

$$\mathbf{x}_i^{(l+1)} = (1 - w_{r,i}^{(l)}) \tilde{\mathbf{x}}_i^{(l+1)} + \boldsymbol{\eta}_i^{(l)}, \quad (2)$$

where $\tilde{\mathbf{x}}_i^{(l+1)}$ is the standard attention update over real tokens, and $\boldsymbol{\eta}_i^{(l)} := \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)}$ aggregates the contribution of random tokens with total attention mass $w_{r,i}^{(l)} := \sum_{j=1}^k \alpha_{i,n+j}^{(l)}$.

The perturbation magnitude is controlled by $w_{r,i}^{(l)}$, with variance scaling on the order of $(w_{r,i}^{(l)})^2$. In practice, since the embedding mean is negligible relative to its standard deviation ($|\mu_E|/\sigma_E \ll 1$), the bias term is dominated by the stochastic component, and the dynamics are well-approximated by a gradient descent process with approximately zero-mean structured noise (see Appendix D.2).

This yields a Langevin-like stochastic process that enables exploration beyond local minima.

Second, RSP modifies the energy landscape. Adding k random keys changes the Hopfield energy as

$$\tilde{E}_i - E_i = -\beta^{-1} \log \left(1 + \frac{B_i}{A_i} \right) \leq 0, \quad (3)$$

where $A_i = \sum_j \exp(\beta \mathbf{k}_j^\top \mathbf{q}_i)$ and $B_i = \sum_j \exp(\beta \mathbf{k}_{r,j}^\top \mathbf{q}_i) > 0$. The magnitude scales with $w_{r,i}^{(l)} / (1 - w_{r,i}^{(l)})$, so tokens assigning more attention to random keys experience larger energy perturbations.

While this does not directly reduce barrier height, it alters the local energy landscape, and together with stochastic perturbations increases the probability of crossing energy barriers between basins.

Both mechanisms are governed by $w_{r,i}^{(l)}$, which controls the overall strength of exploration.

Why performance is preserved: dual decay of $w_{r,i}^{(l)}$. The attention decay bound is:

$$w_{r,i}^{(l)} \leq \frac{k}{k + n \cdot \exp(\Delta_i^{(l)} / \sqrt{d})}, \quad (4)$$

where $\Delta_i^{(l)} := \min_{j \in [n]} \mathbf{q}_i^\top \mathbf{k}_j - \max_{j' \in [k]} \mathbf{q}_i^\top \mathbf{k}_{r,j'}$, is the real-vs-random logit gap and n is the number of real tokens. Since the bias, stochastic scale, and energy modification are all controlled by $w_{r,i}^{(l)}$, this bound governs the overall perturbation strength.

Layer-wise decay. If deeper layers increase the real-vs-random logit gap $\Delta_i^{(l)}$, then $w_{r,i}^{(l)}$ correspondingly decays toward zero. Empirically, as deeper layers develop task-specific representations, $\Delta_i^{(l)}$ tends to increase, making RSP’s perturbation strongest in early layers and negligible in deeper layers.

Sequence-wise decay. In autoregressive generation, each newly generated token adds a real key to the KV cache, so n increases at every generation step while k remains fixed. The growing n in Eq. (4) causes $w_{r,i}^{(l)}$ to decrease monotonically, progressively diluting RSP’s influence as the sequence grows.

This dual decay provides a natural annealing schedule: RSP enables exploration in the early stage of both the internal forward pass and the generation sequence, while the direct influence of the perturbation progressively diminishes as the model accumulates real tokens and converges to task-specific representations.

Empirical verification. We empirically verify this dual decay on Qwen2.5-Math-7B with suffix injection over 500 MATH-500 problems. For each problem, we measure the per-token attention mass directed to the RSP embeddings and to the question tokens within the same sequence, with the question tokens serving as a baseline that factors out sequence-length effects. Figure 4 shows the expected contrast: the question tokens, which remain task-relevant throughout reasoning, receive attention whose distribution varies across layers and reasoning positions without a consistent trend, while the RSP embeddings, which carry no task-relevant content and act as noise, receive attention that decreases along both the layer and sequence axes, directly instantiating the layer-wise and sequence-wise dilution predicted by Eq. (4).

4.2 Output Distribution Shift and Exploration

We study how Random Soft Prompting (RSP) alters the output distribution of an autoregressive language model and how this leads to improved exploration.

In autoregressive decoding, generation proceeds via prefix-conditioned distributions $p(y_t | x, y_{<t})$. RSP does not modify a single global distribution, but perturbs the sequence of conditional distributions across decoding steps. In Section 4.1, we analyzed this from an internal, layer-wise and sequence-wise perspective, showing that RSP introduces structured perturbations through attention that naturally decay across layers. We now examine how these perturbations manifest at the output level.

Limitation of temperature scaling. Temperature scaling transforms probabilities as $p_i^{(\tau)} \propto p_i^{1/\tau}$ and preserves token ranking, i.e., $p_i > p_j \iff p_i^{(\tau)} > p_j^{(\tau)}$. Hence it cannot introduce new token orderings at a fixed decoding state.

Effect of RSP. RSP induces stochastic perturbations to the hidden state, which we locally approximate as $z_i(h) \approx z_i + a_i^\top h$, where h is a random soft prompt.




Figure 4: Per-token attention mass on Qwen2.5-Math-7B (suffix, 500 MATH-500 problems, each with an independently sampled RSP). The layer and reasoning-position axes are each partitioned into five quantile bins; every cell reports the 500-sample mean per-token attention mass directed to (a) the RSP embeddings and (b) the question tokens within that region. The horizontal axis spans reasoning position from early (left) to late (right) generation steps, and the vertical axis spans layer depth from shallow (top) to deep (bottom). Attention to the RSP embeddings in (a) decays along both axes — concentrating at shallow layers and early reasoning positions — while attention to the question tokens in (b) remains comparatively uniform, visually confirming the layer-wise and sequence-wise dilution predicted by Eq. (4).

295 **Proposition 1 (Stochastic ranking).** If there exists a token pair (i, j) such that $a_i \neq a_j$, then

$$0 < \Pr_h(z_i(h) > z_j(h)) < 1.$$

296 Thus, unlike temperature scaling, RSP randomizes token rankings at each decoding step.

297 **Proposition 2 (Support expansion).** Let $\text{TopK}_{\text{Temp}}(s_t)$ be the top- K tokens under temperature
 298 scaling at state s_t . Under the same condition, there exists $i \notin \text{TopK}_{\text{Temp}}(s_t)$ such that

$$\Pr_h(i \in \text{TopK}(p^{\text{RSP}}(\cdot | s_t, h))) > 0.$$

299 **Implication.** RSP expands the effective support of the output distribution, enabling tokens that are
 300 effectively inaccessible under temperature scaling. Since autoregressive generation composes these
 301 distributions over time, such perturbations can lead to distinct trajectories, explaining the empirical
 302 gains in exploration.

303 5 Conclusion

304 **Summary.** This work investigated the role of Random Soft Prompts (RSPs), continuous embeddings
 305 sampled from the pretrained token embedding distribution and injected into the model input without
 306 any training. Through empirical and theoretical analyses, we established three findings. First, RSPs
 307 increase early-stage token entropy, top-1 probability reduction, and varentropy while the model’s
 308 output stabilizes in later generation stages (Section 3). Second, RSPs induce stochastic token rankings
 309 and expand the effective support of the output distribution, a property that temperature sampling
 310 fundamentally lacks due to its ranking-preserving nature (Section 4). Third, combining RSPs with
 311 temperature sampling produces more diverse reasoning trajectories than either method alone, as
 312 measured by $\text{Pass}@n$ scaling (Section 3).

313 **Implications for GRPO training.** GRPO computes advantages by comparing multiple sampled
 314 trajectories within each group, making trajectory diversity essential for effective learning signals.
 315 Existing approaches such as DAPO [Yu et al., 2025] and CE-GPPO [Su et al., 2025] address entropy
 316 collapse by modifying the clipping range or reintroducing gradients from clipped tokens, but the
 317 underlying sampling process remains based on temperature sampling, which preserves token rankings
 318 and cannot expand the output support. Our theoretical and empirical results suggest that RSP
 319 can address this limitation at the sampling level: by perturbing the initial output distribution with
 320 independent random embeddings per rollout, RSP provides a complementary source of diversity
 321 that is orthogonal to temperature-based exploration. We expect this to yield more effective learning
 322 signals for GRPO training, particularly in mitigating entropy collapse.

Limitations. Our experiments are conducted on mathematical reasoning benchmarks (MATH-500, GSM8K, AIME24) with a limited set of model families (Qwen2.5-Math, LLaMA-3.1). The energy-theoretic analysis in Section 4 relies on strong assumptions: the Hopfield–Attention equivalence (A1: $\mathbf{K} = \mathbf{V}$) and weight-tied layers (A2), which hold exactly only in Hopfield-type models and deep equilibrium models, respectively. Since real transformer architectures vary widely in their attention mechanisms and layer structures, the theoretical predictions are most directly applicable to models whose attention dynamics closely approximate the Hopfield energy framework. The GRPO application remains a theoretical expectation and has not been empirically validated.

Future work. Extending the experimental evaluation to other domains (e.g., code generation, commonsense reasoning) and larger model scales would test the generality of RSP’s effects. Empirically validating the GRPO application by training with RSP-augmented rollouts and comparing against DAPO and CE-GPPO is a natural next step. Investigating the optimal RSP distribution and token count, which currently require per-model tuning, could lead to more principled injection strategies.

References

- Farhan Ahmed, Yuya Jeremy Ong, and Chad DeLuca. LogitScope: A Framework for Analyzing LLM Uncertainty through Information Metrics. In *ICLR Workshop*, 2026.
- Shaojie Bai, J. Zico Kolter, and Vladlen Koltun. Deep Equilibrium Models. In *NeurIPS*, 2019.
- Nora Belrose, Zach Furman, Logan Smith, Danny Halawi, Igor Ostrovsky, Lev McKinney, Stella Biderman, and Jacob Steinhardt. Eliciting Latent Predictions from Transformers with the Tuned Lens. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. BGE M3-Embedding: Multi-Lingual, Multi-Functionality, Multi-Granularity Text Embeddings Through Self-Knowledge Distillation. *arXiv:2402.03216*, 2024.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training Verifiers to Solve Math Word Problems. *arXiv:2110.14168*, 2021.
- Jeremy M. Cohen, Elan Rosenfeld, and J. Zico Kolter. Certified Adversarial Robustness via Randomized Smoothing. In *ICML*, 2019.
- Timothée Darcet, Maxime Oquab, Julien Mairal, and Piotr Bojanowski. Vision Transformers Need Registers. In *International Conference on Learning Representations (ICLR)*, 2024.
- Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu. A Density-Based Algorithm for Discovering Clusters in Large Spatial Databases with Noise. In *Proceedings of the Second International Conference on Knowledge Discovery and Data Mining (KDD)*, pages 226–231, 1996.
- Meire Fortunato, Mohammad Gheshlaghi Azar, Bilal Piot, Jacob Menick, Ian Osband, Alex Graves, Vlad Mnih, Remi Munos, Demis Hassabis, Olivier Pietquin, Charles Blundell, and Shane Legg. Noisy Networks for Exploration. In *ICLR*, 2018.
- Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. Transformer Feed-Forward Layers Are Key-Value Memories. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2021.
- Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh Nagarajan. Think before You Speak: Training Language Models with Pause Tokens. In *ICLR*, 2024.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, et al. The Llama 3 Herd of Models. *arXiv:2407.21783*, 2024.
- Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian. Training Large Language Models to Reason in a Continuous Latent Space. In *ICLR*, 2025.

369 Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin de Laroussilhe, Andrea
370 Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-Efficient Transfer Learning for NLP.
371 In *ICML*, 2019.

372 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang,
373 and Weizhu Chen. LoRA: Low-Rank Adaptation of Large Language Models. In *ICLR*, 2022.

374 Neel Jain, Ping-yeh Chiang, Yuxin Wen, John Kirchenbauer, Hong-Min Chu, Gowthami Somepalli,
375 Brian R. Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Aniruddha Saha, Micah Goldblum,
376 Jonas Geiping, and Tom Goldstein. NEFTune: Noisy Embeddings Improve Instruction Finetuning.
377 In *ICLR*, 2024.

378 Xinyue Kang, Diwei Shi, and Li Chen. Model Whisper: Steering Vectors Unlock Large Language
379 Models’ Potential in Test-time. In *AAAI*, 2026.

380 Brian Lester, Rami Al-Rfou, and Noah Constant. The Power of Scale for Parameter-Efficient Prompt
381 Tuning. In *EMNLP*, 2021.

382 Xiang Lisa Li and Percy Liang. Prefix-Tuning: Optimizing Continuous Prompts for Generation. In
383 *ACL-IJCNLP*, 2021.

384 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike,
385 John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s Verify Step by Step. In *ICLR*, 2024.

386 Xiao Liu, Yanan Zheng, Zhengxiao Du, Ming Ding, Yujie Qian, Zhilin Yang, and Jie Tang. GPT
387 Understands, Too. *arXiv:2103.10385*, 2023.

388 Aman Madaan and Amir Yazdanbakhsh. Text and Patterns: For Effective Chain of Thought, It Takes
389 Two to Tango. *arXiv:2209.07686*, 2022.

390 Hubert Ramsauer, Bernhard Schöfl, Johannes Lehner, Philipp Seidl, Michael Widrich, Thomas Adler,
391 Lukas Gruber, Markus Holzleitner, Milena Pavlović, Geir Kjetil Sandve, Thomas Unterthiner,
392 Bernhard Nessler, and Sepp Hochreiter. Hopfield Networks Is All You Need. In *ICLR*, 2021.

393 Alexander Robey, Eric Wong, Hamed Hassani, and George J. Pappas. SmoothLLM: Defending Large
394 Language Models Against Jailbreaking Attacks. *arXiv:2310.03684*, 2023.

395 Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang,
396 Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the Limits of
397 Mathematical Reasoning in Open Language Models. *arXiv:2402.03300*, 2024.

398 Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov.
399 Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *Journal of Machine*
400 *Learning Research*, 15:1929–1958, 2014.

401 Zhenpeng Su, Leiyu Pan, Minxuan Lv, Yuntao Li, Wenping Hu, Fuzheng Zhang, Kun Gai, and Guorui
402 Zhou. CE-GPPO: Coordinating Entropy via Gradient-Preserving Clipping Policy Optimization in
403 Reinforcement Learning. *arXiv:2509.20712*, 2025.

404 Shenzhi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen,
405 Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen
406 Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 Rule: High-Entropy Minority Tokens Drive
407 Effective Reinforcement Learning for LLM Reasoning. In *NeurIPS*, 2025.

408 Yige Xu, Xu Guo, Zhiwei Zeng, and Chunyan Miao. SoftCoT: Soft Chain-of-Thought for Efficient
409 Reasoning with LLMs. In *ACL*, 2025.

410 An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu,
411 Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu,
412 Xingzhang Ren, and Zhenru Zhang. Qwen2.5-Math Technical Report: Toward Mathematical
413 Expert Model via Self-Improvement. *arXiv:2409.12122*, 2024.

414 Wengao Ye, Yan Liang, and Lianlei Shan. Thinking on the Fly: Test-Time Reasoning Enhancement
415 via Latent Thought Policy Optimization. In *ICLR*, 2026.

- 416 Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian
417 Fan, Gaohong Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming Sheng,
418 Yuxuan Tong, Chi Zhang, Mofan Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu, Jiaze Chen, Jiangjie
419 Chen, Chengyi Wang, Hongli Yu, Yuxuan Song, Xiangpeng Wei, Hao Zhou, Jingjing Liu, Wei-
420 Ying Ma, Ya-Qin Zhang, Lin Yan, Mu Qiao, Yonghui Wu, and Mingxuan Wang. DAPO: An
421 Open-Source LLM Reinforcement Learning System at Scale. *arXiv:2503.14476*, 2025.
- 422 Guibin Zhang, Muxin Fu, and Shuicheng Yan. MemGen: Weaving Generative Latent Memory for
423 Self-Evolving Agents. *arXiv:2509.24704*, 2025.

424 A Experimental Setup and Full Accuracy Breakdown

425 Table 6 provides the full per-position accuracy breakdown (Prefix, Infix, Suffix) corresponding to
 426 the best-across-positions results summarized in Table 1 in the main text. Table 7 lists the number of
 427 RSP tokens L used for each model–benchmark pair, selected from $\{10, 15, 20\}$. All experiments are
 428 conducted on a single NVIDIA A6000 40GB GPU with greedy decoding, a maximum generation
 429 length of 3,072 tokens for MATH-500 and GSM8K, and 4,096 tokens for AIME24. Batch size is
 430 fixed per model (16 for LLaMA-3.1-8B-Instruct and Qwen2.5-Math-7B variants, 32 for Qwen2.5-
 431 Math-1.5B variants).

Table 6: Accuracy (%) across model configurations, injection positions, and benchmarks. Values in parentheses indicate the difference from the baseline. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Mode	MATH-500	GSM8K	AIME24
<i>Instruct Models</i>				
LLaMA-3.1-8B-Instruct	Baseline	52.20	85.75	<u>6.67</u>
	Prefix	45.00 (−7.2)	76.35 (−9.4)	3.33 (−3.3)
	Infix	49.40 (−2.8)	85.97 (+0.2)	10.00 (+3.3)
	Suffix	<u>49.40</u> (−2.8)	86.73 (+1.0)	10.00 (+3.3)
Qwen2.5-Math-7B-Instruct	Baseline	83.20	95.45	<u>13.33</u>
	Prefix	81.40 (−1.8)	94.92 (−0.5)	13.33 (+0.0)
	Infix	84.20 (+1.0)	95.60 (+0.2)	16.67 (+3.3)
	Suffix	<u>83.40</u> (+0.2)	95.68 (+0.2)	16.67 (+3.3)
Qwen2.5-Math-1.5B-Instruct	Baseline	73.00	85.29	10.00
	Prefix	74.80 (+1.8)	85.29 (+0.0)	<u>13.33</u> (+3.3)
	Infix	75.20 (+2.2)	85.75 (+0.5)	6.67 (−3.3)
	Suffix	74.20 (+1.2)	84.69 (−0.6)	20.00 (+10.0)
<i>Base Models (Raw Text)</i>				
Qwen2.5-Math-7B	Baseline	72.20	84.15	6.67
	Prefix	<u>71.60</u> (−0.6)	84.31 (+0.2)	16.67 (+10.0)
	Infix	63.20 (−9.0)	64.29 (−19.9)	16.67 (+10.0)
	Suffix	55.60 (−16.6)	54.13 (−30.0)	<u>13.33</u> (+6.7)
Qwen2.5-Math-1.5B	Baseline	<u>61.40</u>	77.86	16.67
	Prefix	64.40 (+3.0)	74.30 (−3.6)	10.00 (−6.7)
	Infix	63.20 (+1.8)	73.92 (−3.9)	<u>10.00</u> (−6.7)
	Suffix	54.60 (−6.8)	58.61 (−19.3)	3.33 (−13.3)
<i>Base Models + ChatML (Format Mismatch)</i>				
Qwen2.5-Math-7B	Baseline	52.20	58.30	23.33
	Prefix	59.20 (+7.0)	56.41 (−1.9)	<u>3.33</u> (−20.0)
	Infix	<u>62.00</u> (+9.8)	69.07 (+10.8)	23.33 (+0.0)
	Suffix	70.40 (+18.2)	75.97 (+17.7)	23.33 (+0.0)
Qwen2.5-Math-1.5B	Baseline	34.40	37.45	<u>13.33</u>
	Prefix	63.40 (+29.0)	67.02 (+29.6)	20.00 (+6.7)
	Infix	39.20 (+4.8)	50.42 (+13.0)	6.67 (−6.7)
	Suffix	<u>51.00</u> (+16.6)	<u>50.64</u> (+13.2)	<u>13.33</u> (+0.0)

Table 7: Number of RSP tokens (L) used for each model and benchmark.

Model	MATH-500	GSM8K	AIME24
LLaMA-3.1-8B-Instruct	15	20	15
Qwen2.5-Math-7B-Instruct	20	20	20
Qwen2.5-Math-1.5B-Instruct	20	10	20
Qwen2.5-Math-7B	10	10	20
Qwen2.5-Math-1.5B	10	10	20
Qwen2.5-Math-1.5B (ChatML)	20	20	10
Qwen2.5-Math-7B (ChatML)	20	20	20

432 A.1 RSP Injection Positions and Prompt Templates

433 Below we show the prompt structure for each injection position across all model types. Blue text
434 indicates RSP embeddings, and purple text indicates special tokens.

435 A.1.1 Prefix

436 RSP embeddings are concatenated before the prompt embeddings. No text modification is applied.

437 ChatML (Qwen Instruct / Base + ChatML).

```
[Random Embeddings (L tokens)]
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{>
<|im_start|>user
{question}<|im_end|>
<|im_start|>assistant
```

439 LLaMA Chat Template.

```
[Random Embeddings (L tokens)]
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{>
<|start_header_id|>user<|end_header_id|>
{question}<|eot_id|>
<|start_header_id|>assistant<|end_header_id|>
```

441 Raw Text (Qwen Base).

```
[Random Embeddings (L tokens)]
{question}
Please reason step by step, and put your final answer within \boxed{>
```

443 A.1.2 Infix

444 A description text and special tokens are inserted into the prompt. The special token positions are
445 then replaced with RSP embeddings at the embedding level.

446 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{>
<|im_start|>user
{question}

There are {L} special tokens that contain compressed latent reasoning information that
might be useful for your reasoning. If these tokens are useful for your case, you can
use them as reference. If these tokens are not useful for your case, you can ignore
them and focus back to solving the problem.

Here are the {L} special tokens: <special_token_1>...<special_token_L>
<|im_end|>
<|im_start|>assistant
```

448 LLaMA Chat Template.

```
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{>
<|start_header_id|>user<|end_header_id|>
```

{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens:

```
<reserved_special_token_0>...  
<reserved_special_token_L>  
<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

450

451 Raw Text (Qwen Base).

{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens: $\langle \text{special_token_1} \rangle \dots \langle \text{special_token_L} \rangle$

Please reason step by step, and put your final answer within `\boxed{\}`.

452

453 A.1.3 Suffix

454 For chat-templated models, RSP embeddings are inserted immediately before the end-of-turn token.

455 For raw text models, RSP embeddings are concatenated at the end of the prompt.

456 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system  
Please reason step by step, and put your final answer within \boxed{ }.<|im_end|>  
<|im_start|>user  
{question}[Random Embeddings ( $L$  tokens)]<|im_end|>  
<|im_start|>assistant
```

457

458 LLaMA Chat Template.

```
<|begin_of_text|>  
<|start_header_id|>system<|end_header_id|>  
Cutting Knowledge Date: December 2023  
Today Date: 26 Jul 2024  
Please reason step by step, and put your final answer within \boxed{ }.<|eot_id|>  
<|start_header_id|>user<|end_header_id|>  
{question}[Random Embeddings ( $L$  tokens)]<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

459

460 Raw Text (Qwen Base).

```
{question}  
Please reason step by step, and put your final answer within \boxed{ }. [Random  
Embeddings ( $L$  tokens)]
```

461

462 A.2 Answer Extraction and Grading

463 The final answer is extracted from the model output through a fallback chain. Each step is attempted

464 in order; if extraction fails, the next step is tried.

Answer Extraction Fallback Chain

1. "final answer is \$...\$. I hope" pattern (Minerva format)
2. `\boxed{...}`: last non-empty box is used; empty `\boxed{}` are skipped
3. Text following "the answer is"
4. Text following "final answer is"
5. Last number in the output (regex fallback)

Extracted answers are normalized by removing \$, \left, \right, and unit strings. Grading is performed via symbolic equivalence checking: exact string match, numerical comparison (`math.isclose, rel_tol=1e-4`), and SymPy symbolic simplification.

A.3 Reproduced Prior Work Hyperparameters

All prior methods are reproduced on Qwen2.5-Math-1.5B-Instruct and Qwen2.5-Math-7B-Instruct, evaluated with greedy decoding (temperature = 0) and our unified answer extraction pipeline.

TTSV. Prefix length 20, trained for 20 epochs with AdamW optimizer, batch size 2, gradient accumulation 8 steps, and linear learning rate schedule. Learning rate is 1e-3 for Qwen-1.5B and 1e-5 for Qwen-7B.

LTPO. Table 8 reports the per-benchmark hyperparameters.

Table 8: LTPO hyperparameters.

Parameter	GSM8K	MATH-500	AIME24
tokens	8	8	8
steps	20	20	20
top_k	10	10	10
sigma	20	20	5
sigma_decay	0.95	0.95	0.95
lr	1e-2	5e-3	5e-2

SoftCoT. Projection module trained for 10 epochs. Thought tokens: 32 during training, 4 during evaluation. Batch size 4 for GSM8K, 1 for MATH with gradient accumulation 4. The small (assistant) model is Qwen2.5-Math-1.5B-Instruct in both configurations, while the large model is Qwen2.5-Math-1.5B-Instruct (learning rate 2e-5) or Qwen2.5-Math-7B-Instruct (learning rate 5e-6). For MATH-500 and AIME24 evaluation, we use the projection module trained on GSM8K, as it yields higher accuracy than the one trained on MATH.

A.4 Full Entropy Curves

Figure 5 shows the full normalized (0–100%) generation trajectories for entropy, top-1 probability, and varentropy. All methods converge to similar levels after the initial generation stage, confirming that the distributional changes induced by RSP are localized to the early reasoning phase.




Figure 5: Mean entropy (top), top-1 probability (middle), and varentropy (bottom) over the full normalized generation trajectory (0–100%). Averaged over MATH-500, shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: existing soft prompt methods. All methods converge after the initial stage.

486 B Additional Semantic Diversity Analysis

487 This appendix complements the semantic diversity results reported in Section 3.4. We include
 488 per-problem t-SNE plots across a broader set of problems.



Figure 6: t-SNE projections of BGE-M3 embeddings for 48 out of the 64 evaluated MATH-500 problems (6×8 grid). Each subplot shows 300 baseline (red) and 300 RSP suffix (blue) samples under $\tau = 1.0$ with Qwen2.5-Math-1.5B-Instruct; filled markers denote correct answers and open markers denote incorrect. Subplot titles indicate subject/problem_id. Across most problems, RSP trajectories spread across multiple distinct regions while baseline samples concentrate in a single dominant cluster.

C Token-Level Distribution Metrics

This appendix provides formal definitions for the three metrics used in the Token-level analysis of Section 3.4. Let P_{base} denote the baseline next-token distribution at $\tau = 1.0$ and P_{target} the candidate distribution being compared (e.g., the baseline at $\tau = 2.0$ or RSP at $\tau = 1.0$). All metrics are computed analytically from saved logits at the first generation step.

Spearman rank correlation. Spearman’s ρ is the Pearson correlation coefficient between the token ranks of two distributions:

$$\rho = \text{Pearson}(\text{rank}(P_{\text{target}}), \text{rank}(P_{\text{base}})). \quad (5)$$

Because temperature scaling $P^{(\tau)} \propto P^{1/\tau}$ is a strictly monotone transform, the rank order of tokens is preserved exactly, so $\rho = 1$ for any $\tau > 0$ as a mathematical identity. Any value $\rho < 1$ therefore indicates rank reorganization beyond what temperature scaling can produce.

Mass outside top- K . Let $\text{TopK}(P_{\text{base}})$ be the set of K tokens with the highest probability under P_{base} . The probability mass that the candidate distribution places *outside* this set is

$$\text{MassOutside}_K = 1 - \sum_{t \in \text{TopK}(P_{\text{base}})} P_{\text{target}}(t). \quad (6)$$

This directly measures support expansion: how much probability the candidate assigns to tokens that are not in the baseline’s preferred set. Reported values use $K = 10$.

Jensen–Shannon divergence. The symmetrized, bounded version of Kullback–Leibler divergence:

$$\text{JS}(P, Q) = \frac{1}{2} \text{KL}(P \parallel M) + \frac{1}{2} \text{KL}(Q \parallel M), \quad M = \frac{1}{2}(P + Q). \quad (7)$$

We use base-2 logarithms so that $\text{JS} \in [0, 1]$. Tokens absent from the stored top-100 are assigned a sentinel logit (-10^9) before softmax, which is negligible for our setting since the top-5 covers more than 99.9% of the probability mass.

D Proofs for Energy-Theoretic Analysis

D.1 Deterministic Trapping

Lemma 1 (Descent lemma). *If E is L_E -smooth and $\eta \leq 1/L_E$, then $E(\boldsymbol{\sigma}^{\text{new}}) \leq E(\boldsymbol{\sigma}) - \frac{\eta}{2} \|\nabla E(\boldsymbol{\sigma})\|^2$.*

Proof. By L_E -smoothness, $E(\boldsymbol{\sigma} - \eta \nabla E) \leq E(\boldsymbol{\sigma}) - \eta \|\nabla E\|^2 + \frac{L_E \eta^2}{2} \|\nabla E\|^2 = E(\boldsymbol{\sigma}) - \eta(1 - \frac{L_E \eta}{2}) \|\nabla E\|^2$. Since $\eta \leq 1/L_E$, we have $L_E \eta/2 \leq 1/2$, giving $E(\boldsymbol{\sigma}^{\text{new}}) \leq E(\boldsymbol{\sigma}) - \frac{\eta}{2} \|\nabla E\|^2$. $\square$

Theorem 1 (Deterministic trapping). *Without perturbation, the forward pass converges to the local minimum $\boldsymbol{\sigma}^*$ in the basin of attraction of $\boldsymbol{\sigma}^{(0)}$, which may satisfy $E(\boldsymbol{\sigma}^*) > E(\boldsymbol{\sigma}_g^*)$.*

Proof. Lemma 1 guarantees monotone energy decrease at each layer: $E(\boldsymbol{\sigma}^{(l+1)}) \leq E(\boldsymbol{\sigma}^{(l)})$. Since E is bounded below (the quadratic term dominates), the sequence $\{E(\boldsymbol{\sigma}^{(l)})\}$ converges, implying $\|\nabla E(\boldsymbol{\sigma}^{(l)})\| \rightarrow 0$ and convergence to a critical point. Deterministic dynamics follow a unique trajectory from $\boldsymbol{\sigma}^{(0)}$ and cannot cross energy barriers to reach minima in other basins. $\square$

D.2 Properties of the RSP-Induced Perturbation

Decomposition of attention output. The attention update with n real keys and k random keys can be decomposed as

$$\mathbf{x}_i^{(l+1)} = (1 - w_{r,i}^{(l)}) \tilde{\mathbf{x}}_i^{(l+1)} + \boldsymbol{\eta}_i^{(l)}, \quad (8)$$

where

$$\tilde{\mathbf{x}}_i^{(l+1)} = \sum_{j=1}^n \tilde{\alpha}_{ij}^{(l)} \mathbf{v}_j^{(l)}, \quad \tilde{\alpha}_{ij}^{(l)} = \frac{\alpha_{ij}^{(l)}}{1 - w_{r,i}^{(l)}}, \quad (9)$$

and

$$\boldsymbol{\eta}_i^{(l)} = \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)}. \quad (10)$$

We now analyze the properties of the perturbation term $\boldsymbol{\eta}_i^{(l)}$.

Conditional expectation. Conditioned on the attention weights, we have

$$\mathbb{E}[\boldsymbol{\eta}_i^{(l)} \mid \alpha^{(l)}] = \mu_E w_{r,i}^{(l)}, \quad (11)$$

where $w_{r,i}^{(l)} := \sum_{j=1}^k \alpha_{i,n+j}^{(l)}$.

Conditional covariance. Assuming the random vectors are independent with covariance $\sigma_E^2 \mathbf{I}$,

$$\text{Cov}(\boldsymbol{\eta}_i^{(l)} \mid \alpha^{(l)}) = \sigma_E^2 \sum_{j=1}^k (\alpha_{i,n+j}^{(l)})^2 \mathbf{I}. \quad (12)$$

Using $\sum_j \alpha_j^2 \leq (\sum_j \alpha_j)^2$, we obtain

$$\text{Cov}(\boldsymbol{\eta}_i^{(l)} \mid \alpha^{(l)}) \preceq \sigma_E^2 (w_{r,i}^{(l)})^2 \mathbf{I}. \quad (13)$$

Bias-to-variance ratio. The ratio between the mean and standard deviation is bounded by

$$\frac{\|\mathbb{E}[\boldsymbol{\eta}_i^{(l)} \mid \alpha^{(l)}]\|}{\sqrt{\text{Var}(\boldsymbol{\eta}_i^{(l)} \mid \alpha^{(l)})}} \leq \frac{|\mu_E|}{\sigma_E}. \quad (14)$$

Empirically, $|\mu_E|/\sigma_E \ll 1$, implying that the stochastic component dominates the bias.

Table 9: Embedding matrix statistics. $|\mu_E|/\sigma_E < 0.02$ across all models.

Model	μ_E	σ_E	$ \mu_E /\sigma_E$
LLaMA-3.1-8B-Instruct	0.00002	0.01057	0.002
Qwen2.5-Math-7B-Instruct	-0.00002	0.01647	0.001
Qwen2.5-Math-1.5B-Instruct	0.00019	0.01585	0.012
Qwen2.5-Math-7B	-0.00003	0.01842	0.001
Qwen2.5-Math-1.5B	0.00038	0.02990	0.013

Empirical verification. Table 9 reports the embedding matrix statistics for all models used in our experiments.

Implication. Therefore, the perturbation behaves as approximately zero-mean noise whose magnitude is controlled by $w_{r,i}^{(l)}$, justifying the Langevin-like interpretation in the main text.

D.3 Attention Decay Bound

Theorem 2 (Attention decay). *The total attention on random tokens satisfies*

$$w_{r,i}^{(l)} \leq \frac{k}{k + n \cdot \exp(\Delta_i^{(l)}/\sqrt{d})}, \quad (15)$$

where $\Delta_i^{(l)} := \min_{j \in [n]} \mathbf{q}_i^\top \mathbf{k}_j - \max_{j' \in [k]} \mathbf{q}_i^\top \mathbf{k}_{r,j'}$ is the real-vs-random logit gap.

Proof. Let $s_j = \mathbf{q}_i^\top \mathbf{k}_j / \sqrt{d}$ for real keys and $s_{r,j'} = \mathbf{q}_i^\top \mathbf{k}_{r,j'} / \sqrt{d}$ for random keys. By definition of $\Delta_i^{(l)}$, each real key satisfies $s_j \geq \Delta_i^{(l)}/\sqrt{d} + s_{r,\max}$ where $s_{r,\max} = \max_{j'} s_{r,j'}$. Therefore $\exp(s_j) \geq \exp(\Delta_i^{(l)}/\sqrt{d}) \cdot \exp(s_{r,\max})$ for every real key. Summing over n real keys:

$$\sum_{j=1}^n \exp(s_j) \geq n \cdot \exp(\Delta_i^{(l)}/\sqrt{d}) \cdot \exp(s_{r,\max}) \geq n \cdot \exp(\Delta_i^{(l)}/\sqrt{d}) \cdot \frac{1}{k} \sum_{j'=1}^k \exp(s_{r,j'}). \quad (16)$$

The total attention on random tokens is $w_{r,i}^{(l)} = \sum_{j'} \exp(s_{r,j'}) / (\sum_j \exp(s_j) + \sum_{j'} \exp(s_{r,j'}))$. Substituting the lower bound on the real-key sum and letting $R = \sum_{j'} \exp(s_{r,j'})$:

$$w_{r,i}^{(l)} = \frac{R}{\sum_j \exp(s_j) + R} \leq \frac{R}{n \cdot \exp(\Delta_i^{(l)}/\sqrt{d}) \cdot R/k + R} = \frac{k}{k + n \cdot \exp(\Delta_i^{(l)}/\sqrt{d})}. \quad (17)$$

□

D.4 Proofs for Output Distribution Shift

D.4.1 Autoregressive Formulation

In autoregressive decoding, the joint distribution over a sequence $y_{1:T}$ factorizes as

$$\Pr(y_{1:T} \mid x) = \prod_{t=1}^T p(y_t \mid x, y_{<t}),$$

where each step is governed by a prefix-conditioned distribution. Any perturbation introduced by RSP affects generation through a sequence of local changes to $p(y_t \mid x, y_{<t})$, rather than a single global distribution.

D.4.2 Local Linearization of Logits under RSP

Let h denote the random soft prompt. We model the perturbed logits as a function $z_i(h)$, assuming local smoothness in a neighborhood of $h = 0$:

$$z_i(h) = z_i(0) + \nabla_h z_i(0)^\top h + r_i(h),$$

where $r_i(h) = O(\|h\|^2)$. Defining $a_i := \nabla_h z_i(0)$, we obtain the local linear approximation $z_i(h) \approx z_i + a_i^\top h$. This approximation is valid for sufficiently small $\|h\|$ and is used to characterize the direction of perturbation rather than exact logit values.

556 **D.4.3 Proof of Proposition 1 (Stochastic Ranking)**

557 *If there exists a token pair (i, j) such that $a_i \neq a_j$, then $0 < \Pr_h(z_i(h) > z_j(h)) < 1$.*

558 *Proof.* Under the local linear approximation, $z_i(h) - z_j(h) = (z_i - z_j) + (a_i - a_j)^\top h$. Let
 559 $b_{ij} := a_i - a_j \neq 0$. If $h \sim \mathcal{N}(0, \sigma^2 I)$, then $b_{ij}^\top h \sim \mathcal{N}(0, \sigma^2 \|b_{ij}\|^2)$, a non-degenerate Gaussian.
 560 The perturbed logit difference $\Delta_{ij}(h) := (z_i - z_j) + b_{ij}^\top h$ therefore has non-zero probability mass
 561 on both sides of zero: $0 < \Pr_h(\Delta_{ij}(h) > 0) < 1$. $\square$

562 **D.4.4 On Proposition 2 (Support Expansion)**

563 Proposition 2 states that tokens outside the baseline top- K set can enter the top- K set under RSP
 564 with positive probability. Proposition 1 establishes that pairwise ranking between tokens becomes
 565 stochastic. Entering the top- K set requires surpassing multiple tokens simultaneously, i.e., conditions
 566 of the form $z_i(h) > z_j(h)$ for all $j \in J$ where J is a subset of higher-ranked tokens. Under the local
 567 linear model, these inequalities define a region in h -space. If this region has non-zero measure under
 568 the distribution of h , then top- K entry occurs with positive probability. We interpret Proposition 2 as a
 569 support-expansion statement: RSP enables tokens that are effectively inaccessible under deterministic
 570 ranking to become reachable under stochastic perturbations.

571 **D.4.5 Trajectory-Level Implications**

572 Since autoregressive generation composes conditional distributions across steps, perturbations at
 573 early steps can lead to distinct trajectories. Let $\mathcal{R}_{k,K}(x)$ denote the set of trajectories containing
 574 at least one token outside the baseline top- K set. If RSP increases the accessibility of such tokens
 575 at early steps, then the probability of sampling trajectories in $\mathcal{R}_{k,K}(x)$ increases under repeated
 576 sampling. This links local ranking perturbations to global trajectory diversity.

E Attention Mass Analysis

This appendix formalizes the per-token attention mass reported in Figure 4 and the exclusion criteria applied to the reasoning span and the layer axis. For each of the 500 MATH-500 problems, a fresh RSP $\mathbf{H} \in \mathbb{R}^{L \times d}$ with $L = 10$ is sampled independently, so the reported heatmaps average over both problem variability and RSP-vector variability.

E.1 Per-Token Attention Mass

For each problem, we measure attention on the concatenated sequence [prompt; $\mathbf{H}$; generation], where the generation is the trajectory produced by the same model under matching RSP injection. Let $\alpha_{i,j}^{(\ell,h)}$ denote the post-softmax attention weight at layer ℓ , head h , from query position i to key position j , satisfying $\sum_j \alpha_{i,j}^{(\ell,h)} = 1$ under the causal mask. Let Q and R denote the index sets of question and RSP tokens with sizes $|Q|$ and $|R| = L$, and let H be the number of heads. The per-token attention mass that query i directs to each group at layer ℓ is

$$\text{PTM}_Q[\ell, i] = \frac{1}{|Q|H} \sum_{h=1}^H \sum_{j \in Q} \alpha_{i,j}^{(\ell,h)}, \quad \text{PTM}_R[\ell, i] = \frac{1}{|R|H} \sum_{h=1}^H \sum_{j \in R} \alpha_{i,j}^{(\ell,h)}. \quad (18)$$

Normalization by $|Q|$ and $|R|$ controls for group size, so both quantities express how much attention a single question (resp. RSP) token receives on average under the same softmax denominator. Question tokens thereby serve as a size-matched baseline that absorbs sequence-length effects common to both groups.

E.2 Heatmap Aggregation

We partition the layer indices and the reasoning position indices each into five equal-size bins $\{B_L^{(b)}\}_{b=1}^5$ and $\{B_P^{(b)}\}_{b=1}^5$. For sample s and target group $\bullet \in \{Q, R\}$, the value in cell (b_L, b_P) is

$$M_{\bullet}^{(s)}[b_L, b_P] = \frac{1}{|B_L^{(b_L)}| \cdot |B_P^{(b_P)}|} \sum_{\ell \in B_L^{(b_L)}} \sum_{i \in B_P^{(b_P)}} \text{PTM}_{\bullet}[\ell, i]. \quad (19)$$

Figure 4 reports the sample average of Eq. (19) over the 500 valid samples.

E.3 Excluding the `\boxed{...}` Span

The reasoning position range terminates strictly before the final `\boxed{...}` span. Chain-of-thought outputs decompose into a *text* (content) component and a *pattern* (template) component that contribute independently to downstream performance [Madaan and Yazdanbakhsh, 2022]. The `\boxed{...}` wrapper is a canonical pattern: a short, stereotyped span whose role is to complete the answer format rather than to extend reasoning. Including it would therefore inject a template-driven attention signature unrelated to reasoning dynamics.

E.4 Excluding the Final Two Layers

The layer axis further excludes the last two transformer layers. Two well-documented properties of late layers, none specific to our setup, motivate this choice.

Output-projection specialization. Late-layer hidden states lie in an approximately linear relationship with vocabulary logits and therefore function as a progressive linear readout rather than as representation-transforming blocks [Belrose et al., 2023]. Consistently, late-layer feed-forward modules have been characterized as output-distribution shaping mechanisms [Geva et al., 2021]. Attention in these layers therefore reflects output formation rather than reasoning computation.

Register-token hypothesis. Uninformative patch tokens in Vision Transformers have been shown to be repurposed as *registers*: storage slots that absorb attention in deeper layers [Darcet et al., 2024]. Because RSPs share a similar profile as continuous tokens without pretrained semantic content, a

615 similar repurposing in late transformer layers cannot be ruled out. Excluding the final two layers
 616 removes this potential register-induced confound from the reasoning-phase signal.

617 The two lines of evidence above identify known late-layer phenomena — linear readout and potential
 618 register repurposing — whose attention signatures are orthogonal to reasoning-phase dynamics.
 619 Removing the final two layers preempts their contamination of the reasoning-phase signal, yielding a
 620 principled, literature-grounded exclusion rather than a model-specific tuning decision.

621 E.5 Additional Suffix Heatmaps

622 Figure 7 reports the same per-token RSP attention analysis under suffix injection for Qwen2.5-Math-
 623 1.5B-Instruct and LLaMA-3.1-8B-Instruct. For the late-layer reasons discussed above, the pattern
 624 does not generalize uniformly across every cell; nevertheless, the overall trend is consistent across
 625 both models: question-token attention varies broadly across layers, whereas RSP attention shows the
 626 layer-wise and sequence-wise decay predicted by Eq. (4).




Figure 7: Per-token RSP attention mass under suffix injection for the remaining two models (500 MATH-500 samples each). Axes and preprocessing match Figure 4: x-axis is reasoning position binned into five quantiles, y-axis is layer depth binned into five quantiles (top = shallow), color encodes the 500-sample mean per-token attention assigned to RSP tokens, and tokens inside `\boxed{}` spans and the last two layers are excluded.

627 **F GRPO Formulation**

628 For a given prompt q , the behavior policy $\pi_{\theta_{\text{old}}}$ samples a group of G responses $\{o_i\}_{i=1}^G$ with
 629 corresponding rewards $\{R_i\}_{i=1}^G$. The advantage of each response is computed by normalizing the
 630 group-level rewards:

$$\hat{A}_{i,t} = \frac{R_i - \text{mean}(\{R_i\}_{i=1}^G)}{\text{std}(\{R_i\}_{i=1}^G)}. \quad (20)$$

631 The policy is updated via a clipped objective with a KL penalty:

$$\mathcal{J}_{\text{GRPO}}(\theta) = \mathbb{E} \left[\frac{1}{G} \sum_{i=1}^G \frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \left(\min \left(r_{i,t} \hat{A}_{i,t}, \text{clip}(r_{i,t}, 1-\varepsilon, 1+\varepsilon) \hat{A}_{i,t} \right) - \beta D_{\text{KL}}(\pi_{\theta} \parallel \pi_{\text{ref}}) \right) \right], \quad (21)$$

632 where $r_{i,t} = \pi_{\theta}(o_{i,t} \mid q, o_{i,<t}) / \pi_{\theta_{\text{old}}}(o_{i,t} \mid q, o_{i,<t})$ is the importance sampling ratio, ε is the clipping
 633 range, and β controls the strength of the KL penalty against the reference policy π_{ref} .

634 G Broader Impacts

635 This work provides empirical and theoretical analysis of how random embedding injection affects
636 the reasoning behavior of large language models. The contributions are primarily foundational,
637 with potential downstream implications for reinforcement learning-based post-training of reasoning
638 models such as GRPO.

639 **Positive impacts.** RSP-induced trajectory diversity could improve the sample efficiency of RL-
640 based LLM post-training by providing richer learning signals within group-relative advantage estima-
641 tion. This has the potential to reduce compute costs associated with alignment and reasoning-oriented
642 fine-tuning, making such methods more accessible to research groups with limited resources.

643 **Potential concerns.** Methods that expand the effective output support of LLMs may increase the
644 reachability of low-probability tokens, which includes both desirable alternative reasoning paths
645 and potentially unsafe or off-distribution outputs. Practitioners applying RSP in production systems
646 should combine it with existing safety filtering and alignment mechanisms rather than treating it as a
647 standalone diversity source.

648 **Limitations on impact scope.** Since this work focuses on analysis rather than deploying a new
649 model or dataset, the direct societal impact is limited. The broader implications described above are
650 contingent on future work integrating RSP into applied training pipelines.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The authors should answer **[Yes]** if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- 855 • The paper should disclose whether the full research project required more compute
856 than the experiments reported in the paper (e.g., preliminary or failed experiments that
857 didn't make it into the paper).

858 **9. Code of ethics**

859 Question: Does the research conducted in the paper conform, in every respect, with the
860 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

861 Answer: **[TODO]**

862 Justification: **[TODO]**

863 Guidelines:

- 864 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of
865 Ethics.
866 • If the authors answer [No], they should explain the special circumstances that require a
867 deviation from the Code of Ethics.
868 • The authors should make sure to preserve anonymity (e.g., if there is a special consid-
869 eration due to laws or regulations in their jurisdiction).

870 **10. Broader impacts**

871 Question: Does the paper discuss both potential positive societal impacts and negative
872 societal impacts of the work performed?

873 Answer: **[TODO]**

874 Justification: **[TODO]**

875 Guidelines:

- 876 • The answer [N/A] means that there is no societal impact of the work performed.
877 • If the authors answer [N/A] or [No], they should explain why their work has no societal
878 impact or why the paper does not address societal impact.
879 • Examples of negative societal impacts include potential malicious or unintended uses
880 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations
881 (e.g., deployment of technologies that could make decisions that unfairly impact specific
882 groups), privacy considerations, and security considerations.
883 • The conference expects that many papers will be foundational research and not tied
884 to particular applications, let alone deployments. However, if there is a direct path to
885 any negative applications, the authors should point it out. For example, it is legitimate
886 to point out that an improvement in the quality of generative models could be used to
887 generate Deepfakes for disinformation. On the other hand, it is not needed to point out
888 that a generic algorithm for optimizing neural networks could enable people to train
889 models that generate Deepfakes faster.
890 • The authors should consider possible harms that could arise when the technology is
891 being used as intended and functioning correctly, harms that could arise when the
892 technology is being used as intended but gives incorrect results, and harms following
893 from (intentional or unintentional) misuse of the technology.
894 • If there are negative societal impacts, the authors could also discuss possible mitigation
895 strategies (e.g., gated release of models, providing defenses in addition to attacks,
896 mechanisms for monitoring misuse, mechanisms to monitor how a system learns from
897 feedback over time, improving the efficiency and accessibility of ML).

898 **11. Safeguards**

899 Question: Does the paper describe safeguards that have been put in place for responsible
900 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
901 image generators, or scraped datasets)?

902 Answer: **[TODO]**

903 Justification: **[TODO]**

904 Guidelines:

- 905 • The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.