
From Noise to Diversity: Random Embedding Injection in LLM Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

Recent *soft prompt* methods, which prepend or insert trainable continuous embeddings into LLM inputs as virtual tokens, optimize these embeddings to enhance mathematical reasoning. However, the mechanism by which embedding injection itself affects model behavior, independent of optimization, remains unexplored. We investigate Random Soft Prompts (RSPs), continuous vectors sampled from an isotropic Gaussian fitted to the empirical mean and variance of the model’s token embedding table and injected without any training. Empirically, RSPs increase early-stage token entropy while generation stabilizes thereafter, and combining RSPs with temperature sampling yields higher trajectory diversity as measured by Pass@ n . Theoretically, we provide two complementary analyses: an energy-based framework that interprets RSP as an idealized stochastic perturbation of a Hopfield-style energy landscape with attention-mediated decay across layers and generation steps, and a distributional analysis showing that RSP induces stochastic token rankings and expands the output support beyond temperature sampling. We further discuss the implications for GRPO training, where trajectory diversity is critical for learning signal quality.

1 Introduction

As large language models (LLMs) scale to billions of parameters, full fine-tuning becomes increasingly expensive. This has motivated parameter-efficient methods that adapt frozen models by introducing a small number of trainable parameters [Hu et al., 2022, Houlsby et al., 2019]. Among these, *soft prompts* prepend or insert learnable continuous embeddings into the input sequence, steering model behavior through the attention mechanism [Li and Liang, 2021, Lester et al., 2021]. This paradigm has been actively adopted for enhancing mathematical reasoning in LLMs [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Hao et al., 2025, Zhang et al., 2025]. Despite their diverse designs, these methods share a common structure: they inject continuous embeddings that lie outside the pretrained vocabulary and optimize them to improve downstream performance.

However, we find that optimization is not a necessary condition for these effects. We observe that injecting randomly initialized, untrained continuous embeddings can also shift the output distribution. For instance, applying a chat template to a base model not fine-tuned with such format causes a prompt-format mismatch that degrades reasoning capability; appending random embeddings to such mismatched inputs mitigates this degradation. This suggests that injecting out-of-vocabulary continuous embeddings itself, independent of any learned content, plays a non-trivial role in shaping model behavior. Existing studies primarily evaluate *soft prompt* methods through end-task accuracy, and the mechanisms by which embedding injection influences the output distribution remain insufficiently explored.

In this work, we investigate the role of Random Soft Prompts (RSPs), continuous vectors sampled from an isotropic Gaussian fitted to the empirical mean and variance of the model’s token embedding table and injected without any training. Empirically, RSPs increase early-stage token entropy while the output stabilizes in later generation stages, and combining RSPs with temperature sampling yields more diverse reasoning trajectories as measured by Pass@ n scaling. Theoretically, we analyze RSPs from two perspectives: an energy-based framework showing that RSP introduces controlled stochastic perturbations with natural decay through the attention mechanism, and a distributional analysis showing that RSP induces stochastic token rankings and expands the output support beyond what temperature sampling can achieve. We further discuss the implications of these findings for Group Relative Policy Optimization (GRPO) [Shao et al., 2024], where trajectory diversity is essential for effective learning signals.

Our contributions are as follows:

- We empirically analyze how RSPs affect LLM generation, showing that RSPs increase early token entropy and produce more diverse reasoning trajectories when combined with temperature sampling.
- We provide theoretical analyses from two perspectives: an energy-based framework explaining why RSP can perturb the output distribution without degrading performance, and a distributional analysis showing that RSP induces stochastic token rankings and expands the output support beyond the reach of temperature sampling.
- We discuss the implications of RSP-induced diversity for GRPO training and argue that RSP provides exploration gains complementary to existing temperature-based approaches.

2 Background

2.1 Soft Prompt Methods for LLM Reasoning

Soft prompts emerged as a parameter-efficient alternative to full fine-tuning. Prefix-Tuning [Li and Liang, 2021] introduced the paradigm by prepending trainable continuous vectors to every transformer layer, Prompt Tuning [Lester et al., 2021] simplified this to input-layer-only tuning and showed that it matches full fine-tuning at sufficient scale, and P-tuning [Liu et al., 2023] extended the approach with an LSTM-based prompt encoder to model inter-position dependencies.

More recently, *soft prompts* have been actively applied to LLM reasoning. TTSV [Kang et al., 2026] optimizes prefix embeddings at test time via entropy minimization, steering the model toward high-certainty states to activate latent reasoning capabilities. SoftCoT [Xu et al., 2025] replaces explicit chain-of-thought tokens with soft tokens generated by an assistant model, leveraging the stronger representational capacity of continuous representations over discrete language. LTPO [Ye et al., 2026] optimizes latent thought embeddings per test instance through policy gradient with a confidence-based intrinsic reward. COCONUT [Hao et al., 2025] feeds the model’s own hidden states back as input embeddings, enabling reasoning in a continuous latent space. MemGen [Zhang et al., 2025] demonstrates that embedding vectors can serve as experience memory by constructing latent token sequences that enrich the model’s reasoning. Pause tokens [Goyal et al., 2024] insert learnable tokens into the input to provide additional computation steps.

These approaches all inject out-of-vocabulary continuous embeddings optimized for task-specific objectives. Because evaluation centers on end-task accuracy, the contribution of the injection itself and that of the optimization are not separated, and the mechanisms by which such out-of-vocabulary continuous embeddings shape model behavior remain a separate and underexplored question.

2.2 Noise Injection in Neural Networks

Neural networks have incorporated noise at various stages. During training, dropout [Srivastava et al., 2014] randomly deactivates hidden units to prevent overfitting, and NEFTune [Jain et al., 2024] adds noise to token embeddings during instruction fine-tuning, reporting improvements in instruction following. These methods use noise as a regularizer and do not apply perturbations at inference time.

At inference time, perturbation-based approaches primarily target robustness and safety. Randomized smoothing [Cohen et al., 2019] injects Gaussian noise into inputs to obtain certified robustness

Table 1: Accuracy (%) across model configurations and benchmarks. RSP reports the best result across injection positions (Prefix, Infix, Suffix); full per-position breakdown is in Appendix A. Values in parentheses indicate the difference from the baseline.

Model	MATH-500		GSM8K		AIME24	
	Baseline	RSP	Baseline	RSP	Baseline	RSP
<i>Instruct Models</i>						
LLaMA-3.1-8B-Inst	52.20	49.40 (−2.8)	85.75	86.73 (+1.0)	6.67	10.00 (+3.3)
Qwen2.5-Math-7B-Inst	83.20	84.20 (+1.0)	95.45	95.68 (+0.2)	13.33	16.67 (+3.3)
Qwen2.5-Math-1.5B-Inst	73.00	75.20 (+2.2)	85.29	85.75 (+0.5)	10.00	20.00 (+10.0)
<i>Base Models (Raw Text)</i>						
Qwen2.5-Math-7B	72.20	71.60 (−0.6)	84.15	84.31 (+0.2)	6.67	16.67 (+10.0)
Qwen2.5-Math-1.5B	61.40	64.40 (+3.0)	77.86	74.30 (−3.6)	16.67	10.00 (−6.7)
<i>Base Models + ChatML (Format Mismatch)</i>						
Qwen2.5-Math-7B	52.20	70.40 (+18.2)	58.30	75.97 (+17.7)	23.33	23.33 (+0.0)
Qwen2.5-Math-1.5B	34.40	63.40 (+29.0)	37.45	67.02 (+29.6)	13.33	20.00 (+6.7)

guarantees. In the LLM setting, SmoothLLM [Robey et al., 2023] perturbs input prompts at the character level to defend against adversarial jailbreaking. These methods require multiple noisy forward passes with output aggregation, and target prediction stability rather than output diversity.

In reinforcement learning, NoisyNet [Fortunato et al., 2018] injects learnable noise into network weights to improve exploration, demonstrating that the location and structure of noise injection affect exploration quality. This motivation is closest to RSPs, which similarly target diversity, but differ in injection target (input embeddings vs. network weights) and require no learned noise parameters.

RSPs differ from these approaches in the following respects. **(1)** RSPs operate solely at inference without any training, whereas dropout and NEFTune inject noise during training, modifying the resulting model parameters accordingly. **(2)** RSPs preserve the original input and append fresh embedding vectors in a single forward pass, whereas robustness methods corrupt the input and aggregate over multiple noisy passes. **(3)** RSPs use purely random vectors sampled from the pretrained embedding statistics, whereas NoisyNet learns noise parameters through optimization.

3 Experimental Analysis

3.1 Random Soft Prompt: Definition and Setup

Let $\mathbf{W}_E \in \mathbb{R}^{V \times d}$ denote the token embedding matrix of a pretrained LLM, where V is the vocabulary size and d is the hidden dimension. We define an RSP $\mathbf{H} = [h_1, \dots, h_L] \in \mathbb{R}^{L \times d}$ as a sequence of L continuous vectors, where each vector is independently sampled from:

$$h_j \sim \mathcal{N}(\mu_E, \sigma_E^2 \mathbf{I}), \quad j = 1, \dots, L, \quad (1)$$

with $\mu_E = \frac{1}{Vd} \sum_{i,k} [\mathbf{W}_E]_{i,k}$ and $\sigma_E = \text{std}(\mathbf{W}_E)$. Let $\mathbf{X} = [x_1, \dots, x_T] \in \mathbb{R}^{T \times d}$ denote the embedding sequence of the original input. The RSP-augmented input $\tilde{\mathbf{X}}$ is constructed by concatenating $\mathbf{H}$ with $\mathbf{X}$ at one of three positions without modifying the original prompt text: *prefix* ($[\mathbf{H}; \mathbf{X}]$), *infix* ($\mathbf{H}$ inserted within $\mathbf{X}$), or *suffix* ($[\mathbf{X}; \mathbf{H}]$), following the injection strategies adopted in prior work [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Zhang et al., 2025]. Sampling a new $\mathbf{H}$ for each batch ensures that the results are not dependent on any realization of the random embeddings.

We evaluate on three mathematical reasoning benchmarks: MATH-500 [Lightman et al., 2024], GSM8K [Cobbe et al., 2021], and AIME24, using LLaMA-3.1-8B-Instruct [Grattafiori et al., 2024] and the Qwen2.5-Math model family [Yang et al., 2024] across seven configurations including instruct models, base models, and base models with mismatched chat templates. All experiments use greedy decoding to isolate the effect of RSPs from sampling stochasticity. Only the final answer is evaluated through a fallback-based extraction pipeline. Full experimental details are provided in Appendix A.

Table 2: Comparison with existing soft prompt methods (TTSV, SoftCoT, LTPO) under unified evaluation. Baseline and RSP values are from Table 1. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Benchmark	Baseline	RSP	TTSV	SoftCoT	LTPO
Qwen2.5-Math-7B-Instruct	MATH-500	83.20	84.20	<u>83.80</u>	82.80	83.00
	GSM8K	<u>95.45</u>	95.68	95.30	95.00	94.84
	AIME24	13.33	<u>16.67</u>	23.30	<u>16.67</u>	13.33
Qwen2.5-Math-1.5B-Instruct	MATH-500	73.00	<u>75.20</u>	<u>75.20</u>	76.00	72.40
	GSM8K	85.29	85.75	<u>85.70</u>	84.46	84.53
	AIME24	<u>10.00</u>	20.00	6.70	6.67	<u>10.00</u>

3.2 Does RSP Degrade Reasoning Performance?

We first examine how injecting random embeddings affects the model’s reasoning performance. Table 1 reports the accuracy for each model configuration, using the best RSP result across injection positions (Prefix, Infix, Suffix) and $L \in \{10, 15, 20\}$ for each benchmark.

For instruct models and base models with their native prompt formats, RSP injection maintains or slightly shifts baseline performance within a few percentage points. Most strikingly, under prompt-format mismatch — where base models receive ChatML templates they were not trained on — RSP yields pronounced improvements of up to +29%p. These gains are difficult to attribute to simple perturbation effects, as noise that merely raises model uncertainty would further degrade performance.

We compare RSPs with prior soft prompt methods (TTSV, SoftCoT, LTPO) that optimize their embeddings for distinct objectives, under a unified evaluation protocol reported in Table 2. All methods are re-evaluated with a shared answer extraction pipeline that applies fallback strategies, enabling a format-agnostic comparison despite each method’s distinct prompt format and grading scheme. The results place RSPs on par with the optimized methods without requiring any training.

3.3 How Does RSP Affect the Output Distribution?

We analyze how RSP injection changes the model’s output distribution during generation on Qwen2.5-Math-1.5B-Instruct with MATH-500, measuring per-token entropy, top-1 probability, and varentropy across the generation process. Figure 1 and Table 3 compare these metrics for the first 5% of generation steps across RSP variants and prior soft prompt methods.




Figure 1: Mean entropy, top-1 probability, and varentropy during the first 5% of generation steps (Qwen2.5-Math-1.5B-Instruct, MATH-500). Shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: existing soft prompt methods (LTPO, SoftCoT, TTSV).

Table 3: Mean entropy, top-1 probability, and varentropy for the first 5% of generation steps (Qwen2.5-Math-1.5B-Instruct, MATH-500).

Metric	Baseline	Prefix	Infix	Suffix	LTPO	SoftCoT	TTSV
Entropy	0.1332	0.1366	0.1402	0.1941	0.1662	0.4218	0.1359
Top-1 Prob	0.9535	0.9523	0.9525	0.9373	0.9437	0.9156	0.9534
Varentropy	0.2181	0.2207	0.2463	0.4010	0.2953	2.2234	0.2174

These results show that RSP modifies the early-stage output distribution across all injection positions: entropy increases (the token distribution flattens), top-1 probability decreases (commitment to a dominant token weakens), and varentropy increases (indicating multimodal distributions [Ahmed et al., 2026]). High-entropy positions are known to act as critical forking points that determine reasoning trajectories [Wang et al., 2025], so this shift moves the model toward a state where multiple reasoning paths can coexist. The change is localized to the initial reasoning stage; in the middle and final portions of generation, all three metrics converge to baseline levels.

Although existing soft prompt methods are optimized for different objectives, they share a common pattern of increasing early-stage entropy. Notably, TTSV, which optimizes soft prompts using a reward that reduces the mean entropy of reasoning trajectories, still exhibits early entropy comparable to the baseline. This suggests that elevating early-stage entropy is an inherent effect of injecting non-pretrained continuous embeddings, independent of the optimization objective.

3.4 Does RSP Induce Trajectory Diversity?

Section 3.3 empirically established that RSP reshapes the early-stage output distribution. We now examine whether this shift translates into reasoning diversity at three complementary levels of analysis: token rankings, semantic content of trajectories, and task-level outcomes. All three analyses share the setting: Qwen2.5-Math-1.5B-Instruct, MATH-500, suffix injection, and $L = 10$.

Token-level: distribution beyond temperature. Under autoregressive decoding, any trajectory-level divergence between RSP and the baseline must originate from the first generation step, a critical forking point for reasoning trajectories [Wang et al., 2025]. Using the baseline at $\tau=1.0$ as reference, we compare RSP at $\tau=1.0$ against the baseline at $\tau=2.0$ (the high-temperature ceiling of what flattening alone can achieve) on three complementary metrics: Spearman rank correlation ρ to detect rank reorganization, the probability mass placed outside the baseline’s top-10 to measure support expansion, and Jensen–Shannon divergence to quantify the overall magnitude of distributional change (formal definitions in Appendix C).

Table 4: First-token distribution metrics measured against the baseline at $\tau=1.0$, comparing the baseline at $\tau=2.0$ with RSP at $\tau=1.0$.

Metric	$\tau=2.0$	RSP
Spearman ρ	1.0	0.651
Mass outside top-10	5.15%	27.64%
JS divergence	0.086	0.231

Table 4 reports the contrast. Even at $\tau=2.0$, only 5.15% of probability mass falls outside the baseline’s top-10; RSP at $\tau=1.0$ reaches 27.64%. The Spearman rank correlation tells the story more starkly: because temperature scaling $p_i^{(\tau)} \propto p_i^{1/\tau}$ is a monotone transform that preserves token ranking ($p_i > p_j \iff p_i^{(\tau)} > p_j^{(\tau)}$) for any $\tau > 0$, ρ relative to the baseline is exactly 1 as a mathematical identity. RSP, by contrast, reranks the token distribution itself, reducing ρ to 0.651—an effect temperature cannot produce at any τ . The Jensen–Shannon divergence (0.231) quantifies the overall magnitude of this distributional change. Since $\rho < 1$ rules out rank-preserving flattening as the sole cause, the high divergence reflects a structural reshape of the distribution rather than a uniform flattening of the same ranking.

Trajectory-level: semantic diversity. Token-level reranking raises a second question: do first-token perturbations produce semantically distinct reasoning paths, or do independently perturbed trajectories converge to the same solution strategy? We randomly sample 64 problems from MATH-500 and generate 300 independent trajectories per problem under two conditions, Baseline and RSP, with temperature $\tau = 1.0$. The trajectories are embedded via BGE-M3 [Chen et al., 2024] and compared on three metrics: mean pairwise cosine distance within each problem, inter-cluster distance between DBSCAN [Ester et al., 1996] centroids, and intra-cluster distance. Table 5 shows that RSP substantially increases the inter-cluster distance while leaving the intra-cluster distance unchanged: clusters become more separated while reasoning within each cluster remains consistent. RSP also wins on pairwise distance in 56 of 64 problems (87.5%).

Table 5: Semantic diversity metrics averaged across the 64 problems.

Metric	Baseline	RSP
Pairwise cos. dist.	0.0612	0.0879
Inter-cluster dist.	0.0183	0.0982
Intra-cluster dist.	0.0526	0.0528

Figure 2(a) illustrates the qualitative difference for the representative problem: the RSP samples (blue) cover the region occupied by baseline samples (red) while extending into clusters that the




Figure 2: Semantic diversity analysis on 64 randomly sampled MATH-500 problems with 300 trajectories per condition (Qwen2.5-Math-1.5B-Instruct, $\tau = 1.0$, BGE-M3 embeddings). (a) t-SNE projection for one representative problem (test/prealgebra/951): baseline (red) concentrates around one region while RSP suffix (blue) spreads across distinct clusters; filled markers denote correct answers and open markers denote incorrect. (b) Per-problem mean pairwise cosine distance; points above the diagonal indicate RSP produces more diverse trajectories than baseline. (c) Per-problem DBSCAN cluster count; RSP yields more distinct clusters than baseline on most problems.

baseline does not reach; the expanded regions also contain correct solutions (filled markers), showing that RSP’s exploration is not limited to incorrect alternatives but reaches new reasoning paths that still yield the correct answer. Figures 2(b) and 2(c) show that this pattern holds across the 64 problems: on most problems RSP produces more clusters and larger pairwise cosine distance than baseline. The combination of large inter-cluster increase and unchanged intra-cluster distance mirrors the token-level finding: RSP alters which cluster a trajectory enters while preserving the local continuation within each cluster. Additional per-problem t-SNE plots are provided in Figure 6 (Appendix B).

Outcome-level: Pass@ n scaling. Finally, we evaluate whether the token- and trajectory-level diversity translates into task-level outcome diversity through Pass@ n , the probability that at least one of n sampled solutions is correct. We compare four conditions with $n = 16$ samples per problem on MATH-500 and AIME24: (i) *Baseline*, $\tau = 1.0$: temperature sampling only; (ii) *RSP (fixed seed)*, $\tau = 1.0$: a single RSP shared across all samples; (iii) *RSP (indep. seed), greedy*: a different RSP per sample without temperature; and (iv) *RSP (indep. seed)*, $\tau = 1.0$: a different RSP per sample combined with temperature sampling.




Figure 3: Pass@ n scaling on MATH-500 (left) and AIME24 (right) with Qwen2.5-Math-1.5B-Instruct, $n = 16$ samples per problem. Baseline: temperature sampling only. Fixed seed: single RSP shared across samples combined with temperature. Indep. seed: a different RSP per sample, with or without temperature.

Condition (iv) achieves the highest Pass@ n on both benchmarks, reaching 92.80% on MATH-500 and 36.67% on AIME24 at $n = 16$. Condition (ii) shows no improvement over the baseline, confirming that the diversity gain originates from varying the random embeddings across samples, not from a single fixed perturbation. Pass@1 remains comparable to the baseline on MATH-500, and on the more challenging AIME24, (iv) achieves higher Pass@1 than the baseline. Independent RSPs combined with temperature sampling produce the most diverse reasoning trajectories while maintaining or improving single-sample accuracy.

209 **Synthesis across the three levels.** The three analyses converge on a consistent picture. At the token
 210 level, RSP induces support expansion concentrated at the first token and decaying monotonically
 211 across steps. At the trajectory level, this localized perturbation translates into distinct reasoning
 212 strategies (large inter-cluster separation) while preserving local coherence within each strategy
 213 (unchanged intra-cluster distance). At the outcome level, the resulting trajectories cover more correct
 214 solutions without sacrificing single-sample accuracy.

215 4 Theoretical Analysis: Why RSP Works

216 Our theoretical goal is narrower than a theorem about the full nonlinear transformer. We want
 217 a mechanism that explains the three empirical regularities from Section 3: random continuous
 218 prompts can alter the model’s earliest reasoning decisions even though they are not optimized for
 219 the task; their influence is strongest at the beginning of decoding and weakens later; and this early
 220 perturbation can increase trajectory diversity without permanently destabilizing generation. The
 221 common control variable throughout the section is $w_{r,i}^{(l)}$, the attention mass assigned to the random
 222 tokens. The argument therefore has three parts. First, we characterize what kind of perturbation can
 223 open reasoning branches without inserting a hand-crafted prior, and explain why centered isotropic
 224 randomness is the relevant choice. Second, we show that attention injects the perturbation through
 225 $w_{r,i}^{(l)}$ and simultaneously places an annealing pressure on it as real-token evidence grows. Third, we
 226 show why this pre-softmax perturbation can change token rankings in a way that temperature scaling
 227 cannot.

228 4.1 Why Random Rather Than a Fixed Token?

229 A generic exploration prompt for reasoning should satisfy three structural requirements. It should
 230 not systematically pre-commit the model to one semantic direction before the input-dependent
 231 computation begins. It should vary across trials, otherwise repeated decoding would trace the same
 232 perturbation every time. And because the branch-sensitive direction is unknown in advance and
 233 depends on the prompt, the decoding state, and the token comparison, the perturbation should not
 234 privilege some directions while leaving others almost untouched. Optimized soft prompts deliberately
 235 encode a task prior [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Zhang et al., 2025]; our
 236 question is the complementary one: what perturbation should one use when no such prior is assumed?

237 A fixed token or a fixed continuous prompt fails this requirement set immediately. It may steer one
 238 trajectory, but if the same vector is reused across runs then there is no across-trial exploration at all.
 239 Resampling discrete vocabulary tokens restores variability, yet each draw still comes with its own
 240 lexical direction and there is no reason for the resulting perturbations to be centered or isotropic. We
 241 formalize the choice we make instead.

242 **Definition 1** (Random Soft Prompt). *Let $\mathbf{E} \in \mathbb{R}^{|\mathcal{V}| \times d}$ denote the pretrained input-embedding table,*
 243 *and define the empirical entrywise mean and variance*

$$\bar{\mu}_E := \frac{1}{|\mathcal{V}|d} \sum_{v \in \mathcal{V}} \sum_{m=1}^d E_{v,m}, \quad \sigma_E^2 := \frac{1}{|\mathcal{V}|d} \sum_{v \in \mathcal{V}} \sum_{m=1}^d (E_{v,m} - \bar{\mu}_E)^2, \quad \mu_E := \bar{\mu}_E \mathbf{1}_d.$$

244 *In the implementation studied here, an RSP of length L is a sequence $\mathbf{H} = (h_1, \dots, h_L)$ with*

$$h_l = \mu_E + \sigma_E \xi_l, \quad \xi_l \stackrel{iid}{\sim} \mathcal{N}(0, \mathbf{I}_d).$$

245 *Equivalently, the centered prompt $\bar{h}_l := h_l - \mu_E$ satisfies $\bar{h}_l \sim \mathcal{N}(0, \sigma_E^2 \mathbf{I}_d)$. Whenever we later*
 246 *analyze repeated-sampling protocols such as Pass@N, we additionally assume the draws used across*
 247 *those compared trials are independent.*

248 The proofs below use only the centered law of $\bar{h}_l$, not the particular form of the centering vector μ_E .
 249 The scalar fit $\mu_E = \bar{\mu}_E \mathbf{1}_d$ therefore matches the implementation without changing the subsequent
 250 arguments. The structural ingredients are centeredness, isotropic covariance, and fresh resampling
 251 whenever multiple RSP trials are compared. The two embedding-derived statistics in Definition 1 play
 252 distinct roles. The scalar σ_E^2 borrows the *scale* of the pretrained embedding distribution so that the
 253 perturbation operates in the magnitude regime the model expects. The covariance $\sigma_E^2 \mathbf{I}_d$, in contrast,
 254 deliberately discards the directional structure of the embedding distribution; Theorem 1 below shows

that this is the unique covariance that maximizes the worst-case directional variance under a fixed trace budget. The downstream attention mechanism then maps this prompt-space distribution to value space via the value projection $\mathbf{W}_V^{(l)}$; the corresponding variance facts are stated explicitly below.

Lemma 1 (Semantic neutrality). *Let $\bar{h} \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_d)$ and $\hat{h} := \bar{h}/\|\bar{h}\|$. For any fixed unit vectors $u_1, \dots, u_M \in \mathbb{S}^{d-1}$ and any $\delta \in (0, 1)$,*

$$\Pr \left[\max_{r \in [M]} |\hat{h}^\top u_r| \leq \sqrt{\frac{2 \log(2M/\delta)}{d-1}} \right] \geq 1 - \delta.$$

In contrast, for any nonzero token embedding e_s , choosing the anchor direction $u = \hat{e}_s := e_s/\|e_s\|$ yields $|\hat{e}_s^\top u| = 1$.

Lemma 1 gives the precise sense in which the perturbation is neutral. The claim is not that a random vector is literally free of every incidental semantic correlation. The relevant and weaker statement is that it does not systematically align with any pre-selected finite family of semantic alternatives. Centering then removes deterministic first-order drift: under the local surrogate below, the expected perturbation of any logit gap is zero when $\mathbb{E}[\bar{h}] = 0$.

To formalize the exploration part, let $\bar{h}$ be the centered prompt vector and consider the local first-order surrogate

$$z_i(\bar{h}) \approx z_i + a_i^\top \bar{h}, \quad a_i := \nabla_{\bar{h}} z_i(0),$$

valid for sufficiently small $\|\bar{h}\|$ (Appendix E); we use this surrogate throughout the section to identify the direction of perturbation rather than to predict its exact magnitude. For a token pair (i, j) , the perturbed logit gap becomes

$$\Delta_{ij}(\bar{h}) \approx \Delta_{ij} + b_{ij}^\top \bar{h}, \quad b_{ij} := a_i - a_j.$$

The relevant question is therefore not “does the perturbation have variance?” but “how much variance does it retain along the unknown direction b_{ij} that matters for this branching decision?” The theorem below answers that question at the level of second moments.

Theorem 1 (Direction-agnostic coverage). *Let D be any zero-mean perturbation distribution on $\mathbb{R}^d$ with covariance Σ_D and variance budget $\text{tr}(\Sigma_D) \leq \rho^2$. Define the worst directional variance*

$$\mathcal{E}(D) := \min_{\|b\|=1} \text{Var}_{h \sim D}(b^\top h).$$

Then

$$\mathcal{E}(D) = \lambda_{\min}(\Sigma_D) \leq \frac{\rho^2}{d},$$

with equality if and only if $\Sigma_D = (\rho^2/d)\mathbf{I}_d$. Hence isotropic covariance is the unique way to equalize first-order perturbation variance across all directions while maximizing the least-covered one.

Theorem 1 is not meant to cast prompting as an adversarial game; its role is simpler. When the branch-sensitive direction is unknown at sample time, isotropy is the unique second-order description that neither favors nor neglects any direction. A fixed prompt has $\Sigma_D = 0$ and therefore creates no across-trial exploration at all. An anisotropic perturbation leaves at least one direction less explored than the isotropic optimum. The Gaussian choice in RSP is the simplest centered, rotation-symmetric family that realizes this covariance while matching the scale of pretrained embeddings. This is the sense in which RSP is not just “some noise”: it is a resampled, centered, direction-agnostic perturbation.

A natural alternative one might consider is to draw the perturbation from a distribution that matches the empirical geometry of the pretrained embedding table $\mathbf{W}_E$, rather than discarding it. Theorem 1 dismisses this alternative as soon as the embedding geometry is known to be anisotropic.

Corollary 1 (Embedding-matched noise is suboptimal in worst-direction coverage). *Let Σ_E be the empirical covariance of the rows of $\mathbf{W}_E$, and let $\rho^2 := \text{tr}(\Sigma_E)$. Consider two zero-mean perturbations on $\mathbb{R}^d$ with the same trace budget ρ^2 : the isotropic RSP law D_{iso} with covariance $(\rho^2/d)\mathbf{I}_d$, and the embedding-matched law D_{emb} with covariance Σ_E . Then*

$$\mathcal{E}(D_{\text{iso}}) = \frac{\rho^2}{d} \geq \lambda_{\min}(\Sigma_E) = \mathcal{E}(D_{\text{emb}}),$$

with strict inequality whenever Σ_E is non-isotropic.

The corollary is one application of Theorem 1 and requires no separate proof. Its empirical bite comes from the fact that pretrained LLM token-embedding matrices are documented to be highly anisotropic, concentrating in a low-dimensional “cone” inside $\mathbb{R}^d$ [Mu and Viswanath, 2018, Ethayarajh, 2019]. In that regime $\lambda_{\min}(\Sigma_E)$ is orders of magnitude smaller than ρ^2/d , so the worst-covered direction under embedding-matched noise is nearly unexplored—precisely the directions a generic exploration prompt should not abandon. Discarding the directional geometry of $\mathbf{W}_E$ while keeping its scale is therefore not a simplification of an embedding-matched alternative but a quantitatively better choice for direction-agnostic exploration.

Lemma 1 and Theorem 1 therefore play complementary roles. The lemma rules out semantic pre-commitment to any fixed anchor set, while the theorem rules out directional under-coverage in the local branch-sensitive geometry. A perturbation that satisfies only the first condition would be neutral but too weak in some directions; a perturbation that satisfies only the second would still inject a hand-chosen semantic prior. RSP is useful because it satisfies both at once.

4.2 How Attention Turns a Neutral Perturbation into Controlled Exploration

The previous subsection explains why centered isotropic randomness is the right kind of perturbation. The next question is temporal: why is the perturbation strong enough to matter at the beginning of decoding but weak enough not to derail the rest of the reasoning chain? The formal answer uses three ingredients. First, in an idealized Hopfield-style view of attention, each transformer block acts like one step of an iterative attractor process; in that picture the unperturbed path is strongly constrained by its initialization, which identifies a structural source of exploration deficit. Second, the attention output admits an exact decomposition into a real-token update and a random-token contribution. Third, the same attention mass that injects the random contribution is structurally forced downward as real-token evidence accumulates, and the perturbation propagates through the feed-forward sublayer with controlled amplification.

Iterative attractor view. Following the Modern Hopfield correspondence [Ramsauer et al., 2021], one attention update with key matrix $\mathbf{K}$ and value matrix $\mathbf{V}$ is equivalent to a gradient step on the energy

$$E_{\mathbf{K}}(\sigma) = -\beta^{-1} \log \sum_j \exp(\beta \mathbf{k}_j^\top \sigma) + \frac{1}{2} \|\sigma\|^2$$

when $\mathbf{K} = \mathbf{V}$. Under three idealized assumptions (A1) $\mathbf{K} = \mathbf{V}$, (A2) *weight-tied attention*, (A3) L_E -smoothness of $E_{\mathbf{K}}$ with $L_E < 2$, Appendix E.1 shows that the unperturbed iterates produce a strictly monotone descent in $E_{\mathbf{K}}$ and the sublevel set $\{\sigma : E_{\mathbf{K}}(\sigma) \leq E_{\mathbf{K}}(\sigma^{(0)})\}$ is forward-invariant; under the additional separation condition $\text{sep}(C(\sigma^{(0)})) > 0$ together with the per-step gradient bound of Corollary 3, the trajectory also remains in the connected component containing $\sigma^{(0)}$.

(A1) and (A2) are not literally satisfied by real transformers, so the Hopfield view here is an idealized structural picture rather than a literal statement about every transformer block. Its role is to identify why a deterministic forward pass can be exploration-poor: under A1–A3 it stays inside a fixed sublevel set, and under the extra separation/small-step condition it also remains in the connected component selected by initialization. We use this picture to motivate why injecting an isotropic stochastic perturbation at every block is the relevant intervention, not to claim formal escape rates for real networks.

Throughout the rest of the section we work with the actual transformer attention. The random values are obtained by applying the value projection to the centered prompt vectors, so

$$\mathbf{v}_{r_j}^{(l)} = \mathbf{W}_V^{(l)} \bar{\mathbf{h}}_j, \quad \text{Cov}(\mathbf{v}_{r_j}^{(l)}) = \sigma_E^2 \mathbf{W}_V^{(l)} \mathbf{W}_V^{(l)\top}.$$

Proposition 5 in Appendix D states the resulting directional-variance bounds explicitly. In particular, if $\mathbf{W}_V^{(l)}$ is full rank then the projected perturbation remains non-degenerate in every value-space direction, and if $\mathbf{W}_V^{(l)}$ is well-conditioned then the prompt-space direction-agnostic coverage carries to value space up to layer-dependent constants.

Exact decomposition. When k random vectors are inserted as additional keys and values, the attention update can be written as

$$\mathbf{x}_i^{(l+1)} = (1 - w_{r,i}^{(l)}) \tilde{\mathbf{x}}_i^{(l+1)} + \boldsymbol{\eta}_i^{(l)}, \quad (2)$$

where $\tilde{\mathbf{x}}_i^{(l+1)}$ is the update produced by the real tokens alone, $\boldsymbol{\eta}_i^{(l)} := \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)}$ is the contribution of the random values, and

$$w_{r,i}^{(l)} := \sum_{j=1}^k \alpha_{i,n+j}^{(l)}$$

is the total attention mass assigned to the random tokens. Appendix E.2 proves the decomposition and the deterministic bound

$$\|\boldsymbol{\eta}_i^{(l)}\| \leq w_{r,i}^{(l)} \max_{j \in [k]} \|\mathbf{v}_{r_j}^{(l)}\|.$$

This is the structural core of the mechanism: the same scalar $w_{r,i}^{(l)}$ controls how much of the original update is retained and how large the random contribution can be. If $w_{r,i}^{(l)}$ is large, the perturbation can materially change the update; if $w_{r,i}^{(l)}$ is small, the computation collapses back toward the real-token update.

Natural annealing through attention. The second structural fact is that $w_{r,i}^{(l)}$ is not arbitrary. It satisfies

$$w_{r,i}^{(l)} \leq \frac{k}{k + n \cdot \exp(\Delta_i^{(l)} / \sqrt{d})}, \quad (3)$$

where $\Delta_i^{(l)} := \min_{j \in [n]} \mathbf{q}_i^\top \mathbf{k}_j - \max_{j' \in [k]} \mathbf{q}_i^\top \mathbf{k}_{r_{j'}}$ is the real-vs-random logit gap and n is the number of real tokens. This bound yields a clean *sequence-wise dilution effect*: for a fixed gap, every newly generated real token tightens the maximum random-token attention compatible with that state. A similar *layer-wise* decay is expected whenever deeper layers enlarge the real-vs-random gap $\Delta_i^{(l)}$. We intentionally treat that layer-wise statement as a model-side hypothesis rather than as a universal theorem, because it depends on how the model sharpens real-token retrieval with depth. The resulting prediction is simple: the perturbation should be strongest exactly when the model has the least context and many continuations remain open, and it should weaken as evidence accumulates.

Theorem 2 (Sequence-wise annealing under a fixed logit gap). *Fix k and $\Delta_i^{(l)}$. The upper bound in Eq. (3) is strictly decreasing in the number n of real tokens. Consequently, cache growth alone imposes a monotone decrease on the largest random-token attention compatible with that logit gap.*

Theorem 2 proves a structural dilution effect, not monotone decay of the realized $w_{r,i}^{(l)}$ on every decoding path, because $\Delta_i^{(l)}$ also evolves with state. The layer-wise half is likewise model-specific and must be checked empirically.

Feed-forward propagation under local Lipschitz control. The decomposition above governs the attention sublayer. A real transformer block applies a feed-forward sublayer $F^{(l)}$ after attention with residual connections, so the cross-block perturbation depends on how $F^{(l)}$ propagates $\boldsymbol{\eta}_i^{(l)}$. If $F^{(l)}$ is $L^{(l)}$ -Lipschitz on a bounded set containing both the perturbed and unperturbed post-attention states, then Appendix E.4 shows

$$\|\mathbf{x}_i^{(l+1, \text{RSP})} - \mathbf{x}_i^{(l+1, \text{base})}\| \leq (1 + L^{(l)}) (\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|),$$

which together with $\|\boldsymbol{\eta}_i^{(l)}\| \leq w_{r,i}^{(l)} \max_j \|\mathbf{v}_{r_j}^{(l)}\|$ shows that the per-block perturbation amplitude scales linearly with $w_{r,i}^{(l)}$. The bounded-set assumption is plausible in practice because LayerNorm constrains the post-attention norm, so the perturbed and unperturbed states traversed during a single forward pass lie in a common bounded region. Iterating across the fixed finite stack yields the explicit cumulative bound of Corollary 7. Under standard architectural assumptions this bound is finite for a fixed model depth, and the later-block injections themselves shrink as $w_{r,i}^{(l)}$ decreases. Stronger uniform-in-depth non-expansion would require additional contraction assumptions, which we do not invoke here.

Empirical verification. We test exactly these model-dependent steps on Qwen2.5-Math-7B under suffix injection. If $w_{r,i}^{(l)}$ really acts as an annealed exploration coefficient, then the realized attention to RSP tokens should decay across reasoning positions and layers, whereas attention to question tokens should not exhibit the same monotone pattern because they remain task-relevant throughout decoding. Figure 4 shows precisely this contrast. The figure therefore serves a specific logical role: it validates the model-side sharpening behavior that Eq. (3) by itself does not prove.




Figure 4: Per-token attention mass on Qwen2.5-Math-7B (suffix, 500 MATH-500 problems, each with an independently sampled RSP). The layer and reasoning-position axes are each partitioned into five quantile bins; every cell reports the 500-sample mean per-token attention mass directed to (a) the RSP embeddings and (b) the question tokens within that region. The horizontal axis spans reasoning position from early (left) to late (right) generation steps, and the vertical axis spans layer depth from shallow (top) to deep (bottom). Attention to the RSP embeddings in (a) decays along both axes, while attention to the question tokens in (b) remains comparatively uniform.

4.3 Output Distribution Shift and Exploration

The practical consequence of the previous mechanism is branch opening. To connect that mechanism to observable decoding behavior, we now return to the output distribution. The relevant question is not whether RSP merely flattens probabilities, but whether it can make the decoder consider tokens that temperature scaling can never promote into contention.

Why temperature is insufficient. Temperature scaling transforms probabilities as $p_i^{(\tau)} \propto p_i^{1/\tau}$ and preserves token ranking, i.e., $p_i > p_j \iff p_i^{(\tau)} > p_j^{(\tau)}$. It can soften or sharpen confidence, but it cannot create a new ordering among candidate tokens at a fixed decoding state.

Why RSP changes rankings. RSP acts before softmax, through the hidden state. At a fixed decoding state, temperature rescales existing logit gaps, whereas RSP perturbs the gaps themselves. Using the same first-order local surrogate as above, define

$$z_i^{\text{lin}}(\bar{h}) := z_i + a_i^\top \bar{h}.$$

The next two propositions are statements about these surrogate logits, not the exact transformer logits; this scope is sufficient for the comparison to temperature because temperature is rank-preserving regardless of approximation.

The next two propositions use only one additional support property of the centered prompt law: every nonempty open subset of $\mathbb{R}^d$ receives positive probability. The Gaussian instantiation of Definition 1 satisfies this automatically, but the proofs themselves do not require Gaussian tails once this support condition is in place.

Proposition 1 (Stochastic ranking under the local affine surrogate). *If there exists a token pair (i, j) such that $a_i \neq a_j$, and if the centered prompt law assigns positive probability to every nonempty open subset of $\mathbb{R}^d$, then*

$$0 < \Pr_{\bar{h}}(z_i^{\text{lin}}(\bar{h}) > z_j^{\text{lin}}(\bar{h})) < 1.$$

Thus, unlike temperature scaling, the local affine surrogate of RSP randomizes token rankings at each decoding step.

Proposition 1 is the minimal statement: pairwise order can flip. Exploration requires a stronger conclusion, namely that tokens outside the baseline candidate set can actually enter it.

411 **Proposition 2** (Top- K entry under a nonempty surrogate entry region). *Let $J_K(s_t)$ be the baseline*
 412 *top- K set at decoding state s_t , and let $i \notin J_K(s_t)$. Define*

$$\mathcal{G}_{i,K} := \{\bar{h} \in \mathbb{R}^d : \#\{j \neq i : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h})\} \leq K - 1\}.$$

413 *If $\mathcal{G}_{i,K}$ contains a nonempty open set and the centered prompt law assigns positive probability to*
 414 *every nonempty open subset of $\mathbb{R}^d$, then*

$$\Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0.$$

415 Proposition 2 is intentionally conditional: it does not claim that every excluded token admits such a
 416 top- K entry region. It states that whenever the first-order model contains an open region in prompt
 417 space where a previously excluded token moves into the top- K , the centered perturbation reaches
 418 that region with positive probability. The open-region hypothesis can be stated in two equivalent
 419 ways. First, $\mathcal{G}_{i,K}$ is the union $\bigcup_S \mathcal{O}_S$ over subsets $S \subseteq \{j \neq i\}$ with $|S| \leq K - 1$, where each $\mathcal{O}_S$
 420 is the strict-inequality polyhedron in which precisely the indices in S outrank token i (Appendix E);
 421 each $\mathcal{O}_S$ is open in $\mathbb{R}^d$ as a finite intersection of open half-spaces, so $\mathcal{G}_{i,K}$ is nonempty whenever
 422 some $\mathcal{O}_S$ is nonempty. Second, sufficient strict-margin condition: the existence of any $\bar{h}_0 \in \mathbb{R}^d$
 423 at which token i outranks all but at most $K - 1$ other tokens with a strict margin (no tie) yields
 424 $\bar{h}_0 \in \mathcal{O}_S$ for $S = \{j : z_j + a_j^\top \bar{h}_0 > z_i + a_i^\top \bar{h}_0\}$, and the open polyhedron $\mathcal{O}_S$ contains an open
 425 ball around $\bar{h}_0$. The complementary case—no such $\bar{h}_0$ exists—is exactly the situation in which
 426 token i is dominated by the baseline top- K in the affine logit surrogate, and Proposition 2 makes
 427 no claim there. Section 3.4 verifies empirically that nonempty entry regions exist for a substantial
 428 fraction of decoding states. Gaussianity is one sufficient way to realize the required support condition;
 429 temperature cannot reach $\mathcal{G}_{i,K}$ because it never changes token order at all. The mechanism here is
 430 branch opening, not mere probability flattening.

431 **Corollary 2** (Sample complexity of top- K admission). *Let $\mu_i := \Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0$ denote the*
 432 *per-sample admission probability for token i . Within N independent fresh RSP draws, the probability*
 433 *of at least one admission is at least*

$$1 - (1 - \mu_i)^N \geq 1 - \exp(-N\mu_i).$$

434 Corollary 2 closes the loop from per-sample positive probability to multi-trial exploration. Even
 435 when μ_i is small, $N \cdot \mu_i = \Theta(1)$ samples suffice to admit a previously excluded token with constant
 436 probability. The independence assumption first enters here: the single-run rank-flip and top- K entry
 437 statements above do not require it. This is the formal bridge from the static guarantee of Proposition 2
 438 to the trajectory-level Pass@ N phenomena reported in Section 3.4.

439 **From rank flips to trajectory diversity to accuracy gain.** Together, Propositions 1 and 2 close the
 440 logical chain from Section 3.3 to Section 3.4, up to the local-surrogate assumption above. Early RSP
 441 perturbations alter logit gaps before softmax, which can admit previously excluded tokens into the
 442 candidate set. Because autoregressive decoding compounds early choices, one such early admission is
 443 enough to push the trajectory into a different reasoning cluster even if later perturbations are already
 444 small. RSP therefore does not need to perturb the entire chain of thought uniformly: it only needs to
 445 alter the first few branching decisions before attention-mediated attenuation suppresses further noise,
 446 which is why early exploration and late stabilization coexist in the same mechanism.

447 The connection from trajectory diversity to accuracy gain requires a property of reasoning tasks
 448 themselves: the existence of multiple valid trajectories. We formalize this as a separate proposition
 449 rather than absorb it into the perturbation analysis.

450 **Proposition 3** (Exploration-meets-multi-path accuracy gain). *Fix a problem x , and consider the pro-*
 451 *cedure that runs the deterministic baseline once and additionally draws N fresh RSPs $\bar{h}^{(1)}, \dots, \bar{h}^{(N)}$,*
 452 *with the model decoded once per sample. Let $\mathcal{T}^*(x)$ denote the (nonempty) set of trajectories*
 453 *reaching the correct answer and $p_{\text{base}}(x) \in \{0, 1\}$ the deterministic-baseline correctness indicator.*
 454 *Assume (M1) for each $n \in \{1, \dots, N\}$, the marginal probability that the n -th RSP run produces a*
 455 *trajectory in $\mathcal{T}^*(x)$ is at least $p_{\min}(x) > 0$ (the same lower bound for every n), and (M2) the N RSP*
 456 *runs are mutually independent (independent prompt draws and, if decoding is stochastic, independent*
 457 *decoding randomness across runs). Then*

$$\Pr(\text{Pass}@ (N+1) \text{ correct on } x) \geq \max\{p_{\text{base}}(x), 1 - (1 - p_{\min}(x))^N\} \geq 1 - e^{-Np_{\min}(x)}.$$

458 *The expected accuracy gain over the deterministic baseline is at least*

$$\frac{1}{|\mathcal{D}|} \sum_{x \in \mathcal{D}} (1 - p_{\text{base}}(x)) (1 - \exp(-N p_{\text{min}}(x))).$$

459 *Proof.* The Pass@($N+1$) event holds if the baseline run reaches $\mathcal{T}^*(x)$ (probability $p_{\text{base}}(x) \in$
 460 $\{0, 1\}$) or if any of the N RSP runs does. By (M2), the failure probabilities of the N RSP runs
 461 multiply, so the probability that at least one RSP run reaches $\mathcal{T}^*(x)$ is $1 - (1 - p_{\text{min}}(x))^N$ by (M1).
 462 Taking the maximum gives the first lower bound; $1 - x \leq e^{-x}$ gives the second. Restricting to
 463 $\{x : p_{\text{base}}(x) = 0\}$ yields the expected gain. $\square$

464 Proposition 3 converts the local guarantees of Propositions 1 and 2 into a global accuracy claim, but
 465 only via the auxiliary assumption that $p_{\text{min}}(x) > 0$. The next lemma formalizes a sufficient condition
 466 under which the perturbation theory of this section actually delivers $p_{\text{min}}(x) > 0$, separating what is
 467 proved (positive measure of an entry region) from what is assumed (extension of a perturbed prefix to
 468 a correct completion).

469 **Lemma 2** (Bridge from open entry region to one-run success). *Specialize Proposition 2 to $K = 1$*
 470 *(greedy decoding) and assume the same support condition on the centered prompt law. Suppose there*
 471 *exists a nonempty open ball $B \subseteq \mathbb{R}^d$ such that for every $\bar{h} \in B$ both of the following hold under a*
 472 *single RSP draw at the start of decoding: (B1) the RSP-induced trajectory under prompt $\bar{h}$ reaches*
 473 *some decoding state $s_{t^*}(\bar{h})$ at which an excluded token $i^*(\bar{h}) \notin J_1(s_{t^*})$ satisfies $\bar{h} \in \mathcal{G}_{i^*,1}(s_{t^*})$, so*
 474 *greedy decoding selects $i^*(\bar{h})$ at step t^* ; (B2) the resulting full trajectory belongs to $\mathcal{T}^*(x)$. Then*
 475 $p_{\text{min}}(x) \geq \Pr(\bar{h} \in B) > 0$.

476 *Proof.* By definition, $p_{\text{min}}(x)$ is the per-sample probability that the RSP-induced full trajectory under
 477 one draw of $\bar{h}$ reaches $\mathcal{T}^*(x)$. Conditions (B1) and (B2) state exactly that this event holds for every
 478 $\bar{h} \in B$. Since B is a nonempty open subset of $\mathbb{R}^d$ and the centered prompt law assigns positive
 479 probability to every nonempty open set, $\Pr(\bar{h} \in B) > 0$, and the lower bound follows. $\square$

480 Lemma 2 converts the abstract assumption $p_{\text{min}}(x) > 0$ of Proposition 3 into the existence of an
 481 open ball B satisfying two simultaneous properties: (B1) the perturbation theory of this section—
 482 specifically Proposition 2 at $K = 1$ —supplies the open-set structure that triggers top-1 entry once
 483 a candidate decoding state s_{t^*} is fixed; (B2) the resulting full trajectory continues into a correct
 484 completion. The prefix-preservation aspect (the trajectory under $\bar{h}$ actually reaches a state s_{t^*} with
 485 the desired entry-region structure) and the completion property (B2) are both model- and dataset-
 486 dependent and treated as empirical hypotheses. Section 3.4 plays a load-bearing role here: it verifies
 487 that RSP-induced trajectories include both improvements and regressions over the baseline with net
 488 gain emerging under repeated sampling, which is precisely the empirical signature predicted by such
 489 an open ball B existing for some but not all problems. The mechanism we have isolated is exploration
 490 meeting multi-path structure, not steering toward correctness.

491 Taken together, the causal chain is now explicit. Randomness matters because RSP is resampled,
 492 centered, and direction-agnostic; attention turns that perturbation into an early exploration coefficient
 493 and then anneals it; the resulting early rank flips open trajectories that temperature scaling alone
 494 cannot reach; and the multi-path structure of reasoning converts trajectory diversity into accuracy gain
 495 under repeated sampling. Section 3 therefore reports four observable faces of a single mechanism
 496 rather than four disconnected phenomena.

497 5 Conclusion

498 **Summary.** This work investigated the role of Random Soft Prompts (RSPs), continuous embeddings
 499 sampled from an isotropic Gaussian fitted to the empirical mean and variance of the pretrained
 500 embedding table and injected into the model input without any training. Through empirical and
 501 theoretical analyses, we established three findings. First, RSPs increase early-stage token entropy,
 502 top-1 probability reduction, and varentropy while the model’s output stabilizes in later generation
 503 stages (Section 3). Second, RSPs induce stochastic token rankings and expand the effective support of

the output distribution, a property that temperature sampling fundamentally lacks due to its ranking-preserving nature (Section 4). Third, combining RSPs with temperature sampling produces more diverse reasoning trajectories than either method alone, as measured by Pass@ n scaling (Section 3).

Implications for GRPO training. GRPO computes advantages by comparing multiple sampled trajectories within each group, making trajectory diversity essential for effective learning signals. Existing approaches such as DAPO [Yu et al., 2025] and CE-GPPO [Su et al., 2025] address entropy collapse by modifying the clipping range or reintroducing gradients from clipped tokens, but the underlying sampling process remains based on temperature sampling, which preserves token rankings and cannot expand the output support. Our theoretical and empirical results suggest that RSP can address this limitation at the sampling level: by perturbing the initial output distribution with freshly resampled random embeddings across compared rollouts, RSP provides a complementary source of diversity that is orthogonal to temperature-based exploration. We expect this to yield more effective learning signals for GRPO training, particularly in mitigating entropy collapse.

Limitations. Our experiments are conducted on mathematical reasoning benchmarks (MATH-500, GSM8K, AIME24) with a limited set of model families (Qwen2.5-Math, LLaMA-3.1). The energy-theoretic analysis in Section 4 relies on strong assumptions: the Hopfield–Attention equivalence (A1: $\mathbf{K} = \mathbf{V}$) and weight-tied layers (A2), which hold exactly only in Hopfield-type models and deep equilibrium models, respectively. Since real transformer architectures vary widely in their attention mechanisms and layer structures, the theoretical predictions are most directly applicable to models whose attention dynamics closely approximate the Hopfield energy framework. The GRPO application remains a theoretical expectation and has not been empirically validated.

Future work. Extending the experimental evaluation to other domains (e.g., code generation, commonsense reasoning) and larger model scales would test the generality of RSP’s effects. Empirically validating the GRPO application by training with RSP-augmented rollouts and comparing against DAPO and CE-GPPO is a natural next step. Investigating the optimal RSP distribution and token count, which currently require per-model tuning, could lead to more principled injection strategies.

References

- Farhan Ahmed, Yuya Jeremy Ong, and Chad DeLuca. LogitScope: A Framework for Analyzing LLM Uncertainty through Information Metrics. In *ICLR Workshop*, 2026.
- Nora Belrose, Zach Furman, Logan Smith, Danny Halawi, Igor Ostrovsky, Lev McKinney, Stella Biderman, and Jacob Steinhardt. Eliciting Latent Predictions from Transformers with the Tuned Lens. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Stéphane Boucheron, Gábor Lugosi, and Pascal Massart. *Concentration Inequalities: A Nonasymptotic Theory of Independence*. Oxford University Press, 2013.
- Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. BGE M3-Embedding: Multi-Lingual, Multi-Functionality, Multi-Granularity Text Embeddings Through Self-Knowledge Distillation. *arXiv:2402.03216*, 2024.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training Verifiers to Solve Math Word Problems. *arXiv:2110.14168*, 2021.
- Jeremy M. Cohen, Elan Rosenfeld, and J. Zico Kolter. Certified Adversarial Robustness via Randomized Smoothing. In *ICML*, 2019.
- Timothée Darcet, Maxime Oquab, Julien Mairal, and Piotr Bojanowski. Vision Transformers Need Registers. In *International Conference on Learning Representations (ICLR)*, 2024.
- Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu. A Density-Based Algorithm for Discovering Clusters in Large Spatial Databases with Noise. In *Proceedings of the Second International Conference on Knowledge Discovery and Data Mining (KDD)*, pages 226–231, 1996.

551 Kawin Ethayarajh. How Contextual are Contextualized Word Representations? Comparing the
552 Geometry of BERT, ELMo, and GPT-2 Embeddings. In *Proceedings of the 2019 Conference on*
553 *Empirical Methods in Natural Language Processing (EMNLP)*, 2019.

554 Meire Fortunato, Mohammad Gheshlaghi Azar, Bilal Piot, Jacob Menick, Ian Osband, Alex Graves,
555 Vlad Mnih, Remi Munos, Demis Hassabis, Olivier Pietquin, Charles Blundell, and Shane Legg.
556 Noisy Networks for Exploration. In *ICLR*, 2018.

557 Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. Transformer Feed-Forward Layers Are
558 Key-Value Memories. In *Proceedings of the 2021 Conference on Empirical Methods in Natural*
559 *Language Processing (EMNLP)*, 2021.

560 Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh
561 Nagarajan. Think before You Speak: Training Language Models with Pause Tokens. In *ICLR*,
562 2024.

563 Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, et al. The
564 Llama 3 Herd of Models. *arXiv:2407.21783*, 2024.

565 Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian.
566 Training Large Language Models to Reason in a Continuous Latent Space. In *ICLR*, 2025.

567 Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin de Laroussilhe, Andrea
568 Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-Efficient Transfer Learning for NLP.
569 In *ICML*, 2019.

570 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang,
571 and Weizhu Chen. LoRA: Low-Rank Adaptation of Large Language Models. In *ICLR*, 2022.

572 Neel Jain, Ping-yeh Chiang, Yuxin Wen, John Kirchenbauer, Hong-Min Chu, Gowthami Somepalli,
573 Brian R. Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Aniruddha Saha, Micah Goldblum,
574 Jonas Geiping, and Tom Goldstein. NEFTune: Noisy Embeddings Improve Instruction Finetuning.
575 In *ICLR*, 2024.

576 Xinyue Kang, Diwei Shi, and Li Chen. Model Whisper: Steering Vectors Unlock Large Language
577 Models’ Potential in Test-time. In *AAAI*, 2026.

578 Hyunjik Kim, George Papamakarios, and Andriy Mnih. The Lipschitz Constant of Self-Attention.
579 *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021.

580 Brian Lester, Rami Al-Rfou, and Noah Constant. The Power of Scale for Parameter-Efficient Prompt
581 Tuning. In *EMNLP*, 2021.

582 Xiang Lisa Li and Percy Liang. Prefix-Tuning: Optimizing Continuous Prompts for Generation. In
583 *ACL-IJCNLP*, 2021.

584 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike,
585 John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s Verify Step by Step. In *ICLR*, 2024.

586 Xiao Liu, Yanan Zheng, Zhengxiao Du, Ming Ding, Yujie Qian, Zhilin Yang, and Jie Tang. GPT
587 Understands, Too. *arXiv:2103.10385*, 2023.

588 Aman Madaan and Amir Yazdanbakhsh. Text and Patterns: For Effective Chain of Thought, It Takes
589 Two to Tango. *arXiv:2209.07686*, 2022.

590 Jiaqi Mu and Pramod Viswanath. All-but-the-Top: Simple and effective postprocessing for word
591 representations. In *International Conference on Learning Representations (ICLR)*, 2018.

592 Yurii Nesterov. *Lectures on Convex Optimization*. Springer, 2 edition, 2018.

593 Hubert Ramsauer, Bernhard Schöfl, Johannes Lehner, Philipp Seidl, Michael Widrich, Thomas Adler,
594 Lukas Gruber, Markus Holzleitner, Milena Pavlović, Geir Kjetil Sandve, Thomas Unterthiner,
595 Bernhard Nessler, and Sepp Hochreiter. Hopfield Networks Is All You Need. In *ICLR*, 2021.

596 Alexander Robey, Eric Wong, Hamed Hassani, and George J. Pappas. SmoothLLM: Defending Large
597 Language Models Against Jailbreaking Attacks. *arXiv:2310.03684*, 2023.

598 Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang,
599 Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the Limits of
600 Mathematical Reasoning in Open Language Models. *arXiv:2402.03300*, 2024.

601 Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov.
602 Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *Journal of Machine*
603 *Learning Research*, 15:1929–1958, 2014.

604 Zhenpeng Su, Leiyu Pan, Minxuan Lv, Yuntao Li, Wenping Hu, Fuzheng Zhang, Kun Gai, and Guorui
605 Zhou. CE-GPPO: Coordinating Entropy via Gradient-Preserving Clipping Policy Optimization in
606 Reinforcement Learning. *arXiv:2509.20712*, 2025.

607 Roman Vershynin. *High-Dimensional Probability: An Introduction with Applications in Data Science*.
608 Cambridge University Press, 2018.

609 Shenzi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen,
610 Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen
611 Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 Rule: High-Entropy Minority Tokens Drive
612 Effective Reinforcement Learning for LLM Reasoning. In *NeurIPS*, 2025.

613 Ruibin Xiong, Yunchang Yang, Di He, Kai Zheng, Shuxin Zheng, Chen Xing, Huishuai Zhang,
614 Yanyan Lan, Liwei Wang, and Tie-Yan Liu. On Layer Normalization in the Transformer Architec-
615 ture. In *International Conference on Machine Learning (ICML)*, 2020.

616 Yige Xu, Xu Guo, Zhiwei Zeng, and Chunyan Miao. SoftCoT: Soft Chain-of-Thought for Efficient
617 Reasoning with LLMs. In *ACL*, 2025.

618 An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu,
619 Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu,
620 Xingzhang Ren, and Zhenru Zhang. Qwen2.5-Math Technical Report: Toward Mathematical
621 Expert Model via Self-Improvement. *arXiv:2409.12122*, 2024.

622 Wengao Ye, Yan Liang, and Lianlei Shan. Thinking on the Fly: Test-Time Reasoning Enhancement
623 via Latent Thought Policy Optimization. In *ICLR*, 2026.

624 Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian
625 Fan, Gaohong Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming Sheng,
626 Yuxuan Tong, Chi Zhang, Mofan Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu, Jiaze Chen, Jiangjie
627 Chen, Chengyi Wang, Hongli Yu, Yuxuan Song, Xiangpeng Wei, Hao Zhou, Jingjing Liu, Wei-
628 Ying Ma, Ya-Qin Zhang, Lin Yan, Mu Qiao, Yonghui Wu, and Mingxuan Wang. DAPO: An
629 Open-Source LLM Reinforcement Learning System at Scale. *arXiv:2503.14476*, 2025.

630 Guibin Zhang, Muxin Fu, and Shuicheng Yan. MemGen: Weaving Generative Latent Memory for
631 Self-Evolving Agents. *arXiv:2509.24704*, 2025.

632 A Experimental Setup and Full Accuracy Breakdown

633 Table 6 provides the full per-position accuracy breakdown (Prefix, Infix, Suffix) corresponding to
634 the best-across-positions results summarized in Table 1 in the main text. Table 7 lists the number of
635 RSP tokens L used for each model–benchmark pair, selected from $\{10, 15, 20\}$. All experiments are
636 conducted on a single NVIDIA A6000 40GB GPU with greedy decoding, a maximum generation
637 length of 3,072 tokens for MATH-500 and GSM8K, and 4,096 tokens for AIME24. Batch size is
638 fixed per model (16 for LLaMA-3.1-8B-Instruct and Qwen2.5-Math-7B variants, 32 for Qwen2.5-
639 Math-1.5B variants).

Table 6: Accuracy (%) across model configurations, injection positions, and benchmarks. Values in parentheses indicate the difference from the baseline. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Mode	MATH-500	GSM8K	AIME24
<i>Instruct Models</i>				
LLaMA-3.1-8B-Instruct	Baseline	52.20	85.75	<u>6.67</u>
	Prefix	45.00 (−7.2)	76.35 (−9.4)	3.33 (−3.3)
	Infix	49.40 (−2.8)	85.97 (+0.2)	10.00 (+3.3)
	Suffix	<u>49.40</u> (−2.8)	86.73 (+1.0)	10.00 (+3.3)
Qwen2.5-Math-7B-Instruct	Baseline	83.20	95.45	<u>13.33</u>
	Prefix	81.40 (−1.8)	94.92 (−0.5)	13.33 (+0.0)
	Infix	84.20 (+1.0)	95.60 (+0.2)	16.67 (+3.3)
	Suffix	<u>83.40</u> (+0.2)	95.68 (+0.2)	16.67 (+3.3)
Qwen2.5-Math-1.5B-Instruct	Baseline	73.00	<u>85.29</u>	10.00
	Prefix	74.80 (+1.8)	85.29 (+0.0)	<u>13.33</u> (+3.3)
	Infix	75.20 (+2.2)	85.75 (+0.5)	6.67 (−3.3)
	Suffix	74.20 (+1.2)	84.69 (−0.6)	20.00 (+10.0)
<i>Base Models (Raw Text)</i>				
Qwen2.5-Math-7B	Baseline	72.20	84.15	6.67
	Prefix	<u>71.60</u> (−0.6)	84.31 (+0.2)	16.67 (+10.0)
	Infix	63.20 (−9.0)	64.29 (−19.9)	16.67 (+10.0)
	Suffix	55.60 (−16.6)	54.13 (−30.0)	<u>13.33</u> (+6.7)
Qwen2.5-Math-1.5B	Baseline	<u>61.40</u>	77.86	16.67
	Prefix	64.40 (+3.0)	74.30 (−3.6)	10.00 (−6.7)
	Infix	63.20 (+1.8)	73.92 (−3.9)	<u>10.00</u> (−6.7)
	Suffix	54.60 (−6.8)	58.61 (−19.3)	3.33 (−13.3)
<i>Base Models + ChatML (Format Mismatch)</i>				
Qwen2.5-Math-7B	Baseline	52.20	58.30	23.33
	Prefix	59.20 (+7.0)	56.41 (−1.9)	<u>3.33</u> (−20.0)
	Infix	<u>62.00</u> (+9.8)	69.07 (+10.8)	23.33 (+0.0)
	Suffix	70.40 (+18.2)	75.97 (+17.7)	23.33 (+0.0)
Qwen2.5-Math-1.5B	Baseline	34.40	37.45	<u>13.33</u>
	Prefix	63.40 (+29.0)	67.02 (+29.6)	20.00 (+6.7)
	Infix	39.20 (+4.8)	50.42 (+13.0)	6.67 (−6.7)
	Suffix	<u>51.00</u> (+16.6)	<u>50.64</u> (+13.2)	<u>13.33</u> (+0.0)

Table 7: Number of RSP tokens (L) used for each model and benchmark.

Model	MATH-500	GSM8K	AIME24
LLaMA-3.1-8B-Instruct	15	20	15
Qwen2.5-Math-7B-Instruct	20	20	20
Qwen2.5-Math-1.5B-Instruct	20	10	20
Qwen2.5-Math-7B	10	10	20
Qwen2.5-Math-1.5B	10	10	20
Qwen2.5-Math-1.5B (ChatML)	20	20	10
Qwen2.5-Math-7B (ChatML)	20	20	20

640 A.1 RSP Injection Positions and Prompt Templates

641 Below we show the prompt structure for each injection position across all model types. Blue text
642 indicates RSP embeddings, and purple text indicates special tokens.

643 A.1.1 Prefix

644 RSP embeddings are concatenated before the prompt embeddings. No text modification is applied.

645 ChatML (Qwen Instruct / Base + ChatML).

```
[Random Embeddings (L tokens)]
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{<|im_end|>
<|im_start|>user
{question}<|im_end|>
<|im_start|>assistant
```

647 LLaMA Chat Template.

```
[Random Embeddings (L tokens)]
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{<|eot_id|>
<|start_header_id|>user<|end_header_id|>
{question}<|eot_id|>
<|start_header_id|>assistant<|end_header_id|>
```

649 Raw Text (Qwen Base).

```
[Random Embeddings (L tokens)]
{question}
Please reason step by step, and put your final answer within \boxed{<|eot_id|>
```

651 A.1.2 Infix

652 A description text and special tokens are inserted into the prompt. The special token positions are
653 then replaced with RSP embeddings at the embedding level.

654 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{<|im_end|>
<|im_start|>user
{question}

There are {L} special tokens that contain compressed latent reasoning information that
might be useful for your reasoning. If these tokens are useful for your case, you can
use them as reference. If these tokens are not useful for your case, you can ignore
them and focus back to solving the problem.

Here are the {L} special tokens: <special_token_1>...<special_token_L>
<|im_end|>
<|im_start|>assistant
```

656 LLaMA Chat Template.

```
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{<|eot_id|>
<|start_header_id|>user<|end_header_id|>
```

{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens:

```
<reserved_special_token_0>...  
<reserved_special_token_L>  
<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

658

659 Raw Text (Qwen Base).

{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens: $\langle \text{special_token_1} \rangle \dots \langle \text{special_token_L} \rangle$

Please reason step by step, and put your final answer within `\boxed{\}`.

660

661 A.1.3 Suffix

662 For chat-templated models, RSP embeddings are inserted immediately before the end-of-turn token.

663 For raw text models, RSP embeddings are concatenated at the end of the prompt.

664 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system  
Please reason step by step, and put your final answer within \boxed{<|im_end|>  
<|im_start|>user  
{question}[Random Embeddings (L tokens)]<|im_end|>  
<|im_start|>assistant
```

665

666 LLaMA Chat Template.

```
<|begin_of_text|>  
<|start_header_id|>system<|end_header_id|>  
Cutting Knowledge Date: December 2023  
Today Date: 26 Jul 2024  
Please reason step by step, and put your final answer within \boxed{<|eot_id|>  
<|start_header_id|>user<|end_header_id|>  
{question}[Random Embeddings (L tokens)]<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

667

668 Raw Text (Qwen Base).

```
{question}  
Please reason step by step, and put your final answer within \boxed{<|eot_id|>[Random  
Embeddings (L tokens)]
```

669

670 A.2 Answer Extraction and Grading

671 The final answer is extracted from the model output through a fallback chain. Each step is attempted

672 in order; if extraction fails, the next step is tried.

Answer Extraction Fallback Chain

1. "final answer is \$...\$. I hope" pattern (Minerva format)
2. `\boxed{...}`: last non-empty box is used; empty `\boxed{}` are skipped
3. Text following "the answer is"
4. Text following "final answer is"
5. Last number in the output (regex fallback)

Extracted answers are normalized by removing \$, \left, \right, and unit strings. Grading is performed via symbolic equivalence checking: exact string match, numerical comparison (`math.isclose, rel_tol=1e-4`), and SymPy symbolic simplification.

A.3 Reproduced Prior Work Hyperparameters

All prior methods are reproduced on Qwen2.5-Math-1.5B-Instruct and Qwen2.5-Math-7B-Instruct, evaluated with greedy decoding (temperature = 0) and our unified answer extraction pipeline.

TTSV. Prefix length 20, trained for 20 epochs with AdamW optimizer, batch size 2, gradient accumulation 8 steps, and linear learning rate schedule. Learning rate is 1e-3 for Qwen-1.5B and 1e-5 for Qwen-7B.

LTPO. Table 8 reports the per-benchmark hyperparameters.

Table 8: LTPO hyperparameters.

Parameter	GSM8K	MATH-500	AIME24
tokens	8	8	8
steps	20	20	20
top_k	10	10	10
sigma	20	20	5
sigma_decay	0.95	0.95	0.95
lr	1e-2	5e-3	5e-2

SoftCoT. Projection module trained for 10 epochs. Thought tokens: 32 during training, 4 during evaluation. Batch size 4 for GSM8K, 1 for MATH with gradient accumulation 4. The small (assistant) model is Qwen2.5-Math-1.5B-Instruct in both configurations, while the large model is Qwen2.5-Math-1.5B-Instruct (learning rate 2e-5) or Qwen2.5-Math-7B-Instruct (learning rate 5e-6). For MATH-500 and AIME24 evaluation, we use the projection module trained on GSM8K, as it yields higher accuracy than the one trained on MATH.

A.4 Full Entropy Curves

Figure 5 shows the full normalized (0–100%) generation trajectories for entropy, top-1 probability, and varentropy. All methods converge to similar levels after the initial generation stage, confirming that the distributional changes induced by RSP are localized to the early reasoning phase.




Figure 5: Mean entropy (top), top-1 probability (middle), and varentropy (bottom) over the full normalized generation trajectory (0–100%). Averaged over MATH-500, shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: existing soft prompt methods. All methods converge after the initial stage.

694 B Additional Semantic Diversity Analysis

695 This appendix complements the semantic diversity results reported in Section 3.4. We include
 696 per-problem t-SNE plots across a broader set of problems.



Figure 6: t-SNE projections of BGE-M3 embeddings for 48 out of the 64 evaluated MATH-500 problems (6×8 grid). Each subplot shows 300 baseline (red) and 300 RSP suffix (blue) samples under $\tau = 1.0$ with Qwen2.5-Math-1.5B-Instruct; filled markers denote correct answers and open markers denote incorrect. Subplot titles indicate subject/problem_id. Across most problems, RSP trajectories spread across multiple distinct regions while baseline samples concentrate in a single dominant cluster.

697 C Token-Level Distribution Metrics

698 This appendix provides formal definitions for the three metrics used in the Token-level analysis of
 699 Section 3.4. Let P_{base} denote the baseline next-token distribution at $\tau = 1.0$ and P_{target} the candidate
 700 distribution being compared (e.g., the baseline at $\tau = 2.0$ or RSP at $\tau = 1.0$). All metrics are
 701 computed analytically from saved logits at the first generation step.

702 **Spearman rank correlation.** Spearman’s ρ is the Pearson correlation coefficient between the token
 703 ranks of two distributions:

$$\rho = \text{Pearson}(\text{rank}(P_{\text{target}}), \text{rank}(P_{\text{base}})). \quad (4)$$

704 Because temperature scaling $P^{(\tau)} \propto P^{1/\tau}$ is a strictly monotone transform, the rank order of tokens
 705 is preserved exactly, so $\rho = 1$ for any $\tau > 0$ as a mathematical identity. Any value $\rho < 1$ therefore
 706 indicates rank reorganization beyond what temperature scaling can produce.

707 **Mass outside top- K .** Let $\text{TopK}(P_{\text{base}})$ be the set of K tokens with the highest probability under
 708 P_{base} . The probability mass that the candidate distribution places *outside* this set is

$$\text{MassOutside}_K = 1 - \sum_{t \in \text{TopK}(P_{\text{base}})} P_{\text{target}}(t). \quad (5)$$

709 This directly measures support expansion: how much probability the candidate assigns to tokens that
 710 are not in the baseline’s preferred set. Reported values use $K = 10$.

711 **Jensen–Shannon divergence.** The symmetrized, bounded version of Kullback–Leibler divergence:

$$\text{JS}(P, Q) = \frac{1}{2} \text{KL}(P \parallel M) + \frac{1}{2} \text{KL}(Q \parallel M), \quad M = \frac{1}{2}(P + Q). \quad (6)$$

712 We use base-2 logarithms so that $\text{JS} \in [0, 1]$. Tokens absent from the stored top-100 are assigned a
 713 sentinel logit (-10^9) before softmax, which is negligible for our setting since the top-5 covers more
 714 than 99.9% of the probability mass.

D Why Random Prompts Are Semantically Neutral and Direction-Agnostic

This appendix formalizes the “why random” claim used in Section 4.1. The point is not that random prompts contain hidden task information. The point is structural: if the perturbation is meant to act as a generic exploration device rather than as a learned prompt prior, then it should be resampled across runs, centered so that it does not induce a fixed first-order drift, and direction-agnostic so that it does not under-serve whichever local branch direction matters for the current prompt.

D.1 Semantic Neutrality

Let $\bar{h} \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_d)$ and $\hat{h} := \bar{h}/\|\bar{h}\|$ whenever $\bar{h} \neq 0$. Let $u_1, \dots, u_M \in \mathbb{S}^{d-1}$ be any fixed unit vectors representing semantic anchor directions.

Lemma 3 (Semantic neutrality). *For any $\delta \in (0, 1)$,*

$$\Pr \left[\max_{r \in [M]} |\hat{h}^\top u_r| \leq \sqrt{\frac{2 \log(2M/\delta)}{d-1}} \right] \geq 1 - \delta.$$

Hence the normalized direction of an isotropic Gaussian prompt is simultaneously nearly orthogonal to every fixed finite collection of semantic anchors with high probability.

Proof. Because $\bar{h}$ is isotropic Gaussian, its normalized direction $\hat{h}$ is uniformly distributed on the sphere $\mathbb{S}^{d-1}$. For a fixed unit vector $u \in \mathbb{S}^{d-1}$, the marginal $\hat{h}^\top u$ satisfies the standard spherical concentration bound [Vershynin, 2018, Theorem 3.4.6][Boucheron et al., 2013, Section 3.2],

$$\Pr(|\hat{h}^\top u| \geq t) \leq 2 \exp\left(-\frac{(d-1)t^2}{2}\right) \quad \text{for } t \in (0, 1),$$

which follows from the spherical-cap area bound $\sigma\{\hat{h} \in \mathbb{S}^{d-1} : \hat{h}^\top u \geq t\} \leq \frac{1}{2}(1-t^2)^{(d-1)/2}$ together with $(1-t^2)^{(d-1)/2} \leq \exp(-(d-1)t^2/2)$ derived from $\log(1-x) \leq -x$ for $x \in [0, 1)$. Applying this bound to each u_r and using a union bound,

$$\Pr \left[\max_{r \in [M]} |\hat{h}^\top u_r| \geq t \right] \leq 2M \exp\left(-\frac{(d-1)t^2}{2}\right).$$

Setting

$$t = \sqrt{\frac{2 \log(2M/\delta)}{d-1}}$$

gives the claim. If the displayed value of t exceeds 1, the statement is trivial because $|\hat{h}^\top u_r| \leq 1$ always. $\square$

The contrast with a token embedding is immediate. If $e_s \neq 0$ is a fixed token embedding and $\hat{e}_s := e_s/\|e_s\|$, then choosing the anchor direction $u = \hat{e}_s$ gives

$$|\hat{e}_s^\top u| = 1.$$

A token prompt is therefore maximally aligned with at least one semantic direction, whereas an isotropic random prompt is nearly orthogonal to every fixed finite set of such directions. This is the sense in which tokens inject semantic commitment while random prompts act as semantically uncommitted perturbations.

D.2 Centered Perturbations Do Not Impose Systematic First-Order Drift

Let Δ_{ij} be the baseline logit gap between tokens i and j , and write the first-order surrogate as

$$\Delta_{ij}(\bar{h}) = \Delta_{ij} + b_{ij}^\top \bar{h}.$$

Proposition 4 (No systematic first-order drift). *If $\mathbb{E}[\bar{h}] = 0$, then*

$$\mathbb{E}[\Delta_{ij}(\bar{h})] = \Delta_{ij}$$

for every token pair (i, j) . If a deterministic prompt h_0 is reused across runs, then

$$\Delta_{ij}(h_0) - \Delta_{ij} = b_{ij}^\top h_0$$

is a fixed directional bias.

747 *Proof.* The centered case is immediate from linearity of expectation:

$$\mathbb{E}[\Delta_{ij}(\bar{h})] = \Delta_{ij} + b_{ij}^\top \mathbb{E}[\bar{h}] = \Delta_{ij}.$$

748 The deterministic case follows by substitution. $\square$

749 The proposition does not say that a centered perturbation leaves every *realization* unchanged; it says
 750 there is no systematic first-order push toward one branch across runs. This is exactly the behavior
 751 desired from a generic exploration device. A fixed token or fixed soft prompt may still be useful, but
 752 it acts as a persistent prompt prior rather than as an unbiased exploration mechanism.

753 D.3 Direction-Agnostic Coverage

754 The second requirement is directional coverage. Let D be any zero-mean perturbation distribution
 755 in $\mathbb{R}^d$ with covariance Σ_D and variance budget $\text{tr}(\Sigma_D) \leq \rho^2$. For a unit vector b , the first-order
 756 perturbation energy available along b is $\text{Var}_{h \sim D}(b^\top h) = b^\top \Sigma_D b$. This is the quantity that appears
 757 in the local linear logit surrogate used in the main text.

758 **Theorem 3** (Direction-agnostic coverage). *Define*

$$\mathcal{E}(D) := \min_{\|b\|=1} \text{Var}_{h \sim D}(b^\top h).$$

759 *Then*

$$\mathcal{E}(D) = \lambda_{\min}(\Sigma_D) \leq \frac{\rho^2}{d},$$

760 *with equality if and only if $\Sigma_D = (\rho^2/d)\mathbf{I}_d$.*

761 *Proof.* For any symmetric covariance matrix Σ_D ,

$$\min_{\|b\|=1} b^\top \Sigma_D b = \lambda_{\min}(\Sigma_D),$$

762 so $\mathcal{E}(D) = \lambda_{\min}(\Sigma_D)$. Let the eigenvalues of Σ_D be $\lambda_1, \dots, \lambda_d$. Then

$$\lambda_{\min}(\Sigma_D) \leq \frac{1}{d} \sum_{m=1}^d \lambda_m = \frac{\text{tr}(\Sigma_D)}{d} \leq \frac{\rho^2}{d}.$$

763 Equality holds if and only if all eigenvalues are equal, i.e., $\Sigma_D = (\rho^2/d)\mathbf{I}_d$. $\square$

764 This theorem is a statement about second moments, not about Gaussianity by itself. Many distributions
 765 could realize isotropic covariance. We use a Gaussian in RSP because it is the simplest centered,
 766 rotation-symmetric family that matches the scale of pretrained embeddings. The theorem clarifies the
 767 structural advantage: a fixed prompt has $\Sigma_D = 0$ and therefore no across-run exploration, whereas
 768 any anisotropic perturbation leaves at least one branch-sensitive direction less explored than the
 769 isotropic optimum.

770 **Proposition 5** (Covariance pushforward under the value projection). *Let $\bar{h}$ be a centered random*
 771 *vector with covariance $\sigma_E^2 \mathbf{I}_d$, and let $v = W\bar{h}$ for a matrix $W \in \mathbb{R}^{d_v \times d}$. Then*

$$\text{Cov}(v) = \sigma_E^2 W W^\top.$$

772 *Consequently, for every unit vector $u \in \mathbb{S}^{d_v-1}$,*

$$\sigma_E^2 \lambda_{\min}(W W^\top) \leq \text{Var}(u^\top v) = \sigma_E^2 u^\top W W^\top u \leq \sigma_E^2 \lambda_{\max}(W W^\top).$$

773 *In particular, if W is full rank then the projected perturbation has strictly positive variance in every*
 774 *output-space direction.*

775 *Proof.* Since $v = W\bar{h}$ and $\mathbb{E}[\bar{h}] = 0$,

$$\text{Cov}(v) = \mathbb{E}[v v^\top] = W \mathbb{E}[\bar{h} \bar{h}^\top] W^\top = \sigma_E^2 W W^\top.$$

776 For any unit vector u ,

$$\text{Var}(u^\top v) = u^\top \text{Cov}(v) u = \sigma_E^2 u^\top W W^\top u,$$

777 and the Rayleigh quotient bounds give the stated inequalities. $\square$

E Proofs for Attention-Mediated Annealing

E.1 Hopfield-Energy Descent and Basin Trapping

We restate the Modern Hopfield correspondence in the form used in the main text. For a key matrix $\mathbf{K} = (\mathbf{k}_1, \dots, \mathbf{k}_n)$ at a fixed query position, the Hopfield energy is

$$E_{\mathbf{K}}(\sigma) := -\beta^{-1} \log \sum_{j=1}^n \exp(\beta \mathbf{k}_j^\top \sigma) + \frac{1}{2} \|\sigma\|^2.$$

A direct calculation gives the gradient

$$\nabla E_{\mathbf{K}}(\sigma) = \sigma - \mathbf{K}^\top \text{softmax}(\beta \mathbf{K} \sigma).$$

Under (A1) $\mathbf{K} = \mathbf{V}$, the attention update with this $\mathbf{K}$ is therefore $\sigma^{\text{new}} = \sigma - \nabla E_{\mathbf{K}}(\sigma) = \mathbf{K}^\top \text{softmax}(\beta \mathbf{K} \sigma)$, which is one gradient step on $E_{\mathbf{K}}$ at unit step size [Ramsauer et al., 2021].

Lemma 4 (Descent inequality). *If $E_{\mathbf{K}}$ is L_E -smooth on a neighborhood of σ , then the unit-step gradient update $\sigma^{\text{new}} := \sigma - \nabla E_{\mathbf{K}}(\sigma)$ satisfies*

$$E_{\mathbf{K}}(\sigma^{\text{new}}) \leq E_{\mathbf{K}}(\sigma) - \left(1 - \frac{L_E}{2}\right) \|\nabla E_{\mathbf{K}}(\sigma)\|^2.$$

Proof. The standard L_E -smooth quadratic upper bound [Nesterov, 2018, Lemma 1.2.3] states $E(\sigma + v) \leq E(\sigma) + \nabla E(\sigma)^\top v + (L_E/2) \|v\|^2$ for any v in the smoothness neighborhood. Applying this to $v = -\nabla E(\sigma)$ gives $E(\sigma - \nabla E) \leq E(\sigma) - \|\nabla E\|^2 + (L_E/2) \|\nabla E\|^2 = E(\sigma) - (1 - L_E/2) \|\nabla E\|^2$. $\square$

Theorem 4 (Deterministic descent and sublevel-set invariance). *Assume (A1) $\mathbf{K} = \mathbf{V}$, (A2) the attention component is weight-tied across blocks, and (A3) $E_{\mathbf{K}}$ is L_E -smooth on a neighborhood of the unperturbed trajectory with $L_E < 2$. Then the unit-step iterate $\sigma^{(l+1)} = \sigma^{(l)} - \nabla E_{\mathbf{K}}(\sigma^{(l)})$ is strictly monotone-descending in $E_{\mathbf{K}}$ at every non-stationary point, and the sublevel set $\{\sigma : E_{\mathbf{K}}(\sigma) \leq E_{\mathbf{K}}(\sigma^{(0)})\}$ is forward-invariant.*

Proof. By Lemma 4, $E_{\mathbf{K}}(\sigma^{(l+1)}) \leq E_{\mathbf{K}}(\sigma^{(l)}) - (1 - L_E/2) \|\nabla E_{\mathbf{K}}(\sigma^{(l)})\|^2$. Under (A3), the coefficient $1 - L_E/2 > 0$, so $E_{\mathbf{K}}(\sigma^{(l+1)}) \leq E_{\mathbf{K}}(\sigma^{(l)})$ with strict inequality unless $\nabla E_{\mathbf{K}}(\sigma^{(l)}) = 0$, which proves forward invariance. $\square$

Corollary 3 (Connected-component containment for the unit-step iterate). *Let $S := \{\sigma : E_{\mathbf{K}}(\sigma) \leq E_{\mathbf{K}}(\sigma^{(0)})\}$ and let $C(\sigma^{(0)})$ denote the connected component of S containing $\sigma^{(0)}$. Define*

$$\text{sep}(C(\sigma^{(0)})) := \inf\{\|x - y\| : x \in C(\sigma^{(0)}), y \in S \setminus C(\sigma^{(0)})\},$$

with the convention $\text{sep} = +\infty$ if S is connected. Assume $\text{sep}(C(\sigma^{(0)})) > 0$. Then the unit-step iterate $\sigma^{(l+1)} = \sigma^{(l)} - \nabla E_{\mathbf{K}}(\sigma^{(l)})$ remains in $C(\sigma^{(0)})$ for all $l \geq 0$ whenever

$$\|\nabla E_{\mathbf{K}}(\sigma^{(l)})\| < \text{sep}(C(\sigma^{(0)})) \quad \text{for every } l.$$

Proof. By Theorem 4, every iterate lies in S . Suppose by induction $\sigma^{(l)} \in C(\sigma^{(0)})$. The displacement $\sigma^{(l+1)} - \sigma^{(l)} = -\nabla E_{\mathbf{K}}(\sigma^{(l)})$ has Euclidean length $\|\nabla E_{\mathbf{K}}(\sigma^{(l)})\|$. A jump to a distinct component $C' \subseteq S$ would require this length to be at least $\inf\{\|x - y\| : x \in C(\sigma^{(0)}), y \in C'\} \geq \text{sep}(C(\sigma^{(0)}))$, which the hypothesis excludes. Hence $\sigma^{(l+1)} \in C(\sigma^{(0)})$. $\square$

Positivity of sep. $\text{sep}(C) > 0$ holds whenever the components of S are spatially separated by a region of strictly higher energy than $E_{\mathbf{K}}(\sigma^{(0)})$, which we treat as a problem-dependent assumption rather than a structural property: it requires that the relevant components of the Hopfield energy do not share closure boundary points at level $E_{\mathbf{K}}(\sigma^{(0)})$. If $\text{sep} = 0$ the corollary becomes vacuous; in that case the trajectory is still in S by Theorem 4, but is not formally trapped in $C(\sigma^{(0)})$.

Caveats and scope. Theorem 4 depends on three assumptions that are individually mild but jointly idealized for real transformers. (A1) $\mathbf{K} = \mathbf{V}$ collapses the key and value projections; this is the standard Modern Hopfield assumption [Ramsauer et al., 2021] and isolates the energy-iteration aspect of attention from the additional value-projection degree of freedom. (A2) weight-tying treats the attention component as a single repeated map; without it, each block has its own energy and the trajectory does not minimize a single fixed function across blocks. (A3) is a smoothness condition on $E_{\mathbf{K}}$. For the Hopfield energy used here, $\nabla^2 E_{\mathbf{K}}(\sigma) = \mathbf{I} - \beta \mathbf{K}^\top (\text{diag}(p_\sigma) - p_\sigma p_\sigma^\top) \mathbf{K}$ where $p_\sigma = \text{softmax}(\beta \mathbf{K} \sigma)$, so the operator-norm bound on $\nabla^2 E_{\mathbf{K}}$ depends on the inverse-temperature β and on $\|\mathbf{K}\|$. (A3) holds whenever $\beta \|\mathbf{K}\|^2$ is sufficiently small, which is the standard regime in which the Hopfield iterate converges in a small number of steps. The theorem therefore characterizes an idealized tendency of attention-only iteration in this regime; it is not a literal property of every transformer block. In particular, sublevel-set confinement follows from Theorem 4, whereas connected-component confinement additionally requires the separation and per-step gradient assumptions of Corollary 3. Its role in the main text is to supply an idealized source of exploration deficit that an injected stochastic perturbation addresses; the empirical link is the cross-layer attention decay observed in Figure 4 and the trajectory diversity gain reported in Section 3.4.

828 E.2 Exact Attention Decomposition and Magnitude Control

Proposition 6 (Attention decomposition). *Let $\alpha_{ij}^{(l)}$ be the attention weights for query position i at layer l , with n real tokens and k random tokens. Define*

$$w_{r,i}^{(l)} := \sum_{j=1}^k \alpha_{i,n+j}^{(l)}$$

831 *and, for $j \in [n]$,*

$$\tilde{\alpha}_{ij}^{(l)} := \frac{\alpha_{ij}^{(l)}}{1 - w_{r,i}^{(l)}}.$$

832 *Then*

$$\mathbf{x}_i^{(l+1)} = (1 - w_{r,i}^{(l)}) \tilde{\mathbf{x}}_i^{(l+1)} + \boldsymbol{\eta}_i^{(l)},$$

833 *where*

$$\tilde{\mathbf{x}}_i^{(l+1)} := \sum_{j=1}^n \tilde{\alpha}_{ij}^{(l)} \mathbf{v}_j^{(l)} \quad \text{and} \quad \boldsymbol{\eta}_i^{(l)} := \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)}.$$

834 *Proof.* The attention output is

$$\mathbf{x}_i^{(l+1)} = \sum_{j=1}^n \alpha_{ij}^{(l)} \mathbf{v}_j^{(l)} + \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)}.$$

835 By definition, $\sum_{j=1}^n \alpha_{ij}^{(l)} = 1 - w_{r,i}^{(l)}$, so

$$\sum_{j=1}^n \alpha_{ij}^{(l)} \mathbf{v}_j^{(l)} = (1 - w_{r,i}^{(l)}) \sum_{j=1}^n \frac{\alpha_{ij}^{(l)}}{1 - w_{r,i}^{(l)}} \mathbf{v}_j^{(l)} = (1 - w_{r,i}^{(l)}) \tilde{\mathbf{x}}_i^{(l+1)}.$$

836 Substituting this identity gives the decomposition. □

837 **Corollary 4** (Magnitude scales with random-token attention). *For every realization of the random*
838 *values,*

$$\|\boldsymbol{\eta}_i^{(l)}\| \leq w_{r,i}^{(l)} \max_{j \in [k]} \|\mathbf{v}_{r_j}^{(l)}\|.$$

839 *Proof.* By the triangle inequality,

$$\|\boldsymbol{\eta}_i^{(l)}\| = \left\| \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)} \right\| \leq \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \|\mathbf{v}_{r_j}^{(l)}\| \leq w_{r,i}^{(l)} \max_{j \in [k]} \|\mathbf{v}_{r_j}^{(l)}\|.$$

840 □

Corollary 4 is the only quantitative fact needed in the main text. No independence assumption between attention weights and random keys is required. Once $w_{r,i}^{(l)}$ is small, the random contribution must be small as well.

E.3 Attention Decay Bound

Theorem 5 (Attention decay). *The total attention on random tokens satisfies*

$$w_{r,i}^{(l)} \leq \frac{k}{k + n \cdot \exp(\Delta_i^{(l)} / \sqrt{d})}, \quad (7)$$

where $\Delta_i^{(l)} := \min_{j \in [n]} \mathbf{q}_i^\top \mathbf{k}_j - \max_{j' \in [k]} \mathbf{q}_i^\top \mathbf{k}_{r,j'}$ is the real-vs-random logit gap.

Proof. Let $s_j = \mathbf{q}_i^\top \mathbf{k}_j / \sqrt{d}$ for real keys and $s_{r,j'} = \mathbf{q}_i^\top \mathbf{k}_{r,j'} / \sqrt{d}$ for random keys. By definition of $\Delta_i^{(l)}$, each real key satisfies $s_j \geq \Delta_i^{(l)} / \sqrt{d} + s_{r,\max}$ where $s_{r,\max} = \max_{j'} s_{r,j'}$. Therefore $\exp(s_j) \geq \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot \exp(s_{r,\max})$ for every real key. Summing over n real keys gives

$$\sum_{j=1}^n \exp(s_j) \geq n \cdot \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot \exp(s_{r,\max}) \geq n \cdot \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot \frac{1}{k} \sum_{j'=1}^k \exp(s_{r,j'}).$$

Let $R := \sum_{j'=1}^k \exp(s_{r,j'})$. Then

$$w_{r,i}^{(l)} = \frac{R}{\sum_{j=1}^n \exp(s_j) + R} \leq \frac{R}{n \cdot \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot R/k + R} = \frac{k}{k + n \cdot \exp(\Delta_i^{(l)} / \sqrt{d})}.$$

□

Corollary 5 (Sequence-wise annealing under a fixed logit gap). *For fixed k and $\Delta_i^{(l)}$, the upper bound*

$$f(n) = \frac{k}{k + n \exp(\Delta_i^{(l)} / \sqrt{d})}$$

is strictly decreasing in n . Consequently, cache growth alone decreases the maximum perturbation mass compatible with that gap.

Proof. Differentiate $f(n)$ with respect to n :

$$f'(n) = -\frac{k \exp(\Delta_i^{(l)} / \sqrt{d})}{\left(k + n \exp(\Delta_i^{(l)} / \sqrt{d})\right)^2} < 0.$$

Thus $f(n)$ is strictly decreasing. □

Corollary 6 (Vanishing random contribution under annealing). *If $n \rightarrow \infty$ along decoding while k and $\Delta_i^{(l)}$ remain fixed, then $w_{r,i}^{(l)} \rightarrow 0$. Moreover, for any fixed realization of the k random values,*

$$\|\boldsymbol{\eta}_i^{(l)}\| \leq w_{r,i}^{(l)} \max_{j \in [k]} \|\mathbf{v}_{r,j}^{(l)}\| \rightarrow 0.$$

Since the implementation samples finite-dimensional vectors and k is finite, the displayed maximum is almost surely finite.

Proof. Corollary 5 implies $w_{r,i}^{(l)} \rightarrow 0$. The bound on $\|\boldsymbol{\eta}_i^{(l)}\|$ is Corollary 4. Since the multiplier $\max_{j \in [k]} \|\mathbf{v}_{r,j}^{(l)}\|$ is finite for any fixed realization of finitely many sampled values, the product converges to zero. □

865 E.4 Feed-Forward Propagation under Local Lipschitz Control

866 A transformer block applies an attention sublayer followed by a feed-forward sublayer with residual
867 connections:

$$\mathbf{x}_i^{(l,\text{attn})} = \mathbf{x}_i^{(l)} + \text{Attn}^{(l)}(\mathbf{x}_i^{(l)}), \quad \mathbf{x}_i^{(l+1)} = \mathbf{x}_i^{(l,\text{attn})} + F^{(l)}(\mathbf{x}_i^{(l,\text{attn})}).$$

868 Let $\mathbf{x}_i^{(l,\text{attn},\text{RSP})}$ and $\mathbf{x}_i^{(l,\text{attn},\text{base})}$ denote the post-attention states under RSP and the unperturbed forward
869 pass at the same input $\mathbf{x}_i^{(l)}$.

870 **Lemma 5** (Per-block FFN propagation). *The post-attention difference satisfies*

$$\|\mathbf{x}_i^{(l,\text{attn},\text{RSP})} - \mathbf{x}_i^{(l,\text{attn},\text{base})}\| \leq \|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|.$$

871 If, in addition, $F^{(l)}$ is $L^{(l)}$ -Lipschitz on a bounded set containing both $\mathbf{x}_i^{(l,\text{attn},\text{RSP})}$ and $\mathbf{x}_i^{(l,\text{attn},\text{base})}$,
872 then

$$\|\mathbf{x}_i^{(l+1,\text{RSP})} - \mathbf{x}_i^{(l+1,\text{base})}\| \leq (1 + L^{(l)})(\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|).$$

873 *Proof.* The unperturbed attention sublayer produces $\mathbf{x}_i^{(l,\text{attn},\text{base})} = \mathbf{x}_i^{(l)} + \tilde{\mathbf{x}}_i^{(l+1)}$, while the RSP
874 attention sublayer produces $\mathbf{x}_i^{(l,\text{attn},\text{RSP})} = \mathbf{x}_i^{(l)} + (1 - w_{r,i}^{(l)})\tilde{\mathbf{x}}_i^{(l+1)} + \boldsymbol{\eta}_i^{(l)}$ by Proposition 6. Subtracting,

$$\mathbf{x}_i^{(l,\text{attn},\text{RSP})} - \mathbf{x}_i^{(l,\text{attn},\text{base})} = \boldsymbol{\eta}_i^{(l)} - w_{r,i}^{(l)} \tilde{\mathbf{x}}_i^{(l+1)},$$

875 and the triangle inequality gives the first bound. For the second bound, write the post-FFN difference
876 as

$$[\mathbf{x}_i^{(l,\text{attn},\text{RSP})} + F^{(l)}(\mathbf{x}_i^{(l,\text{attn},\text{RSP})})] - [\mathbf{x}_i^{(l,\text{attn},\text{base})} + F^{(l)}(\mathbf{x}_i^{(l,\text{attn},\text{base})})].$$

877 By Lipschitz continuity, $\|F^{(l)}(\mathbf{x}_i^{(l,\text{attn},\text{RSP})}) - F^{(l)}(\mathbf{x}_i^{(l,\text{attn},\text{base})})\| \leq L^{(l)} \|\mathbf{x}_i^{(l,\text{attn},\text{RSP})} - \mathbf{x}_i^{(l,\text{attn},\text{base})}\|$.
878 The triangle inequality combined with the first bound gives the second. $\square$

879 **Corollary 7** (Cross-block propagation control). Define $\Delta^{(l)} := \mathbf{x}_i^{(l,\text{RSP})} - \mathbf{x}_i^{(l,\text{base})}$ with $\Delta^{(0)} = \mathbf{0}$
880 (matched inputs at layer 0), and let $L_{\text{block}}^{(l)}$ denote the local Lipschitz constant of the block map
881 $B_{\text{base}}^{(l)} : \mathbf{x} \mapsto \mathbf{x} + \text{Attn}^{(l)}(\mathbf{x}) + F^{(l)}(\mathbf{x} + \text{Attn}^{(l)}(\mathbf{x}))$ at layer l . Then

$$\|\Delta^{(l+1)}\| \leq (1 + L_{\text{block}}^{(l)}) \|\Delta^{(l)}\| + (1 + L^{(l)})(\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|).$$

882 Iterating across L blocks with $\Delta^{(0)} = \mathbf{0}$ yields

$$\|\Delta^{(L+1)}\| \leq \sum_{l=0}^L \prod_{l'=l+1}^L (1 + L_{\text{block}}^{(l')}) \cdot (1 + L^{(l)})(\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|).$$

883 *Proof.* Decompose the cross-block discrepancy at layer $l + 1$ into two contributions:

$$\Delta^{(l+1)} = \underbrace{B_{\text{RSP}}^{(l)}(\mathbf{x}^{(l,\text{RSP})}) - B_{\text{RSP}}^{(l)}(\mathbf{x}^{(l,\text{base})})}_{\text{(I) propagation of prior } \Delta^{(l)}} + \underbrace{B_{\text{RSP}}^{(l)}(\mathbf{x}^{(l,\text{base})}) - B_{\text{base}}^{(l)}(\mathbf{x}^{(l,\text{base})})}_{\text{(II) RSP-induced perturbation injected at block } l},$$

884 where $B_{\text{RSP}}^{(l)}$ replaces $\text{Attn}^{(l)}$ with the RSP-augmented attention. Term (I) is bounded via the block
885 Lipschitz constant: $\|(I)\| \leq L_{\text{block}}^{(l)} \|\Delta^{(l)}\|$. With residuals included one gets the factor $(1 + L_{\text{block}}^{(l)})$
886 when also propagating the residual term $\Delta^{(l)}$ itself. Term (II) is exactly the per-block injection from
887 Lemma 5, bounded by $(1 + L^{(l)})(\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|)$. Combining via the triangle inequality
888 gives the recursion. The summed bound follows by unrolling the recursion: with $\Delta^{(0)} = \mathbf{0}$, induction
889 on L shows that $\|\Delta^{(L+1)}\|$ is at most the displayed sum. $\square$

Boundedness of $L^{(l)}$ for standard architectures. Modern transformer blocks combine LayerNorm, GELU or SwiGLU activations, and linear projections. LayerNorm has Lipschitz constant on the order of $\|\gamma\|/\inf\|x\|$ on every set bounded away from the origin, where γ is the LayerNorm scale parameter [Xiong et al., 2020]. GELU and SwiGLU are smooth and therefore locally Lipschitz on bounded activation regions, and linear projections have Lipschitz constant equal to the spectral norm of the weight matrix. The composed self-attention sublayer has a Lipschitz constant that depends on the spectral norms of $\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V$ and the inverse-temperature β as analyzed by Kim et al. [2021]. On the bounded activation regions traversed during a single forward pass, the per-block Lipschitz constants in Corollary 7 are therefore finite. For a fixed model depth this yields a finite cumulative bound on $\|\Delta^{(L+1)}\|$; stronger depth-uniform non-expansion would require additional contraction assumptions not imposed here. This is the formal sense in which the injected perturbation stays controlled across a finite forward pass, while the size of each new injection is set by $w_{r,i}^{(l)}$.

902 E.5 Proofs for Output Distribution Shift

903 The main text makes two output-level claims. Proposition 1 establishes pairwise rank randomness, 904 and Proposition 2 upgrades that statement to top- K entry. Both are statements about a local first-order 905 surrogate of the perturbed logits.

906 E.5.1 Autoregressive Formulation

907 In autoregressive decoding, the joint distribution over a sequence $y_{1:T}$ factorizes as

$$\Pr(y_{1:T} \mid x) = \prod_{t=1}^T p(y_t \mid x, y_{<t}),$$

908 where each step is governed by a prefix-conditioned distribution. Any perturbation introduced by 909 RSP affects generation through a sequence of local changes to $p(y_t \mid x, y_{<t})$, rather than a single 910 global distribution.

911 E.5.2 Local Linearization of Logits under RSP

912 Let $\bar{h}$ denote the centered random soft prompt. We model the perturbed logits as a function $z_i(\bar{h})$, 913 assuming local smoothness in a neighborhood of $\bar{h} = 0$:

$$z_i(\bar{h}) = z_i(0) + \nabla_{\bar{h}} z_i(0)^\top \bar{h} + r_i(\bar{h}),$$

914 where $r_i(\bar{h}) = O(\|\bar{h}\|^2)$. Defining $a_i := \nabla_{\bar{h}} z_i(0)$, we obtain the local affine surrogate

$$z_i^{\text{lin}}(\bar{h}) := z_i + a_i^\top \bar{h}.$$

915 This approximation is valid for sufficiently small $\|\bar{h}\|$ and is used as a first-order surrogate for the 916 direction of perturbation rather than as an exact global model of the transformer logits.

917 For the two propositions below we isolate the only support property that is actually used. Let D 918 denote the centered prompt law. We assume

$$(S1) \quad D(U) > 0 \quad \text{for every nonempty open set } U \subseteq \mathbb{R}^d.$$

919 The Gaussian instantiation in Section 4.1 satisfies (S1), but the proofs below require only this support 920 condition.

921 E.5.3 Proof of Proposition 1 (Stochastic Ranking)

922 *If there exists a token pair (i, j) such that $a_i \neq a_j$, and if the centered prompt law satisfies (S1), then* 923 $0 < \Pr_{\bar{h}}(z_i^{\text{lin}}(\bar{h}) > z_j^{\text{lin}}(\bar{h})) < 1$.

924 *Proof.* Under the local affine surrogate,

$$z_i^{\text{lin}}(\bar{h}) - z_j^{\text{lin}}(\bar{h}) = (z_i - z_j) + (a_i - a_j)^\top \bar{h}.$$

925 Let $b_{ij} := a_i - a_j \neq 0$. The perturbed logit difference

$$\Delta_{ij}(\bar{h}) := (z_i - z_j) + b_{ij}^\top \bar{h}$$

926 vanishes on the affine hyperplane

$$H_0 := \{\bar{h} \in \mathbb{R}^d : (z_i - z_j) + b_{ij}^\top \bar{h} = 0\}.$$

927 Because $b_{ij} \neq 0$, H_0 has codimension one and partitions $\mathbb{R}^d$ into the two nonempty open half-spaces

$$H_+ := \{\bar{h} : (z_i - z_j) + b_{ij}^\top \bar{h} > 0\}, \quad H_- := \{\bar{h} : (z_i - z_j) + b_{ij}^\top \bar{h} < 0\}.$$

928 By (S1), both H_+ and H_- have positive probability, which yields

$$0 < \Pr_{\bar{h}}(\Delta_{ij}(\bar{h}) > 0) < 1.$$

929

□

930 **E.5.4 Proof of Proposition 2 (Top-K Entry)**

931 Let $J_K(s_t)$ be the baseline top- K set at state s_t , and let $i \notin J_K(s_t)$. Define

$$\mathcal{G}_{i,K} := \{\bar{h} \in \mathbb{R}^d : \#\{j \neq i : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h})\} \leq K - 1\}.$$

932 If $\mathcal{G}_{i,K}$ contains a nonempty open set and the centered prompt law satisfies (S1), then

$$\Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0.$$

933 *Proof.* Under the local affine surrogate, token i belongs to the surrogate top- K set if and only if at
 934 most $K - 1$ tokens outrank it, i.e., if and only if $\bar{h} \in \mathcal{G}_{i,K}$. For each subset S of token indices with
 935 $i \notin S$ and $|S| \leq K - 1$, define

$$\mathcal{O}_S := \{\bar{h} : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h}) \text{ for all } j \in S, z_i^{\text{lin}}(\bar{h}) > z_j^{\text{lin}}(\bar{h}) \text{ for all } j \notin S \cup \{i\}\}.$$

936 Each $\mathcal{O}_S$ is an intersection of open half-spaces and is therefore open. Moreover,

$$\mathcal{G}_{i,K} = \bigcup_{S: i \notin S, |S| \leq K-1} \mathcal{O}_S,$$

937 so $\mathcal{G}_{i,K}$ is a finite union of open sets and is itself open. By assumption it contains a nonempty open
 938 set. Condition (S1) therefore implies

$$\Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0,$$

939 which is exactly the event that token i enters the surrogate top- K set. □

940 **E.5.5 Proof of Corollary 2 (Sample Complexity)**

941 Let $\mu_i = \Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0$. Within N independent RSP samples $\bar{h}_1, \dots, \bar{h}_N$, the probability of at
 942 least one admission is at least $1 - (1 - \mu_i)^N \geq 1 - \exp(-N\mu_i)$.

943 *Proof.* The samples $\{\bar{h}_n\}_{n=1}^N$ are independent, so

$$\Pr(\text{no admission within } N \text{ samples}) = \prod_{n=1}^N \Pr(\bar{h}_n \notin \mathcal{G}_{i,K}) = (1 - \mu_i)^N.$$

944 The complement gives the first inequality. The second inequality is standard: $1 - \mu_i \leq \exp(-\mu_i)$,
 945 hence $(1 - \mu_i)^N \leq \exp(-N\mu_i)$. □

946 **E.5.6 Trajectory-Level Implications**

947 Since autoregressive generation composes conditional distributions across steps, perturbations at
 948 early steps can lead to distinct trajectories. Let $\mathcal{R}_{k,K}(x)$ denote the set of trajectories containing
 949 at least one token outside the baseline top- K set. If RSP increases the accessibility of such tokens
 950 at early steps, then the probability of sampling trajectories in $\mathcal{R}_{k,K}(x)$ increases under repeated
 951 sampling. The key point is temporal compounding: one early top- K entry changes the next decoding
 952 state, which changes the next candidate set, and so on. This is the formal bridge from local rank
 953 perturbations to global trajectory diversity.

F Attention Mass Analysis

This appendix formalizes the per-token attention mass reported in Figure 4 and the exclusion criteria applied to the reasoning span and the layer axis. For each of the 500 MATH-500 problems, a fresh RSP $\mathbf{H} \in \mathbb{R}^{L \times d}$ with $L = 10$ is sampled independently, so the reported heatmaps average over both problem variability and RSP-vector variability.

F.1 Per-Token Attention Mass

For each problem, we measure attention on the concatenated sequence [prompt; $\mathbf{H}$; generation], where the generation is the trajectory produced by the same model under matching RSP injection. Let $\alpha_{i,j}^{(\ell,h)}$ denote the post-softmax attention weight at layer ℓ , head h , from query position i to key position j , satisfying $\sum_j \alpha_{i,j}^{(\ell,h)} = 1$ under the causal mask. Let Q and R denote the index sets of question and RSP tokens with sizes $|Q|$ and $|R| = L$, and let H be the number of heads. The per-token attention mass that query i directs to each group at layer ℓ is

$$\text{PTM}_Q[\ell, i] = \frac{1}{|Q|H} \sum_{h=1}^H \sum_{j \in Q} \alpha_{i,j}^{(\ell,h)}, \quad \text{PTM}_R[\ell, i] = \frac{1}{|R|H} \sum_{h=1}^H \sum_{j \in R} \alpha_{i,j}^{(\ell,h)}. \quad (8)$$

Normalization by $|Q|$ and $|R|$ controls for group size, so both quantities express how much attention a single question (resp. RSP) token receives on average under the same softmax denominator. Question tokens thereby serve as a size-matched baseline that absorbs sequence-length effects common to both groups.

F.2 Heatmap Aggregation

We partition the layer indices and the reasoning position indices each into five equal-size bins $\{B_L^{(b)}\}_{b=1}^5$ and $\{B_P^{(b)}\}_{b=1}^5$. For sample s and target group $\bullet \in \{Q, R\}$, the value in cell (b_L, b_P) is

$$M_\bullet^{(s)}[b_L, b_P] = \frac{1}{|B_L^{(b_L)}| \cdot |B_P^{(b_P)}|} \sum_{\ell \in B_L^{(b_L)}} \sum_{i \in B_P^{(b_P)}} \text{PTM}_\bullet[\ell, i]. \quad (9)$$

Figure 4 reports the sample average of Eq. (9) over the 500 valid samples.

F.3 Excluding the `\boxed{\dots}` Span

The reasoning position range terminates strictly before the final `\boxed{\dots}` span. Chain-of-thought outputs decompose into a *text* (content) component and a *pattern* (template) component that contribute independently to downstream performance [Madaan and Yazdanbakhsh, 2022]. The `\boxed{\dots}` wrapper is a canonical pattern: a short, stereotyped span whose role is to complete the answer format rather than to extend reasoning. Including it would therefore inject a template-driven attention signature unrelated to reasoning dynamics.

F.4 Excluding the Final Two Layers

The layer axis further excludes the last two transformer layers. Two well-documented properties of late layers, none specific to our setup, motivate this choice.

Output-projection specialization. Late-layer hidden states lie in an approximately linear relationship with vocabulary logits and therefore function as a progressive linear readout rather than as representation-transforming blocks [Belrose et al., 2023]. Consistently, late-layer feed-forward modules have been characterized as output-distribution shaping mechanisms [Geva et al., 2021]. Attention in these layers therefore reflects output formation rather than reasoning computation.

Register-token hypothesis. Uninformative patch tokens in Vision Transformers have been shown to be repurposed as *registers*: storage slots that absorb attention in deeper layers [Darcet et al., 2024]. Because RSPs share a similar profile as continuous tokens without pretrained semantic content, a

similar repurposing in late transformer layers cannot be ruled out. Excluding the final two layers removes this potential register-induced confound from the reasoning-phase signal.

The two lines of evidence above identify known late-layer phenomena — linear readout and potential register repurposing — whose attention signatures are orthogonal to reasoning-phase dynamics. Removing the final two layers preempts their contamination of the reasoning-phase signal, yielding a principled, literature-grounded exclusion rather than a model-specific tuning decision.

F.5 Additional Suffix Heatmaps

Figure 7 reports the same per-token RSP attention analysis under suffix injection for Qwen2.5-Math-1.5B-Instruct and LLaMA-3.1-8B-Instruct. For the late-layer reasons discussed above, the pattern does not generalize uniformly across every cell; nevertheless, the overall trend is consistent across both models: question-token attention varies broadly across layers, whereas RSP attention shows the layer-wise and sequence-wise decay predicted by Eq. (3).




Figure 7: Per-token RSP attention mass under suffix injection for the remaining two models (500 MATH-500 samples each). Axes and preprocessing match Figure 4: x-axis is reasoning position binned into five quantiles, y-axis is layer depth binned into five quantiles (top = shallow), color encodes the 500-sample mean per-token attention assigned to RSP tokens, and tokens inside `\boxed{}` spans and the last two layers are excluded.

1004 **G GRPO Formulation**

1005 For a given prompt q , the behavior policy $\pi_{\theta_{\text{old}}}$ samples a group of G responses $\{o_i\}_{i=1}^G$ with
 1006 corresponding rewards $\{R_i\}_{i=1}^G$. The advantage of each response is computed by normalizing the
 1007 group-level rewards:

$$\hat{A}_{i,t} = \frac{R_i - \text{mean}(\{R_i\}_{i=1}^G)}{\text{std}(\{R_i\}_{i=1}^G)}. \quad (10)$$

1008 The policy is updated via a clipped objective with a KL penalty:

$$\mathcal{J}_{\text{GRPO}}(\theta) = \mathbb{E} \left[\frac{1}{G} \sum_{i=1}^G \frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \left(\min \left(r_{i,t} \hat{A}_{i,t}, \text{clip}(r_{i,t}, 1-\varepsilon, 1+\varepsilon) \hat{A}_{i,t} \right) - \beta D_{\text{KL}}(\pi_{\theta} \parallel \pi_{\text{ref}}) \right) \right], \quad (11)$$

1009 where $r_{i,t} = \pi_{\theta}(o_{i,t} \mid q, o_{i,<t}) / \pi_{\theta_{\text{old}}}(o_{i,t} \mid q, o_{i,<t})$ is the importance sampling ratio, ε is the clipping
 1010 range, and β controls the strength of the KL penalty against the reference policy π_{ref} .

1011 H Broader Impacts

1012 This work provides empirical and theoretical analysis of how random embedding injection affects
1013 the reasoning behavior of large language models. The contributions are primarily foundational,
1014 with potential downstream implications for reinforcement learning-based post-training of reasoning
1015 models such as GRPO.

1016 **Positive impacts.** RSP-induced trajectory diversity could improve the sample efficiency of RL-
1017 based LLM post-training by providing richer learning signals within group-relative advantage estima-
1018 tion. This has the potential to reduce compute costs associated with alignment and reasoning-oriented
1019 fine-tuning, making such methods more accessible to research groups with limited resources.

1020 **Potential concerns.** Methods that expand the effective output support of LLMs may increase the
1021 reachability of low-probability tokens, which includes both desirable alternative reasoning paths
1022 and potentially unsafe or off-distribution outputs. Practitioners applying RSP in production systems
1023 should combine it with existing safety filtering and alignment mechanisms rather than treating it as a
1024 standalone diversity source.

1025 **Limitations on impact scope.** Since this work focuses on analysis rather than deploying a new
1026 model or dataset, the direct societal impact is limited. The broader implications described above are
1027 contingent on future work integrating RSP into applied training pipelines.

1028 NeurIPS Paper Checklist

1029 The checklist is designed to encourage best practices for responsible machine learning research,
1030 addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove
1031 the checklist: **The papers not including the checklist will be desk rejected.** The checklist should
1032 follow the references and follow the (optional) supplemental material. The checklist does NOT count
1033 towards the page limit.

1034 Please read the checklist guidelines carefully for information on how to answer these questions. For
1035 each question in the checklist:

- 1036 • You should answer [Yes], [No], or [N/A].
- 1037 • [N/A] means either that the question is Not Applicable for that particular paper or the
1038 relevant information is Not Available.
- 1039 • Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

1040 **The checklist answers are an integral part of your paper submission.** They are visible to the
1041 reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it
1042 (after eventual revisions) with the final version of your paper, and its final version will be published
1043 with the paper.

1044 The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation.
1045 While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a
1046 proper justification is given (e.g., error bars are not reported because it would be too computationally
1047 expensive” or “we were unable to find the license for the dataset we used”). In general, answering
1048 [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we
1049 acknowledge that the true answer is often more nuanced, so please just use your best judgment and
1050 write a justification to elaborate. All supporting evidence can appear either in the main paper or the
1051 supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification
1052 please point to the section(s) where related material for the question can be found.

1053 IMPORTANT, please:

- 1054 • **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- 1055 • **Keep the checklist subsection headings, questions/answers and guidelines below.**
- 1056 • **Do not modify the questions and only use the provided macros for your answers.**

1057 1. Claims

1058 Question: Do the main claims made in the abstract and introduction accurately reflect the
1059 paper’s contributions and scope?

1060 Answer: **[TODO]**

1061 Justification: **[TODO]**

1062 Guidelines:

- 1063 • The answer [N/A] means that the abstract and introduction do not include the claims
1064 made in the paper.
- 1065 • The abstract and/or introduction should clearly state the claims made, including the
1066 contributions made in the paper and important assumptions and limitations. A [No] or
1067 [N/A] answer to this question will not be perceived well by the reviewers.
- 1068 • The claims made should match theoretical and experimental results, and reflect how
1069 much the results can be expected to generalize to other settings.
- 1070 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
1071 are not attained by the paper.

1072 2. Limitations

1073 Question: Does the paper discuss the limitations of the work performed by the authors?

1074 Answer: **[TODO]**

1075 Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The authors should answer **[Yes]** if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- 1232 • The paper should disclose whether the full research project required more compute
1233 than the experiments reported in the paper (e.g., preliminary or failed experiments that
1234 didn't make it into the paper).

1235 **9. Code of ethics**

1236 Question: Does the research conducted in the paper conform, in every respect, with the
1237 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

1238 Answer: **[TODO]**

1239 Justification: **[TODO]**

1240 Guidelines:

- 1241 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of
1242 Ethics.
1243 • If the authors answer [No], they should explain the special circumstances that require a
1244 deviation from the Code of Ethics.
1245 • The authors should make sure to preserve anonymity (e.g., if there is a special consid-
1246 eration due to laws or regulations in their jurisdiction).

1247 **10. Broader impacts**

1248 Question: Does the paper discuss both potential positive societal impacts and negative
1249 societal impacts of the work performed?

1250 Answer: **[TODO]**

1251 Justification: **[TODO]**

1252 Guidelines:

- 1253 • The answer [N/A] means that there is no societal impact of the work performed.
1254 • If the authors answer [N/A] or [No], they should explain why their work has no societal
1255 impact or why the paper does not address societal impact.
1256 • Examples of negative societal impacts include potential malicious or unintended uses
1257 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations
1258 (e.g., deployment of technologies that could make decisions that unfairly impact specific
1259 groups), privacy considerations, and security considerations.
1260 • The conference expects that many papers will be foundational research and not tied
1261 to particular applications, let alone deployments. However, if there is a direct path to
1262 any negative applications, the authors should point it out. For example, it is legitimate
1263 to point out that an improvement in the quality of generative models could be used to
1264 generate Deepfakes for disinformation. On the other hand, it is not needed to point out
1265 that a generic algorithm for optimizing neural networks could enable people to train
1266 models that generate Deepfakes faster.
1267 • The authors should consider possible harms that could arise when the technology is
1268 being used as intended and functioning correctly, harms that could arise when the
1269 technology is being used as intended but gives incorrect results, and harms following
1270 from (intentional or unintentional) misuse of the technology.
1271 • If there are negative societal impacts, the authors could also discuss possible mitigation
1272 strategies (e.g., gated release of models, providing defenses in addition to attacks,
1273 mechanisms for monitoring misuse, mechanisms to monitor how a system learns from
1274 feedback over time, improving the efficiency and accessibility of ML).

1275 **11. Safeguards**

1276 Question: Does the paper describe safeguards that have been put in place for responsible
1277 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
1278 image generators, or scraped datasets)?

1279 Answer: **[TODO]**

1280 Justification: **[TODO]**

1281 Guidelines:

- 1282 • The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

1334 Guidelines:

1335 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
1336 with human subjects.

1337 • Including this information in the supplemental material is fine, but if the main contribu-
1338 tion of the paper involves human subjects, then as much detail as possible should be
1339 included in the main paper.

1340 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,
1341 or other labor should be paid at least the minimum wage in the country of the data
1342 collector.

1343 **15. Institutional review board (IRB) approvals or equivalent for research with human**
1344 **subjects**

1345 Question: Does the paper describe potential risks incurred by study participants, whether
1346 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
1347 approvals (or an equivalent approval/review based on the requirements of your country or
1348 institution) were obtained?

1349 Answer: **[TODO]**

1350 Justification: **[TODO]**

1351 Guidelines:

1352 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
1353 with human subjects.

1354 • Depending on the country in which research is conducted, IRB approval (or equivalent)
1355 may be required for any human subjects research. If you obtained IRB approval, you
1356 should clearly state this in the paper.

1357 • We recognize that the procedures for this may vary significantly between institutions
1358 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
1359 guidelines for their institution.

1360 • For initial submissions, do not include any information that would break anonymity (if
1361 applicable), such as the institution conducting the review.

1362 **16. Declaration of LLM usage**

1363 Question: Does the paper describe the usage of LLMs if it is an important, original, or
1364 non-standard component of the core methods in this research? Note that if the LLM is used
1365 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
1366 scientific rigor, or originality of the research, declaration is not required.

1367 Answer: **[TODO]**

1368 Justification: **[TODO]**

1369 Guidelines:

1370 • The answer [N/A] means that the core method development in this research does not
1371 involve LLMs as any important, original, or non-standard components.

1372 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
1373 be described.