
From Noise to Diversity: Random Embedding Injection in LLM Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

Recent *soft prompt* methods, which prepend or insert trainable continuous embeddings into LLM inputs as virtual tokens, optimize these embeddings to enhance mathematical reasoning. However, the mechanism by which embedding injection itself affects model behavior, independent of optimization, remains unexplored. We investigate Random Soft Prompts (RSPs), continuous vectors sampled from an isotropic Gaussian fitted to the entrywise mean and variance of the pretrained embedding table and injected without any training. Theoretically, we show that the RSP contribution at each attention layer is exactly captured by a scalar random-token attention mass that is structurally diluted as autoregressive cache growth accumulates real tokens, and that under a local affine analysis of the perturbed logits RSPs can open previously excluded tokens into the effective top- K candidate set in a way temperature sampling provably cannot. Empirically, RSPs increase early-stage token entropy while generation stabilizes thereafter, and combining RSPs with temperature sampling yields higher trajectory diversity as measured by Pass@ N . We further discuss the implications for GRPO training, where trajectory diversity is critical for learning signal quality.

1 Introduction

As large language models (LLMs) scale to billions of parameters, full fine-tuning becomes increasingly expensive. This has motivated parameter-efficient methods that adapt frozen models by introducing a small number of trainable parameters [Hu et al., 2022, Houlsby et al., 2019]. Among these, *soft prompts* prepend or insert learnable continuous embeddings into the input sequence, steering model behavior through the attention mechanism [Li and Liang, 2021, Lester et al., 2021]. This paradigm has been actively adopted for enhancing mathematical reasoning in LLMs [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Hao et al., 2025, Zhang et al., 2026]. Despite their diverse designs, these methods share a common structure: they inject continuous embeddings that lie outside the pretrained vocabulary and optimize them to improve downstream performance.

However, we find that optimization is not a necessary condition for these effects. We observe that injecting randomly initialized, untrained continuous embeddings can also shift the output distribution. For instance, applying a chat template to a base model not fine-tuned with such format causes a prompt-format mismatch that degrades reasoning capability, but appending random embeddings to such mismatched inputs mitigates this degradation. If simple noise were merely perturbing the model, accuracy should worsen instead. This suggests that injecting out-of-vocabulary continuous embeddings itself, independent of any learned content, plays a non-trivial role in shaping model behavior. Existing studies primarily evaluate *soft prompt* methods through end-task accuracy, and the mechanisms by which embedding injection influences the output distribution remain insufficiently explored.

In this work, we investigate the role of Random Soft Prompts (RSPs), continuous vectors sampled from an isotropic Gaussian fitted to the entrywise mean and variance of the pretrained embedding table and injected without any training. The role of randomness is not to imitate the embedding distribution; it is to ensure that repeated rollouts receive *independent* perturbations, which is what exploration requires. Theoretically, we identify an attention-level mechanism in which the RSP contribution is exactly captured by a scalar random-token attention mass that is structurally diluted by autoregressive cache growth, and we show that under a local affine analysis of the perturbed logits RSPs can open previously excluded tokens into the effective top- K candidate set in a way temperature sampling provably cannot. Empirically, RSPs increase early-stage token entropy while the output stabilizes in later generation stages, and combining RSPs with temperature sampling yields more diverse reasoning trajectories as measured by Pass@ N scaling. We further discuss the implications of these findings for Group Relative Policy Optimization (GRPO) [Shao et al., 2024], where trajectory diversity is essential for effective learning signals.

Our contributions are as follows:

- We give a transformer-level mechanism in which the RSP contribution is controlled by a scalar attention mass with structural cache-driven dilution, and a distributional analysis under a local affine surrogate showing that *independent isotropic resampling*—not a fixed perturbation—can admit previously excluded tokens into the effective top- K candidate set in a way temperature sampling cannot.
- We empirically analyze how RSPs affect LLM generation, showing that RSPs increase early token entropy and produce more diverse reasoning trajectories when combined with temperature sampling, and that the accumulation gain disappears when one RSP is shared across samples.
- We discuss the implications of RSP-induced diversity for GRPO training and argue that RSP, being rank-changing while temperature is rank-preserving, provides exploration gains complementary to existing temperature-based approaches.

2 Background

2.1 Soft Prompt Methods for LLM Reasoning

Soft prompts emerged as a parameter-efficient alternative to full fine-tuning. Prefix-Tuning [Li and Liang, 2021] introduced the paradigm by prepending trainable continuous vectors to every transformer layer, Prompt Tuning [Lester et al., 2021] simplified this to input-layer-only tuning and showed that it matches full fine-tuning at sufficient scale, and P-tuning [Liu et al., 2024] extended the approach with an LSTM-based prompt encoder to model inter-position dependencies.

More recently, *soft prompts* have been actively applied to LLM reasoning. TTSV [Kang et al., 2026] optimizes prefix embeddings at test time via entropy minimization, steering the model toward high-certainty states to activate latent reasoning capabilities. SoftCoT [Xu et al., 2025] replaces explicit chain-of-thought tokens with soft tokens generated by an assistant model, leveraging the stronger representational capacity of continuous representations over discrete language. LTPO [Ye et al., 2026] optimizes latent thought embeddings per test instance through policy gradient with a confidence-based intrinsic reward. COCONUT [Hao et al., 2025] feeds the model’s own hidden states back as input embeddings, enabling reasoning in a continuous latent space. MemGen [Zhang et al., 2026] demonstrates that embedding vectors can serve as experience memory by constructing latent token sequences that enrich the model’s reasoning. Pause tokens [Goyal et al., 2024] insert learnable tokens into the input to provide additional computation steps.

These approaches all inject out-of-vocabulary continuous embeddings optimized for task-specific objectives. Because evaluation centers on end-task accuracy, the contribution of the injection itself and that of the optimization are not separated, and the mechanisms by which such out-of-vocabulary continuous embeddings shape model behavior remain a separate and underexplored question.

2.2 Noise Injection in Neural Networks

Neural networks have incorporated noise at various stages. During training, dropout [Srivastava et al., 2014] randomly deactivates hidden units to prevent overfitting, and NEFTune [Jain et al., 2024]

adds noise to token embeddings during instruction fine-tuning, reporting improvements in instruction following. These methods use noise as a regularizer and do not apply perturbations at inference time. At inference time, perturbation-based approaches primarily target robustness and safety. Randomized smoothing [Cohen et al., 2019] injects Gaussian noise into inputs to obtain certified robustness guarantees. In the LLM setting, SmoothLLM [Robey et al., 2025] perturbs input prompts at the character level to defend against adversarial jailbreaking. These methods require multiple noisy forward passes with output aggregation, and target prediction stability rather than output diversity. In reinforcement learning, NoisyNet [Fortunato et al., 2018] injects learnable noise into network weights to improve exploration, demonstrating that the location and structure of noise injection affect exploration quality. This motivation is closest to RSPs, which similarly target diversity, but differ in injection target (input embeddings vs. network weights) and require no learned noise parameters. RSPs differ from these approaches in the following respects. (1) RSPs operate solely at inference without any training, whereas dropout and NEFTune inject noise during training, modifying the resulting model parameters accordingly. (2) RSPs preserve the original input and append fresh embedding vectors in a single forward pass, whereas robustness methods corrupt the input and aggregate over multiple noisy passes. (3) RSPs use purely random vectors sampled from the pretrained embedding statistics, whereas NoisyNet learns noise parameters through optimization.

3 Random Soft Prompts and Why They Work

This section states the mechanism used in the paper. A *rollout* means one full autoregressive decode under one sampled RSP. Within a rollout, one sampled RSP is fixed, can perturb early next-token decisions through attention, and then fades as the autoregressive cache grows. Across rollouts, independent resampling can accumulate those branch changes into Pass@N gains. We keep only the main statements and intuition here; derivations are deferred to Appendix D.

3.1 Random Soft Prompt

We start with the object being analyzed. Let $\mathbf{W}_E \in \mathbb{R}^{V \times d}$ denote the pretrained token-embedding matrix, where V is the vocabulary size and d is the hidden dimension. Write its entrywise mean and standard deviation as $\mu_E := \frac{1}{Vd} \sum_{i,k} [\mathbf{W}_E]_{i,k}$ and $\sigma_E := \text{std}(\mathbf{W}_E)$. An RSP of length L is a sequence of continuous vectors $\mathbf{H} = [h_1, \dots, h_L] \in \mathbb{R}^{L \times d}$ drawn independently from

$$h_j \sim \mathcal{N}(\mu_E \mathbf{1}_d, \sigma_E^2 \mathbf{I}_d), \quad j = 1, \dots, L. \quad (1)$$

The centered form $\bar{h}_j := h_j - \mu_E \mathbf{1}_d \sim \mathcal{N}(0, \sigma_E^2 \mathbf{I}_d)$ is a zero-mean isotropic Gaussian. The prompt can be injected as a prefix $[\mathbf{H}; \mathbf{X}]$, infix, or suffix $[\mathbf{X}; \mathbf{H}]$ [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Zhang et al., 2026]. For repeated-sampling protocols, prompt draws and any decoding randomness across compared trials are assumed independent.

The role of randomness is not to imitate the empirical embedding distribution. It is to make repeated rollouts from the same text prompt receive *independent* perturbations. That independence is what exploration requires.

3.2 Exploration: why the first few tokens move

A single RSP can matter early because it enters the transformer through one scalar that can still be nontrivial before the cache fills. Consider self-attention layer ℓ with n real tokens and L RSP tokens, attention weights $\alpha_{ij}^{(\ell)}$, and value vectors $\mathbf{v}_j^{(\ell)}$ and $\mathbf{v}_{r_j}^{(\ell)}$ on the real and RSP sides. Define the RSP attention mass

$$w_{r,i}^{(\ell)} := \sum_{j'=1}^L \alpha_{i,n+j'}^{(\ell)}.$$

The attention output decomposes exactly as

$$\mathbf{x}_i^{(\ell+1)} = (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{x}}_i^{(\ell+1)} + \boldsymbol{\eta}_i^{(\ell)}, \quad \tilde{\mathbf{x}}_i^{(\ell+1)} := \sum_{j=1}^n \frac{\alpha_{ij}^{(\ell)}}{1 - w_{r,i}^{(\ell)}} \mathbf{v}_j^{(\ell)}, \quad \boldsymbol{\eta}_i^{(\ell)} := \sum_{j'=1}^L \alpha_{i,n+j'}^{(\ell)} \mathbf{v}_{r_{j'}}^{(\ell)}. \quad (2)$$

129 This is an algebraic identity, not an approximation. Because $\|\boldsymbol{\eta}_i^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_{j'} \|\mathbf{v}_{r,j'}^{(\ell)}\|$, the same
 130 scalar $w_{r,i}^{(\ell)} \in [0, 1]$ controls both how much real-token signal is attenuated and how large the random
 131 contribution can be.

132 The only remaining question is whether this scalar can stay meaningfully above zero at early steps.
 133 Let $\mathbf{q}_i^{(\ell)}$ be the query and $\mathbf{k}_j^{(\ell)}, \mathbf{k}_{r,j'}^{(\ell)}$ the real and RSP keys at layer ℓ .

134 **Proposition 1** (Lower envelope for early exploration). *Define the reverse logit gap $\Delta_i'^{(\ell)} :=$
 135 $\max_{j \in [n]} (\mathbf{q}_i^{(\ell)})^\top \mathbf{k}_j^{(\ell)} - \min_{j' \in [L]} (\mathbf{q}_i^{(\ell)})^\top \mathbf{k}_{r,j'}^{(\ell)}$. Then in scaled dot-product attention,*

$$w_{r,i}^{(\ell)} \geq \frac{L}{L + n \exp(\Delta_i'^{(\ell)} / \sqrt{d})}.$$

136 The proof is a symmetric softmax bound given in Appendix D. The point of Proposition 1 is modest
 137 but useful: if the worst random-token score is not too far below the best real-token score, then early
 138 decoding still assigns the random tokens a strictly positive amount of attention mass. That is enough
 139 for a single RSP draw to open an alternative early branch within one rollout.

140 3.3 Annealing: why the effect fades

141 The same scalar that enables early exploration is also what makes the perturbation fade. Define the
 142 forward logit gap

$$\Delta_i^{(\ell)} := \min_{j \in [n]} (\mathbf{q}_i^{(\ell)})^\top \mathbf{k}_j^{(\ell)} - \max_{j' \in [L]} (\mathbf{q}_i^{(\ell)})^\top \mathbf{k}_{r,j'}^{(\ell)}, \quad \Delta_i^{(\ell)} \leq \Delta_i'^{(\ell)}.$$

143 **Theorem 1** (Upper envelope under annealing). *In scaled dot-product attention,*

$$w_{r,i}^{(\ell)} \leq \frac{L}{L + n \exp(\Delta_i^{(\ell)} / \sqrt{d})}.$$

144 At fixed L and $\Delta_i^{(\ell)}$, the right-hand side is strictly decreasing in n and converges to zero as $n \rightarrow \infty$.

145 Combined with Proposition 1, this gives the two-sided sandwich

$$\frac{L}{L + n \exp(\Delta_i'^{(\ell)} / \sqrt{d})} \leq w_{r,i}^{(\ell)} \leq \frac{L}{L + n \exp(\Delta_i^{(\ell)} / \sqrt{d})}. \quad (3)$$

146 The theorem does *not* say that the realized mass must decrease step by step: queries, keys, and hidden
 147 states co-evolve, so both gaps change over time. What it does say is simpler and more useful: cache
 148 growth alone narrows the largest RSP attention mass compatible with a fixed gap. Since Eq. (2)
 149 already bounds the random contribution by the same scalar, the perturbation is also bounded by that
 150 fixed-gap envelope. Appendix D extends this site-local statement to finite-depth propagation under
 151 local Lipschitz control.




Figure 1: Per-token attention mass on Qwen2.5-Math-7B (suffix, 500 MATH-500 problems, each with an independently sampled RSP). (a) RSP-token attention decreases along the reasoning-position axis, consistent with Eq. (3). The additional decrease along the layer axis is a separate empirical observation suggesting that deeper layers sharpen the real-vs-random gap; it does not follow from Theorem 1 alone. (b) Question-token attention remains comparatively uniform. The reported quantity is the per-token mass of Appendix E, equal to $w_{r,i}^{(\ell)} / L$ averaged over heads and samples.

3.4 Performance gain: why independent reruns help

The early and late envelopes explain what one rollout can do. The remaining question is performance: why can *fresh* reruns help more than repeating one fixed perturbation? The answer needs one distributional fact about the prompt law and one separate task-side fact about productive branches.

The first fact is directional. The attention identities above do not require isotropy, but branch opening benefits from a perturbation law that does not under-serve whichever local direction happens to matter for the current prompt.

Theorem 2 (Direction-agnostic minimax coverage). *Let D be any zero-mean distribution on $\mathbb{R}^d$ with covariance Σ_D and trace budget $\text{tr}(\Sigma_D) \leq \rho^2$. Then*

$$\min_{\|b\|=1} \text{Var}_{h \sim D}(b^\top h) = \lambda_{\min}(\Sigma_D) \leq \rho^2/d,$$

with equality if and only if $\Sigma_D = (\rho^2/d)\mathbf{I}_d$.

Equality does not require Gaussianity; isotropic covariance is the key property. We choose a Gaussian because it adds the second property we need: positive probability on every nonempty open set. The prompt law therefore does not privilege one direction in second moment, yet it still leaves every local branch-opening region reachable (Proposition 5).

We now formalize branch opening on a first-order surrogate. Let

$$z_i^{\text{lin}}(\bar{h}) := z_i + a_i^\top \bar{h}, \quad a_i := \nabla_{\bar{h}} z_i(0),$$

be the affine approximation of the next-token logits around $\bar{h} = 0$. For a baseline-excluded token i , define the top- K entry region

$$\mathcal{G}_{i,K} := \{\bar{h} \in \mathbb{R}^d : \#\{j \neq i : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h})\} \leq K - 1\}.$$

Proposition 2 (Conditional top- K entry). *If $\mathcal{G}_{i,K}$ contains a nonempty open set, then the isotropic Gaussian prompt law assigns it positive probability, so*

$$\mu_i := \Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0.$$

The mechanism described by Proposition 2 is deliberately limited: it says that a branch can open, not that the opened branch is useful. To connect single-draw branch opening to task accuracy, we add two task-side assumptions. For a problem x , let (M1) each independent RSP run reach a correct trajectory with marginal probability at least $p_{\min}(x) > 0$, and let (M2) the N runs be mutually independent.

Proposition 3 (Multi-path Pass@N). *Under (M1)–(M2), with N independent RSP runs,*

$$\Pr(\text{Pass@N on } x) \geq 1 - (1 - p_{\min}(x))^N \geq 1 - \exp(-N p_{\min}(x)).$$

This division of labor is the main interpretive point. The mechanism is responsible only for *admitting* baseline-excluded candidates into the local candidate set. Whether an admitted branch is *productive* is task-dependent, which is why the assumption $p_{\min}(x) > 0$ is separate. It is also why sharing one RSP across all samples removes the accumulation effect: without independent resampling, the same branch is revisited rather than diversified. Section 4 tests exactly this signature.

4 Empirical Validation

We now examine how the explanation in §3 appears in actual LLM generation through three scenes. First, we ask whether RSP preserves baseline accuracy without breaking it, and recovers it in some settings (§4.2). Second, we check whether the early diversification followed by late stabilization—predicted by the two-sided envelope—is observed empirically (§4.3). Finally, we contrast Pass@N between sharing one RSP across samples versus drawing a fresh RSP per sample, directly probing whether *independence* is the key (§4.4).

Table 1: Accuracy (%) across model configurations and benchmarks. RSP reports the best result across injection positions (Prefix, Infix, Suffix). Full per-position breakdown is in Appendix A. Values in parentheses indicate the difference from the baseline.

Model	MATH-500		GSM8K		AIME24	
	Baseline	RSP	Baseline	RSP	Baseline	RSP
<i>Instruct Models</i>						
LLaMA-3.1-8B-Inst	52.20	49.40 (−2.8)	85.75	86.73 (+1.0)	6.67	10.00 (+3.3)
Qwen2.5-Math-7B-Inst	83.20	84.20 (+1.0)	95.45	95.68 (+0.2)	13.33	16.67 (+3.3)
Qwen2.5-Math-1.5B-Inst	73.00	75.20 (+2.2)	85.29	85.75 (+0.5)	10.00	20.00 (+10.0)
<i>Base Models (Raw Text)</i>						
Qwen2.5-Math-7B	72.20	71.60 (−0.6)	84.15	84.31 (+0.2)	6.67	16.67 (+10.0)
Qwen2.5-Math-1.5B	61.40	64.40 (+3.0)	77.86	74.30 (−3.6)	16.67	10.00 (−6.7)
<i>Base Models + ChatML (Format Mismatch)</i>						
Qwen2.5-Math-7B	52.20	70.40 (+18.2)	58.30	75.97 (+17.7)	23.33	23.33 (+0.0)
Qwen2.5-Math-1.5B	34.40	63.40 (+29.0)	37.45	67.02 (+29.6)	13.33	20.00 (+6.7)

Table 2: Comparison with existing soft prompt methods (TTSV, SoftCoT, LTPO) under unified evaluation. Baseline and RSP values are from Table 1. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Benchmark	Baseline	RSP	TTSV	SoftCoT	LTPO
Qwen2.5-Math-7B-Instruct	MATH-500	83.20	84.20	83.80	82.80	83.00
	GSM8K	<u>95.45</u>	95.68	<u>95.30</u>	95.00	94.84
	AIME24	13.33	<u>16.67</u>	23.30	<u>16.67</u>	13.33
Qwen2.5-Math-1.5B-Instruct	MATH-500	73.00	<u>75.20</u>	<u>75.20</u>	76.00	72.40
	GSM8K	85.29	85.75	<u>85.70</u>	84.46	84.53
	AIME24	<u>10.00</u>	20.00	6.70	6.67	<u>10.00</u>

4.1 Experimental Setup

We evaluate on three mathematical reasoning benchmarks, MATH-500 [Lightman et al., 2024], GSM8K [Cobbe et al., 2021], and AIME24, using LLaMA-3.1-8B-Instruct [Grattafiori et al., 2024] and the Qwen2.5-Math model family [Yang et al., 2024] across instruct models, base models, and base models with mismatched chat templates—seven configurations in total. All experiments use greedy decoding to isolate the effect of RSPs from sampling stochasticity, and the answer-extraction pipeline uses fallbacks to score only the final answer. RSPs follow the definition of §3.1 and are concatenated at one of three positions (prefix, infix, suffix) with a freshly drawn **H** for each batch. Full experimental details and per-position results are in Appendix A.

4.2 How Does Untrained RSP Compare to Baselines & Optimized Soft Prompts?

We first examine how injecting random embeddings affects the model’s reasoning performance. Table 1 reports the accuracy for each model configuration, using the best RSP result across injection positions (Prefix, Infix, Suffix) and $L \in \{10, 15, 20\}$ for each benchmark.

For instruct models and base models with their native prompt formats, RSP injection maintains or slightly shifts baseline performance within a few percentage points. Most strikingly, under prompt-format mismatch — where base models receive ChatML templates they were not trained on — RSP yields pronounced improvements of up to +29 pp. These gains are difficult to attribute to simple perturbation effects, as noise that merely raises model uncertainty would further degrade performance.

We compare RSPs with prior soft prompt methods (TTSV, SoftCoT, LTPO) that optimize their embeddings for distinct objectives, under a unified evaluation protocol reported in Table 2. All methods are re-evaluated with the same unified extraction pipeline (Appendix A.2), enabling a format-agnostic comparison despite each method’s distinct prompt format and grading scheme. The results place RSPs on par with the optimized methods without requiring any training.

4.3 How Does RSP Affect the Output Distribution?

We analyze how RSP injection changes the model’s output distribution on Qwen2.5-Math-1.5B-Instruct with MATH-500, measuring per-token entropy, top-1 probability, and varentropy across the generation process. Figure 2 compares these metrics for the first 5% of generation steps across RSP variants and prior soft prompt methods (LTPO, SoftCoT, TTSV). Full statistics are in Appendix A.4.




Figure 2: Mean entropy, top-1 probability, and varentropy during the first 5% of generation steps (Qwen2.5-Math-1.5B-Instruct, MATH-500). Shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: prior soft prompt methods (LTPO, SoftCoT, TTSV).

These results show that RSP modifies the early-stage output distribution across all injection positions: entropy increases (the token distribution flattens), top-1 probability decreases (commitment to a dominant token weakens), and varentropy increases (indicating multimodal distributions [Ahmed et al., 2026]). High-entropy positions are known to act as critical forking points that determine reasoning trajectories [Wang et al., 2025], so this shift moves the model toward a state where multiple reasoning paths can coexist. The change is localized to the initial reasoning stage. In the middle and final portions of generation, all three metrics converge to baseline levels.

Although existing soft prompt methods are optimized for different objectives, they share a common pattern of increasing early-stage entropy. Notably, TTSV, which optimizes soft prompts using a reward that reduces the mean entropy of reasoning trajectories, still exhibits early entropy comparable to the baseline. This suggests that elevating early-stage entropy is an inherent effect of injecting non-pretrained continuous embeddings, independent of the optimization objective.

4.4 Does RSP Induce Trajectory Diversity?

Section 4.3 empirically established that RSP reshapes the early-stage output distribution. We now examine whether this shift translates into reasoning diversity at three complementary levels of analysis: token rankings, semantic content of trajectories, and task-level outcomes. The three analyses converge on a directional picture: RSP’s first-token perturbation propagates into semantically more diverse reasoning trajectories, which in turn yields greater output diversity at the task level. All three analyses share the setting: Qwen2.5-Math-1.5B-Instruct, MATH-500, suffix injection, and $L = 10$.

Token-level: distribution beyond temperature. We test whether RSP modifies the first-token distribution in a qualitatively different way from temperature scaling, which preserves token rankings while only flattening probabilities. Under autoregressive decoding, any trajectory-level divergence between RSP and the baseline must originate from the first generation step, a critical forking point for reasoning trajectories [Wang et al., 2025]. Using the baseline at $\tau=1.0$ as reference, we compare RSP at $\tau=1.0$ against the baseline at $\tau=2.0$ (the high-temperature ceiling of what flattening alone can achieve) on three complementary metrics: Spearman rank correlation ρ to detect rank reorganization, the probability mass placed outside the baseline’s top-10 to measure support expansion, and Jensen–Shannon divergence to quantify the overall magnitude of distributional change (formal definitions in Appendix B).

Table 3: First-token distribution metrics measured against the baseline at $\tau=1.0$, comparing the baseline at $\tau=2.0$ with RSP at $\tau=1.0$.

Metric	$\tau=2.0$	RSP
Spearman ρ	1.0	0.651
Mass outside top-10	5.15%	27.64%
JS divergence	0.086	0.231

Table 3 reports the contrast. Even at $\tau=2.0$, only 5.15% of probability mass falls outside the baseline’s top-10. RSP at $\tau=1.0$ reaches 27.64%. The Spearman rank correlation tells the story: because temperature scaling $p_i^{(\tau)} \propto p_i^{1/\tau}$ is a monotone transform that preserves token ranking ($p_i > p_j \iff p_i^{(\tau)} > p_j^{(\tau)}$) for any $\tau > 0$, ρ relative to the baseline is exactly 1. RSP, by contrast, reranks the token distribution itself, reducing ρ to 0.651—an effect temperature cannot produce at any τ . The Jensen–Shannon divergence (0.231) quantifies the overall magnitude of this distributional change. Since $\rho < 1$ rules out rank-preserving flattening as the sole cause, the high divergence reflects a structural reshape of the distribution rather than a uniform flattening of the same ranking.

Trajectory-level: semantic diversity. Token-level reranking raises a second question: do first-token perturbations produce semantically diverse reasoning trajectories, or do independently perturbed trajectories converge to semantically similar ones? We randomly sample 64 problems from MATH-500 and generate 300 independent trajectories per problem under Baseline and RSP at temperature $\tau = 1.0$. Each trajectory is encoded into a dense semantic vector by BGE-M3 [Chen et al., 2024], a sentence-level multilingual embedding model, and we compute three complementary metrics over these embeddings: *pairwise cosine distance* between trajectories captures the overall semantic spread of the trajectory set; *inter-cluster distance* between centroids identified by DBSCAN [Ester et al., 1996], a density-based clustering algorithm that groups nearby trajectories without prespecifying the number of clusters, captures the separation between distinct trajectory groups; and *intra-cluster distance* captures semantic coherence within each group.

Table 4 reveals three concurrent patterns. Pairwise cosine distance rises by about $1.4\times$, indicating that the trajectory set as a whole becomes more semantically diverse. Inter-cluster distance jumps by about $5.4\times$, showing that the trajectories split into substantially more separated groups. Intra-cluster distance is essentially unchanged, indicating that semantic coherence within each group remains comparable to the baseline. RSP also yields larger pairwise distance than baseline on 56 of 64 problems (87.5%).

Outcome-level: Pass@ n scaling. Finally, we evaluate whether the token- and trajectory-level diversity translates into task-level outcome diversity through Pass@ n [Chen et al., 2021], the probability that at least one of n sampled solutions is correct. We compare four conditions with $n = 16$ samples per problem on MATH-500 and AIME24: (i) *Baseline*, $\tau=1.0$: temperature sampling only; (ii) *RSP (fixed seed)*, $\tau=1.0$: a single RSP shared across all samples; (iii) *RSP (indep. seed), greedy*: a different RSP per sample without temperature; and (iv) *RSP (indep. seed)*, $\tau=1.0$: a different RSP per sample combined with temperature sampling.




Figure 3: Pass@ n scaling on (a) MATH-500 and (b) AIME24 with Qwen2.5-Math-1.5B-Instruct, $n = 16$ samples per problem. Baseline: temperature sampling only. Fixed seed: single RSP shared across samples combined with temperature. Indep. seed: a different RSP per sample, with or without temperature.

Condition (iv) achieves the highest Pass@ n on both benchmarks, reaching 92.80% on MATH-500 and 36.67% on AIME24 at $n = 16$. Condition (ii) shows no improvement over the baseline, suggesting that the diversity gain likely depends on varying the random embeddings across samples rather than a single fixed perturbation. Pass@1 remains comparable to the baseline on MATH-500, and on the more challenging AIME24, (iv) achieves higher Pass@1 than the baseline. Together these

observations suggest that combining independent RSPs with temperature sampling may yield greater trajectory diversity while preserving single-sample accuracy.

5 Conclusion

We studied Random Soft Prompts (RSPs), isotropic Gaussian vectors fitted to the entrywise mean and variance of the pretrained embedding table and injected without training. Theoretically, the RSP contribution at each attention layer is exactly captured by a scalar attention mass that cache growth structurally dilutes, and under a local affine analysis RSPs can admit previously excluded tokens into the effective top- K candidate set in a way temperature scaling cannot (§3). Empirically, RSPs raise early-stage entropy while later generation stabilizes, and the accumulation gain disappears when one RSP is shared across samples (§4). Because RSP is rank-changing while temperature is rank-preserving, the two are complementary, suggesting a sampling-level remedy for entropy collapse in GRPO [Shao et al., 2024] beyond DAPO [Yu et al., 2025] and CE-GPPO [Su et al., 2026]; training-time validation, extension beyond mathematical reasoning, and a principled choice of RSP length and injection position—together with the task-side bound $p_{\min}(x) > 0$ of Proposition 3—are natural next steps.

References

- Farhan Ahmed, Yuya Jeremy Ong, and Chad DeLuca. LogitScope: A Framework for Analyzing LLM Uncertainty through Information Metrics. In *ICLR Workshop*, 2026.
- Nora Belrose, Igor Ostrovsky, Lev McKinney, Zach Furman, Logan Smith, Danny Halawi, Stella Biderman, and Jacob Steinhardt. Eliciting Latent Predictions from Transformers with the Tuned Lens. *arXiv:2303.08112*, 2023.
- Jianlyu Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. M3-Embedding: Multi-Linguality, Multi-Functionality, Multi-Granularity Text Embeddings Through Self-Knowledge Distillation. In *Findings of ACL*, 2024.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating Large Language Models Trained on Code. *arXiv:2107.03374*, 2021.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training Verifiers to Solve Math Word Problems. *arXiv:2110.14168*, 2021.
- Jeremy M. Cohen, Elan Rosenfeld, and J. Zico Kolter. Certified Adversarial Robustness via Randomized Smoothing. In *ICML*, 2019.
- Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu. A Density-Based Algorithm for Discovering Clusters in Large Spatial Databases with Noise. In *KDD*, 1996.
- Meire Fortunato, Mohammad Gheshlaghi Azar, Bilal Piot, Jacob Menick, Ian Osband, Alex Graves, Vlad Mnih, Remi Munos, Demis Hassabis, Olivier Pietquin, Charles Blundell, and Shane Legg. Noisy Networks for Exploration. In *ICLR*, 2018.
- Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. Transformer Feed-Forward Layers Are Key-Value Memories. In *EMNLP*, 2021.
- Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh Nagarajan. Think before You Speak: Training Language Models with Pause Tokens. In *ICLR*, 2024.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurelien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Roziere, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya

336 Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang
 337 Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song,
 338 Danielle Pintz, Danny Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan,
 339 Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina
 340 Lobanova, Emily Dinan, Eric Michael Smith, Filip Radenovic, Francisco Guzmán, Frank Zhang,
 341 Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, Graeme Nail, Gregoire
 342 Mialon, Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron,
 343 Iliyan Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang,
 344 Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelmer
 345 van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang,
 346 Jiawen Liu, Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua
 347 Johnstun, Joshua Saxe, Junteng Jia, Kalyan Vasuden Alwala, Karthik Prasad, Kartikeya Upasani,
 348 Kate Plawiak, Ke Li, Kenneth Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz
 349 Malik, Kuenley Chiu, Kunal Bhalla, Kushal Lakhotia, Lauren Rantala-Yearly, Laurens van der
 350 Maaten, Lawrence Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo,
 351 Lukas Blecher, Lukas Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat
 352 Singh, Manohar Paluri, Marcin Kardas, Maria Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya
 353 Pavlova, Melanie Kambadur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman
 354 Goyal, Narjes Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri Chatterji, Ning Zhang,
 355 Olivier Duchenne, Onur Çelebi, Patrick Alrassy, Pengchuan Zhang, Pengwei Li, Petar Vasic,
 356 Peter Weng, Prajwal Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Puxin Xu,
 357 Qing He, Qingxiao Dong, Ragavan Srinivasan, Raj Ganapathy, Ramon Calderer, Ricardo Silveira
 358 Cabral, Robert Stojnic, Roberta Raileanu, Rohan Maheswari, Rohit Girdhar, Rohit Patel, Romain
 359 Sauvestre, Ronnie Polidoro, Roshan Sumbaly, Ross Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar
 360 Hosseini, Sahana Chennabasappa, Sanjay Singh, Sean Bell, Seohyun Sonia Kim, Sergey Edunov,
 361 Shao-liang Nie, Sharan Narang, Sharath Raparthy, Sheng Shen, Shengye Wan, Shruti Bhosale,
 362 Shun Zhang, Simon Vandenhende, Soumya Batra, Spencer Whitman, Sten Sootla, Stephane
 363 Collot, Suchin Gururangan, Sydney Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha,
 364 Thomas Georgiou, Thomas Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal
 365 Karn, Vedanuj Goswami, Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Gouget,
 366 Virginie Do, Vish Vogeti, Vitor Albiero, Vladan Petrovic, Weiwei Chu, Wenhan Xiong, Wenyan
 367 Fu, Whitney Meers, Xavier Martinet, Xiaodong Wang, Xiaofang Wang, Xiaoqing Ellen Tan, Xide
 368 Xia, Xinfeng Xie, Xuchao Jia, Xuwei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei,
 369 Yi Wen, Yiwen Song, Yuchen Zhang, Yue Li, Yuning Mao, Zacharie Delpierre Coudert, Zheng Yan,
 370 Zhengxing Chen, Zoe Papakipos, Aaditya Singh, Aayushi Srivastava, Abha Jain, Adam Kelsey,
 371 Adam Shajnfeld, Adithya Gangidi, Adolfo Victoria, Ahuva Goldstand, Ajay Menon, Ajay Sharma,
 372 Alex Boesenber, Alexei Baevski, Allie Feinstein, Amanda Kallet, Amit Sangani, Amos Teo,
 373 Anam Yunus, Andrei Lupu, Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew
 374 Poulton, Andrew Ryan, Ankit Ramchandani, Annie Dong, Annie Franco, Anuj Goyal, Aparajita
 375 Saraf, Arkabandhu Chowdhury, Ashley Gabriel, Ashwin Bharambe, Assaf Eisenman, Azadeh
 376 Yazdan, Beau James, Ben Maurer, Benjamin Leonhardi, Bernie Huang, Beth Loyd, Beto De Paola,
 377 Bhargavi Paranjape, Bing Liu, Bo Wu, Boyu Ni, Braden Hancock, Bram Wasti, Brandon Spence,
 378 Brani Stojkovic, Brian Gamido, Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu,
 379 Changan Wang, Changkyu Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris
 380 Tindal, Christoph Feichtenhofer, Cynthia Gao, Damon Civin, Dana Beaty, Daniel Kreymer, Daniel
 381 Li, David Adkins, David Xu, Davide Testuggine, Delia David, Devi Parikh, Diana Liskovich,
 382 Didem Foss, Dingkan Wang, Duc Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine
 383 Montgomery, Eleonora Presani, Emily Hahn, Emily Wood, Eric-Tuan Le, Erik Brinkman, Esteban
 384 Arcaute, Evan Dunbar, Evan Smothers, Fei Sun, Felix Kreuk, Feng Tian, Filippos Kokkinos, Firat
 385 Ozgenel, Francesco Caggioni, Frank Kanayet, Frank Seide, Gabriela Medina Florez, Gabriella
 386 Schwarz, Gada Badeer, Georgia Swee, Gil Halpern, Grant Herman, Grigory Sizov, Guangyi Zhang,
 387 Guna Lakshminarayanan, Hakan Inan, Hamid Shojanazeri, Han Zou, Hannah Wang, Hanwen Zha,
 388 Haroun Habeeb, Harrison Rudolph, Helen Suk, Henry Aspegren, Hunter Goldmann, Hongyuan
 389 Zhan, Ibrahim Damlaj, Igor Molybog, Igor Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai
 390 Gat, Jake Weissman, James Geboski, James Kohli, Janice Lam, Japhet Asher, Jean-Baptiste Gaya,
 391 Jeff Marcus, Jeff Tang, Jennifer Chan, Jenny Zhen, Jeremy Reizenstein, Jeremy Teboul, Jessica
 392 Zhong, Jian Jin, Jingyi Yang, Joe Cummings, Jon Carvill, Jon Shepard, Jonathan McPhie, Jonathan
 393 Torres, Josh Ginsburg, Junjie Wang, Kai Wu, Kam Hou U, Karan Saxena, Kartikay Khandelwal,
 394 Katayoun Zand, Kathy Matosich, Kaushik Veeraraghavan, Kelly Michelena, Keqian Li, Kiran

395 Jagadeesh, Kun Huang, Kunal Chawla, Kyle Huang, Lailin Chen, Lakshya Garg, Lavender A,
396 Leandro Silva, Lee Bell, Lei Zhang, Liangpeng Guo, Licheng Yu, Liron Moshkovich, Luca
397 Wehrstedt, Madian Khabsa, Manav Avalani, Manish Bhatt, Martynas Mankus, Matan Hasson,
398 Matthew Lennie, Matthias Reso, Maxim Groshev, Maxim Naumov, Maya Lathi, Meghan Keneally,
399 Miao Liu, Michael L. Seltzer, Michal Valko, Michelle Restrepo, Mihir Patel, Mik Vyatskov,
400 Mikayel Samvelyan, Mike Clark, Mike Macey, Mike Wang, Miquel Jubert Hermoso, Mo Metanat,
401 Mohammad Rastegari, Munish Bansal, Nandhini Santhanam, Natascha Parks, Natasha White,
402 Navyata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier, Nikhil Mehta, Nikolay Pavlovich
403 Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, Olivia Hart, Omkar Salpekar, Ozlem
404 Kalinli, Parkin Kent, Parth Parekh, Paul Saab, Pavan Balaji, Pedro Rittner, Philip Bontrager,
405 Pierre Roux, Piotr Dollar, Polina Zvyagina, Prashant Ratanchandani, Pritish Yuvraj, Qian Liang,
406 Rachad Alao, Rachel Rodriguez, Rafi Ayub, Raghotham Murthy, Raghu Nayani, Rahul Mitra,
407 Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin Battey, Rocky Wang, Russ
408 Howes, Rutu Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu, Samyak Datta, Sara Chugh,
409 Sara Hunt, Sargun Dhillon, Sasha Sidorov, Satadru Pan, Saurabh Mahajan, Saurabh Verma, Seiji
410 Yamamoto, Sharadh Ramaswamy, Shaun Lindsay, Shaun Lindsay, Sheng Feng, Shenghao Lin,
411 Shengxin Cindy Zha, Shishir Patil, Shiva Shankar, Shuqiang Zhang, Shuqiang Zhang, Sinong Wang,
412 Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala, Stephanie Max, Stephen Chen, Steve Kehoe,
413 Steve Satterfield, Sudarshan Govindaprasad, Sumit Gupta, Summer Deng, Sungmin Cho, Sunny
414 Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman, Tal Remez, Tamar Glaser, Tamara
415 Best, Thilo Koehler, Thomas Robinson, Tianhe Li, Tianjun Zhang, Tim Matthews, Timothy Chou,
416 Tzook Shaked, Varun Vontimitta, Victoria Ajayi, Victoria Montanez, Vijai Mohan, Vinay Satish
417 Kumar, Vishal Mangla, Vlad Ionescu, Vlad Poenaru, Vlad Tiberiu Mihailescu, Vladimir Ivanov,
418 Wei Li, Wenchen Wang, Wenwen Jiang, Wes Bouaziz, Will Constable, Xiaocheng Tang, Xiaojian
419 Wu, Xiaolan Wang, Xilun Wu, Xinbo Gao, Yaniv Kleinman, Yanjun Chen, Ye Hu, Ye Jia, Ye Qi,
420 Yenda Li, Yilin Zhang, Ying Zhang, Yossi Adi, Youngjin Nam, Yu Wang, Yu Zhao, Yuchen Hao,
421 Yundi Qian, Yunlu Li, Yuzi He, Zach Rait, Zachary DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhenyu
422 Yang, Zhiwei Zhao, and Zhiyu Ma. The Llama 3 Herd of Models. *arXiv:2407.21783*, 2024.

423 Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian.
424 Training Large Language Models to Reason in a Continuous Latent Space. In *ICLR*, 2025.

425 Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin de Laroussilhe, Andrea
426 Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-Efficient Transfer Learning for NLP.
427 In *ICML*, 2019.

428 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang,
429 and Weizhu Chen. LoRA: Low-Rank Adaptation of Large Language Models. In *ICLR*, 2022.

430 Neel Jain, Ping-yeh Chiang, Yuxin Wen, John Kirchenbauer, Hong-Min Chu, Gowthami Somepalli,
431 Brian R. Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Aniruddha Saha, Micah Goldblum,
432 Jonas Geiping, and Tom Goldstein. NEFTune: Noisy Embeddings Improve Instruction Finetuning.
433 In *ICLR*, 2024.

434 Xinyue Kang, Diwei Shi, and Li Chen. Model Whisper: Steering Vectors Unlock Large Language
435 Models’ Potential in Test-time. In *AAAI*, 2026.

436 Hyunjik Kim, George Papamakarios, and Andriy Mnih. The Lipschitz Constant of Self-Attention.
437 *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021.

438 Brian Lester, Rami Al-Rfou, and Noah Constant. The Power of Scale for Parameter-Efficient Prompt
439 Tuning. In *EMNLP*, 2021.

440 Xiang Lisa Li and Percy Liang. Prefix-Tuning: Optimizing Continuous Prompts for Generation. In
441 *ACL-IJCNLP*, 2021.

442 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike,
443 John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s Verify Step by Step. In *ICLR*, 2024.

444 Xiao Liu, Yanan Zheng, Zhengxiao Du, Ming Ding, Yujie Qian, Zhilin Yang, and Jie Tang. GPT
445 Understands, Too. *AI Open*, 5:208–215, 2024.

446 Aman Madaan and Amir Yazdanbakhsh. Text and Patterns: For Effective Chain of Thought, It Takes
447 Two to Tango. *arXiv:2209.07686*, 2022.

448 Alexander Robey, Eric Wong, Hamed Hassani, and George J. Pappas. SmoothLLM: Defending Large
449 Language Models Against Jailbreaking Attacks. *Transactions on Machine Learning Research*,
450 2025.

451 Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang,
452 Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the Limits of
453 Mathematical Reasoning in Open Language Models. *arXiv:2402.03300*, 2024.

454 Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov.
455 Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *Journal of Machine*
456 *Learning Research*, 15:1929–1958, 2014.

457 Zhenpeng Su, Leiyu Pan, Minxuan Lv, Yuntao Li, Wenping Hu, Fuzheng Zhang, Kun Gai, and Guorui
458 Zhou. CE-GPPO: Coordinating Entropy via Gradient-Preserving Clipping Policy Optimization in
459 Reinforcement Learning. In *ACL*, 2026.

460 Shenchi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen,
461 Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen
462 Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 Rule: High-Entropy Minority Tokens Drive
463 Effective Reinforcement Learning for LLM Reasoning. In *NeurIPS*, 2025.

464 Ruibin Xiong, Yunchang Yang, Di He, Kai Zheng, Shuxin Zheng, Chen Xing, Huishuai Zhang,
465 Yanyan Lan, Liwei Wang, and Tie-Yan Liu. On Layer Normalization in the Transformer Architec-
466 ture. In *International Conference on Machine Learning (ICML)*, 2020.

467 Yige Xu, Xu Guo, Zhiwei Zeng, and Chunyan Miao. SoftCoT: Soft Chain-of-Thought for Efficient
468 Reasoning with LLMs. In *ACL*, 2025.

469 An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu,
470 Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu,
471 Xingzhang Ren, and Zhenru Zhang. Qwen2.5-Math Technical Report: Toward Mathematical
472 Expert Model via Self-Improvement. *arXiv:2409.12122*, 2024.

473 Wengao Ye, Yan Liang, and Lianlei Shan. Thinking on the Fly: Test-Time Reasoning Enhancement
474 via Latent Thought Policy Optimization. In *ICLR*, 2026.

475 Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian
476 Fan, Gaohong Liu, Juncai Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming
477 Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Ru Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu,
478 Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Yuxuan Song, Xiangpeng Wei, Hao Zhou,
479 Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Yonghui Wu, and Mingxuan Wang. DAPO:
480 An Open-Source LLM Reinforcement Learning System at Scale. In *NeurIPS*, 2025.

481 Guibin Zhang, Muxin Fu, and Shuicheng Yan. MemGen: Weaving Generative Latent Memory for
482 Self-Evolving Agents. In *ICLR*, 2026.

483 A Experimental Setup and Full Accuracy Breakdown

484 Table 5 provides the full per-position accuracy breakdown (Prefix, Infix, Suffix) corresponding to
 485 the best-across-positions results summarized in Table 1 in the main text. Table 6 lists the number of
 486 RSP tokens L used for each model–benchmark pair, selected from $\{10, 15, 20\}$. All experiments are
 487 conducted on a single NVIDIA A6000 40GB GPU with greedy decoding, a maximum generation
 488 length of 3,072 tokens for MATH-500 and GSM8K, and 4,096 tokens for AIME24. Batch size is
 489 fixed per model (16 for LLaMA-3.1-8B-Instruct and Qwen2.5-Math-7B variants, 32 for Qwen2.5-
 490 Math-1.5B variants).

Table 5: Accuracy (%) across model configurations, injection positions, and benchmarks. Values in parentheses indicate the difference from the baseline. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Mode	MATH-500	GSM8K	AIME24
<i>Instruct Models</i>				
LLaMA-3.1-8B-Instruct	Baseline	52.20	85.75	<u>6.67</u>
	Prefix	45.00 (−7.2)	76.35 (−9.4)	3.33 (−3.3)
	Infix	49.40 (−2.8)	85.97 (+0.2)	10.00 (+3.3)
	Suffix	<u>49.40</u> (−2.8)	86.73 (+1.0)	10.00 (+3.3)
Qwen2.5-Math-7B-Instruct	Baseline	83.20	95.45	<u>13.33</u>
	Prefix	81.40 (−1.8)	94.92 (−0.5)	13.33 (+0.0)
	Infix	84.20 (+1.0)	95.60 (+0.2)	16.67 (+3.3)
	Suffix	<u>83.40</u> (+0.2)	95.68 (+0.2)	16.67 (+3.3)
Qwen2.5-Math-1.5B-Instruct	Baseline	73.00	<u>85.29</u>	10.00
	Prefix	74.80 (+1.8)	85.29 (+0.0)	<u>13.33</u> (+3.3)
	Infix	75.20 (+2.2)	85.75 (+0.5)	6.67 (−3.3)
	Suffix	74.20 (+1.2)	84.69 (−0.6)	20.00 (+10.0)
<i>Base Models (Raw Text)</i>				
Qwen2.5-Math-7B	Baseline	72.20	84.15	6.67
	Prefix	<u>71.60</u> (−0.6)	84.31 (+0.2)	16.67 (+10.0)
	Infix	63.20 (−9.0)	64.29 (−19.9)	16.67 (+10.0)
	Suffix	55.60 (−16.6)	54.13 (−30.0)	<u>13.33</u> (+6.7)
Qwen2.5-Math-1.5B	Baseline	<u>61.40</u>	77.86	16.67
	Prefix	64.40 (+3.0)	74.30 (−3.6)	10.00 (−6.7)
	Infix	63.20 (+1.8)	73.92 (−3.9)	<u>10.00</u> (−6.7)
	Suffix	54.60 (−6.8)	58.61 (−19.3)	3.33 (−13.3)
<i>Base Models + ChatML (Format Mismatch)</i>				
Qwen2.5-Math-7B	Baseline	52.20	58.30	23.33
	Prefix	59.20 (+7.0)	56.41 (−1.9)	3.33 (−20.0)
	Infix	<u>62.00</u> (+9.8)	69.07 (+10.8)	23.33 (+0.0)
	Suffix	70.40 (+18.2)	75.97 (+17.7)	23.33 (+0.0)
Qwen2.5-Math-1.5B	Baseline	34.40	37.45	<u>13.33</u>
	Prefix	63.40 (+29.0)	67.02 (+29.6)	20.00 (+6.7)
	Infix	39.20 (+4.8)	50.42 (+13.0)	6.67 (−6.7)
	Suffix	<u>51.00</u> (+16.6)	<u>50.64</u> (+13.2)	<u>13.33</u> (+0.0)

Table 6: Number of RSP tokens (L) used for each model and benchmark.

Model	MATH-500	GSM8K	AIME24
LLaMA-3.1-8B-Instruct	15	20	15
Qwen2.5-Math-7B-Instruct	20	20	20
Qwen2.5-Math-1.5B-Instruct	20	10	20
Qwen2.5-Math-7B	10	10	20
Qwen2.5-Math-1.5B	10	10	20
Qwen2.5-Math-1.5B (ChatML)	20	20	10
Qwen2.5-Math-7B (ChatML)	20	20	20

491 A.1 RSP Injection Positions and Prompt Templates

492 Below we show the prompt structure for each injection position across all model types. Blue text
493 indicates RSP embeddings, and purple text indicates special tokens.

494 A.1.1 Prefix

495 RSP embeddings are concatenated before the prompt embeddings. No text modification is applied.

496 ChatML (Qwen Instruct / Base + ChatML).

```
[Random Embeddings (L tokens)]
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{>
<|im_start|>user
{question}<|im_end|>
<|im_start|>assistant
```

498 LLaMA Chat Template.

```
[Random Embeddings (L tokens)]
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{>
<|start_header_id|>user<|end_header_id|>
{question}<|eot_id|>
<|start_header_id|>assistant<|end_header_id|>
```

500 Raw Text (Qwen Base).

```
[Random Embeddings (L tokens)]
{question}
Please reason step by step, and put your final answer within \boxed{>
```

502 A.1.2 Infix

503 A description text and special tokens are inserted into the prompt. The special token positions are
504 then replaced with RSP embeddings at the embedding level.

505 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{>
<|im_start|>user
{question}

There are {L} special tokens that contain compressed latent reasoning information that
might be useful for your reasoning. If these tokens are useful for your case, you can
use them as reference. If these tokens are not useful for your case, you can ignore
them and focus back to solving the problem.

Here are the {L} special tokens: <special_token_1>...<special_token_L>
<|im_end|>
<|im_start|>assistant
```

507 LLaMA Chat Template.

```
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{>
<|start_header_id|>user<|end_header_id|>
```

{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens:

```
<reserved_special_token_0>...  
<reserved_special_token_L>  
<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

509

510 Raw Text (Qwen Base).

{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens: $\langle \text{special_token_1} \rangle \dots \langle \text{special_token_L} \rangle$

Please reason step by step, and put your final answer within `\boxed{\}`.

511

512 A.1.3 Suffix

513 For chat-templated models, RSP embeddings are inserted immediately before the end-of-turn token.

514 For raw text models, RSP embeddings are concatenated at the end of the prompt.

515 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system  
Please reason step by step, and put your final answer within \boxed{<|im_end|>  
<|im_start|>user  
{question}[Random Embeddings (L tokens)]<|im_end|>  
<|im_start|>assistant
```

516

517 LLaMA Chat Template.

```
<|begin_of_text|>  
<|start_header_id|>system<|end_header_id|>  
Cutting Knowledge Date: December 2023  
Today Date: 26 Jul 2024  
Please reason step by step, and put your final answer within \boxed{<|eot_id|>  
<|start_header_id|>user<|end_header_id|>  
{question}[Random Embeddings (L tokens)]<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

518

519 Raw Text (Qwen Base).

```
{question}  
Please reason step by step, and put your final answer within \boxed{<|eot_id|>  
[Random Embeddings (L tokens)]
```

520

521 A.2 Answer Extraction and Grading

522 The final answer is extracted from the model output through a fallback chain. Each step is attempted
523 in order; if extraction fails, the next step is tried.

Answer Extraction Fallback Chain

1. "final answer is \$...\$. I hope" pattern (Minerva format)
2. `\boxed{...}`: last non-empty box is used; empty `\boxed{}` are skipped
3. Text following "the answer is"
4. Text following "final answer is"
5. Last number in the output (regex fallback)

Extracted answers are normalized by removing \$, \left, \right, and unit strings. Grading is performed via symbolic equivalence checking: exact string match, numerical comparison (`math.isclose, rel_tol=1e-4`), and SymPy symbolic simplification.

A.3 Reproduced Prior Work Hyperparameters

All prior methods are reproduced on Qwen2.5-Math-1.5B-Instruct and Qwen2.5-Math-7B-Instruct, evaluated with greedy decoding (temperature = 0) and our unified answer extraction pipeline.

TTSV. Prefix length 20, trained for 20 epochs with AdamW optimizer, batch size 2, gradient accumulation 8 steps, and linear learning rate schedule. Learning rate is 1e-3 for Qwen-1.5B and 1e-5 for Qwen-7B.

LTPO. Table 7 reports the per-benchmark hyperparameters.

Table 7: LTPO hyperparameters.

Parameter	GSM8K	MATH-500	AIME24
tokens	8	8	8
steps	20	20	20
top_k	10	10	10
sigma	20	20	5
sigma_decay	0.95	0.95	0.95
lr	1e-2	5e-3	5e-2

SoftCoT. Projection module trained for 10 epochs. Thought tokens: 32 during training, 4 during evaluation. Batch size 4 for GSM8K, 1 for MATH with gradient accumulation 4. The small (assistant) model is Qwen2.5-Math-1.5B-Instruct in both configurations, while the large model is Qwen2.5-Math-1.5B-Instruct (learning rate 2e-5) or Qwen2.5-Math-7B-Instruct (learning rate 5e-6). For MATH-500 and AIME24 evaluation, we use the projection module trained on GSM8K, as it yields higher accuracy than the one trained on MATH.

A.4 Full Entropy Curves

Figure 4 shows the full normalized (0–100%) generation trajectories for entropy, top-1 probability, and varentropy, and Table 8 reports the corresponding scalar statistics including overall means, first-5% means, and the average number of generated tokens per problem. All methods converge to similar levels after the initial generation stage, confirming that the distributional changes induced by RSP are localized to the early reasoning phase.

Table 8: Per-step statistics for Qwen2.5-Math-1.5B-Instruct on MATH-500 across the full generation and the first 5% of generation steps, together with the average number of generated tokens per problem. Top-1 Prob (overall) is reported as a weighted average over the first-10%/middle-last-10% segment means with weights 0.1/0.8/0.1. Extension of Figure 2 (main text) and Figure 4.

Metric	Baseline	Prefix	Infix	Suffix	LTPO	SoftCoT	TTSV
Tokens (mean)	575.7	558.4	549.9	573.3	592.3	594.4	577.4
Entropy (first 5%)	0.1332	0.1366	0.1402	0.1941	0.1662	0.4218	0.1359
Entropy (overall)	0.0871	0.0881	0.0899	0.1049	0.0909	0.1059	0.0864
Top-1 Prob (first 5%)	0.9535	0.9523	0.9525	0.9373	0.9437	0.9156	0.9534
Top-1 Prob (overall)	0.9684	0.9681	0.9678	0.9647	0.9683	0.9651	0.9685
Varentropy (first 5%)	0.2181	0.2207	0.2463	0.4010	0.2953	2.2234	0.2174
Varentropy (overall)	0.1353	0.1375	0.1422	0.2038	0.1452	0.2404	0.1361




Figure 4: Mean entropy (top), top-1 probability (middle), and varentropy (bottom) over the full normalized generation trajectory (0–100%). Averaged over MATH-500, shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: existing soft prompt methods. All methods converge after the initial stage.

547 B Token-Level Distribution Metrics

548 This appendix provides formal definitions for the three metrics used in the Token-level analysis of
 549 Section 4.4. Let P_{base} denote the baseline next-token distribution at $\tau = 1.0$ and P_{target} the candidate
 550 distribution being compared (e.g., the baseline at $\tau = 2.0$ or RSP at $\tau = 1.0$). All metrics are
 551 computed analytically from saved logits at the first generation step.

552 **Spearman rank correlation.** Spearman’s ρ is the Pearson correlation coefficient between the token
 553 ranks of two distributions:

$$\rho = \text{Pearson}(\text{rank}(P_{\text{target}}), \text{rank}(P_{\text{base}})). \quad (4)$$

554 Because temperature scaling $P^{(\tau)} \propto P^{1/\tau}$ is a strictly monotone transform, the rank order of tokens
 555 is preserved exactly, so $\rho = 1$ for any $\tau > 0$ as a mathematical identity. Any value $\rho < 1$ therefore
 556 indicates rank reorganization beyond what temperature scaling can produce.

557 **Mass outside top- K .** Let $\text{TopK}(P_{\text{base}})$ be the set of K tokens with the highest probability under
 558 P_{base} . The probability mass that the candidate distribution places *outside* this set is

$$\text{MassOutside}_K = 1 - \sum_{t \in \text{TopK}(P_{\text{base}})} P_{\text{target}}(t). \quad (5)$$

559 This directly measures support expansion: how much probability the candidate assigns to tokens that
 560 are not in the baseline’s preferred set. Reported values use $K = 10$.

561 **Jensen–Shannon divergence.** The symmetrized, bounded version of Kullback–Leibler divergence:

$$\text{JS}(P, Q) = \frac{1}{2} \text{KL}(P \parallel M) + \frac{1}{2} \text{KL}(Q \parallel M), \quad M = \frac{1}{2}(P + Q). \quad (6)$$

562 We use base-2 logarithms so that $\text{JS} \in [0, 1]$. Tokens absent from the stored top-100 are assigned a
 563 sentinel logit (-10^9) before softmax, which is negligible for our setting since the top-5 covers more
 564 than 99.9% of the probability mass.

C Prompt-Law Properties Used by the Theory

This appendix isolates the only properties of the RSP law used in the main text: centeredness, direction-agnostic covariance under a variance budget, and positive probability on every nonempty open set. We avoid stronger semantic claims because Section 3.4 needs only these geometric and measure-theoretic facts.

C.1 Centered Perturbations Do Not Impose Systematic First-Order Drift

Let Δ_{ij} be the baseline logit gap between tokens i and j , and write the first-order surrogate as

$$\Delta_{ij}(\bar{h}) = \Delta_{ij} + b_{ij}^\top \bar{h}.$$

Proposition 4 (No systematic first-order drift). *If $\mathbb{E}[\bar{h}] = 0$, then*

$$\mathbb{E}[\Delta_{ij}(\bar{h})] = \Delta_{ij}$$

for every token pair (i, j) . If a deterministic prompt h_0 is reused across runs, then

$$\Delta_{ij}(h_0) - \Delta_{ij} = b_{ij}^\top h_0$$

is a fixed directional bias.

Proof. The centered case is immediate from linearity of expectation:

$$\mathbb{E}[\Delta_{ij}(\bar{h})] = \Delta_{ij} + b_{ij}^\top \mathbb{E}[\bar{h}] = \Delta_{ij}.$$

The deterministic case follows by substitution. $\square$

This proposition rules out a run-averaged first-order bias. It does not say that individual prompt draws leave the logits unchanged.

C.2 Proof of Theorem 2

For a unit vector b , the first-order perturbation energy available along b is $\text{Var}_{h \sim D}(b^\top h) = b^\top \Sigma_D b$. The worst-case directional energy is therefore the minimum Rayleigh quotient of Σ_D .

Proof. For any symmetric covariance matrix Σ_D ,

$$\min_{\|b\|=1} b^\top \Sigma_D b = \lambda_{\min}(\Sigma_D).$$

Let the eigenvalues of Σ_D be $\lambda_1, \dots, \lambda_d$. Then

$$\lambda_{\min}(\Sigma_D) \leq \frac{1}{d} \sum_{m=1}^d \lambda_m = \frac{\text{tr}(\Sigma_D)}{d} \leq \frac{\rho^2}{d}.$$

Equality holds if and only if all eigenvalues are equal, i.e., $\Sigma_D = (\rho^2/d)\mathbf{I}_d$. $\square$

The proof uses only covariance. We choose a Gaussian law for RSP because it adds the open-set reachability property proved next.

C.3 Open-Set Reachability of Isotropic Gaussian Prompts

Proposition 5 (Open-set reachability). *Let $\bar{h} \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_d)$ with $\sigma > 0$. Then every nonempty open set $U \subseteq \mathbb{R}^d$ satisfies*

$$\Pr(\bar{h} \in U) > 0.$$

Proof. The Gaussian density

$$(2\pi\sigma^2)^{-d/2} \exp\left(-\frac{\|\bar{h}\|^2}{2\sigma^2}\right)$$

is strictly positive at every point of $\mathbb{R}^d$. Every nonempty open set contains a ball of positive Lebesgue measure, and integrating a strictly positive density over that ball gives positive probability. $\square$

This is the only support property used in Proposition 2.

D Proofs for the Transformer-Level Mechanism

This appendix follows the same order as the main theory section: exploration, annealing, and performance gain. The prompt-law facts used before these proofs are collected in Appendix C.

D.1 Exploration: attention decomposition and perturbation size

Proposition 6 (Attention decomposition). *Let $\alpha_{ij}^{(\ell)}$ be the attention weights for query position i at layer ℓ , with n real tokens and L random tokens. Define*

$$w_{r,i}^{(\ell)} := \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)}$$

and, for $j \in [n]$,

$$\tilde{\alpha}_{ij}^{(\ell)} := \frac{\alpha_{ij}^{(\ell)}}{1 - w_{r,i}^{(\ell)}}.$$

Then

$$\mathbf{x}_i^{(\ell+1)} = (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{x}}_i^{(\ell+1)} + \boldsymbol{\eta}_i^{(\ell)},$$

where

$$\tilde{\mathbf{x}}_i^{(\ell+1)} := \sum_{j=1}^n \tilde{\alpha}_{ij}^{(\ell)} \mathbf{v}_j^{(\ell)} \quad \text{and} \quad \boldsymbol{\eta}_i^{(\ell)} := \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)}.$$

Proof. The attention output is

$$\mathbf{x}_i^{(\ell+1)} = \sum_{j=1}^n \alpha_{ij}^{(\ell)} \mathbf{v}_j^{(\ell)} + \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)}.$$

By definition, $\sum_{j=1}^n \alpha_{ij}^{(\ell)} = 1 - w_{r,i}^{(\ell)}$, so

$$\sum_{j=1}^n \alpha_{ij}^{(\ell)} \mathbf{v}_j^{(\ell)} = (1 - w_{r,i}^{(\ell)}) \sum_{j=1}^n \frac{\alpha_{ij}^{(\ell)}}{1 - w_{r,i}^{(\ell)}} \mathbf{v}_j^{(\ell)} = (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{x}}_i^{(\ell+1)}.$$

Substituting this identity gives the decomposition. $\square$

Corollary 1 (Magnitude scales with random-token attention). *For every realization of the random values,*

$$\|\boldsymbol{\eta}_i^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_{j \in [L]} \|\mathbf{v}_{r_j}^{(\ell)}\|.$$

Proof. By the triangle inequality,

$$\|\boldsymbol{\eta}_i^{(\ell)}\| = \left\| \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)} \right\| \leq \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \|\mathbf{v}_{r_j}^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_{j \in [L]} \|\mathbf{v}_{r_j}^{(\ell)}\|.$$

Corollary 1 is the only quantitative fact needed in the main text. No independence assumption between attention weights and random keys is required. Once $w_{r,i}^{(\ell)}$ is small, the random contribution must be small as well. $\square$

613 D.2 Annealing: upper and lower envelopes

614 **Proof of Theorem 1 (upper envelope).** Write the scaled scores as $s_j^{(\ell)} := (\mathbf{q}_i^{(\ell)})^\top \mathbf{k}_j^{(\ell)} / \sqrt{d}$ for
 615 $j \in [n]$ and $s_{r,j'}^{(\ell)} := (\mathbf{q}_i^{(\ell)})^\top \mathbf{k}_{r,j'}^{(\ell)} / \sqrt{d}$ for $j' \in [L]$, and let $s_{r,\max}^{(\ell)} := \max_{j'} s_{r,j'}^{(\ell)}$ and $R^{(\ell)} :=$
 616 $\sum_{j'=1}^L \exp(s_{r,j'}^{(\ell)})$. The forward-gap definition gives, for every $j \in [n]$,

$$s_j^{(\ell)} \geq s_{r,\max}^{(\ell)} + \Delta_i^{(\ell)} / \sqrt{d}.$$

617 Since $\exp(s_{r,\max}^{(\ell)}) \geq R^{(\ell)} / L$ (max dominates average), summing over the n real keys yields

$$\sum_{j=1}^n \exp(s_j^{(\ell)}) \geq n \exp(\Delta_i^{(\ell)} / \sqrt{d}) \exp(s_{r,\max}^{(\ell)}) \geq \frac{n \exp(\Delta_i^{(\ell)} / \sqrt{d})}{L} R^{(\ell)}.$$

618 Plugging this into the softmax-normalized RSP mass,

$$w_{r,i}^{(\ell)} = \frac{R^{(\ell)}}{\sum_{j=1}^n \exp(s_j^{(\ell)}) + R^{(\ell)}} \leq \frac{R^{(\ell)}}{(n \exp(\Delta_i^{(\ell)} / \sqrt{d}) / L) R^{(\ell)} + R^{(\ell)}} = \frac{L}{L + n \exp(\Delta_i^{(\ell)} / \sqrt{d})}.$$

619 The right-hand side $f(n) := L / (L + n \exp(\Delta_i^{(\ell)} / \sqrt{d}))$ is strictly decreasing in n at fixed L and $\Delta_i^{(\ell)}$,
 620 since $f'(n) = -L \exp(\Delta_i^{(\ell)} / \sqrt{d}) / (L + n \exp(\Delta_i^{(\ell)} / \sqrt{d}))^2 < 0$, and $f(n) \rightarrow 0$ as $n \rightarrow \infty$. $\square$

621 **Proof of Proposition 1 (lower envelope).** The argument is symmetric. Let $s_{r,\min}^{(\ell)} := \min_{j'} s_{r,j'}^{(\ell)}$.
 622 The reverse-gap definition gives, for every $j \in [n]$,

$$s_j^{(\ell)} \leq s_{r,\min}^{(\ell)} + \Delta_i'^{(\ell)} / \sqrt{d}.$$

623 Since $\exp(s_{r,\min}^{(\ell)}) \leq R^{(\ell)} / L$ (min is dominated by average),

$$\sum_{j=1}^n \exp(s_j^{(\ell)}) \leq \frac{n \exp(\Delta_i'^{(\ell)} / \sqrt{d})}{L} R^{(\ell)},$$

624 and therefore

$$w_{r,i}^{(\ell)} = \frac{R^{(\ell)}}{\sum_{j=1}^n \exp(s_j^{(\ell)}) + R^{(\ell)}} \geq \frac{L}{L + n \exp(\Delta_i'^{(\ell)} / \sqrt{d})}.$$

625 D.3 Annealing: corollaries on cache growth

626 The two envelopes yield the corollaries below.

627 **Corollary 2** (Sequence-wise annealing of the fixed-gap envelope). *For fixed L and $\Delta_i^{(\ell)}$, the upper*
 628 *bound*

$$f(n) = \frac{L}{L + n \exp(\Delta_i^{(\ell)} / \sqrt{d})}$$

629 *is strictly decreasing in n . Cache growth alone therefore narrows the largest RSP attention mass*
 630 *compatible with that gap.*

631 During autoregressive decoding $\Delta_i^{(\ell)}$ also moves because queries, keys, and hidden states co-evolve.
 632 The accurate reading is not “the realized mass decreases at every step” but “the envelope compatible
 633 with a fixed gap shrinks as the cache grows.”

634 *Proof.* Differentiate $f(n)$ with respect to n :

$$f'(n) = -\frac{L \exp(\Delta_i^{(\ell)} / \sqrt{d})}{\left(L + n \exp(\Delta_i^{(\ell)} / \sqrt{d})\right)^2} < 0.$$

635 Thus $f(n)$ is strictly decreasing. $\square$

636 **Corollary 3** (Vanishing random contribution under annealing). *If $n \rightarrow \infty$ while L and $\Delta_i^{(\ell)}$ remain*
 637 *fixed, $w_{r,i}^{(\ell)} \rightarrow 0$, and for every realization of the random values*

$$\|\boldsymbol{\eta}_i^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_{j \in [L]} \|\mathbf{v}_{r_j}^{(\ell)}\| \rightarrow 0.$$

638 *Proof.* $w_{r,i}^{(\ell)} \rightarrow 0$ by Corollary 2. Apply Corollary 1 to bound $\|\boldsymbol{\eta}_i^{(\ell)}\|$. The maximum norm is finite
 639 for any realization of finitely many sampled values. $\square$

640 D.4 Annealing: finite-depth propagation under local Lipschitz control

641 A transformer block applies an attention sublayer followed by a feed-forward sublayer with residual
 642 connections:

$$\mathbf{x}_i^{(\ell, \text{attn})} = \mathbf{x}_i^{(\ell)} + \text{Attn}^{(\ell)}(\mathbf{x}_i^{(\ell)}), \quad \mathbf{x}_i^{(\ell+1)} = \mathbf{x}_i^{(\ell, \text{attn})} + F^{(\ell)}(\mathbf{x}_i^{(\ell, \text{attn})}).$$

643 Let $\mathbf{x}_i^{(\ell, \text{attn}, \text{RSP})}$ and $\mathbf{x}_i^{(\ell, \text{attn}, \text{base})}$ denote the post-attention states under RSP and the unperturbed forward
 644 pass at the same input $\mathbf{x}_i^{(\ell)}$. To avoid clashing with the RSP length L , we write Lipschitz constants of
 645 the FFN sublayer as $\kappa_B^{(\ell)}$ and of the full block as $\kappa_B^{(\ell)}$, and the model depth as Λ .

646 **Lemma 1** (Per-block FFN propagation). *The post-attention difference satisfies*

$$\|\mathbf{x}_i^{(\ell, \text{attn}, \text{RSP})} - \mathbf{x}_i^{(\ell, \text{attn}, \text{base})}\| \leq \|\boldsymbol{\eta}_i^{(\ell)}\| + w_{r,i}^{(\ell)} \|\tilde{\mathbf{x}}_i^{(\ell+1)}\|.$$

647 *If, in addition, $F^{(\ell)}$ is $\kappa_B^{(\ell)}$ -Lipschitz on a bounded set containing both $\mathbf{x}_i^{(\ell, \text{attn}, \text{RSP})}$ and $\mathbf{x}_i^{(\ell, \text{attn}, \text{base})}$,*
 648 *then*

$$\|\mathbf{x}_i^{(\ell+1, \text{RSP})} - \mathbf{x}_i^{(\ell+1, \text{base})}\| \leq (1 + \kappa_B^{(\ell)}) (\|\boldsymbol{\eta}_i^{(\ell)}\| + w_{r,i}^{(\ell)} \|\tilde{\mathbf{x}}_i^{(\ell+1)}\|).$$

649 *Proof.* The unperturbed attention sublayer produces $\mathbf{x}_i^{(\ell, \text{attn}, \text{base})} = \mathbf{x}_i^{(\ell)} + \tilde{\mathbf{x}}_i^{(\ell+1)}$, while the RSP
 650 attention sublayer produces $\mathbf{x}_i^{(\ell, \text{attn}, \text{RSP})} = \mathbf{x}_i^{(\ell)} + (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{x}}_i^{(\ell+1)} + \boldsymbol{\eta}_i^{(\ell)}$ by Proposition 6. Subtracting,

$$\mathbf{x}_i^{(\ell, \text{attn}, \text{RSP})} - \mathbf{x}_i^{(\ell, \text{attn}, \text{base})} = \boldsymbol{\eta}_i^{(\ell)} - w_{r,i}^{(\ell)} \tilde{\mathbf{x}}_i^{(\ell+1)},$$

651 and the triangle inequality gives the first bound. For the second, write the post-FFN difference as

$$[\mathbf{x}_i^{(\ell, \text{attn}, \text{RSP})} + F^{(\ell)}(\mathbf{x}_i^{(\ell, \text{attn}, \text{RSP})})] - [\mathbf{x}_i^{(\ell, \text{attn}, \text{base})} + F^{(\ell)}(\mathbf{x}_i^{(\ell, \text{attn}, \text{base})})].$$

652 By Lipschitz continuity, $\|F^{(\ell)}(\mathbf{x}_i^{(\ell, \text{attn}, \text{RSP})}) - F^{(\ell)}(\mathbf{x}_i^{(\ell, \text{attn}, \text{base})})\| \leq \kappa_B^{(\ell)} \|\mathbf{x}_i^{(\ell, \text{attn}, \text{RSP})} - \mathbf{x}_i^{(\ell, \text{attn}, \text{base})}\|$.
 653 The triangle inequality combined with the first bound gives the second. $\square$

654 **Corollary 4** (Cross-block propagation control). *Define $\Delta^{(\ell)} := \mathbf{x}_i^{(\ell, \text{RSP})} - \mathbf{x}_i^{(\ell, \text{base})}$ with $\Delta^{(0)} = \mathbf{0}$*
 655 *(matched inputs at layer 0), and let $\kappa_B^{(\ell)}$ denote a common local Lipschitz constant valid for both the*
 656 *baseline block map $B_{\text{base}}^{(\ell)} : \mathbf{x} \mapsto \mathbf{x} + \text{Attn}^{(\ell)}(\mathbf{x}) + F^{(\ell)}(\mathbf{x} + \text{Attn}^{(\ell)}(\mathbf{x}))$ and the RSP-augmented*
 657 *block map $B_{\text{RSP}}^{(\ell)}$ (which replaces $\text{Attn}^{(\ell)}$ with the attention computation that includes the random*
 658 *tokens) on a ball containing both trajectories. Such a constant exists for standard transformer blocks*
 659 *because Attn, LayerNorm, and the FFN sublayer are each locally Lipschitz on any bounded set*
 660 *[Xiong et al., 2020, Kim et al., 2021], and the random-key augmentation only adds a finite number of*
 661 *additional softmax inputs at every layer. Then*

$$\|\Delta^{(\ell+1)}\| \leq (1 + \kappa_B^{(\ell)}) \|\Delta^{(\ell)}\| + (1 + \kappa_B^{(\ell)}) (\|\boldsymbol{\eta}_i^{(\ell)}\| + w_{r,i}^{(\ell)} \|\tilde{\mathbf{x}}_i^{(\ell+1)}\|).$$

662 *Iterating across Λ blocks with $\Delta^{(0)} = \mathbf{0}$ yields*

$$\|\Delta^{(\Lambda)}\| \leq \sum_{\ell=0}^{\Lambda-1} \prod_{\ell'=\ell+1}^{\Lambda-1} (1 + \kappa_B^{(\ell')}) \cdot (1 + \kappa_B^{(\ell)}) (\|\boldsymbol{\eta}_i^{(\ell)}\| + w_{r,i}^{(\ell)} \|\tilde{\mathbf{x}}_i^{(\ell+1)}\|).$$

663 *Proof.* Decompose the cross-block discrepancy at layer $\ell + 1$ into two contributions:

$$\Delta^{(\ell+1)} = \underbrace{B_{\text{RSP}}^{(\ell)}(\mathbf{x}^{(\ell, \text{RSP})}) - B_{\text{RSP}}^{(\ell)}(\mathbf{x}^{(\ell, \text{base})})}_{\text{(I) propagation of prior } \Delta^{(\ell)}} + \underbrace{B_{\text{RSP}}^{(\ell)}(\mathbf{x}^{(\ell, \text{base})}) - B_{\text{base}}^{(\ell)}(\mathbf{x}^{(\ell, \text{base})})}_{\text{(II) RSP-induced injection at block } \ell},$$

664 where $B_{\text{RSP}}^{(\ell)}$ replaces $\text{Attn}^{(\ell)}$ with the RSP-augmented attention. Term (I) is bounded via the block
665 Lipschitz constant: $\|(I)\| \leq \kappa_{\text{B}}^{(\ell)} \|\Delta^{(\ell)}\|$. With residuals included one obtains the factor $(1 + \kappa_{\text{B}}^{(\ell)})$
666 when also propagating $\Delta^{(\ell)}$ itself. Term (II) is exactly the per-block injection from Lemma 1,
667 bounded by $(1 + \kappa^{(\ell)})(\|\boldsymbol{\eta}_i^{(\ell)}\| + w_{r,i}^{(\ell)} \|\bar{\mathbf{x}}_i^{(\ell+1)}\|)$. The triangle inequality gives the recursion, and
668 unrolling with $\Delta^{(0)} = \mathbf{0}$ yields the displayed sum. $\square$

669 **Boundedness of $\kappa^{(\ell)}$ for standard architectures.** Modern transformer blocks combine LayerNorm,
670 GELU or SwiGLU activations, and linear projections. LayerNorm has Lipschitz constant on the
671 order of $\|\gamma\| / \inf \|x\|$ on every set bounded away from the origin, where γ is the LayerNorm scale
672 parameter [Xiong et al., 2020]. GELU and SwiGLU are smooth and therefore locally Lipschitz on
673 bounded activation regions, and linear projections have Lipschitz constant equal to the spectral norm
674 of the weight matrix. The composed self-attention sublayer has a Lipschitz constant determined by
675 the spectral norms of $\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V$ and the inverse-temperature β as analyzed by Kim et al. [2021].
676 On the bounded activation regions traversed during a single forward pass, the per-block Lipschitz
677 constants in Corollary 4 are therefore finite. For a fixed depth Λ this yields a finite cumulative bound
678 on $\|\Delta^{(\Lambda)}\|$; stronger depth-uniform non-expansion would require additional contraction assumptions
679 not imposed here. This is the formal sense in which the injected perturbation stays controlled across
680 a finite forward pass, while the size of each new injection is set by $w_{r,i}^{(\ell)}$.

681 D.5 Performance gain: from local admission to Pass@N

682 The main text §3.4 uses a first-order surrogate to state two claims: a baseline-excluded token can
683 enter the local top- K set, and independent reruns can turn such local openings into Pass@ N gains.
684 The proofs below use only one support condition on the centered prompt law D :

$$(S1) \quad D(U) > 0 \quad \text{for every nonempty open set } U \subseteq \mathbb{R}^d.$$

685 Proposition 5 shows that the isotropic Gaussian law of §3.4 satisfies (S1).

686 D.5.1 Autoregressive Formulation

687 In autoregressive decoding, the joint distribution over a sequence $y_{1:T}$ factorizes as

$$\Pr(y_{1:T} \mid x) = \prod_{t=1}^T p(y_t \mid x, y_{<t}),$$

688 where each step is governed by a prefix-conditioned distribution. Any perturbation introduced by
689 RSP affects generation through a sequence of local changes to $p(y_t \mid x, y_{<t})$, rather than a single
690 global distribution.

691 D.5.2 Local Linearization of Logits under RSP

692 Throughout this subsection, $\bar{h} \in \mathbb{R}^d$ denotes *one* centered prompt vector from the RSP definition
693 (Eq. (1) in §3.1), treated as a per-token surrogate. The actual implementation uses a length- L prompt
694 $\mathbf{H} = (h_1, \dots, h_L)$ with L independent draws; the arguments below isolate how one local perturbation
695 can shift the logits. We model the perturbed logits as a function $z_i(\bar{h})$, assuming local smoothness in
696 a neighborhood of $\bar{h} = 0$:

$$z_i(\bar{h}) = z_i(0) + \nabla_h z_i(0)^\top \bar{h} + r_i(\bar{h}),$$

697 where $r_i(\bar{h}) = O(\|\bar{h}\|^2)$. Defining $a_i := \nabla_h z_i(0)$, we obtain the local affine surrogate

$$z_i^{\text{lin}}(\bar{h}) := z_i + a_i^\top \bar{h}.$$

698 This approximation is used as a first-order surrogate for branch opening, not as an exact global model
699 of the transformer logits.

700 **D.5.3 Proof of Proposition 2**

701 Let

$$\mathcal{G}_{i,K} := \{\bar{h} \in \mathbb{R}^d : \#\{j \neq i : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h})\} \leq K - 1\}.$$

702 *Proof.* Assume $\mathcal{G}_{i,K}$ contains a nonempty open set U . By (S1), $D(U) > 0$. Since $U \subseteq \mathcal{G}_{i,K}$,

$$\Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) \geq D(U) > 0.$$

703 This is exactly the event that token i enters the surrogate top- K set.

704 A useful sufficient condition is the existence of some $\bar{h}_0$ at which token i strictly outranks all but
 705 at most $K - 1$ other tokens in the affine surrogate. By continuity, a neighborhood of $\bar{h}_0$ is then
 706 contained in $\mathcal{G}_{i,K}$. $\square$

707 **D.5.4 Proof of Proposition 3**

708 This proposition does not derive $p_{\min}(x) > 0$ from the mechanism. It only converts a task-side lower
 709 bound and independence into Pass@ N growth.

710 *Proof.* Let A_n be the event that the n -th independent RSP run reaches a correct trajectory on problem
 711 x , and write $p_n(x) := \Pr(A_n) \geq p_{\min}(x)$ as in (M1). Under (M2),

$$\Pr\left(\bigcap_{n=1}^N A_n^c\right) = \prod_{n=1}^N (1 - p_n(x)) \leq \prod_{n=1}^N (1 - p_{\min}(x)) = (1 - p_{\min}(x))^N.$$

712 Taking complements gives

$$\Pr(\text{Pass@}N \text{ on } x) = \Pr\left(\bigcup_{n=1}^N A_n\right) \geq 1 - (1 - p_{\min}(x))^N.$$

713 The bound $1 - x \leq e^{-x}$ for $x \in [0, 1]$ then yields

$$1 - (1 - p_{\min}(x))^N \geq 1 - e^{-N p_{\min}(x)}.$$

714 $\square$

E Attention Mass Analysis

This appendix formalizes the per-token attention mass reported in Figure 1 and the exclusion criteria applied to the reasoning span and the layer axis. For each of the 500 MATH-500 problems, a fresh RSP $\mathbf{H} \in \mathbb{R}^{L \times d}$ with $L = 10$ is sampled independently, so the reported heatmaps average over both problem variability and RSP-vector variability.

E.1 Per-Token Attention Mass

For each problem, we measure attention on the concatenated sequence [prompt; $\mathbf{H}$; generation], where the generation is the trajectory produced by the same model under matching RSP injection. Let $\alpha_{i,j}^{(\ell,h)}$ denote the post-softmax attention weight at layer ℓ , head h , from query position i to key position j , satisfying $\sum_j \alpha_{i,j}^{(\ell,h)} = 1$ under the causal mask. Let Q and R denote the index sets of question and RSP tokens with sizes $|Q|$ and $|R| = L$, and let H be the number of heads. The per-token attention mass that query i directs to each group at layer ℓ is

$$\text{PTM}_Q[\ell, i] = \frac{1}{|Q|H} \sum_{h=1}^H \sum_{j \in Q} \alpha_{i,j}^{(\ell,h)}, \quad \text{PTM}_R[\ell, i] = \frac{1}{|R|H} \sum_{h=1}^H \sum_{j \in R} \alpha_{i,j}^{(\ell,h)}. \quad (7)$$

Normalization by $|Q|$ and $|R|$ controls for group size, so both quantities express how much attention a single question (resp. RSP) token receives on average under the same softmax denominator. Question tokens thereby serve as a size-matched baseline that absorbs sequence-length effects common to both groups.

E.2 Heatmap Aggregation

We partition the layer indices and the reasoning position indices each into five equal-size bins $\{B_L^{(b)}\}_{b=1}^5$ and $\{B_P^{(b)}\}_{b=1}^5$. For sample s and target group $\bullet \in \{Q, R\}$, the value in cell (b_L, b_P) is

$$M_\bullet^{(s)}[b_L, b_P] = \frac{1}{|B_L^{(b_L)}| \cdot |B_P^{(b_P)}|} \sum_{\ell \in B_L^{(b_L)}} \sum_{i \in B_P^{(b_P)}} \text{PTM}_\bullet[\ell, i]. \quad (8)$$

Figure 1 reports the sample average of Eq. (8) over the 500 valid samples.

E.3 Excluding the `\boxed{\dots}` Span

The reasoning position range terminates strictly before the final `\boxed{\dots}` span. Chain-of-thought outputs decompose into a *text* (content) component and a *pattern* (template) component that contribute independently to downstream performance [Madaan and Yazdanbakhsh, 2022]. The `\boxed{\dots}` wrapper is a canonical pattern: a short, stereotyped span whose role is to complete the answer format rather than to extend reasoning. Including it would therefore inject a template-driven attention signature unrelated to reasoning dynamics.

E.4 Excluding the Final Two Layers

The layer axis further excludes the last two transformer layers. This choice is motivated by a standard late-layer effect rather than by any model-specific tuning decision.

Output-projection specialization. Late-layer hidden states lie in an approximately linear relationship with vocabulary logits and therefore function as a progressive linear readout rather than as representation-transforming blocks [Belrose et al., 2023]. Consistently, late-layer feed-forward modules have been characterized as output-distribution shaping mechanisms [Geva et al., 2021]. Attention in these layers therefore reflects output formation rather than reasoning computation.

Excluding the final two layers keeps the heatmap focused on reasoning-phase dynamics instead of readout-phase dynamics. This exclusion is not load-bearing for the main claim: the sequence-direction attenuation emphasized in the main text is also visible without it, and Theorem 1 concerns sequence-direction attenuation rather than layer-wise monotonicity.

754 E.5 Additional Suffix Heatmaps

755 Figure 5 reports the same per-token RSP attention analysis under suffix injection for Qwen2.5-Math-
 756 1.5B-Instruct and LLaMA-3.1-8B-Instruct. For the late-layer reasons discussed above, the pattern
 757 does not generalize uniformly across every cell; nevertheless, the overall trend is consistent across
 758 both models: question-token attention varies broadly across layers, whereas RSP attention shows the
 759 layer-wise and sequence-wise decay predicted by Theorem 1.




Figure 5: Per-token RSP attention mass under suffix injection for the remaining two models (500 MATH-500 samples each). Axes and preprocessing match Figure 1: x-axis is reasoning position binned into five quantiles, y-axis is layer depth binned into five quantiles (top = shallow), color encodes the 500-sample mean per-token attention assigned to RSP tokens, and tokens inside `\boxed{}` spans and the last two layers are excluded.

760 **F GRPO Formulation**

761 For a given prompt q , the behavior policy $\pi_{\theta_{\text{old}}}$ samples a group of G responses $\{o_i\}_{i=1}^G$ with
 762 corresponding rewards $\{R_i\}_{i=1}^G$. The advantage of each response is computed by normalizing the
 763 group-level rewards:

$$\hat{A}_{i,t} = \frac{R_i - \text{mean}(\{R_i\}_{i=1}^G)}{\text{std}(\{R_i\}_{i=1}^G)}. \quad (9)$$

764 The policy is updated via a clipped objective with a KL penalty:

$$\mathcal{J}_{\text{GRPO}}(\theta) = \mathbb{E} \left[\frac{1}{G} \sum_{i=1}^G \frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \left(\min \left(r_{i,t} \hat{A}_{i,t}, \text{clip}(r_{i,t}, 1-\varepsilon, 1+\varepsilon) \hat{A}_{i,t} \right) - \beta D_{\text{KL}}(\pi_{\theta} \parallel \pi_{\text{ref}}) \right) \right], \quad (10)$$

765 where $r_{i,t} = \pi_{\theta}(o_{i,t} \mid q, o_{i,<t}) / \pi_{\theta_{\text{old}}}(o_{i,t} \mid q, o_{i,<t})$ is the importance sampling ratio, ε is the clipping
 766 range, and β controls the strength of the KL penalty against the reference policy π_{ref} .

767 G Broader Impacts

768 This work provides empirical and theoretical analysis of how random embedding injection affects
769 the reasoning behavior of large language models. The contributions are primarily foundational,
770 with potential downstream implications for reinforcement learning-based post-training of reasoning
771 models such as GRPO.

772 **Positive impacts.** RSP-induced trajectory diversity could improve the sample efficiency of RL-
773 based LLM post-training by providing richer learning signals within group-relative advantage estima-
774 tion. This has the potential to reduce compute costs associated with alignment and reasoning-oriented
775 fine-tuning, making such methods more accessible to research groups with limited resources.

776 **Potential concerns.** Methods that expand the effective output support of LLMs may increase the
777 reachability of low-probability tokens, which includes both desirable alternative reasoning paths
778 and potentially unsafe or off-distribution outputs. Practitioners applying RSP in production systems
779 should combine it with existing safety filtering and alignment mechanisms rather than treating it as a
780 standalone diversity source.

781 **Limitations on impact scope.** Since this work focuses on analysis rather than deploying a new
782 model or dataset, the direct societal impact is limited. The broader implications described above are
783 contingent on future work integrating RSP into applied training pipelines.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.
- Keep the checklist subsection headings, questions/answers and guidelines below.
- Do not modify the questions and only use the provided macros for your answers.

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The authors should answer **[Yes]** if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

1090 Guidelines:

1091 • The answer [N/A] means that the paper does not involve crowdsourcing nor research

1092 with human subjects.

1093 • Including this information in the supplemental material is fine, but if the main contribu-

1094 tion of the paper involves human subjects, then as much detail as possible should be

1095 included in the main paper.

1096 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,

1097 or other labor should be paid at least the minimum wage in the country of the data

1098 collector.

1099 **15. Institutional review board (IRB) approvals or equivalent for research with human**

1100 **subjects**

1101 Question: Does the paper describe potential risks incurred by study participants, whether

1102 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)

1103 approvals (or an equivalent approval/review based on the requirements of your country or

1104 institution) were obtained?

1105 Answer: [TODO]

1106 Justification: [TODO]

1107 Guidelines:

1108 • The answer [N/A] means that the paper does not involve crowdsourcing nor research

1109 with human subjects.

1110 • Depending on the country in which research is conducted, IRB approval (or equivalent)

1111 may be required for any human subjects research. If you obtained IRB approval, you

1112 should clearly state this in the paper.

1113 • We recognize that the procedures for this may vary significantly between institutions

1114 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the

1115 guidelines for their institution.

1116 • For initial submissions, do not include any information that would break anonymity (if

1117 applicable), such as the institution conducting the review.

1118 **16. Declaration of LLM usage**

1119 Question: Does the paper describe the usage of LLMs if it is an important, original, or

1120 non-standard component of the core methods in this research? Note that if the LLM is used

1121 only for writing, editing, or formatting purposes and does *not* impact the core methodology,

1122 scientific rigor, or originality of the research, declaration is not required.

1123 Answer: [TODO]

1124 Justification: [TODO]

1125 Guidelines:

1126 • The answer [N/A] means that the core method development in this research does not

1127 involve LLMs as any important, original, or non-standard components.

1128 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not

1129 be described.