
From Noise to Diversity: Random Embedding Injection in LLM Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

Recent *soft prompt* methods, which prepend or insert trainable continuous embeddings into LLM inputs as virtual tokens, optimize these embeddings to enhance mathematical reasoning. However, the mechanism by which embedding injection itself affects model behavior, independent of optimization, remains unexplored. We investigate Random Soft Prompts (RSPs), continuous vectors sampled from the model’s token embedding distribution and injected without any training. Empirically, RSPs increase early-stage token entropy while generation stabilizes thereafter, and combining RSPs with temperature sampling yields higher trajectory diversity as measured by Pass@ n . Theoretically, we provide two complementary analyses: an energy-based framework showing that RSP introduces Langevin-like perturbations on the Hopfield energy landscape with natural dual decay across layers and generation steps, and a distributional analysis showing that RSP induces stochastic token rankings and expands the output support beyond temperature sampling. We further discuss the implications for GRPO training, where trajectory diversity is critical for learning signal quality.

1 Introduction

As large language models (LLMs) scale to billions of parameters, full fine-tuning becomes increasingly expensive. This has motivated parameter-efficient methods that adapt frozen models by introducing a small number of trainable parameters [Hu et al., 2022, Housby et al., 2019]. Among these, *soft prompts* prepend or insert learnable continuous embeddings into the input sequence, steering model behavior through the attention mechanism [Li and Liang, 2021, Lester et al., 2021]. This paradigm has been actively adopted for enhancing mathematical reasoning in LLMs [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Hao et al., 2025, Zhang et al., 2026]. Despite their diverse designs, these methods share a common structure: they inject continuous embeddings that lie outside the pretrained vocabulary and optimize them to improve downstream performance.

However, we find that optimization is not a necessary condition for these effects. We observe that injecting randomly initialized, untrained continuous embeddings can also shift the output distribution. For instance, applying a chat template to a base model not fine-tuned with such format causes a prompt-format mismatch that degrades reasoning capability, but appending random embeddings to such mismatched inputs mitigates this degradation. This suggests that injecting out-of-vocabulary continuous embeddings itself, independent of any learned content, plays a non-trivial role in shaping model behavior. Existing studies primarily evaluate *soft prompt* methods through end-task accuracy, and the mechanisms by which embedding injection influences the output distribution remain insufficiently explored.

In this work, we investigate the role of Random Soft Prompts (RSPs), continuous vectors sampled from the Gaussian distribution of the model’s token embedding matrix and injected without any

training. Empirically, RSPs increase early-stage token entropy while the output stabilizes in later generation stages, and combining RSPs with temperature sampling yields more diverse reasoning trajectories as measured by Pass@ n scaling. Theoretically, we analyze RSPs from two perspectives: an energy-based framework showing that RSP introduces controlled stochastic perturbations with natural decay through the attention mechanism, and a distributional analysis showing that RSP induces stochastic token rankings and expands the output support beyond what temperature sampling can achieve. We further discuss the implications of these findings for Group Relative Policy Optimization (GRPO) [Shao et al., 2024], where trajectory diversity is essential for effective learning signals.

Our contributions are as follows:

- We empirically analyze how RSPs affect LLM generation, showing that RSPs increase early token entropy and produce more diverse reasoning trajectories when combined with temperature sampling.
- We provide theoretical analyses from two perspectives: an energy-based framework explaining why RSP can perturb the output distribution without degrading performance, and a distributional analysis showing that RSP induces stochastic token rankings and expands the output support beyond the reach of temperature sampling.
- We discuss the implications of RSP-induced diversity for GRPO training and argue that RSP provides exploration gains complementary to existing temperature-based approaches.

2 Background

2.1 Soft Prompt Methods for LLM Reasoning

Soft prompts emerged as a parameter-efficient alternative to full fine-tuning. Prefix-Tuning [Li and Liang, 2021] introduced the paradigm by prepending trainable continuous vectors to every transformer layer, Prompt Tuning [Lester et al., 2021] simplified this to input-layer-only tuning and showed that it matches full fine-tuning at sufficient scale, and P-tuning [Liu et al., 2024] extended the approach with an LSTM-based prompt encoder to model inter-position dependencies.

More recently, *soft prompts* have been actively applied to LLM reasoning. TTSV [Kang et al., 2026] optimizes prefix embeddings at test time via entropy minimization, steering the model toward high-certainty states to activate latent reasoning capabilities. SoftCoT [Xu et al., 2025] replaces explicit chain-of-thought tokens with soft tokens generated by an assistant model, leveraging the stronger representational capacity of continuous representations over discrete language. LTPO [Ye et al., 2026] optimizes latent thought embeddings per test instance through policy gradient with a confidence-based intrinsic reward. COCONUT [Hao et al., 2025] feeds the model’s own hidden states back as input embeddings, enabling reasoning in a continuous latent space. MemGen [Zhang et al., 2026] demonstrates that embedding vectors can serve as experience memory by constructing latent token sequences that enrich the model’s reasoning. Pause tokens [Goyal et al., 2024] insert learnable tokens into the input to provide additional computation steps.

These approaches all inject out-of-vocabulary continuous embeddings optimized for task-specific objectives. Because evaluation centers on end-task accuracy, the contribution of the injection itself and that of the optimization are not separated, and the mechanisms by which such out-of-vocabulary continuous embeddings shape model behavior remain a separate and underexplored question.

2.2 Noise Injection in Neural Networks

Neural networks have incorporated noise at various stages. During training, dropout [Srivastava et al., 2014] randomly deactivates hidden units to prevent overfitting, and NEFTune [Jain et al., 2024] adds noise to token embeddings during instruction fine-tuning, reporting improvements in instruction following. These methods use noise as a regularizer and do not apply perturbations at inference time.

At inference time, perturbation-based approaches primarily target robustness and safety. Randomized smoothing [Cohen et al., 2019] injects Gaussian noise into inputs to obtain certified robustness guarantees. In the LLM setting, SmoothLLM [Robey et al., 2025] perturbs input prompts at the character level to defend against adversarial jailbreaking. These methods require multiple noisy forward passes with output aggregation, and target prediction stability rather than output diversity.

Table 1: Accuracy (%) across model configurations and benchmarks. RSP reports the best result across injection positions (Prefix, Infix, Suffix). Full per-position breakdown is in Appendix A. Values in parentheses indicate the difference from the baseline.

Model	MATH-500		GSM8K		AIME24	
	Baseline	RSP	Baseline	RSP	Baseline	RSP
<i>Instruct Models</i>						
LLaMA-3.1-8B-Inst	52.20	49.40 (−2.8)	85.75	86.73 (+1.0)	6.67	10.00 (+3.3)
Qwen2.5-Math-7B-Inst	83.20	84.20 (+1.0)	95.45	95.68 (+0.2)	13.33	16.67 (+3.3)
Qwen2.5-Math-1.5B-Inst	73.00	75.20 (+2.2)	85.29	85.75 (+0.5)	10.00	20.00 (+10.0)
<i>Base Models (Raw Text)</i>						
Qwen2.5-Math-7B	72.20	71.60 (−0.6)	84.15	84.31 (+0.2)	6.67	16.67 (+10.0)
Qwen2.5-Math-1.5B	61.40	64.40 (+3.0)	77.86	74.30 (−3.6)	16.67	10.00 (−6.7)
<i>Base Models + ChatML (Format Mismatch)</i>						
Qwen2.5-Math-7B	52.20	70.40 (+18.2)	58.30	75.97 (+17.7)	23.33	23.33 (+0.0)
Qwen2.5-Math-1.5B	34.40	63.40 (+29.0)	37.45	67.02 (+29.6)	13.33	20.00 (+6.7)

In reinforcement learning, NoisyNet [Fortunato et al., 2018] injects learnable noise into network weights to improve exploration, demonstrating that the location and structure of noise injection affect exploration quality. This motivation is closest to RSPs, which similarly target diversity, but differ in injection target (input embeddings vs. network weights) and require no learned noise parameters.

RSPs differ from these approaches in the following respects. **(1)** RSPs operate solely at inference without any training, whereas dropout and NEFTune inject noise during training, modifying the resulting model parameters accordingly. **(2)** RSPs preserve the original input and append fresh embedding vectors in a single forward pass, whereas robustness methods corrupt the input and aggregate over multiple noisy passes. **(3)** RSPs use purely random vectors sampled from the pretrained embedding statistics, whereas NoisyNet learns noise parameters through optimization.

3 Experimental Analysis

3.1 Random Soft Prompts: Definition and Experimental Setup

Let $\mathbf{W}_E \in \mathbb{R}^{V \times d}$ denote the token embedding matrix of a pretrained LLM, where V is the vocabulary size and d is the hidden dimension. We define an RSP $\mathbf{H} = [h_1, \dots, h_L] \in \mathbb{R}^{L \times d}$ as a sequence of L continuous vectors, where each vector is independently sampled from an isotropic Gaussian distribution whose mean and variance are matched to those of the pretrained token embedding matrix:

$$h_j \sim \mathcal{N}(\mu_E, \sigma_E^2 \mathbf{I}), \quad j = 1, \dots, L, \quad (1)$$

with $\mu_E = \frac{1}{Vd} \sum_{i,k} [\mathbf{W}_E]_{i,k}$ and $\sigma_E = \text{std}(\mathbf{W}_E)$. Let $\mathbf{X} = [x_1, \dots, x_T] \in \mathbb{R}^{T \times d}$ denote the embedding sequence of the original input. The RSP-augmented input $\tilde{\mathbf{X}}$ is constructed by concatenating $\mathbf{H}$ with $\mathbf{X}$ at one of three positions without modifying the original prompt text: *prefix* ($[\mathbf{H}; \mathbf{X}]$), *infix* ($\mathbf{H}$ inserted within $\mathbf{X}$), or *suffix* ($[\mathbf{X}; \mathbf{H}]$), following the injection strategies adopted in prior work [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Zhang et al., 2026]. Sampling a new $\mathbf{H}$ for each batch ensures that the results are not dependent on any realization of the random embeddings.

We evaluate on mathematical reasoning benchmarks: MATH-500 [Lightman et al., 2024], GSM8K [Cobbe et al., 2021], and AIME24, using LLaMA-3.1-8B-Instruct [Grattafiori et al., 2024] and the Qwen2.5-Math model family [Yang et al., 2024] across instruct models, base models, and base models with mismatched chat templates. All experiments use greedy decoding to isolate the effect of RSPs from sampling stochasticity. Full experimental details are provided in Appendix A.

3.2 How Does Untrained RSP Compare to Baselines & Optimized Soft Prompts?

We first examine how injecting random embeddings affects the model’s reasoning performance. Table 1 reports the accuracy for each model configuration, using the best RSP result across injection positions (Prefix, Infix, Suffix) and $L \in \{10, 15, 20\}$ for each benchmark.

Table 2: Comparison with existing soft prompt methods (TTSV, SoftCoT, LTPO) under unified evaluation. Baseline and RSP values are from Table 1. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Benchmark	Baseline	RSP	TTSV	SoftCoT	LTPO
Qwen2.5-Math-7B-Instruct	MATH-500	83.20	84.20	<u>83.80</u>	82.80	83.00
	GSM8K	<u>95.45</u>	95.68	95.30	95.00	94.84
	AIME24	13.33	<u>16.67</u>	23.30	<u>16.67</u>	13.33
Qwen2.5-Math-1.5B-Instruct	MATH-500	73.00	<u>75.20</u>	<u>75.20</u>	76.00	72.40
	GSM8K	85.29	85.75	<u>85.70</u>	84.46	84.53
	AIME24	<u>10.00</u>	20.00	6.70	6.67	<u>10.00</u>

For instruct models and base models with their native prompt formats, RSP injection maintains or slightly shifts baseline performance within a few percentage points. Most strikingly, under prompt-format mismatch — where base models receive ChatML templates they were not trained on — RSP yields pronounced improvements of up to +29 pp. These gains are difficult to attribute to simple perturbation effects, as noise that merely raises model uncertainty would further degrade performance.

We compare RSPs with prior soft prompt methods (TTSV, SoftCoT, LTPO) that optimize their embeddings for distinct objectives, under a unified evaluation protocol reported in Table 2. All methods are re-evaluated with the same unified extraction pipeline (Appendix A.2), enabling a format-agnostic comparison despite each method’s distinct prompt format and grading scheme. The results place RSPs on par with the optimized methods without requiring any training.

3.3 How Does RSP Affect the Output Distribution?

We analyze how RSP injection changes the model’s output distribution on Qwen2.5-Math-1.5B-Instruct with MATH-500, measuring per-token entropy, top-1 probability, and varentropy across the generation process. Figure 1 compares these metrics for the first 5% of generation steps across RSP variants and prior soft prompt methods (LTPO, SoftCoT, TTSV). Full statistics are in Appendix A.4.




Figure 1: Mean entropy, top-1 probability, and varentropy during the first 5% of generation steps (Qwen2.5-Math-1.5B-Instruct, MATH-500). Shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: prior soft prompt methods (LTPO, SoftCoT, TTSV).

These results show that RSP modifies the early-stage output distribution across all injection positions: entropy increases (the token distribution flattens), top-1 probability decreases (commitment to a dominant token weakens), and varentropy increases (indicating multimodal distributions [Ahmed et al., 2026]). High-entropy positions are known to act as critical forking points that determine reasoning trajectories [Wang et al., 2025], so this shift moves the model toward a state where multiple reasoning paths can coexist. The change is localized to the initial reasoning stage. In the middle and final portions of generation, all three metrics converge to baseline levels.

Although existing soft prompt methods are optimized for different objectives, they share a common pattern of increasing early-stage entropy. Notably, TTSV, which optimizes soft prompts using a reward that reduces the mean entropy of reasoning trajectories, still exhibits early entropy comparable to the baseline. This suggests that elevating early-stage entropy is an inherent effect of injecting non-pretrained continuous embeddings, independent of the optimization objective.

3.4 Does RSP Induce Trajectory Diversity?

Section 3.3 empirically established that RSP reshapes the early-stage output distribution. We now examine whether this shift translates into reasoning diversity at three complementary levels of analysis: token rankings, semantic content of trajectories, and task-level outcomes. The three analyses converge on a directional picture: RSP’s first-token perturbation propagates into semantically more diverse reasoning trajectories, which in turn yields greater output diversity at the task level. All three analyses share the setting: Qwen2.5-Math-1.5B-Instruct, MATH-500, suffix injection, and $L = 10$.

Token-level: distribution beyond temperature. We test whether RSP modifies the first-token distribution in a qualitatively different way from temperature scaling, which preserves token rankings while only flattening probabilities. Under autoregressive decoding, any trajectory-level divergence between RSP and the baseline must originate from the first generation step, a critical forking point for reasoning trajectories [Wang et al., 2025]. Using the baseline at $\tau=1.0$ as reference, we compare RSP at $\tau=1.0$ against the baseline at $\tau=2.0$

(the high-temperature ceiling of what flattening alone can achieve) on three complementary metrics: Spearman rank correlation ρ to detect rank reorganization, the probability mass placed outside the baseline’s top-10 to measure support expansion, and Jensen–Shannon divergence to quantify the overall magnitude of distributional change (formal definitions in Appendix B).

Table 3 reports the contrast. Even at $\tau=2.0$, only 5.15% of probability mass falls outside the baseline’s top-10. RSP at $\tau=1.0$ reaches 27.64%. The Spearman rank correlation tells the story: because temperature scaling $p_i^{(\tau)} \propto p_i^{1/\tau}$ is a monotone transform that preserves token ranking ($p_i > p_j \iff p_i^{(\tau)} > p_j^{(\tau)}$) for any $\tau > 0$, ρ relative to the baseline is exactly 1. RSP, by contrast, reranks the token distribution itself, reducing ρ to 0.651—an effect temperature cannot produce at any τ . The Jensen–Shannon divergence (0.231) quantifies the overall magnitude of this distributional change. Since $\rho < 1$ rules out rank-preserving flattening as the sole cause, the high divergence reflects a structural reshape of the distribution rather than a uniform flattening of the same ranking.

Table 3: First-token distribution metrics measured against the baseline at $\tau=1.0$, comparing the baseline at $\tau=2.0$ with RSP at $\tau=1.0$.

Metric	$\tau=2.0$	RSP
Spearman ρ	1.0	0.651
Mass outside top-10	5.15%	27.64%
JS divergence	0.086	0.231

Trajectory-level: semantic diversity. Token-level reranking raises a second question: do first-token perturbations produce semantically diverse reasoning trajectories, or do independently perturbed trajectories converge to semantically similar ones? We randomly sample 64 problems from MATH-500 and generate 300 independent trajectories per problem under Baseline and RSP at temperature $\tau = 1.0$. Each trajectory is encoded into a dense semantic vector by BGE-M3 [Chen et al., 2024], a sentence-level multilingual embedding model, and we compute three complementary metrics over these embeddings: *pairwise cosine distance* between trajectories captures the overall semantic spread of the trajectory set; *inter-cluster distance* between centroids identified by DBSCAN [Ester et al., 1996], a density-based clustering algorithm that groups nearby trajectories without prespecifying the number of clusters, captures the separation between distinct trajectory groups; and *intra-cluster distance* captures semantic coherence within each group.

Table 4 reveals three concurrent patterns. Pairwise cosine distance rises by about $1.4\times$, indicating that the trajectory set as a whole becomes more semantically diverse. Inter-cluster distance jumps by about $5.4\times$, showing that the trajectories split into substantially more separated groups. Intra-cluster distance is essentially unchanged, indicating that semantic coherence within each group remains comparable to the baseline. RSP also yields larger pairwise distance than baseline on 56 of 64 problems (87.5%).

Table 4: Semantic diversity metrics (64 problems, 300 trajectories per condition).

Metric	Baseline	RSP
Pairwise cos. dist.	0.0612	0.0879
Inter-cluster dist.	0.0183	0.0982
Intra-cluster dist.	0.0526	0.0528

Outcome-level: Pass@ n scaling. Finally, we evaluate whether the token- and trajectory-level diversity translates into task-level outcome diversity through Pass@ n [Chen et al., 2021], the probability that at least one of n sampled solutions is correct. We compare four conditions with $n = 16$ samples per problem on MATH-500 and AIME24: (i) *Baseline*, $\tau=1.0$: temperature sampling only;

198 (ii) *RSP (fixed seed)*, $\tau = 1.0$: a single RSP shared across all samples; (iii) *RSP (indep. seed)*, *greedy*:
 199 a different RSP per sample without temperature; and (iv) *RSP (indep. seed)*, $\tau = 1.0$: a different RSP
 200 per sample combined with temperature sampling.




Figure 2: Pass@ n scaling on (a) MATH-500 and (b) AIME24 with Qwen2.5-Math-1.5B-Instruct, $n = 16$ samples per problem. Baseline: temperature sampling only. Fixed seed: single RSP shared across samples combined with temperature. Indep. seed: a different RSP per sample, with or without temperature.

201 Condition (iv) achieves the highest Pass@ n on both benchmarks, reaching 92.80% on MATH-500 and
 202 36.67% on AIME24 at $n = 16$. Condition (ii) shows no improvement over the baseline, suggesting
 203 that the diversity gain likely depends on varying the random embeddings across samples rather
 204 than a single fixed perturbation. Pass@1 remains comparable to the baseline on MATH-500, and
 205 on the more challenging AIME24, (iv) achieves higher Pass@1 than the baseline. Together these
 206 observations suggest that combining independent RSPs with temperature sampling may yield greater
 207 trajectory diversity while preserving single-sample accuracy.

208 4 Theoretical Analysis: RSP as Attention-Mediated Exploration

209 We isolate a transformer-level mechanism explaining the empirical regularities of Section 3: RSP
 210 affects the hidden state before the output softmax, its influence is controlled by the attention mass
 211 assigned to the random tokens, and this influence is structurally diluted as real-token evidence
 212 accumulates. The argument is built from standard self-attention identities; in particular, none of the
 213 formal results in Sections 4.3–4.5 assume $\mathbf{K} = \mathbf{V}$, weight-tied layers, or the existence of a global
 214 energy.

215 4.1 What RSP Is: Scale-Matched Isotropic Perturbation

216 **Definition 1** (Random Soft Prompt). Let $\mathbf{W}_E \in \mathbb{R}^{V \times d}$ be the pretrained input-embedding table,
 217 with μ_E and σ_E the entrywise scalar mean and standard deviation defined in Section 3.1. An RSP
 218 of length k is a sequence $(h_1, \dots, h_k)$ with $h_j = \mu_E \mathbf{1}_d + \sigma_E \xi_j$ and $\xi_j \stackrel{iid}{\sim} \mathcal{N}(0, \mathbf{I}_d)$. The centered
 219 prompt $\tilde{h}_j := h_j - \mu_E \mathbf{1}_d$ satisfies $\tilde{h}_j \sim \mathcal{N}(0, \sigma_E^2 \mathbf{I}_d)$. Whenever we analyze repeated-sampling
 220 protocols such as Pass@ N , we additionally assume the prompt draws and any decoding randomness
 221 used across compared trials are independent.

222 This is not a model of the empirical embedding distribution, which is generally anisotropic [Mu and
 223 Viswanath, 2018, Ethayarajh, 2019]. RSP deliberately discards directional geometry and replaces the
 224 row-covariance of $\mathbf{W}_E$ by the isotropic $\sigma_E^2 \mathbf{I}_d$: σ_E^2 borrows only the *scale*, so that the perturbation
 225 operates in the magnitude regime the model expects, while the identity covariance *deliberately*
 226 discards direction. The reason is that the branch-sensitive direction at a decoding state depends on
 227 prompt, layer, position, and the specific token pair compared; a generic exploration perturbation
 228 should not privilege a small set of embedding directions.

229 4.2 Why Randomness Is Necessary for Exploration

230 *RSP is not random because we believe the embedding table is Gaussian; it is random because*
 231 *exploration requires resampling, centering, and direction-agnostic coverage.* A fixed continuous
 232 prompt can steer a model, but reusing the same vector across rollouts produces the same perturbation
 233 direction every time, so it cannot generate independent alternatives for repeated sampling. Discrete

random tokens are also inadequate, because each vocabulary token carries a lexical and semantic direction inherited from pretraining: discrete sampling injects uncontrolled semantic priors. RSP avoids both by using continuous vectors that are centered and isotropic rather than tied to any specific embedding direction. We formalize the choice through two properties.

Lemma 1 (Semantic neutrality). *Let $\bar{h} \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_d)$ and $\hat{h} := \bar{h}/\|\bar{h}\|$. For any fixed unit vectors $u_1, \dots, u_M \in \mathbb{S}^{d-1}$ and $\delta \in (0, 1)$,*

$$\Pr \left[\max_{r \in [M]} |\hat{h}^\top u_r| \leq \sqrt{\frac{2 \log(2M/\delta)}{d-1}} \right] \geq 1 - \delta.$$

A token embedding does not have this property: any fixed direction $u = e_s/\|e_s\|$ gives $|\hat{e}_s^\top u| = 1$. Lemma 1 therefore says, precisely, that an isotropic random direction is nearly orthogonal to any fixed finite collection of semantic anchors with high probability; it does *not* claim that a single random vector is free of every incidental correlation.

Theorem 1 (Direction-agnostic coverage). *Let D be any zero-mean perturbation distribution on $\mathbb{R}^d$ with covariance Σ_D and trace budget $\text{tr}(\Sigma_D) \leq \rho^2$. Define the worst directional variance $\mathcal{E}(D) := \min_{\|b\|=1} \text{Var}_{h \sim D}(b^\top h)$. Then*

$$\mathcal{E}(D) = \lambda_{\min}(\Sigma_D) \leq \frac{\rho^2}{d},$$

with equality if and only if $\Sigma_D = (\rho^2/d)\mathbf{I}_d$.

Theorem 1 is a minimax statement: among all zero-mean perturbations with the same total variance budget, isotropic covariance uniquely maximizes the least-covered direction. Together with Lemma 1 this rules out the most natural alternatives: a fixed prompt has $\Sigma_D = 0$ and produces no exploration; PCA-aligned noise is anisotropic and underexplores the orthogonal complement; vocabulary-token sampling is both anisotropic and semantically committing. The same theorem also rules out the apparently natural alternative of *matching* the empirical embedding covariance.

Corollary 1 (Embedding-matched noise is suboptimal). *Let Σ_E be the row-covariance of $\mathbf{W}_E$ around the same scalar-broadcast center $\mu_E \mathbf{1}_d$ used in Definition 1, with $\text{tr}(\Sigma_E) = d\sigma_E^2$. Compare two zero-mean perturbations on $\mathbb{R}^d$ with the common trace budget $\rho^2 := d\sigma_E^2$: the isotropic RSP law D_{iso} with covariance $\sigma_E^2 \mathbf{I}_d$, and the embedding-matched law D_{emb} with covariance Σ_E . Then $\mathcal{E}(D_{\text{iso}}) = \sigma_E^2 \geq \lambda_{\min}(\Sigma_E) = \mathcal{E}(D_{\text{emb}})$, with strict inequality whenever Σ_E is non-isotropic.*

The corollary requires no separate proof beyond Theorem 1. To the extent that pretrained embedding tables inherit the documented anisotropy of word and contextual representations [Mu and Viswanath, 2018, Ethayarajh, 2019], $\lambda_{\min}(\Sigma_E)$ is strictly smaller than σ_E^2 , so an embedding-matched alternative leaves at least one direction less explored than the isotropic comparator. Discarding the directional geometry of $\mathbf{W}_E$ while keeping its scale is therefore not a simplification but a deliberately better choice for direction-agnostic exploration.

4.3 Exact Transformer-Level Mechanism

We now show that the effect of RSP in a standard transformer is controlled by the attention mass assigned to the random tokens. None of the statements in this subsection assume $\mathbf{K} = \mathbf{V}$, weight-tied layers, or the existence of a global energy function; they follow from standard self-attention identities. (A complementary interpretation as iterative retrieval on a Hopfield-style energy landscape [Ramsauer et al., 2021] is given in Appendix D; the formal results below do not depend on that view.)

Attention decomposition. Consider a self-attention layer ℓ with n real tokens and k RSP tokens. Let $\alpha_{ij}^{(\ell)}$ denote the attention weight from query position i to token j . Define the random-token attention mass

$$w_{r,i}^{(\ell)} = \sum_{j=1}^k \alpha_{i,n+j}^{(\ell)}, \quad (2)$$

and write $\mathbf{v}_j^{(\ell)}$ and $\mathbf{v}_{r_j}^{(\ell)}$ for the value vectors of the j -th real and the j -th RSP token, respectively. The attention output decomposes *exactly* as

$$\mathbf{x}_i^{(\ell+1)} = (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{x}}_i^{(\ell+1)} + \boldsymbol{\eta}_i^{(\ell)}, \quad (3)$$

276 where $\tilde{\mathbf{x}}_i^{(\ell+1)} = \sum_{j=1}^n \frac{\alpha_{ij}^{(\ell)}}{1-w_{r,i}^{(\ell)}} \mathbf{v}_j^{(\ell)}$ is the renormalized attention output over the real tokens alone
 277 and $\boldsymbol{\eta}_i^{(\ell)} = \sum_{j=1}^k \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)}$ is the RSP contribution. By the triangle inequality, $\|\boldsymbol{\eta}_i^{(\ell)}\| \leq$
 278 $w_{r,i}^{(\ell)} \max_{j \in [k]} \|\mathbf{v}_{r_j}^{(\ell)}\|$. Hence the same scalar $w_{r,i}^{(\ell)}$ controls both how much real-token attention
 279 is displaced and how large the random contribution can be. Eq. (3) is an algebraic identity, not an
 280 approximation; the proof and the deterministic bound on $\boldsymbol{\eta}_i^{(\ell)}$ are in Appendix E.1.

281 **Sequence-wise dilution.** With $\Delta_i^{(\ell)} := \min_{j \in [n]} \mathbf{q}_i^\top \mathbf{k}_j - \max_{j' \in [k]} \mathbf{q}_i^\top \mathbf{k}_{r_{j'}}$, the real-vs-random
 282 pre-scaling logit gap, the random-token mass satisfies

$$w_{r,i}^{(\ell)} \leq \frac{k}{k + n \exp(\Delta_i^{(\ell)} / \sqrt{d})}. \quad (4)$$

283 For fixed k and $\Delta_i^{(\ell)}$ the right-hand side is strictly decreasing in n , so autoregressive cache growth
 284 alone imposes a monotone decrease on the largest random-token attention compatible with that gap.
 285 This is a structural dilution effect, not monotone decay of the realized $w_{r,i}^{(\ell)}$ on every path: $\Delta_i^{(\ell)}$ itself
 286 evolves with state, and layer-wise decay (deeper layers widen $\Delta_i^{(\ell)}$ as real-token retrieval sharpens)
 287 is a model-side hypothesis verified empirically by Figure 3.

288 **Block-level propagation.** A full transformer block adds residual connections, normalization, and a
 289 feed-forward sublayer. Under a common local Lipschitz constant for the full block map (covering the
 290 attention sublayer, LayerNorm, and the FFN) on a ball containing both the perturbed and unperturbed
 291 trajectories—which holds in practice because each component is locally Lipschitz on bounded sets
 292 [Xiong et al., 2020, Kim et al., 2021]—Appendix E.3 shows that the per-block perturbation amplitude
 293 scales linearly with $w_{r,i}^{(\ell)}$, so iterating across a finite depth gives a finite cumulative bound. Combined
 294 with Eq. (4), this explains why the perturbation enters strongly at early blocks and shrinks as evidence
 295 accumulates, without requiring a persistent perturbation throughout the reasoning chain.

296 **Empirical verification.** We test these model-dependent steps on Qwen2.5-Math-7B under suffix
 297 injection. If $w_{r,i}^{(\ell)}$ acts as an annealed exploration coefficient, then the realized attention to RSP tokens
 298 should decay across reasoning positions and layers, while attention to question tokens should not
 299 exhibit the same monotone pattern because they remain task-relevant. Figure 3 shows this contrast.




Figure 3: Per-token attention mass on Qwen2.5-Math-7B (suffix, 500 MATH-500 problems, each with an independently sampled RSP). Layer and reasoning-position axes are partitioned into five quantile bins; every cell reports the 500-sample mean per-token attention mass directed to (a) the RSP embeddings and (b) the question tokens within that region. Attention to the RSP embeddings in (a) decays along both axes, while attention to the question tokens in (b) remains comparatively uniform.

300 4.4 Output Support Expansion Beyond Temperature

301 Temperature sampling is monotone, $p_i^{(\tau)} \propto p_i^{1/\tau}$, so $p_i > p_j \iff p_i^{(\tau)} > p_j^{(\tau)}$ for any $\tau > 0$:
 302 it cannot change token rankings at a fixed decoding state. RSP acts before the output softmax.
 303 Throughout this subsection $\bar{h} \in \mathbb{R}^d$ denotes one centered prompt vector from Definition 1, treated as
 304 a per-token surrogate; the multi-token compounding through softmax-weighted attention is handled

empirically by Section 3.4. The two propositions below are stated for the local affine surrogate $z_i^{\text{lin}}(\bar{h}) := z_i + a_i^\top \bar{h}$ with $a_i := \nabla_{\bar{h}} z_i(0)$, valid for small $\|\bar{h}\|$ (Appendix E); they identify the structure of pre-softmax rank flips under that surrogate, while the multi-token, fully nonlinear effect is reported empirically in Section 3.4.

Proposition 1 (Stochastic ranking under the local affine surrogate). *If there exists a token pair (i, j) with $a_i \neq a_j$ and the centered prompt law assigns positive probability to every nonempty open subset of $\mathbb{R}^d$, then $0 < \Pr_{\bar{h}}(z_i^{\text{lin}}(\bar{h}) > z_j^{\text{lin}}(\bar{h})) < 1$.*

Proposition 2 (Top- K entry under the local affine surrogate). *Let $J_K(s_t)$ be the baseline top- K token set at decoding state s_t and $i \notin J_K(s_t)$. Define the surrogate entry region*

$$\mathcal{G}_{i,K} := \{\bar{h} \in \mathbb{R}^d : \#\{j \neq i : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h})\} \leq K - 1\}.$$

A sufficient condition for $\mathcal{G}_{i,K}$ to contain a nonempty open set is the existence of $\bar{h}_0 \in \mathbb{R}^d$ at which token i outranks all but at most $K - 1$ other tokens with a strict margin (no tie); the affine surrogate is continuous, so an open ball around such $\bar{h}_0$ lies in $\mathcal{G}_{i,K}$ (Appendix E). If this condition holds and the centered prompt law assigns positive probability to every nonempty open subset of $\mathbb{R}^d$, then $\Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0$.

Proposition 2 is intentionally conditional: it does not claim every excluded token admits an entry region. The complementary case in which token i is dominated by the baseline top- K in the affine surrogate is exactly the case in which the proposition makes no claim. Temperature scaling reaches no $\mathcal{G}_{i,K}$ at all because it never changes token order; the mechanism here is branch opening, not probability flattening.

4.5 From Independent Perturbations to Pass@ N

Let $\mu_i := \Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0$ be the per-sample top- K admission probability for token i . Independent draws compound geometrically.

Corollary 2 (Sample complexity of admission). *With N independent fresh RSP draws, the probability that i is admitted into the top- K at least once is $1 - (1 - \mu_i)^N \geq 1 - \exp(-N\mu_i)$.*

For task accuracy, the connection to multi-trial sampling needs an additional task-dependent ingredient.

Proposition 3 (Exploration-meets-multi-path accuracy gain). *Fix a problem x . Run the deterministic baseline once and draw N fresh independent RSPs $\bar{h}^{(1)}, \dots, \bar{h}^{(N)}$, decoding once per sample. Let $\mathcal{T}^*(x)$ be the (nonempty) set of correct trajectories and $p_{\text{base}}(x) \in \{0, 1\}$ the baseline correctness indicator. Assume (M1) for each n , the marginal probability that the n -th RSP run produces a trajectory in $\mathcal{T}^*(x)$ is at least $p_{\min}(x) > 0$, and (M2) the N RSP runs are mutually independent. Then*

$$\Pr(\text{Pass}@ (N+1) \text{ correct on } x) \geq \max\{p_{\text{base}}(x), 1 - (1 - p_{\min}(x))^N\} \geq 1 - e^{-Np_{\min}(x)}.$$

Assumption (M1) is task-side: it asserts that the per-run probability of reaching a correct trajectory is bounded below by some $p_{\min}(x) > 0$ (a sufficient way to obtain such a bound is the existence of a positive-measure prompt region leading to a correct completion, but Proposition 3 only requires the marginal lower bound). Section 3.4 verifies the empirical signature predicted by this assumption: RSP-induced trajectories include both improvements and regressions, with net gain emerging under repeated sampling. Proposition 3 does not claim RSP improves single-sample accuracy; it states that when a problem admits multiple valid trajectories and at least one alternative branch has positive probability, repeated independent RSP draws turn that positive probability into Pass@ N improvement.

RSP is resampled, centered, and direction-agnostic; attention injects it through a scalar mass that cache growth structurally dilutes; the early pre-softmax perturbation changes token rankings beyond what temperature can; and multi-path reasoning converts trajectory diversity into Pass@ N gain. Section 3 reports four faces of this single mechanism.

5 Conclusion

Summary. This work investigated the role of Random Soft Prompts (RSPs), continuous embeddings sampled from the pretrained token embedding distribution and injected into the model input without

any training. Through empirical and theoretical analyses, we established three findings. First, RSPs increase early-stage token entropy, top-1 probability reduction, and varentropy while the model’s output stabilizes in later generation stages (Section 3). Second, RSPs induce stochastic token rankings and expand the effective support of the output distribution, a property that temperature sampling fundamentally lacks due to its ranking-preserving nature (Section 4). Third, combining RSPs with temperature sampling produces more diverse reasoning trajectories than either method alone, as measured by Pass@ n scaling (Section 3).

Implications for GRPO training. GRPO computes advantages by comparing multiple sampled trajectories within each group, making trajectory diversity essential for effective learning signals. Existing approaches such as DAPO [Yu et al., 2025] and CE-GPPO [Su et al., 2026] address entropy collapse by modifying the clipping range or reintroducing gradients from clipped tokens, but the underlying sampling process remains based on temperature sampling, which preserves token rankings and cannot expand the output support. Our theoretical and empirical results suggest that RSP can address this limitation at the sampling level: by perturbing the initial output distribution with independent random embeddings per rollout, RSP provides a complementary source of diversity that is orthogonal to temperature-based exploration. We expect this to yield more effective learning signals for GRPO training, particularly in mitigating entropy collapse.

Limitations. Our experiments are conducted on mathematical reasoning benchmarks (MATH-500, GSM8K, AIME24) with a limited set of model families (Qwen2.5-Math, LLaMA-3.1). The energy-theoretic analysis in Section 4 relies on strong assumptions: the Hopfield–Attention equivalence (A1: $\mathbf{K} = \mathbf{V}$) and weight-tied layers (A2), which hold exactly only in Hopfield-type models and deep equilibrium models, respectively. Since real transformer architectures vary widely in their attention mechanisms and layer structures, the theoretical predictions are most directly applicable to models whose attention dynamics closely approximate the Hopfield energy framework. The GRPO application remains a theoretical expectation and has not been empirically validated.

Future work. Extending the experimental evaluation to other domains (e.g., code generation, commonsense reasoning) and larger model scales would test the generality of RSP’s effects. Empirically validating the GRPO application by training with RSP-augmented rollouts and comparing against DAPO and CE-GPPO is a natural next step. Investigating the optimal RSP distribution and token count, which currently require per-model tuning, could lead to more principled injection strategies.

References

- Farhan Ahmed, Yuya Jeremy Ong, and Chad DeLuca. LogitScope: A Framework for Analyzing LLM Uncertainty through Information Metrics. In *ICLR Workshop*, 2026.
- Nora Belrose, Igor Ostrovsky, Lev McKinney, Zach Furman, Logan Smith, Danny Halawi, Stella Biderman, and Jacob Steinhardt. Eliciting Latent Predictions from Transformers with the Tuned Lens. *arXiv:2303.08112*, 2023.
- Stéphane Boucheron, Gábor Lugosi, and Pascal Massart. *Concentration Inequalities: A Nonasymptotic Theory of Independence*. Oxford University Press, 2013.
- Jianlyu Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. M3-Embedding: Multi-Linguality, Multi-Functionality, Multi-Granularity Text Embeddings Through Self-Knowledge Distillation. In *Findings of ACL*, 2024.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating Large Language Models Trained on Code. *arXiv:2107.03374*, 2021.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training Verifiers to Solve Math Word Problems. *arXiv:2110.14168*, 2021.
- Jeremy M. Cohen, Elan Rosenfeld, and J. Zico Kolter. Certified Adversarial Robustness via Randomized Smoothing. In *ICML*, 2019.

401 Timothée Darcet, Maxime Oquab, Julien Mairal, and Piotr Bojanowski. Vision Transformers Need
402 Registers. In *ICLR*, 2024.

403 Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu. A Density-Based Algorithm for
404 Discovering Clusters in Large Spatial Databases with Noise. In *KDD*, 1996.

405 Kawin Ethayarajh. How Contextual are Contextualized Word Representations? Comparing the
406 Geometry of BERT, ELMo, and GPT-2 Embeddings. In *Proceedings of the 2019 Conference on*
407 *Empirical Methods in Natural Language Processing (EMNLP)*, 2019.

408 Meire Fortunato, Mohammad Gheshlaghi Azar, Bilal Piot, Jacob Menick, Ian Osband, Alex Graves,
409 Vlad Mnih, Remi Munos, Demis Hassabis, Olivier Pietquin, Charles Blundell, and Shane Legg.
410 Noisy Networks for Exploration. In *ICLR*, 2018.

411 Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. Transformer Feed-Forward Layers Are
412 Key-Value Memories. In *EMNLP*, 2021.

413 Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh
414 Nagarajan. Think before You Speak: Training Language Models with Pause Tokens. In *ICLR*,
415 2024.

416 Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad
417 Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela
418 Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem
419 Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurelien Rodriguez, Austen Gregerson, Ava
420 Spataru, Baptiste Roziere, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya
421 Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang
422 Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song,
423 Danielle Pintz, Danny Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan,
424 Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina
425 Lobanova, Emily Dinan, Eric Michael Smith, Filip Radenovic, Francisco Guzmán, Frank Zhang,
426 Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, Graeme Nail, Gregoire
427 Mialon, Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron,
428 Iliyan Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang,
429 Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jëlmer
430 van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang,
431 Jiawen Liu, Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua
432 Johnstun, Joshua Saxe, Junteng Jia, Kalyan Vasuden Alwala, Karthik Prasad, Kartikeya Upasani,
433 Kate Plawiak, Ke Li, Kenneth Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz
434 Malik, Kuenley Chiu, Kunal Bhalla, Kushal Lakhotia, Lauren Rantala-Yeary, Laurens van der
435 Maaten, Lawrence Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo,
436 Lukas Blecher, Lukas Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat
437 Singh, Manohar Paluri, Marcin Kardas, Maria Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya
438 Pavlova, Melanie Kambadur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman
439 Goyal, Narjes Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri Chatterji, Ning Zhang,
440 Olivier Duchenne, Onur Çelebi, Patrick Alrassy, Pengchuan Zhang, Pengwei Li, Petar Vasic,
441 Peter Weng, Prajjwal Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Puxin Xu,
442 Qing He, Qingxiao Dong, Ragavan Srinivasan, Raj Ganapathy, Ramon Calderer, Ricardo Silveira
443 Cabral, Robert Stojnic, Roberta Raileanu, Rohan Maheswari, Rohit Girdhar, Rohit Patel, Romain
444 Sauvestre, Ronnie Polidoro, Roshan Sumbaly, Ross Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar
445 Hosseini, Sahana Chennabasappa, Sanjay Singh, Sean Bell, Seohyun Sonia Kim, Sergey Edunov,
446 Shaoliang Nie, Sharan Narang, Sharath Raparthy, Sheng Shen, Shengye Wan, Shruti Bhosale,
447 Shun Zhang, Simon Vandenhende, Soumya Batra, Spencer Whitman, Sten Sootla, Stéphane
448 Collot, Suchin Gururangan, Sydney Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha,
449 Thomas Georgiou, Thomas Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal
450 Karn, Vedanuj Goswami, Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Gonguet,
451 Virginie Do, Vish Vogeti, Vítor Albiero, Vladan Petrovic, Weiwei Chu, Wenhan Xiong, Wenyin
452 Fu, Whitney Meers, Xavier Martinet, Xiaodong Wang, Xiaofang Wang, Xiaoqing Ellen Tan, Xide
453 Xia, Xinfeng Xie, Xuchao Jia, Xuwei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei,
454 Yi Wen, Yiwen Song, Yuchen Zhang, Yue Li, Yuning Mao, Zacharie Delpratier Coudert, Zheng Yan,
455 Zhengxing Chen, Zoe Papakipos, Aaditya Singh, Aayushi Srivastava, Abha Jain, Adam Kelsey,

456 Adam Shajnfeld, Adithya Gangidi, Adolfo Victoria, Ahuva Goldstand, Ajay Menon, Ajay Sharma,
 457 Alex Boesenberg, Alexei Baevski, Allie Feinstein, Amanda Kallet, Amit Sangani, Amos Teo,
 458 Anam Yunus, Andrei Lupu, Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew
 459 Poulton, Andrew Ryan, Ankit Ramchandani, Annie Dong, Annie Franco, Anuj Goyal, Aparajita
 460 Saraf, Arkabandhu Chowdhury, Ashley Gabriel, Ashwin Bharambe, Assaf Eisenman, Azadeh
 461 Yazdan, Beau James, Ben Maurer, Benjamin Leonhardi, Bernie Huang, Beth Loyd, Beto De Paola,
 462 Bhargavi Paranjape, Bing Liu, Bo Wu, Boyu Ni, Braden Hancock, Bram Wasti, Brandon Spence,
 463 Brani Stojkovic, Brian Gamido, Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu,
 464 Changhan Wang, Changkyu Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris
 465 Tindal, Christoph Feichtenhofer, Cynthia Gao, Damon Civin, Dana Beaty, Daniel Kreymer, Daniel
 466 Li, David Adkins, David Xu, Davide Testuggine, Delia David, Devi Parikh, Diana Liskovich,
 467 Didem Foss, Dingkan Wang, Duc Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine
 468 Montgomery, Eleonora Presani, Emily Hahn, Emily Wood, Eric-Tuan Le, Erik Brinkman, Esteban
 469 Arcaute, Evan Dunbar, Evan Smothers, Fei Sun, Felix Kreuk, Feng Tian, Filippas Kokkinos, Firat
 470 Ozgenel, Francesco Caggioni, Frank Kanayet, Frank Seide, Gabriela Medina Florez, Gabriella
 471 Schwarz, Gada Badeer, Georgia Swee, Gil Halpern, Grant Herman, Grigory Sizov, Guangyi Zhang,
 472 Guna Lakshminarayanan, Hakan Inan, Hamid Shojanazeri, Han Zou, Hannah Wang, Hanwen Zha,
 473 Haroun Habeeb, Harrison Rudolph, Helen Suk, Henry Aspegren, Hunter Goldman, Hongyuan
 474 Zhan, Ibrahim Damlaj, Igor Molybog, Igor Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai
 475 Gat, Jake Weissman, James Geboski, James Kohli, Janice Lam, Japhet Asher, Jean-Baptiste Gaya,
 476 Jeff Marcus, Jeff Tang, Jennifer Chan, Jenny Zhen, Jeremy Reizenstein, Jeremy Teboul, Jessica
 477 Zhong, Jian Jin, Jingyi Yang, Joe Cummings, Jon Carvill, Jon Shepard, Jonathan McPhie, Jonathan
 478 Torres, Josh Ginsburg, Junjie Wang, Kai Wu, Kam Hou U, Karan Saxena, Kartikay Khandelwal,
 479 Katayoun Zand, Kathy Matosich, Kaushik Veeraraghavan, Kelly Michelena, Keqian Li, Kiran
 480 Jagadeesh, Kun Huang, Kunal Chawla, Kyle Huang, Lailin Chen, Lakshya Garg, Lavender A,
 481 Leandro Silva, Lee Bell, Lei Zhang, Liangpeng Guo, Licheng Yu, Liron Moshkovich, Luca
 482 Wehrstedt, Madian Khabza, Manav Avalani, Manish Bhatt, Martynas Mankus, Matan Hasson,
 483 Matthew Lennie, Matthias Reso, Maxim Groshev, Maxim Naumov, Maya Lathi, Meghan Keneally,
 484 Miao Liu, Michael L. Seltzer, Michal Valko, Michelle Restrepo, Mihir Patel, Mik Vyatskov,
 485 Mikayel Samvelyan, Mike Clark, Mike Macey, Mike Wang, Miquel Jubert Hermoso, Mo Metanat,
 486 Mohammad Rastegari, Munish Bansal, Nandhini Santhanam, Natascha Parks, Natasha White,
 487 Navyata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier, Nikhil Mehta, Nikolay Pavlovich
 488 Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, Olivia Hart, Omkar Salpekar, Ozlem
 489 Kalinli, Parkin Kent, Parth Parekh, Paul Saab, Pavan Balaji, Pedro Rittner, Philip Bontrager,
 490 Pierre Roux, Piotr Dollar, Polina Zvyagina, Prashant Ratanchandani, Pritish Yuvraj, Qian Liang,
 491 Rachad Alao, Rachel Rodriguez, Rafi Ayub, Raghotham Murthy, Raghu Nayani, Rahul Mitra,
 492 Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin Battey, Rocky Wang, Russ
 493 Howes, Ruty Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu, Samyak Datta, Sara Chugh,
 494 Sara Hunt, Sargun Dhillon, Sasha Sidorov, Satadru Pan, Saurabh Mahajan, Saurabh Verma, Seiji
 495 Yamamoto, Sharadh Ramaswamy, Shaun Lindsay, Shaun Lindsay, Sheng Feng, Shenghao Lin,
 496 Shengxin Cindy Zha, Shishir Patil, Shiva Shankar, Shuqiang Zhang, Shuqiang Zhang, Sinong Wang,
 497 Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala, Stephanie Max, Stephen Chen, Steve Kehoe,
 498 Steve Satterfield, Sudarshan Govindaprasad, Sumit Gupta, Summer Deng, Sungmin Cho, Sunny
 499 Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman, Tal Remez, Tamar Glaser, Tamara
 500 Best, Thilo Koehler, Thomas Robinson, Tianhe Li, Tianjun Zhang, Tim Matthews, Timothy Chou,
 501 Tzook Shaked, Varun Vontimitta, Victoria Ajayi, Victoria Montanez, Vijai Mohan, Vinay Satish
 502 Kumar, Vishal Mangla, Vlad Ionescu, Vlad Poenaru, Vlad Tiberiu Mihailescu, Vladimir Ivanov,
 503 Wei Li, Wenchen Wang, Wenwen Jiang, Wes Bouaziz, Will Constable, Xiaocong Tang, Xiaojian
 504 Wu, Xiaolan Wang, Xilun Wu, Xinbo Gao, Yaniv Kleinman, Yanjun Chen, Ye Hu, Ye Jia, Ye Qi,
 505 Yenda Li, Yilin Zhang, Ying Zhang, Yossi Adi, Youngjin Nam, Yu Wang, Yu Zhao, Yuchen Hao,
 506 Yundi Qian, Yunlu Li, Yuzi He, Zach Rait, Zachary DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhenyu
 507 Yang, Zhiwei Zhao, and Zhiyu Ma. The Llama 3 Herd of Models. *arXiv:2407.21783*, 2024.

508 Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian.
 509 Training Large Language Models to Reason in a Continuous Latent Space. In *ICLR*, 2025.

510 Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin de Laroussilhe, Andrea
 511 Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-Efficient Transfer Learning for NLP.
 512 In *ICML*, 2019.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-Rank Adaptation of Large Language Models. In *ICLR*, 2022.

Neel Jain, Ping-yeh Chiang, Yuxin Wen, John Kirchenbauer, Hong-Min Chu, Gowthami Somepalli, Brian R. Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Aniruddha Saha, Micah Goldblum, Jonas Geiping, and Tom Goldstein. NEFTune: Noisy Embeddings Improve Instruction Finetuning. In *ICLR*, 2024.

Xinyue Kang, Diwei Shi, and Li Chen. Model Whisper: Steering Vectors Unlock Large Language Models’ Potential in Test-time. In *AAAI*, 2026.

Hyunjik Kim, George Papamakarios, and Andriy Mnih. The Lipschitz Constant of Self-Attention. *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021.

Brian Lester, Rami Al-Rfou, and Noah Constant. The Power of Scale for Parameter-Efficient Prompt Tuning. In *EMNLP*, 2021.

Xiang Lisa Li and Percy Liang. Prefix-Tuning: Optimizing Continuous Prompts for Generation. In *ACL-IJCNLP*, 2021.

Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s Verify Step by Step. In *ICLR*, 2024.

Xiao Liu, Yanan Zheng, Zhengxiao Du, Ming Ding, Yujie Qian, Zhilin Yang, and Jie Tang. GPT Understands, Too. *AI Open*, 5:208–215, 2024.

Aman Madaan and Amir Yazdanbakhsh. Text and Patterns: For Effective Chain of Thought, It Takes Two to Tango. *arXiv:2209.07686*, 2022.

Jiaqi Mu and Pramod Viswanath. All-but-the-Top: Simple and effective postprocessing for word representations. In *International Conference on Learning Representations (ICLR)*, 2018.

Yurii Nesterov. *Lectures on Convex Optimization*. Springer, 2 edition, 2018.

Hubert Ramsauer, Bernhard Schöfl, Johannes Lehner, Philipp Seidl, Michael Widrich, Lukas Gruber, Markus Holzleitner, Thomas Adler, David Kreil, Michael K Kopp, Günter Klambauer, Johannes Brandstetter, and Sepp Hochreiter. Hopfield Networks Is All You Need. In *ICLR*, 2021.

Alexander Robey, Eric Wong, Hamed Hassani, and George J. Pappas. SmoothLLM: Defending Large Language Models Against Jailbreaking Attacks. *Transactions on Machine Learning Research*, 2025.

Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models. *arXiv:2402.03300*, 2024.

Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *Journal of Machine Learning Research*, 15:1929–1958, 2014.

Zhenpeng Su, Leiyu Pan, Minxuan Lv, Yuntao Li, Wenping Hu, Fuzheng Zhang, Kun Gai, and Guorui Zhou. CE-GPPO: Coordinating Entropy via Gradient-Preserving Clipping Policy Optimization in Reinforcement Learning. In *ACL*, 2026.

Roman Vershynin. *High-Dimensional Probability: An Introduction with Applications in Data Science*. Cambridge University Press, 2018.

Shenzhi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen, Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 Rule: High-Entropy Minority Tokens Drive Effective Reinforcement Learning for LLM Reasoning. In *NeurIPS*, 2025.

Ruibin Xiong, Yunchang Yang, Di He, Kai Zheng, Shuxin Zheng, Chen Xing, Huishuai Zhang, Yanyan Lan, Liwei Wang, and Tie-Yan Liu. On Layer Normalization in the Transformer Architecture. In *International Conference on Machine Learning (ICML)*, 2020.

- 560 Yige Xu, Xu Guo, Zhiwei Zeng, and Chunyan Miao. SoftCoT: Soft Chain-of-Thought for Efficient
561 Reasoning with LLMs. In *ACL*, 2025.
- 562 An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu,
563 Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu,
564 Xingzhang Ren, and Zhenru Zhang. Qwen2.5-Math Technical Report: Toward Mathematical
565 Expert Model via Self-Improvement. *arXiv:2409.12122*, 2024.
- 566 Wengao Ye, Yan Liang, and Lianlei Shan. Thinking on the Fly: Test-Time Reasoning Enhancement
567 via Latent Thought Policy Optimization. In *ICLR*, 2026.
- 568 Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian
569 Fan, Gaohong Liu, Juncai Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming
570 Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Ru Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu,
571 Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Yuxuan Song, Xiangpeng Wei, Hao Zhou,
572 Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Yonghui Wu, and Mingxuan Wang. DAPO:
573 An Open-Source LLM Reinforcement Learning System at Scale. In *NeurIPS*, 2025.
- 574 Guibin Zhang, Muxin Fu, and Shuicheng Yan. MemGen: Weaving Generative Latent Memory for
575 Self-Evolving Agents. In *ICLR*, 2026.

576 A Experimental Setup and Full Accuracy Breakdown

577 Table 5 provides the full per-position accuracy breakdown (Prefix, Infix, Suffix) corresponding to
578 the best-across-positions results summarized in Table 1 in the main text. Table 6 lists the number of
579 RSP tokens L used for each model–benchmark pair, selected from $\{10, 15, 20\}$. All experiments are
580 conducted on a single NVIDIA A6000 40GB GPU with greedy decoding, a maximum generation
581 length of 3,072 tokens for MATH-500 and GSM8K, and 4,096 tokens for AIME24. Batch size is
582 fixed per model (16 for LLaMA-3.1-8B-Instruct and Qwen2.5-Math-7B variants, 32 for Qwen2.5-
583 Math-1.5B variants).

Table 5: Accuracy (%) across model configurations, injection positions, and benchmarks. Values in parentheses indicate the difference from the baseline. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Mode	MATH-500	GSM8K	AIME24
<i>Instruct Models</i>				
LLaMA-3.1-8B-Instruct	Baseline	52.20	85.75	<u>6.67</u>
	Prefix	45.00 (−7.2)	76.35 (−9.4)	3.33 (−3.3)
	Infix	49.40 (−2.8)	85.97 (+0.2)	10.00 (+3.3)
	Suffix	<u>49.40</u> (−2.8)	86.73 (+1.0)	10.00 (+3.3)
Qwen2.5-Math-7B-Instruct	Baseline	83.20	95.45	<u>13.33</u>
	Prefix	81.40 (−1.8)	94.92 (−0.5)	13.33 (+0.0)
	Infix	84.20 (+1.0)	95.60 (+0.2)	16.67 (+3.3)
	Suffix	<u>83.40</u> (+0.2)	95.68 (+0.2)	16.67 (+3.3)
Qwen2.5-Math-1.5B-Instruct	Baseline	73.00	<u>85.29</u>	10.00
	Prefix	74.80 (+1.8)	85.29 (+0.0)	<u>13.33</u> (+3.3)
	Infix	75.20 (+2.2)	85.75 (+0.5)	6.67 (−3.3)
	Suffix	74.20 (+1.2)	84.69 (−0.6)	20.00 (+10.0)
<i>Base Models (Raw Text)</i>				
Qwen2.5-Math-7B	Baseline	72.20	84.15	6.67
	Prefix	<u>71.60</u> (−0.6)	84.31 (+0.2)	16.67 (+10.0)
	Infix	63.20 (−9.0)	64.29 (−19.9)	16.67 (+10.0)
	Suffix	55.60 (−16.6)	54.13 (−30.0)	<u>13.33</u> (+6.7)
Qwen2.5-Math-1.5B	Baseline	<u>61.40</u>	77.86	16.67
	Prefix	64.40 (+3.0)	74.30 (−3.6)	10.00 (−6.7)
	Infix	63.20 (+1.8)	73.92 (−3.9)	<u>10.00</u> (−6.7)
	Suffix	54.60 (−6.8)	58.61 (−19.3)	3.33 (−13.3)
<i>Base Models + ChatML (Format Mismatch)</i>				
Qwen2.5-Math-7B	Baseline	52.20	58.30	23.33
	Prefix	59.20 (+7.0)	56.41 (−1.9)	<u>3.33</u> (−20.0)
	Infix	<u>62.00</u> (+9.8)	69.07 (+10.8)	23.33 (+0.0)
	Suffix	70.40 (+18.2)	75.97 (+17.7)	23.33 (+0.0)
Qwen2.5-Math-1.5B	Baseline	34.40	37.45	<u>13.33</u>
	Prefix	63.40 (+29.0)	67.02 (+29.6)	20.00 (+6.7)
	Infix	39.20 (+4.8)	50.42 (+13.0)	6.67 (−6.7)
	Suffix	<u>51.00</u> (+16.6)	<u>50.64</u> (+13.2)	<u>13.33</u> (+0.0)

Table 6: Number of RSP tokens (L) used for each model and benchmark.

Model	MATH-500	GSM8K	AIME24
LLaMA-3.1-8B-Instruct	15	20	15
Qwen2.5-Math-7B-Instruct	20	20	20
Qwen2.5-Math-1.5B-Instruct	20	10	20
Qwen2.5-Math-7B	10	10	20
Qwen2.5-Math-1.5B	10	10	20
Qwen2.5-Math-1.5B (ChatML)	20	20	10
Qwen2.5-Math-7B (ChatML)	20	20	20

584 A.1 RSP Injection Positions and Prompt Templates

585 Below we show the prompt structure for each injection position across all model types. Blue text
586 indicates RSP embeddings, and purple text indicates special tokens.

587 A.1.1 Prefix

588 RSP embeddings are concatenated before the prompt embeddings. No text modification is applied.

589 ChatML (Qwen Instruct / Base + ChatML).

```
[Random Embeddings (L tokens)]
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{>
<|im_start|>user
{question}<|im_end|>
<|im_start|>assistant
```

591 LLaMA Chat Template.

```
[Random Embeddings (L tokens)]
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{>
<|start_header_id|>user<|end_header_id|>
{question}<|eot_id|>
<|start_header_id|>assistant<|end_header_id|>
```

593 Raw Text (Qwen Base).

```
[Random Embeddings (L tokens)]
{question}
Please reason step by step, and put your final answer within \boxed{>
```

595 A.1.2 Infix

596 A description text and special tokens are inserted into the prompt. The special token positions are
597 then replaced with RSP embeddings at the embedding level.

598 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{>
<|im_start|>user
{question}

There are {L} special tokens that contain compressed latent reasoning information that
might be useful for your reasoning. If these tokens are useful for your case, you can
use them as reference. If these tokens are not useful for your case, you can ignore
them and focus back to solving the problem.

Here are the {L} special tokens: <special_token_1>...<special_token_L>
<|im_end|>
<|im_start|>assistant
```

600 LLaMA Chat Template.

```
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{>
<|start_header_id|>user<|end_header_id|>
```

{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens:

```
<reserved_special_token_0>...  
<reserved_special_token_L>  
<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

602

603 Raw Text (Qwen Base).

{question}

There are $\{L\}$ special tokens that contain compressed latent reasoning information that might be useful for your reasoning. If these tokens are useful for your case, you can use them as reference. If these tokens are not useful for your case, you can ignore them and focus back to solving the problem.

Here are the $\{L\}$ special tokens: $\langle \text{special_token_1} \rangle \dots \langle \text{special_token_L} \rangle$

Please reason step by step, and put your final answer within `\boxed{\}`.

604

605 A.1.3 Suffix

606 For chat-templated models, RSP embeddings are inserted immediately before the end-of-turn token.
607 For raw text models, RSP embeddings are concatenated at the end of the prompt.

608 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system  
Please reason step by step, and put your final answer within \boxed{ }.<|im_end|>  
<|im_start|>user  
{question}[Random Embeddings (L tokens)]<|im_end|>  
<|im_start|>assistant
```

609

610 LLaMA Chat Template.

```
<|begin_of_text|>  
<|start_header_id|>system<|end_header_id|>  
Cutting Knowledge Date: December 2023  
Today Date: 26 Jul 2024  
Please reason step by step, and put your final answer within \boxed{ }.<|eot_id|>  
<|start_header_id|>user<|end_header_id|>  
{question}[Random Embeddings (L tokens)]<|eot_id|>  
<|start_header_id|>assistant<|end_header_id|>
```

611

612 Raw Text (Qwen Base).

```
{question}  
Please reason step by step, and put your final answer within \boxed{ }. [Random  
Embeddings (L tokens)]
```

613

614 A.2 Answer Extraction and Grading

615 The final answer is extracted from the model output through a fallback chain. Each step is attempted
616 in order; if extraction fails, the next step is tried.

Answer Extraction Fallback Chain

1. "final answer is \$...\$. I hope" pattern (Minerva format)
2. `\boxed{...}`: last non-empty box is used; empty `\boxed{}` are skipped
3. Text following "the answer is"
4. Text following "final answer is"
5. Last number in the output (regex fallback)

Extracted answers are normalized by removing \$, \left, \right, and unit strings. Grading is performed via symbolic equivalence checking: exact string match, numerical comparison (`math.isclose, rel_tol=1e-4`), and SymPy symbolic simplification.

A.3 Reproduced Prior Work Hyperparameters

All prior methods are reproduced on Qwen2.5-Math-1.5B-Instruct and Qwen2.5-Math-7B-Instruct, evaluated with greedy decoding (temperature = 0) and our unified answer extraction pipeline.

TTSV. Prefix length 20, trained for 20 epochs with AdamW optimizer, batch size 2, gradient accumulation 8 steps, and linear learning rate schedule. Learning rate is 1e-3 for Qwen-1.5B and 1e-5 for Qwen-7B.

LTPO. Table 7 reports the per-benchmark hyperparameters.

Table 7: LTPO hyperparameters.

Parameter	GSM8K	MATH-500	AIME24
tokens	8	8	8
steps	20	20	20
top_k	10	10	10
sigma	20	20	5
sigma_decay	0.95	0.95	0.95
lr	1e-2	5e-3	5e-2

SoftCoT. Projection module trained for 10 epochs. Thought tokens: 32 during training, 4 during evaluation. Batch size 4 for GSM8K, 1 for MATH with gradient accumulation 4. The small (assistant) model is Qwen2.5-Math-1.5B-Instruct in both configurations, while the large model is Qwen2.5-Math-1.5B-Instruct (learning rate 2e-5) or Qwen2.5-Math-7B-Instruct (learning rate 5e-6). For MATH-500 and AIME24 evaluation, we use the projection module trained on GSM8K, as it yields higher accuracy than the one trained on MATH.

A.4 Full Entropy Curves

Figure 4 shows the full normalized (0–100%) generation trajectories for entropy, top-1 probability, and varentropy, and Table 8 reports the corresponding scalar statistics including overall means, first-5% means, and the average number of generated tokens per problem. All methods converge to similar levels after the initial generation stage, confirming that the distributional changes induced by RSP are localized to the early reasoning phase.

Table 8: Per-step statistics for Qwen2.5-Math-1.5B-Instruct on MATH-500 across the full generation and the first 5% of generation steps, together with the average number of generated tokens per problem. Top-1 Prob (overall) is reported as a weighted average over the first-10%/middle-last-10% segment means with weights 0.1/0.8/0.1. Extension of Figure 1 (main text) and Figure 4.

Metric	Baseline	Prefix	Infix	Suffix	LTPO	SoftCoT	TTSV
Tokens (mean)	575.7	558.4	549.9	573.3	592.3	594.4	577.4
Entropy (first 5%)	0.1332	0.1366	0.1402	0.1941	0.1662	0.4218	0.1359
Entropy (overall)	0.0871	0.0881	0.0899	0.1049	0.0909	0.1059	0.0864
Top-1 Prob (first 5%)	0.9535	0.9523	0.9525	0.9373	0.9437	0.9156	0.9534
Top-1 Prob (overall)	0.9684	0.9681	0.9678	0.9647	0.9683	0.9651	0.9685
Varentropy (first 5%)	0.2181	0.2207	0.2463	0.4010	0.2953	2.2234	0.2174
Varentropy (overall)	0.1353	0.1375	0.1422	0.2038	0.1452	0.2404	0.1361




Figure 4: Mean entropy (top), top-1 probability (middle), and varentropy (bottom) over the full normalized generation trajectory (0–100%). Averaged over MATH-500, shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: existing soft prompt methods. All methods converge after the initial stage.

640 B Token-Level Distribution Metrics

641 This appendix provides formal definitions for the three metrics used in the Token-level analysis of
 642 Section 3.4. Let P_{base} denote the baseline next-token distribution at $\tau = 1.0$ and P_{target} the candidate
 643 distribution being compared (e.g., the baseline at $\tau = 2.0$ or RSP at $\tau = 1.0$). All metrics are
 644 computed analytically from saved logits at the first generation step.

645 **Spearman rank correlation.** Spearman’s ρ is the Pearson correlation coefficient between the token
 646 ranks of two distributions:

$$\rho = \text{Pearson}(\text{rank}(P_{\text{target}}), \text{rank}(P_{\text{base}})). \quad (5)$$

647 Because temperature scaling $P^{(\tau)} \propto P^{1/\tau}$ is a strictly monotone transform, the rank order of tokens
 648 is preserved exactly, so $\rho = 1$ for any $\tau > 0$ as a mathematical identity. Any value $\rho < 1$ therefore
 649 indicates rank reorganization beyond what temperature scaling can produce.

650 **Mass outside top- K .** Let $\text{TopK}(P_{\text{base}})$ be the set of K tokens with the highest probability under
 651 P_{base} . The probability mass that the candidate distribution places *outside* this set is

$$\text{MassOutside}_K = 1 - \sum_{t \in \text{TopK}(P_{\text{base}})} P_{\text{target}}(t). \quad (6)$$

652 This directly measures support expansion: how much probability the candidate assigns to tokens that
 653 are not in the baseline’s preferred set. Reported values use $K = 10$.

654 **Jensen–Shannon divergence.** The symmetrized, bounded version of Kullback–Leibler divergence:

$$\text{JS}(P, Q) = \frac{1}{2} \text{KL}(P \parallel M) + \frac{1}{2} \text{KL}(Q \parallel M), \quad M = \frac{1}{2}(P + Q). \quad (7)$$

655 We use base-2 logarithms so that $\text{JS} \in [0, 1]$. Tokens absent from the stored top-100 are assigned a
 656 sentinel logit (-10^9) before softmax, which is negligible for our setting since the top-5 covers more
 657 than 99.9% of the probability mass.

C Why Random Prompts Are Semantically Neutral and Direction-Agnostic

This appendix formalizes the “why random” claim used in Section 4.2. The point is not that random prompts contain hidden task information. The point is structural: if the perturbation is meant to act as a generic exploration device rather than as a learned prompt prior, then it should be resampled across runs, centered so that it does not induce a fixed first-order drift, and direction-agnostic so that it does not under-serve whichever local branch direction matters for the current prompt.

C.1 Semantic Neutrality

Let $\bar{h} \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_d)$ and $\hat{h} := \bar{h}/\|\bar{h}\|$ whenever $\bar{h} \neq 0$. Let $u_1, \dots, u_M \in \mathbb{S}^{d-1}$ be any fixed unit vectors representing semantic anchor directions.

Lemma 2 (Semantic neutrality). *For any $\delta \in (0, 1)$,*

$$\Pr \left[\max_{r \in [M]} |\hat{h}^\top u_r| \leq \sqrt{\frac{2 \log(2M/\delta)}{d-1}} \right] \geq 1 - \delta.$$

Hence the normalized direction of an isotropic Gaussian prompt is simultaneously nearly orthogonal to every fixed finite collection of semantic anchors with high probability.

Proof. Because $\bar{h}$ is isotropic Gaussian, its normalized direction $\hat{h}$ is uniformly distributed on the sphere $\mathbb{S}^{d-1}$. For a fixed unit vector $u \in \mathbb{S}^{d-1}$, the marginal $\hat{h}^\top u$ satisfies the standard spherical concentration bound [Vershynin, 2018, Theorem 3.4.6][Boucheron et al., 2013, Section 3.2],

$$\Pr(|\hat{h}^\top u| \geq t) \leq 2 \exp\left(-\frac{(d-1)t^2}{2}\right) \quad \text{for } t \in (0, 1),$$

which follows from the spherical-cap area bound $\sigma\{\hat{h} \in \mathbb{S}^{d-1} : \hat{h}^\top u \geq t\} \leq \frac{1}{2}(1-t^2)^{(d-1)/2}$ together with $(1-t^2)^{(d-1)/2} \leq \exp(-(d-1)t^2/2)$ derived from $\log(1-x) \leq -x$ for $x \in [0, 1)$. Applying this bound to each u_r and using a union bound,

$$\Pr \left[\max_{r \in [M]} |\hat{h}^\top u_r| \geq t \right] \leq 2M \exp\left(-\frac{(d-1)t^2}{2}\right).$$

Setting

$$t = \sqrt{\frac{2 \log(2M/\delta)}{d-1}}$$

gives the claim. If the displayed value of t exceeds 1, the statement is trivial because $|\hat{h}^\top u_r| \leq 1$ always. $\square$

The contrast with a token embedding is immediate. If $e_s \neq 0$ is a fixed token embedding and $\hat{e}_s := e_s/\|e_s\|$, then choosing the anchor direction $u = \hat{e}_s$ gives

$$|\hat{e}_s^\top u| = 1.$$

A token prompt is therefore maximally aligned with at least one semantic direction, whereas an isotropic random prompt is nearly orthogonal to every fixed finite set of such directions. This is the sense in which tokens inject semantic commitment while random prompts act as semantically uncommitted perturbations.

C.2 Centered Perturbations Do Not Impose Systematic First-Order Drift

Let Δ_{ij} be the baseline logit gap between tokens i and j , and write the first-order surrogate as

$$\Delta_{ij}(\bar{h}) = \Delta_{ij} + b_{ij}^\top \bar{h}.$$

Proposition 4 (No systematic first-order drift). *If $\mathbb{E}[\bar{h}] = 0$, then*

$$\mathbb{E}[\Delta_{ij}(\bar{h})] = \Delta_{ij}$$

for every token pair (i, j) . If a deterministic prompt h_0 is reused across runs, then

$$\Delta_{ij}(h_0) - \Delta_{ij} = b_{ij}^\top h_0$$

is a fixed directional bias.

690 *Proof.* The centered case is immediate from linearity of expectation:

$$\mathbb{E}[\Delta_{ij}(\bar{h})] = \Delta_{ij} + b_{ij}^\top \mathbb{E}[\bar{h}] = \Delta_{ij}.$$

691 The deterministic case follows by substitution. $\square$

692 The proposition does not say that a centered perturbation leaves every *realization* unchanged; it says
 693 there is no systematic first-order push toward one branch across runs. This is exactly the behavior
 694 desired from a generic exploration device. A fixed token or fixed soft prompt may still be useful, but
 695 it acts as a persistent prompt prior rather than as an unbiased exploration mechanism.

696 C.3 Direction-Agnostic Coverage

697 The second requirement is directional coverage. Let D be any zero-mean perturbation distribution
 698 in $\mathbb{R}^d$ with covariance Σ_D and variance budget $\text{tr}(\Sigma_D) \leq \rho^2$. For a unit vector b , the first-order
 699 perturbation energy available along b is $\text{Var}_{h \sim D}(b^\top h) = b^\top \Sigma_D b$. This is the quantity that appears
 700 in the local linear logit surrogate used in the main text.

701 **Theorem 2** (Direction-agnostic coverage). *Define*

$$\mathcal{E}(D) := \min_{\|b\|=1} \text{Var}_{h \sim D}(b^\top h).$$

702 *Then*

$$\mathcal{E}(D) = \lambda_{\min}(\Sigma_D) \leq \frac{\rho^2}{d},$$

703 *with equality if and only if* $\Sigma_D = (\rho^2/d)\mathbf{I}_d$.

704 *Proof.* For any symmetric covariance matrix Σ_D ,

$$\min_{\|b\|=1} b^\top \Sigma_D b = \lambda_{\min}(\Sigma_D),$$

705 so $\mathcal{E}(D) = \lambda_{\min}(\Sigma_D)$. Let the eigenvalues of Σ_D be $\lambda_1, \dots, \lambda_d$. Then

$$\lambda_{\min}(\Sigma_D) \leq \frac{1}{d} \sum_{m=1}^d \lambda_m = \frac{\text{tr}(\Sigma_D)}{d} \leq \frac{\rho^2}{d}.$$

706 Equality holds if and only if all eigenvalues are equal, i.e., $\Sigma_D = (\rho^2/d)\mathbf{I}_d$. $\square$

707 This theorem is a statement about second moments, not about Gaussianity by itself. Many distributions
 708 could realize isotropic covariance. We use a Gaussian in RSP because it is the simplest centered,
 709 rotation-symmetric family that matches the scale of pretrained embeddings. The theorem clarifies the
 710 structural advantage: a fixed prompt has $\Sigma_D = 0$ and therefore no across-run exploration, whereas
 711 any anisotropic perturbation leaves at least one branch-sensitive direction less explored than the
 712 isotropic optimum.

713 **Proposition 5** (Covariance pushforward under the value projection). *Let $\bar{h}$ be a centered random*
 714 *vector with covariance $\sigma_E^2 \mathbf{I}_d$, and let $v = W\bar{h}$ for a matrix $W \in \mathbb{R}^{d_v \times d}$. Then*

$$\text{Cov}(v) = \sigma_E^2 W W^\top.$$

715 *Consequently, for every unit vector $u \in \mathbb{S}^{d_v-1}$,*

$$\sigma_E^2 \lambda_{\min}(W W^\top) \leq \text{Var}(u^\top v) = \sigma_E^2 u^\top W W^\top u \leq \sigma_E^2 \lambda_{\max}(W W^\top).$$

716 *In particular, if W is full rank then the projected perturbation has strictly positive variance in every*
 717 *output-space direction.*

718 *Proof.* Since $v = W\bar{h}$ and $\mathbb{E}[\bar{h}] = 0$,

$$\text{Cov}(v) = \mathbb{E}[v v^\top] = W \mathbb{E}[\bar{h} \bar{h}^\top] W^\top = \sigma_E^2 W W^\top.$$

719 For any unit vector u ,

$$\text{Var}(u^\top v) = u^\top \text{Cov}(v) u = \sigma_E^2 u^\top W W^\top u,$$

720 and the Rayleigh quotient bounds give the stated inequalities. $\square$

D Interpretive Connection to Hopfield Retrieval (Optional)

This appendix is included for completeness and is not load-bearing for any quantitative result in the main text. Section 4.3 of the main paper deliberately uses only standard self-attention identities (the exact decomposition of Eq. (3) and the decay bound of Eq. (4)); the formal claims there do not assume $\mathbf{K} = \mathbf{V}$, weight-tied layers, or the existence of a global energy. The present appendix records the alternative reading—originating from the Modern Hopfield correspondence [Ramsauer et al., 2021]—under which a single attention update with $\mathbf{K} = \mathbf{V}$ becomes one gradient step on a Hopfield energy. This reading provides additional intuition for why a deterministic forward pass can be exploration-poor and has been a useful lens in the related literature, but it is not used in any proof of Section 4.

We restate the Modern Hopfield correspondence in the form used in this interpretive lens. For a key matrix $\mathbf{K} = (\mathbf{k}_1, \dots, \mathbf{k}_n)$ at a fixed query position, the Hopfield energy is

$$E_{\mathbf{K}}(\sigma) := -\beta^{-1} \log \sum_{j=1}^n \exp(\beta \mathbf{k}_j^\top \sigma) + \frac{1}{2} \|\sigma\|^2.$$

A direct calculation gives the gradient

$$\nabla E_{\mathbf{K}}(\sigma) = \sigma - \mathbf{K}^\top \text{softmax}(\beta \mathbf{K} \sigma).$$

Under (A1) $\mathbf{K} = \mathbf{V}$, the attention update with this $\mathbf{K}$ is therefore $\sigma^{\text{new}} = \sigma - \nabla E_{\mathbf{K}}(\sigma) = \mathbf{K}^\top \text{softmax}(\beta \mathbf{K} \sigma)$, which is one gradient step on $E_{\mathbf{K}}$ at unit step size [Ramsauer et al., 2021].

Lemma 3 (Descent inequality). *If $E_{\mathbf{K}}$ is L_E -smooth on a neighborhood of σ , then the unit-step gradient update $\sigma^{\text{new}} := \sigma - \nabla E_{\mathbf{K}}(\sigma)$ satisfies*

$$E_{\mathbf{K}}(\sigma^{\text{new}}) \leq E_{\mathbf{K}}(\sigma) - (1 - \frac{L_E}{2}) \|\nabla E_{\mathbf{K}}(\sigma)\|^2.$$

Proof. The standard L_E -smooth quadratic upper bound [Nesterov, 2018, Lemma 1.2.3] states $E(\sigma + v) \leq E(\sigma) + \nabla E(\sigma)^\top v + (L_E/2) \|v\|^2$ for any v in the smoothness neighborhood. Applying this to $v = -\nabla E(\sigma)$ gives $E(\sigma - \nabla E) \leq E(\sigma) - \|\nabla E\|^2 + (L_E/2) \|\nabla E\|^2 = E(\sigma) - (1 - L_E/2) \|\nabla E\|^2$. $\square$

Theorem 3 (Deterministic descent and sublevel-set invariance). *Assume (A1) $\mathbf{K} = \mathbf{V}$, (A2) the attention component is weight-tied across blocks, and (A3) $E_{\mathbf{K}}$ is L_E -smooth on a neighborhood of the unperturbed trajectory with $L_E < 2$. Then the unit-step iterate $\sigma^{(l+1)} = \sigma^{(l)} - \nabla E_{\mathbf{K}}(\sigma^{(l)})$ is strictly monotone-descending in $E_{\mathbf{K}}$ at every non-stationary point, and the sublevel set $\{\sigma : E_{\mathbf{K}}(\sigma) \leq E_{\mathbf{K}}(\sigma^{(0)})\}$ is forward-invariant.*

Proof. By Lemma 3, $E_{\mathbf{K}}(\sigma^{(l+1)}) \leq E_{\mathbf{K}}(\sigma^{(l)}) - (1 - L_E/2) \|\nabla E_{\mathbf{K}}(\sigma^{(l)})\|^2$. Under (A3), the coefficient $1 - L_E/2 > 0$, so $E_{\mathbf{K}}(\sigma^{(l+1)}) \leq E_{\mathbf{K}}(\sigma^{(l)})$ with strict inequality unless $\nabla E_{\mathbf{K}}(\sigma^{(l)}) = 0$, which proves forward invariance. $\square$

Corollary 3 (Connected-component containment for the unit-step iterate). *Let $S := \{\sigma : E_{\mathbf{K}}(\sigma) \leq E_{\mathbf{K}}(\sigma^{(0)})\}$ and let $C(\sigma^{(0)})$ denote the connected component of S containing $\sigma^{(0)}$. Define*

$$\text{sep}(C(\sigma^{(0)})) := \inf\{\|x - y\| : x \in C(\sigma^{(0)}), y \in S \setminus C(\sigma^{(0)})\},$$

with the convention $\text{sep} = +\infty$ if S is connected. Assume $\text{sep}(C(\sigma^{(0)})) > 0$. Then the unit-step iterate $\sigma^{(l+1)} = \sigma^{(l)} - \nabla E_{\mathbf{K}}(\sigma^{(l)})$ remains in $C(\sigma^{(0)})$ for all $l \geq 0$ whenever

$$\|\nabla E_{\mathbf{K}}(\sigma^{(l)})\| < \text{sep}(C(\sigma^{(0)})) \quad \text{for every } l.$$

Proof. By Theorem 3, every iterate lies in S . Suppose by induction $\sigma^{(l)} \in C(\sigma^{(0)})$. The displacement $\sigma^{(l+1)} - \sigma^{(l)} = -\nabla E_{\mathbf{K}}(\sigma^{(l)})$ has Euclidean length $\|\nabla E_{\mathbf{K}}(\sigma^{(l)})\|$. A jump to a distinct component $C' \subseteq S$ would require this length to be at least $\inf\{\|x - y\| : x \in C(\sigma^{(0)}), y \in C'\} \geq \text{sep}(C(\sigma^{(0)}))$, which the hypothesis excludes. Hence $\sigma^{(l+1)} \in C(\sigma^{(0)})$. $\square$

758 **Positivity of sep.** $\text{sep}(C) > 0$ holds whenever the components of S are spatially separated by a
759 region of strictly higher energy than $E_{\mathbf{K}}(\sigma^{(0)})$, which we treat as a problem-dependent assumption
760 rather than a structural property: it requires that the relevant components of the Hopfield energy do
761 not share closure boundary points at level $E_{\mathbf{K}}(\sigma^{(0)})$. If $\text{sep} = 0$ the corollary becomes vacuous; in
762 that case the trajectory is still in S by Theorem 3, but is not formally trapped in $C(\sigma^{(0)})$.

763 **Caveats and scope.** Theorem 3 depends on three assumptions that are individually mild but jointly
764 idealized for real transformers. (A1) $\mathbf{K} = \mathbf{V}$ collapses the key and value projections; this is the
765 standard Modern Hopfield assumption [Ramsauer et al., 2021] and isolates the energy-iteration aspect
766 of attention from the additional value-projection degree of freedom. (A2) weight-tying treats the
767 attention component as a single repeated map; without it, each block has its own energy and the
768 trajectory does not minimize a single fixed function across blocks. (A3) is a smoothness condition on
769 $E_{\mathbf{K}}$. For the Hopfield energy used here, $\nabla^2 E_{\mathbf{K}}(\sigma) = \mathbf{I} - \beta \mathbf{K}^\top (\text{diag}(p_\sigma) - p_\sigma p_\sigma^\top) \mathbf{K}$ where $p_\sigma =$
770 $\text{softmax}(\beta \mathbf{K} \sigma)$, so the operator-norm bound on $\nabla^2 E_{\mathbf{K}}$ depends on the inverse-temperature β and
771 on $\|\mathbf{K}\|$. (A3) holds whenever $\beta \|\mathbf{K}\|^2$ is sufficiently small, which is the standard regime in which the
772 Hopfield iterate converges in a small number of steps. The theorem therefore characterizes an idealized
773 tendency of attention-only iteration in this regime; it is not a literal property of every transformer
774 block. In particular, sublevel-set confinement follows from Theorem 3, whereas connected-component
775 confinement additionally requires the separation and per-step gradient assumptions of Corollary 3. Its
776 role in the main text is to supply an idealized source of exploration deficit that an injected stochastic
777 perturbation addresses; the empirical link is the cross-layer attention decay observed in Figure 3 and
778 the trajectory diversity gain reported in Section 3.4.

E Proofs for the Transformer-Level Mechanism

This appendix collects the formal proofs underlying Sections 4.3–4.5 of the main text.

E.1 Exact Attention Decomposition and Magnitude Control

Proposition 6 (Attention decomposition). *Let $\alpha_{ij}^{(l)}$ be the attention weights for query position i at layer l , with n real tokens and k random tokens. Define*

$$w_{r,i}^{(l)} := \sum_{j=1}^k \alpha_{i,n+j}^{(l)}$$

and, for $j \in [n]$,

$$\tilde{\alpha}_{ij}^{(l)} := \frac{\alpha_{ij}^{(l)}}{1 - w_{r,i}^{(l)}}.$$

Then

$$\mathbf{x}_i^{(l+1)} = (1 - w_{r,i}^{(l)}) \tilde{\mathbf{x}}_i^{(l+1)} + \boldsymbol{\eta}_i^{(l)},$$

where

$$\tilde{\mathbf{x}}_i^{(l+1)} := \sum_{j=1}^n \tilde{\alpha}_{ij}^{(l)} \mathbf{v}_j^{(l)} \quad \text{and} \quad \boldsymbol{\eta}_i^{(l)} := \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)}.$$

Proof. The attention output is

$$\mathbf{x}_i^{(l+1)} = \sum_{j=1}^n \alpha_{ij}^{(l)} \mathbf{v}_j^{(l)} + \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)}.$$

By definition, $\sum_{j=1}^n \alpha_{ij}^{(l)} = 1 - w_{r,i}^{(l)}$, so

$$\sum_{j=1}^n \alpha_{ij}^{(l)} \mathbf{v}_j^{(l)} = (1 - w_{r,i}^{(l)}) \sum_{j=1}^n \frac{\alpha_{ij}^{(l)}}{1 - w_{r,i}^{(l)}} \mathbf{v}_j^{(l)} = (1 - w_{r,i}^{(l)}) \tilde{\mathbf{x}}_i^{(l+1)}.$$

Substituting this identity gives the decomposition. □

Corollary 4 (Magnitude scales with random-token attention). *For every realization of the random values,*

$$\|\boldsymbol{\eta}_i^{(l)}\| \leq w_{r,i}^{(l)} \max_{j \in [k]} \|\mathbf{v}_{r_j}^{(l)}\|.$$

Proof. By the triangle inequality,

$$\|\boldsymbol{\eta}_i^{(l)}\| = \left\| \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \mathbf{v}_{r_j}^{(l)} \right\| \leq \sum_{j=1}^k \alpha_{i,n+j}^{(l)} \|\mathbf{v}_{r_j}^{(l)}\| \leq w_{r,i}^{(l)} \max_{j \in [k]} \|\mathbf{v}_{r_j}^{(l)}\|.$$

□

Corollary 4 is the only quantitative fact needed in the main text. No independence assumption between attention weights and random keys is required. Once $w_{r,i}^{(l)}$ is small, the random contribution must be small as well.

797 E.2 Attention Decay Bound

798 **Theorem 4** (Attention decay). *The total attention on random tokens satisfies*

$$w_{r,i}^{(l)} \leq \frac{k}{k + n \cdot \exp(\Delta_i^{(l)} / \sqrt{d})}, \quad (8)$$

799 where $\Delta_i^{(l)} := \min_{j \in [n]} \mathbf{q}_i^\top \mathbf{k}_j - \max_{j' \in [k]} \mathbf{q}_i^\top \mathbf{k}_{r_{j'}}$, is the real-vs-random logit gap.

800 *Proof.* Let $s_j = \mathbf{q}_i^\top \mathbf{k}_j / \sqrt{d}$ for real keys and $s_{r_{j'}} = \mathbf{q}_i^\top \mathbf{k}_{r_{j'}} / \sqrt{d}$ for random keys. By definition
801 of $\Delta_i^{(l)}$, each real key satisfies $s_j \geq \Delta_i^{(l)} / \sqrt{d} + s_{r,\max}$ where $s_{r,\max} = \max_{j'} s_{r_{j'}}$. Therefore
802 $\exp(s_j) \geq \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot \exp(s_{r,\max})$ for every real key. Summing over n real keys gives

$$\sum_{j=1}^n \exp(s_j) \geq n \cdot \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot \exp(s_{r,\max}) \geq n \cdot \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot \frac{1}{k} \sum_{j'=1}^k \exp(s_{r_{j'}}).$$

803 Let $R := \sum_{j'=1}^k \exp(s_{r_{j'}})$. Then

$$w_{r,i}^{(l)} = \frac{R}{\sum_{j=1}^n \exp(s_j) + R} \leq \frac{R}{n \cdot \exp(\Delta_i^{(l)} / \sqrt{d}) \cdot R/k + R} = \frac{k}{k + n \cdot \exp(\Delta_i^{(l)} / \sqrt{d})}.$$

804 □

805 **Corollary 5** (Sequence-wise annealing under a fixed logit gap). *For fixed k and $\Delta_i^{(l)}$, the upper*
806 *bound*

$$f(n) = \frac{k}{k + n \exp(\Delta_i^{(l)} / \sqrt{d})}$$

807 *is strictly decreasing in n . Consequently, cache growth alone decreases the maximum perturbation*
808 *mass compatible with that gap.*

809 *Proof.* Differentiate $f(n)$ with respect to n :

$$f'(n) = -\frac{k \exp(\Delta_i^{(l)} / \sqrt{d})}{\left(k + n \exp(\Delta_i^{(l)} / \sqrt{d})\right)^2} < 0.$$

810 Thus $f(n)$ is strictly decreasing. □

811 **Corollary 6** (Vanishing random contribution under annealing). *If $n \rightarrow \infty$ along decoding while k*
812 *and $\Delta_i^{(l)}$ remain fixed, then $w_{r,i}^{(l)} \rightarrow 0$. Moreover, for any fixed realization of the k random values,*

$$\|\boldsymbol{\eta}_i^{(l)}\| \leq w_{r,i}^{(l)} \max_{j \in [k]} \|\mathbf{v}_{r_j}^{(l)}\| \rightarrow 0.$$

813 *Since the implementation samples finite-dimensional vectors and k is finite, the displayed maximum*
814 *is almost surely finite.*

815 *Proof.* Corollary 5 implies $w_{r,i}^{(l)} \rightarrow 0$. The bound on $\|\boldsymbol{\eta}_i^{(l)}\|$ is Corollary 4. Since the multiplier
816 $\max_{j \in [k]} \|\mathbf{v}_{r_j}^{(l)}\|$ is finite for any fixed realization of finitely many sampled values, the product
817 converges to zero. □

818 E.3 Feed-Forward Propagation under Local Lipschitz Control

819 A transformer block applies an attention sublayer followed by a feed-forward sublayer with residual
820 connections:

$$\mathbf{x}_i^{(l,\text{attn})} = \mathbf{x}_i^{(l)} + \text{Attn}^{(l)}(\mathbf{x}_i^{(l)}), \quad \mathbf{x}_i^{(l+1)} = \mathbf{x}_i^{(l,\text{attn})} + F^{(l)}(\mathbf{x}_i^{(l,\text{attn})}).$$

821 Let $\mathbf{x}_i^{(l,\text{attn},\text{RSP})}$ and $\mathbf{x}_i^{(l,\text{attn},\text{base})}$ denote the post-attention states under RSP and the unperturbed forward
822 pass at the same input $\mathbf{x}_i^{(l)}$.

823 **Lemma 4** (Per-block FFN propagation). *The post-attention difference satisfies*

$$\|\mathbf{x}_i^{(l, \text{attn}, \text{RSP})} - \mathbf{x}_i^{(l, \text{attn}, \text{base})}\| \leq \|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|.$$

824 *If, in addition, $F^{(l)}$ is $L^{(l)}$ -Lipschitz on a bounded set containing both $\mathbf{x}_i^{(l, \text{attn}, \text{RSP})}$ and $\mathbf{x}_i^{(l, \text{attn}, \text{base})}$,*
 825 *then*

$$\|\mathbf{x}_i^{(l+1, \text{RSP})} - \mathbf{x}_i^{(l+1, \text{base})}\| \leq (1 + L^{(l)}) (\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|).$$

826 *Proof.* The unperturbed attention sublayer produces $\mathbf{x}_i^{(l, \text{attn}, \text{base})} = \mathbf{x}_i^{(l)} + \tilde{\mathbf{x}}_i^{(l+1)}$, while the RSP
 827 attention sublayer produces $\mathbf{x}_i^{(l, \text{attn}, \text{RSP})} = \mathbf{x}_i^{(l)} + (1 - w_{r,i}^{(l)}) \tilde{\mathbf{x}}_i^{(l+1)} + \boldsymbol{\eta}_i^{(l)}$ by Proposition 6. Subtracting,

$$\mathbf{x}_i^{(l, \text{attn}, \text{RSP})} - \mathbf{x}_i^{(l, \text{attn}, \text{base})} = \boldsymbol{\eta}_i^{(l)} - w_{r,i}^{(l)} \tilde{\mathbf{x}}_i^{(l+1)},$$

828 and the triangle inequality gives the first bound. For the second bound, write the post-FFN difference
 829 as

$$[\mathbf{x}_i^{(l, \text{attn}, \text{RSP})} + F^{(l)}(\mathbf{x}_i^{(l, \text{attn}, \text{RSP})})] - [\mathbf{x}_i^{(l, \text{attn}, \text{base})} + F^{(l)}(\mathbf{x}_i^{(l, \text{attn}, \text{base})})].$$

830 By Lipschitz continuity, $\|F^{(l)}(\mathbf{x}_i^{(l, \text{attn}, \text{RSP})}) - F^{(l)}(\mathbf{x}_i^{(l, \text{attn}, \text{base})})\| \leq L^{(l)} \|\mathbf{x}_i^{(l, \text{attn}, \text{RSP})} - \mathbf{x}_i^{(l, \text{attn}, \text{base})}\|$.
 831 The triangle inequality combined with the first bound gives the second. $\square$

832 **Corollary 7** (Cross-block propagation control). *Define $\Delta^{(l)} := \mathbf{x}_i^{(l, \text{RSP})} - \mathbf{x}_i^{(l, \text{base})}$ with $\Delta^{(0)} = \mathbf{0}$*
 833 *(matched inputs at layer 0), and let $L_{\text{block}}^{(l)}$ denote a common local Lipschitz constant valid for both the*
 834 *baseline block map $B_{\text{base}}^{(l)} : \mathbf{x} \mapsto \mathbf{x} + \text{Attn}^{(l)}(\mathbf{x}) + F^{(l)}(\mathbf{x} + \text{Attn}^{(l)}(\mathbf{x}))$ and the RSP-augmented*
 835 *block map $B_{\text{RSP}}^{(l)}$ (which replaces $\text{Attn}^{(l)}$ with the attention computation that includes the random*
 836 *tokens) on a ball containing both trajectories. Such a constant exists for standard transformer blocks*
 837 *because Attn , LayerNorm , and the FFN sublayer are each locally Lipschitz on any bounded set*
 838 *[Xiong et al., 2020, Kim et al., 2021], and the random-key augmentation only adds a finite number of*
 839 *additional softmax inputs at every layer. Then*

$$\|\Delta^{(l+1)}\| \leq (1 + L_{\text{block}}^{(l)}) \|\Delta^{(l)}\| + (1 + L^{(l)}) (\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|).$$

840 *Iterating across L blocks with $\Delta^{(0)} = \mathbf{0}$ yields*

$$\|\Delta^{(L+1)}\| \leq \sum_{l=0}^L \prod_{l'=l+1}^L (1 + L_{\text{block}}^{(l')}) \cdot (1 + L^{(l)}) (\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|).$$

841 *Proof.* Decompose the cross-block discrepancy at layer $l + 1$ into two contributions:

$$\Delta^{(l+1)} = \underbrace{B_{\text{RSP}}^{(l)}(\mathbf{x}^{(l, \text{RSP})}) - B_{\text{RSP}}^{(l)}(\mathbf{x}^{(l, \text{base})})}_{\text{(I) propagation of prior } \Delta^{(l)}} + \underbrace{B_{\text{RSP}}^{(l)}(\mathbf{x}^{(l, \text{base})}) - B_{\text{base}}^{(l)}(\mathbf{x}^{(l, \text{base})})}_{\text{(II) RSP-induced perturbation injected at block } l},$$

842 where $B_{\text{RSP}}^{(l)}$ replaces $\text{Attn}^{(l)}$ with the RSP-augmented attention. Term (I) is bounded via the block
 843 Lipschitz constant: $\|(I)\| \leq L_{\text{block}}^{(l)} \|\Delta^{(l)}\|$. With residuals included one gets the factor $(1 + L_{\text{block}}^{(l)})$
 844 when also propagating the residual term $\Delta^{(l)}$ itself. Term (II) is exactly the per-block injection from
 845 Lemma 4, bounded by $(1 + L^{(l)}) (\|\boldsymbol{\eta}_i^{(l)}\| + w_{r,i}^{(l)} \|\tilde{\mathbf{x}}_i^{(l+1)}\|)$. Combining via the triangle inequality
 846 gives the recursion. The summed bound follows by unrolling the recursion: with $\Delta^{(0)} = \mathbf{0}$, induction
 847 on L shows that $\|\Delta^{(L+1)}\|$ is at most the displayed sum. $\square$

848 **Boundedness of $L^{(l)}$ for standard architectures.** Modern transformer blocks combine LayerNorm,
 849 GELU or SwiGLU activations, and linear projections. LayerNorm has Lipschitz constant on the
 850 order of $\|\gamma\| / \inf \|x\|$ on every set bounded away from the origin, where γ is the LayerNorm scale
 851 parameter [Xiong et al., 2020]. GELU and SwiGLU are smooth and therefore locally Lipschitz on
 852 bounded activation regions, and linear projections have Lipschitz constant equal to the spectral norm
 853 of the weight matrix. The composed self-attention sublayer has a Lipschitz constant that depends
 854 on the spectral norms of $\mathbf{W}_Q$, $\mathbf{W}_K$, $\mathbf{W}_V$ and the inverse-temperature β as analyzed by Kim et al.
 855 [2021]. On the bounded activation regions traversed during a single forward pass, the per-block
 856 Lipschitz constants in Corollary 7 are therefore finite. For a fixed model depth this yields a finite
 857 cumulative bound on $\|\Delta^{(L+1)}\|$; stronger depth-uniform non-expansion would require additional
 858 contraction assumptions not imposed here. This is the formal sense in which the injected perturbation
 859 stays controlled across a finite forward pass, while the size of each new injection is set by $w_{r,i}^{(l)}$.

860 E.4 Proofs for Output Distribution Shift

861 The main text makes two output-level claims. Proposition 1 establishes pairwise rank randomness,
862 and Proposition 2 upgrades that statement to top- K entry. Both are statements about a local first-order
863 surrogate of the perturbed logits.

864 E.4.1 Autoregressive Formulation

865 In autoregressive decoding, the joint distribution over a sequence $y_{1:T}$ factorizes as

$$\Pr(y_{1:T} \mid x) = \prod_{t=1}^T p(y_t \mid x, y_{<t}),$$

866 where each step is governed by a prefix-conditioned distribution. Any perturbation introduced by
867 RSP affects generation through a sequence of local changes to $p(y_t \mid x, y_{<t})$, rather than a single
868 global distribution.

869 E.4.2 Local Linearization of Logits under RSP

870 Throughout this subsection, $\bar{h} \in \mathbb{R}^d$ denotes *one* centered prompt vector $\bar{h}_l$ from Definition 1, treated
871 as a per-token surrogate. The actual implementation uses a length- L prompt $\mathbf{H} = (h_1, \dots, h_L)$ with
872 L independent draws; the propositions below identify how each individual $\bar{h}_l$ shifts the logits, and
873 the multi-token compounding through softmax-weighted attention is reported empirically in Figure 3
874 and Section 3.4. We model the perturbed logits as a function $z_i(\bar{h})$, assuming local smoothness in a
875 neighborhood of $\bar{h} = 0$:

$$z_i(\bar{h}) = z_i(0) + \nabla_h z_i(0)^\top \bar{h} + r_i(\bar{h}),$$

876 where $r_i(\bar{h}) = O(\|\bar{h}\|^2)$. Defining $a_i := \nabla_h z_i(0)$, we obtain the local affine surrogate

$$z_i^{\text{lin}}(\bar{h}) := z_i + a_i^\top \bar{h}.$$

877 This approximation is valid for sufficiently small $\|\bar{h}\|$ and is used as a first-order surrogate for the
878 direction of perturbation rather than as an exact global model of the transformer logits.

879 For the two propositions below we isolate the only support property that is actually used. Let D
880 denote the centered prompt law. We assume

$$(S1) \quad D(U) > 0 \quad \text{for every nonempty open set } U \subseteq \mathbb{R}^d.$$

881 The Gaussian instantiation in Section 4.2 satisfies (S1), but the proofs below require only this support
882 condition.

883 E.4.3 Proof of Proposition 1 (Stochastic Ranking)

884 *If there exists a token pair (i, j) such that $a_i \neq a_j$, and if the centered prompt law satisfies (S1), then*
885 $0 < \Pr_{\bar{h}}(z_i^{\text{lin}}(\bar{h}) > z_j^{\text{lin}}(\bar{h})) < 1$.

886 *Proof.* Under the local affine surrogate,

$$z_i^{\text{lin}}(\bar{h}) - z_j^{\text{lin}}(\bar{h}) = (z_i - z_j) + (a_i - a_j)^\top \bar{h}.$$

887 Let $b_{ij} := a_i - a_j \neq 0$. The perturbed logit difference

$$\Delta_{ij}(\bar{h}) := (z_i - z_j) + b_{ij}^\top \bar{h}$$

888 vanishes on the affine hyperplane

$$H_0 := \{\bar{h} \in \mathbb{R}^d : (z_i - z_j) + b_{ij}^\top \bar{h} = 0\}.$$

889 Because $b_{ij} \neq 0$, H_0 has codimension one and partitions $\mathbb{R}^d$ into the two nonempty open half-spaces

$$H_+ := \{\bar{h} : (z_i - z_j) + b_{ij}^\top \bar{h} > 0\}, \quad H_- := \{\bar{h} : (z_i - z_j) + b_{ij}^\top \bar{h} < 0\}.$$

890 By (S1), both H_+ and H_- have positive probability, which yields

$$0 < \Pr_{\bar{h}}(\Delta_{ij}(\bar{h}) > 0) < 1.$$

891 □

892 **E.4.4 Proof of Proposition 2 (Top-K Entry)**

893 Let $J_K(s_t)$ be the baseline top- K set at state s_t , and let $i \notin J_K(s_t)$. Define

$$\mathcal{G}_{i,K} := \{\bar{h} \in \mathbb{R}^d : \#\{j \neq i : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h})\} \leq K-1\}.$$

894 If $\mathcal{G}_{i,K}$ contains a nonempty open set and the centered prompt law satisfies (S1), then

$$\Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0.$$

895 *Proof.* Under the local affine surrogate, token i belongs to the surrogate top- K set if and only if at
 896 most $K-1$ tokens outrank it, i.e., if and only if $\bar{h} \in \mathcal{G}_{i,K}$. For each subset S of token indices with
 897 $i \notin S$ and $|S| \leq K-1$, define

$$\mathcal{O}_S := \{\bar{h} : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h}) \text{ for all } j \in S, z_i^{\text{lin}}(\bar{h}) > z_j^{\text{lin}}(\bar{h}) \text{ for all } j \notin S \cup \{i\}\}.$$

898 Each $\mathcal{O}_S$ is an intersection of open half-spaces and is therefore an open subset of $\mathbb{R}^d$. By construction,
 899 every $\bar{h} \in \mathcal{O}_S$ has token i ranked $|S|+1 \leq K$ in the affine surrogate, so $\mathcal{O}_S \subseteq \mathcal{G}_{i,K}$. Hence

$$\bigcup_{S: i \notin S, |S| \leq K-1} \mathcal{O}_S \subseteq \mathcal{G}_{i,K}.$$

900 The reverse inclusion fails only on the measure-zero set of $\bar{h}$ that produce ties, which we do not need:
 901 the assumption that $\mathcal{G}_{i,K}$ contains a nonempty open set is sufficient. Condition (S1) therefore implies

$$\Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0,$$

902 which is exactly the event that token i enters the surrogate top- K set. $\square$

903 **E.4.5 Proof of Corollary 2 (Sample Complexity)**

904 Let $\mu_i = \Pr_{\bar{h}}(\bar{h} \in \mathcal{G}_{i,K}) > 0$. Within N independent RSP samples $\bar{h}_1, \dots, \bar{h}_N$, the probability of at
 905 least one admission is at least $1 - (1 - \mu_i)^N \geq 1 - \exp(-N\mu_i)$.

906 *Proof.* The samples $\{\bar{h}_n\}_{n=1}^N$ are independent, so

$$\Pr(\text{no admission within } N \text{ samples}) = \prod_{n=1}^N \Pr(\bar{h}_n \notin \mathcal{G}_{i,K}) = (1 - \mu_i)^N.$$

907 The complement gives the first inequality. The second inequality is standard: $1 - \mu_i \leq \exp(-\mu_i)$,
 908 hence $(1 - \mu_i)^N \leq \exp(-N\mu_i)$. $\square$

909 **E.4.6 Proof of Proposition 3 (Exploration-meets-multi-path Pass@ N gain)**

910 Fix a problem x and let $p_n(x) \geq p_{\min}(x) > 0$ be the per-run success probability of the n -th RSP run
 911 as in (M1). Under (M1)+(M2),

$$\Pr(\text{Pass@}(N+1) \text{ correct on } x) \geq \max\{p_{\text{base}}(x), 1 - (1 - p_{\min}(x))^N\} \geq 1 - e^{-Np_{\min}(x)}.$$

912 *Proof.* Let A_0 be the event that the deterministic baseline run reaches a trajectory in $\mathcal{T}^*(x)$, and
 913 let A_n ($n = 1, \dots, N$) be the event that the n -th RSP run reaches such a trajectory. Pass@ $(N+1)$
 914 succeeds on x iff at least one of $A_0, A_1, \dots, A_N$ occurs.

915 By assumption $p_{\text{base}}(x) := \Pr(A_0) \in \{0, 1\}$ and $\Pr(A_n) = p_n(x) \geq p_{\min}(x)$ for each $n \geq 1$.
 916 Independence of the N RSP runs (M2) gives

$$\Pr\left(\bigcap_{n=1}^N A_n^c\right) = \prod_{n=1}^N (1 - p_n(x)) \leq \prod_{n=1}^N (1 - p_{\min}(x)) = (1 - p_{\min}(x))^N,$$

917 using $p_n(x) \geq p_{\min}(x)$ in the inequality. Therefore

$$\Pr\left(\bigcup_{n=1}^N A_n\right) = 1 - \Pr\left(\bigcap_{n=1}^N A_n^c\right) \geq 1 - (1 - p_{\min}(x))^N.$$

918 The Pass@ $(N+1)$ event $A_0 \cup A_1 \cup \dots \cup A_N$ has probability at least $\max\{\Pr(A_0), \Pr(\bigcup_{n \geq 1} A_n)\}$,
 919 giving the first claimed inequality. The second inequality follows from the standard bound $1-x \leq e^{-x}$
 920 for $x \in [0, 1]$ applied with $x = p_{\min}(x)$:

$$(1 - p_{\min}(x))^N \leq e^{-Np_{\min}(x)},$$

921 hence $1 - (1 - p_{\min}(x))^N \geq 1 - e^{-Np_{\min}(x)}$. □

922 **E.4.7 Trajectory-Level Compounding**

923 Since autoregressive generation composes conditional distributions across steps, perturbations at
 924 early steps can lead to distinct trajectories. Let $\mathcal{R}_{k,K}(x)$ denote the set of trajectories containing
 925 at least one token outside the baseline top- K set. If RSP increases the accessibility of such tokens
 926 at early steps, then the probability of sampling trajectories in $\mathcal{R}_{k,K}(x)$ increases under repeated
 927 sampling. The key point is temporal compounding: one early top- K entry changes the next decoding
 928 state, which changes the next candidate set, and so on. This is the heuristic bridge from local rank
 929 perturbations to global trajectory diversity that motivates assumption (M1) of Proposition 3.

F Attention Mass Analysis

This appendix formalizes the per-token attention mass reported in Figure 3 and the exclusion criteria applied to the reasoning span and the layer axis. For each of the 500 MATH-500 problems, a fresh RSP $\mathbf{H} \in \mathbb{R}^{L \times d}$ with $L = 10$ is sampled independently, so the reported heatmaps average over both problem variability and RSP-vector variability.

F.1 Per-Token Attention Mass

For each problem, we measure attention on the concatenated sequence [prompt; $\mathbf{H}$; generation], where the generation is the trajectory produced by the same model under matching RSP injection. Let $\alpha_{i,j}^{(\ell,h)}$ denote the post-softmax attention weight at layer ℓ , head h , from query position i to key position j , satisfying $\sum_j \alpha_{i,j}^{(\ell,h)} = 1$ under the causal mask. Let Q and R denote the index sets of question and RSP tokens with sizes $|Q|$ and $|R| = L$, and let H be the number of heads. The per-token attention mass that query i directs to each group at layer ℓ is

$$\text{PTM}_Q[\ell, i] = \frac{1}{|Q|H} \sum_{h=1}^H \sum_{j \in Q} \alpha_{i,j}^{(\ell,h)}, \quad \text{PTM}_R[\ell, i] = \frac{1}{|R|H} \sum_{h=1}^H \sum_{j \in R} \alpha_{i,j}^{(\ell,h)}. \quad (9)$$

Normalization by $|Q|$ and $|R|$ controls for group size, so both quantities express how much attention a single question (resp. RSP) token receives on average under the same softmax denominator. Question tokens thereby serve as a size-matched baseline that absorbs sequence-length effects common to both groups.

F.2 Heatmap Aggregation

We partition the layer indices and the reasoning position indices each into five equal-size bins $\{B_L^{(b)}\}_{b=1}^5$ and $\{B_P^{(b)}\}_{b=1}^5$. For sample s and target group $\bullet \in \{Q, R\}$, the value in cell (b_L, b_P) is

$$M_{\bullet}^{(s)}[b_L, b_P] = \frac{1}{|B_L^{(b_L)}| \cdot |B_P^{(b_P)}|} \sum_{\ell \in B_L^{(b_L)}} \sum_{i \in B_P^{(b_P)}} \text{PTM}_{\bullet}[\ell, i]. \quad (10)$$

Figure 3 reports the sample average of Eq. (10) over the 500 valid samples.

F.3 Excluding the `\boxed{\dots}` Span

The reasoning position range terminates strictly before the final `\boxed{\dots}` span. Chain-of-thought outputs decompose into a *text* (content) component and a *pattern* (template) component that contribute independently to downstream performance [Madaan and Yazdanbakhsh, 2022]. The `\boxed{\dots}` wrapper is a canonical pattern: a short, stereotyped span whose role is to complete the answer format rather than to extend reasoning. Including it would therefore inject a template-driven attention signature unrelated to reasoning dynamics.

F.4 Excluding the Final Two Layers

The layer axis further excludes the last two transformer layers. Two well-documented properties of late layers, none specific to our setup, motivate this choice.

Output-projection specialization. Late-layer hidden states lie in an approximately linear relationship with vocabulary logits and therefore function as a progressive linear readout rather than as representation-transforming blocks [Belrose et al., 2023]. Consistently, late-layer feed-forward modules have been characterized as output-distribution shaping mechanisms [Geva et al., 2021]. Attention in these layers therefore reflects output formation rather than reasoning computation.

Register-token hypothesis. Uninformative patch tokens in Vision Transformers have been shown to be repurposed as *registers*: storage slots that absorb attention in deeper layers [Darcet et al., 2024]. Because RSPs share a similar profile as continuous tokens without pretrained semantic content, a

similar repurposing in late transformer layers cannot be ruled out. Excluding the final two layers removes this potential register-induced confound from the reasoning-phase signal.

The two lines of evidence above identify known late-layer phenomena — linear readout and potential register repurposing — whose attention signatures are orthogonal to reasoning-phase dynamics. Removing the final two layers preempts their contamination of the reasoning-phase signal, yielding a principled, literature-grounded exclusion rather than a model-specific tuning decision.

F.5 Additional Suffix Heatmaps

Figure 5 reports the same per-token RSP attention analysis under suffix injection for Qwen2.5-Math-1.5B-Instruct and LLaMA-3.1-8B-Instruct. For the late-layer reasons discussed above, the pattern does not generalize uniformly across every cell; nevertheless, the overall trend is consistent across both models: question-token attention varies broadly across layers, whereas RSP attention shows the layer-wise and sequence-wise decay predicted by Eq. (4).




Figure 5: Per-token RSP attention mass under suffix injection for the remaining two models (500 MATH-500 samples each). Axes and preprocessing match Figure 3: x-axis is reasoning position binned into five quantiles, y-axis is layer depth binned into five quantiles (top = shallow), color encodes the 500-sample mean per-token attention assigned to RSP tokens, and tokens inside `\boxed{}` spans and the last two layers are excluded.

980 **G GRPO Formulation**

981 For a given prompt q , the behavior policy $\pi_{\theta_{\text{old}}}$ samples a group of G responses $\{o_i\}_{i=1}^G$ with
 982 corresponding rewards $\{R_i\}_{i=1}^G$. The advantage of each response is computed by normalizing the
 983 group-level rewards:

$$\hat{A}_{i,t} = \frac{R_i - \text{mean}(\{R_i\}_{i=1}^G)}{\text{std}(\{R_i\}_{i=1}^G)}. \quad (11)$$

984 The policy is updated via a clipped objective with a KL penalty:

$$\mathcal{J}_{\text{GRPO}}(\theta) = \mathbb{E} \left[\frac{1}{G} \sum_{i=1}^G \frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \left(\min \left(r_{i,t} \hat{A}_{i,t}, \text{clip}(r_{i,t}, 1-\varepsilon, 1+\varepsilon) \hat{A}_{i,t} \right) - \beta D_{\text{KL}}(\pi_{\theta} \parallel \pi_{\text{ref}}) \right) \right], \quad (12)$$

985 where $r_{i,t} = \pi_{\theta}(o_{i,t} \mid q, o_{i,<t}) / \pi_{\theta_{\text{old}}}(o_{i,t} \mid q, o_{i,<t})$ is the importance sampling ratio, ε is the clipping
 986 range, and β controls the strength of the KL penalty against the reference policy π_{ref} .

987 H Broader Impacts

988 This work provides empirical and theoretical analysis of how random embedding injection affects
989 the reasoning behavior of large language models. The contributions are primarily foundational,
990 with potential downstream implications for reinforcement learning-based post-training of reasoning
991 models such as GRPO.

992 **Positive impacts.** RSP-induced trajectory diversity could improve the sample efficiency of RL-
993 based LLM post-training by providing richer learning signals within group-relative advantage estima-
994 tion. This has the potential to reduce compute costs associated with alignment and reasoning-oriented
995 fine-tuning, making such methods more accessible to research groups with limited resources.

996 **Potential concerns.** Methods that expand the effective output support of LLMs may increase the
997 reachability of low-probability tokens, which includes both desirable alternative reasoning paths
998 and potentially unsafe or off-distribution outputs. Practitioners applying RSP in production systems
999 should combine it with existing safety filtering and alignment mechanisms rather than treating it as a
1000 standalone diversity source.

1001 **Limitations on impact scope.** Since this work focuses on analysis rather than deploying a new
1002 model or dataset, the direct societal impact is limited. The broader implications described above are
1003 contingent on future work integrating RSP into applied training pipelines.

1004 NeurIPS Paper Checklist

1005 The checklist is designed to encourage best practices for responsible machine learning research,
1006 addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove
1007 the checklist: **The papers not including the checklist will be desk rejected.** The checklist should
1008 follow the references and follow the (optional) supplemental material. The checklist does NOT count
1009 towards the page limit.

1010 Please read the checklist guidelines carefully for information on how to answer these questions. For
1011 each question in the checklist:

- 1012 • You should answer [Yes], [No], or [N/A].
- 1013 • [N/A] means either that the question is Not Applicable for that particular paper or the
1014 relevant information is Not Available.
- 1015 • Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

1016 **The checklist answers are an integral part of your paper submission.** They are visible to the
1017 reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it
1018 (after eventual revisions) with the final version of your paper, and its final version will be published
1019 with the paper.

1020 The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation.
1021 While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a
1022 proper justification is given (e.g., error bars are not reported because it would be too computationally
1023 expensive” or “we were unable to find the license for the dataset we used”). In general, answering
1024 [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we
1025 acknowledge that the true answer is often more nuanced, so please just use your best judgment and
1026 write a justification to elaborate. All supporting evidence can appear either in the main paper or the
1027 supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification
1028 please point to the section(s) where related material for the question can be found.

1029 IMPORTANT, please:

- 1030 • **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- 1031 • **Keep the checklist subsection headings, questions/answers and guidelines below.**
- 1032 • **Do not modify the questions and only use the provided macros for your answers.**

1033 1. Claims

1034 Question: Do the main claims made in the abstract and introduction accurately reflect the
1035 paper’s contributions and scope?

1036 Answer: **[TODO]**

1037 Justification: **[TODO]**

1038 Guidelines:

- 1039 • The answer [N/A] means that the abstract and introduction do not include the claims
1040 made in the paper.
- 1041 • The abstract and/or introduction should clearly state the claims made, including the
1042 contributions made in the paper and important assumptions and limitations. A [No] or
1043 [N/A] answer to this question will not be perceived well by the reviewers.
- 1044 • The claims made should match theoretical and experimental results, and reflect how
1045 much the results can be expected to generalize to other settings.
- 1046 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
1047 are not attained by the paper.

1048 2. Limitations

1049 Question: Does the paper discuss the limitations of the work performed by the authors?

1050 Answer: **[TODO]**

1051 Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer **[Yes]** if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- 1208 • The paper should disclose whether the full research project required more compute
1209 than the experiments reported in the paper (e.g., preliminary or failed experiments that
1210 didn't make it into the paper).

1211 **9. Code of ethics**

1212 Question: Does the research conducted in the paper conform, in every respect, with the
1213 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

1214 Answer: **[TODO]**

1215 Justification: **[TODO]**

1216 Guidelines:

- 1217 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of
1218 Ethics.
1219 • If the authors answer [No], they should explain the special circumstances that require a
1220 deviation from the Code of Ethics.
1221 • The authors should make sure to preserve anonymity (e.g., if there is a special consid-
1222 eration due to laws or regulations in their jurisdiction).

1223 **10. Broader impacts**

1224 Question: Does the paper discuss both potential positive societal impacts and negative
1225 societal impacts of the work performed?

1226 Answer: **[TODO]**

1227 Justification: **[TODO]**

1228 Guidelines:

- 1229 • The answer [N/A] means that there is no societal impact of the work performed.
1230 • If the authors answer [N/A] or [No], they should explain why their work has no societal
1231 impact or why the paper does not address societal impact.
1232 • Examples of negative societal impacts include potential malicious or unintended uses
1233 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations
1234 (e.g., deployment of technologies that could make decisions that unfairly impact specific
1235 groups), privacy considerations, and security considerations.
1236 • The conference expects that many papers will be foundational research and not tied
1237 to particular applications, let alone deployments. However, if there is a direct path to
1238 any negative applications, the authors should point it out. For example, it is legitimate
1239 to point out that an improvement in the quality of generative models could be used to
1240 generate Deepfakes for disinformation. On the other hand, it is not needed to point out
1241 that a generic algorithm for optimizing neural networks could enable people to train
1242 models that generate Deepfakes faster.
1243 • The authors should consider possible harms that could arise when the technology is
1244 being used as intended and functioning correctly, harms that could arise when the
1245 technology is being used as intended but gives incorrect results, and harms following
1246 from (intentional or unintentional) misuse of the technology.
1247 • If there are negative societal impacts, the authors could also discuss possible mitigation
1248 strategies (e.g., gated release of models, providing defenses in addition to attacks,
1249 mechanisms for monitoring misuse, mechanisms to monitor how a system learns from
1250 feedback over time, improving the efficiency and accessibility of ML).

1251 **11. Safeguards**

1252 Question: Does the paper describe safeguards that have been put in place for responsible
1253 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
1254 image generators, or scraped datasets)?

1255 Answer: **[TODO]**

1256 Justification: **[TODO]**

1257 Guidelines:

- 1258 • The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

1310 Guidelines:

1311 • The answer [N/A] means that the paper does not involve crowdsourcing nor research

1312 with human subjects.

1313 • Including this information in the supplemental material is fine, but if the main contribu-

1314 tion of the paper involves human subjects, then as much detail as possible should be

1315 included in the main paper.

1316 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,

1317 or other labor should be paid at least the minimum wage in the country of the data

1318 collector.

1319 **15. Institutional review board (IRB) approvals or equivalent for research with human**

1320 **subjects**

1321 Question: Does the paper describe potential risks incurred by study participants, whether

1322 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)

1323 approvals (or an equivalent approval/review based on the requirements of your country or

1324 institution) were obtained?

1325 Answer: [TODO]

1326 Justification: [TODO]

1327 Guidelines:

1328 • The answer [N/A] means that the paper does not involve crowdsourcing nor research

1329 with human subjects.

1330 • Depending on the country in which research is conducted, IRB approval (or equivalent)

1331 may be required for any human subjects research. If you obtained IRB approval, you

1332 should clearly state this in the paper.

1333 • We recognize that the procedures for this may vary significantly between institutions

1334 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the

1335 guidelines for their institution.

1336 • For initial submissions, do not include any information that would break anonymity (if

1337 applicable), such as the institution conducting the review.

1338 **16. Declaration of LLM usage**

1339 Question: Does the paper describe the usage of LLMs if it is an important, original, or

1340 non-standard component of the core methods in this research? Note that if the LLM is used

1341 only for writing, editing, or formatting purposes and does *not* impact the core methodology,

1342 scientific rigor, or originality of the research, declaration is not required.

1343 Answer: [TODO]

1344 Justification: [TODO]

1345 Guidelines:

1346 • The answer [N/A] means that the core method development in this research does not

1347 involve LLMs as any important, original, or non-standard components.

1348 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not

1349 be described.