
From Noise to Diversity: Random Embedding Injection in LLM Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

Recent *soft prompt* research has tried to improve reasoning by inserting trained vectors into LLM inputs, yet whether the gain comes from the *learned content* or from the *act of injection itself* has not been carefully separated. We study Random Soft Prompts (RSPs), which drop the training step entirely and append a freshly drawn sequence of random embedding vectors to the input. Each RSP vector is sampled from an isotropic Gaussian fitted to the entrywise mean and variance of the pretrained embedding table; the sequence carries no learned content, and yet matches the accuracy of trained soft prompts on math reasoning benchmarks. The mechanism unfolds in two stages: because attention has to absorb a never-seen-before random position, the distribution over the first few generated tokens flattens and reasoning trajectories *branch*, and as generation continues this influence dilutes naturally so the response *converges toward the goal*. We show that during inference RSPs lift early-stage token diversity and, combined with temperature sampling, widen Pass@ N , the probability that at least one out of N attempts is correct. Beyond inference, we carry the same effect into DAPO training and demonstrate practical gains. Our contributions are: (i) a new finding that purely training-free random embedding injection alone drives latent reasoning; (ii) a theoretical and empirical validation of the underlying mechanism; and (iii) an extension from inference to training.

1 Introduction

As large language models (LLMs) scale, full fine-tuning becomes prohibitively expensive, motivating parameter-efficient methods that adapt frozen models with few extra parameters [Hu et al., 2022, Houshy et al., 2019]. Among these, *soft prompts* prepend or insert learnable continuous embeddings into the input, steering behavior through attention [Li and Liang, 2021, Lester et al., 2021]. The paradigm is actively adopted for LLM mathematical reasoning [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Hao et al., 2025, Zhang et al., 2026]: despite diverse designs, these methods share a common structure of injecting out-of-vocabulary continuous embeddings and *optimizing* them for downstream performance. Such optimization is task- and dataset-specific, and prior work generally interprets the resulting gains as a consequence of the trained values of these embeddings.

However, we observed that optimization is not necessary for these effects: injecting untrained random embeddings can also lift downstream performance. For instance, applying a chat template to a base model not fine-tuned with such format causes a prompt-format mismatch that degrades reasoning, but appending random embeddings to such mismatched inputs mitigates this degradation. If simple noise were merely shaking the model, accuracy should worsen instead. Therefore, the injection itself, independent of any learned content, alters model behavior in non-trivial ways. Existing studies focus on end-task accuracy and have not separated the effects of injection from those of optimization, leaving how embedding injection reshapes the output distribution insufficiently explored.




Figure 1: Conceptual overview of RSP-induced trajectory diversity. The hidden states shown are the final-layer outputs that drive the next-token distribution. (a) Early decoding: RSP perturbation reaches a *branching state* (red; formalized in §3.1), whose top- K differs from the no-RSP baseline and induces a diverse output distribution. (b) Later decoding: as the KV cache grows, the bound on RSP attention mass narrows (Theorem 1); branching events become rare and the hidden state (blue) commits to the branch already opened. (c) Across rollouts, the trajectories opened by early RSP perturbation accumulate into independent paths that converge to goal-directed completions.

We analyze the effect of training-free Random Soft Prompts (RSPs) on LLM reasoning. RSPs are continuous vectors drawn from an isotropic Gaussian fitted to the entrywise mean and variance of the pretrained embedding table, with each rollout receiving an *independent* draw that enables broader exploration. The benefit appears most strongly during early decoding where wider exploration is needed: a single RSP injection induces hidden states that *branch* into diverse trajectories, and as decoding deepens these trajectories *stabilize* into goal-directed completions. Empirically, RSPs lift early-stage entropy and, combined with temperature sampling, expand the Pass@ N trajectory diversity, and we further show that the same exploration benefit carries from inference into DAPO [Yu et al., 2025] training. Our contributions are: (i) the new finding that *training-free* random embedding injection alone drives latent reasoning capabilities, suggesting much of the effect attributed to learned soft prompts may come from the *act of injection itself*; (ii) theoretical and empirical validation of the mechanism; and (iii) an extension from inference to training, demonstrating practical applicability.

2 Background

2.1 Soft Prompt Methods for LLM Reasoning

Prior *soft prompt* methods all share a common structure: they inject out-of-vocabulary continuous embeddings and *optimize* them under task-specific objectives. Because evaluation is anchored to end-task accuracy, the contributions of “injection itself” and “optimization” are not separated, and how plain embedding injection reshapes model behavior remains an open question. As parameter-efficient alternatives to full fine-tuning, Prefix-Tuning [Li and Liang, 2021], Prompt Tuning [Lester et al., 2021], and P-tuning [Liu et al., 2024] reached near-fine-tuning performance using learnable continuous vectors, after which LLM-reasoning-specific variants quickly followed. TTSV [Kang et al., 2026] optimizes *prefix* embeddings at test time via trajectory entropy minimization; SoftCoT [Xu et al., 2025] replaces explicit chain-of-thought tokens with soft tokens generated by an assistant model; LTPO [Ye et al., 2026] optimizes the soft prompt per problem to maximize its own confidence on that problem; COCONUT [Hao et al., 2025] feeds the model’s own hidden states back as input embeddings; MemGen [Zhang et al., 2026] uses embedding vectors as experience memory; and Pause tokens [Goyal et al., 2024] insert learnable tokens to provide additional computation steps.

2.2 Noise Injection in Neural Networks

Neural networks have incorporated noise at several stages and locations. During training, dropout [Srivastava et al., 2014] randomly deactivates *hidden activations*, and NEFTune [Jain et al., 2024] *adds* Gaussian noise to *input embeddings* during instruction fine-tuning to improve instruction following. At inference time, randomized smoothing [Cohen et al., 2019] injects Gaussian noise into *raw inputs* to obtain certified robustness, and SmoothLLM [Robey et al., 2025] perturbs prompts at

71 the *discrete character* level to defend against jailbreaking. In RL, NoisyNet [Fortunato et al., 2018]
 72 adds learnable, trained noise to *network weights* to aid exploration.

73 RSP appends noise to *embeddings*, and differs from these in three respects. (1) It operates without
 74 any training, purely at inference, and does not modify model parameters (unlike dropout, NEFTune,
 75 and NoisyNet). (2) It preserves the original input and appends the embedding within a single forward
 76 pass, so it does not require the “multiple noisy passes plus aggregation” that robustness methods
 77 need. (3) It uses vectors drawn directly from the pretrained embedding statistics, with no learned
 78 parameters at all (NoisyNet, by contrast, learns its noise parameters).

79 3 Random Soft Prompts and Why They Work

80 RSP enters the hidden state through self-attention rather than directly, so within one attention head its
 81 instantaneous influence reduces to one scalar: the *attention mass* w on RSP tokens. We track how
 82 w evolves during decoding. Early on, w is large enough to branch the hidden state into multiple
 83 directions (§3.2); as the KV cache grows, its upper bound narrows and the response converges (§3.3).
 84 Randomness supplies *independent* perturbations across rollouts so each opens a fresh branch (further
 85 proofs in Appendix D).

86 3.1 Random Soft Prompt

87 We start with the object being analyzed. Let $\mathbf{W}_E \in \mathbb{R}^{V \times d}$ denote the pretrained token-embedding
 88 matrix, with V the vocabulary size and d the hidden dimension. Write its entrywise statistics as
 89 $\mu_E := \frac{1}{V \cdot d} \sum_{i,k} [\mathbf{W}_E]_{i,k}$ (the mean over all $V \cdot d$ entries) and $\sigma_E := \text{std}(\mathbf{W}_E)$. An RSP of length
 90 L is a sequence of continuous vectors $\mathbf{H} = [h_1, \dots, h_L] \in \mathbb{R}^{L \times d}$ drawn independently from

$$h_j \sim \mathcal{N}(\mu_E \mathbf{1}_d, \sigma_E^2 \mathbf{I}_d), \quad j = 1, \dots, L. \quad (1)$$

91 The centered form $\bar{h}_j := h_j - \mu_E \mathbf{1}_d \sim \mathcal{N}(0, \sigma_E^2 \mathbf{I}_d)$ is a zero-mean isotropic Gaussian. Letting
 92 $\mathbf{X} \in \mathbb{R}^{T \times d}$ denote the input token-embedding sequence (the T rows of $\mathbf{W}_E$ corresponding to the
 93 prompt tokens), the main text uses the *suffix* form $[\mathbf{X}; \mathbf{H}] \in \mathbb{R}^{(T+L) \times d}$ as the default; other positions
 94 (*prefix*, *infix* [Kang et al., 2026, Xu et al., 2025, Ye et al., 2026, Zhang et al., 2026]) are reported as
 95 ablations in §4.2 and Appendix A. RSP is not a learnable parameter, so it receives no gradient and is
 96 freshly drawn for each rollout.

97 Call decoding step t a *branching event* and the last-layer hidden state $\mathbf{s}^{(t)}$ a *branching state* when
 98 the LM-head top- K differs from that of the no-RSP baseline $\bar{\mathbf{s}}^{(t)}$ (Figure 1, red).¹ §3.2 asks *when* a
 99 branching event is reachable, and §3.3 asks *how* that reachability fades.

100 3.2 Exploration: RSP strengthens early-stage exploration

101 Inserting a random prompt concentrates attention on it, and that concentration is what amplifies
 102 exploration. The effect of a single injection on the hidden state raises two questions — *how much*
 103 and *where* — decided respectively by the scalar w inside one attention head and by the distribution of
 104 $\mathbf{H}$. The isotropic Gaussian addresses both at once.

105 Consider one attention head in self-attention layer ℓ . At query position i , the query attends $n \geq 1$
 106 unmasked real tokens together with $L \geq 1$ RSP tokens, all attention logits finite. Write the attention
 107 weights as $\alpha_{ij}^{(\ell)}$ and the value vectors on the real and RSP sides as $\mathbf{v}_j^{(\ell)}, \mathbf{v}_{r_j}^{(\ell)}$. Defining the total
 108 attention mass on random tokens $w_{r,i}^{(\ell)} := \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)}$, the real-token mass $1 - w_{r,i}^{(\ell)}$ is positive
 109 and the head output $\mathbf{o}_i^{(\ell)}$ splits *exactly* into a renormalized real-token term $\tilde{\mathbf{o}}_i^{(\ell)}$ and an RSP-induced
 110 contribution $\boldsymbol{\eta}_i^{(\ell)}$:

$$\mathbf{o}_i^{(\ell)} = (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{o}}_i^{(\ell)} + \boldsymbol{\eta}_i^{(\ell)}, \quad \tilde{\mathbf{o}}_i^{(\ell)} := \sum_{j=1}^n \frac{\alpha_{ij}^{(\ell)}}{1 - w_{r,i}^{(\ell)}} \mathbf{v}_j^{(\ell)}, \quad \boldsymbol{\eta}_i^{(\ell)} := \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)}. \quad (2)$$

¹ Italic h_j : RSP input space; bold $\mathbf{s}^{(t)}$: hidden state. K is the sampler’s candidate-set size (top- K cutoff or top- p nucleus);
 t for decoding step, ℓ for self-attention layer.

111 *Derivation of Eq. (2).* Write the head output as $\sum_{j=1}^n \alpha_{ij}^{(\ell)} \mathbf{v}_j^{(\ell)} + \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)}$, then factor
 112 $1 - w_{r,i}^{(\ell)} = \sum_{j=1}^n \alpha_{ij}^{(\ell)} > 0$ out of the real-token sum. No approximation beyond the softmax-
 113 normalization identity is used. $\square$

114 This single quantity $w_{r,i}^{(\ell)}$ controls both the attenuation ratio $1 - w_{r,i}^{(\ell)}$ on the real signal and the upper
 115 bound on the random contribution $\|\boldsymbol{\eta}_i^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_j \|\mathbf{v}_{r_j}^{(\ell)}\|$. The *magnitude* of RSP-induced
 116 variation thus reduces to one number in $[0, 1]$. Early in decoding the KV cache is short and this value
 117 is non-negligible, so a single RSP draw can perturb next-token decisions.

118 The scalar w controls *how much*; *where* is decided by the distribution of $\mathbf{H}$. Isolating one position
 119 — $\bar{h} \in \mathbb{R}^d$ the i.i.d. marginal of any $\bar{h}_j$ in Eq. (1), others fixed at 0 — let $z_i(\bar{h})$ be the under-RSP
 120 output logit at vocab token i and write the vocab-logit gap $\Delta_{ij}(\bar{h}) := z_i(\bar{h}) - z_j(\bar{h})$ ($\bar{h} = 0$ is the
 121 centered mean, not the no-RSP state; Appendix C). The transformer’s logit map is non-linear in $\bar{h}$,
 122 but a first-order Taylor expansion around $\bar{h} = 0$ gives the local surrogate $\Delta_{ij}(\bar{h}) \approx \Delta_{ij}(0) + b_{ij}^\top \bar{h}$,
 123 with $b_{ij} := \nabla_{\bar{h}}(z_i - z_j)|_{\bar{h}=0} \in \mathbb{R}^d$ the gradient direction along which i, j swap rank (not unit-norm
 124 in general). Because RSP is training-free, we cannot know in advance which b_{ij} opens a useful
 125 branch, so the distribution must avoid systematically under-perturbing any direction. A deterministic
 126 injection pins to one direction; vocabulary sampling inherits the embedding table’s anisotropy. The
 127 remaining design — maximize the worst-case directional variance at fixed budget — is uniquely
 128 solved by isotropy.

129 **Proposition 1** (Maximin directional coverage). *Let $\mathcal{D}_\rho$ denote the family of zero-mean distributions*
 130 *on $\mathbb{R}^d$ whose covariance satisfies $\text{tr}(\Sigma_D) \leq \rho^2$. For each $D \in \mathcal{D}_\rho$, $\min_{\|b\|=1} \text{Var}_{h \sim D}(b^\top h) =$*
 131 *$\lambda_{\min}(\Sigma_D)$, and $\sup_{D \in \mathcal{D}_\rho} \lambda_{\min}(\Sigma_D) = \rho^2/d$, attained iff $\Sigma_D = (\rho^2/d)\mathbf{I}_d$.*

132 This is a *design criterion* for the prompt law, not a guarantee of correctness; correctness enters
 133 through the task-side $p_{\min}(x)$ assumption of §3.3 and independent resampling.²

134 Gaussian RSP attains the equality with $\Sigma = \sigma_E^2 \mathbf{I}_d$ ($\rho^2 = d\sigma_E^2$) and adds *full support* on $\mathbb{R}^d$
 135 (Appendix C). Isotropy excludes no direction; full support gives positive probability to the open
 136 branching-event region (§3.1) whenever it is non-empty (Appendix D.3.3). Monotone temperature
 137 preserves ranking and cannot reach this region, so RSP’s *rank-changing* branching follows from the
 138 two properties together, splitting the hidden state into multiple branching states early in reasoning.

139 3.3 Annealing: KV cache growth dilutes RSP attention

140 The early branching stabilizes as decoding proceeds: the upper bound on $w_{r,i}^{(\ell)}$ narrows along the
 141 sequence axis with autoregressive cache growth. Each step adds one real-token term while RSP terms
 142 stay fixed; at frozen attention-logit gap, this asymmetry yields the following bound.

143 **Theorem 1** (Attention-mass decay under KV cache growth). *Consider a single attention head with*
 144 *per-head key dimension d_{head} , where query $\mathbf{q}_i$ at position i attends $n \geq 1$ real keys $\mathbf{k}_j$ and $L \geq 1$*
 145 *RSP keys $\mathbf{k}_{r_{j'}}$ under finite logits. Define the pre-scale attention-logit gap $\Delta_i := \min_{j \in [n]} \mathbf{q}_i^\top \mathbf{k}_j -$*
 146 *$\max_{j' \in [L]} \mathbf{q}_i^\top \mathbf{k}_{r_{j'}}$ (distinct from the vocab-logit gap Δ_{ij} of §3.2). In scaled dot-product attention,*

$$w_{r,i} \leq \frac{L}{L + n \exp(\Delta_i / \sqrt{d_{\text{head}}})},$$

147 *and the right-hand side is strictly decreasing in n and tends to 0 at fixed L and Δ_i .*

148 *Proof.* Let $s_j := \mathbf{q}_i^\top \mathbf{k}_j / \sqrt{d_{\text{head}}}$, $s_{r_{j'}} := \mathbf{q}_i^\top \mathbf{k}_{r_{j'}} / \sqrt{d_{\text{head}}}$, $s_{r,\max} := \max_{j'} s_{r_{j'}}$, and $R :=$
 149 $\sum_{j'=1}^L \exp(s_{r_{j'}})$. The gap definition gives $s_j \geq \Delta_i / \sqrt{d_{\text{head}}} + s_{r,\max}$ for every real j , and
 150 $\exp(s_{r,\max}) \geq R/L$ holds because the max dominates the average; combining the two and sum-
 151 ming over the n real keys yields $\sum_{j=1}^n \exp(s_j) \geq (n/L) \exp(\Delta_i / \sqrt{d_{\text{head}}}) R$. Substituting into

²The maximin is stated in input space; isotropy is not claimed to survive transformer layers. Any logit direction pulls back via the model Jacobian to an input-space b , and isotropy guarantees only that no such pulled-back direction has zero projected variance.

152 $w_{r,i} = R/(\sum_j \exp(s_j) + R)$ gives

$$w_{r,i} \leq \frac{R}{(n/L) \exp(\Delta_i/\sqrt{d_{\text{head}}}) R + R} = \frac{L}{L + n \exp(\Delta_i/\sqrt{d_{\text{head}}})}.$$

153 The derivative of the right-hand side in n is $-L \exp(\Delta_i/\sqrt{d_{\text{head}}})/(L + n \exp(\Delta_i/\sqrt{d_{\text{head}}}))^2 < 0$, so
 154 the bound is strictly decreasing and tends to 0 as $n \rightarrow \infty$. $\square$

155 Applied at each layer ℓ to $w_{r,i}^{(\ell)}$ with gap $\Delta_i^{(\ell)}$, L caps the maximum influence while n grows each
 156 step. The theorem guarantees only sequence-axis attenuation at fixed (layer, query, gap); the gap
 157 itself sharpens at deeper layers, but that effect is empirical and outside its scope. The accurate reading
 158 is “the attainable upper bound at a fixed gap decreases monotonically.” Because Eq. (2) ties the
 159 random contribution to the same scalar, branching-event reachability inherits this envelope — as the
 160 cache grows, events become rare and the response commits to the branch already opened, an implicit
 161 explore-then-exploit. Figure 2 confirms both axes.




Figure 2: Per-token attention mass on Qwen2.5-Math-7B (*suffix*, 500 MATH-500 problems, independent RSP per problem). (a) RSP-token attention decreases along the reasoning-position axis, matching Theorem 1’s KV cache-growth bound; the additional layer-axis decrease is an empirical phenomenon (deeper layers sharpen the real-vs-random gap) outside Theorem 1’s scope. (b) Question-token attention stays comparatively uniform. Reported quantity: per-token mass of Appendix E, equal to $w_{r,i}^{(\ell)}/L$ averaged over heads and samples.

162 Lifting single-shot branching to task accuracy across N rollouts requires two assumptions: (M1)
 163 one execution reaches a correct trajectory with marginal probability $p_{\min}(x) > 0$, and (M2) the
 164 N executions are mutually independent. Under both, $\Pr(\text{Pass}@N \text{ on } x) \geq 1 - (1 - p_{\min}(x))^N$
 165 (Appendix D.3.4). §3.2’s isotropy and full support deliver (M1); *independent resampling* delivers
 166 (M2). Sharing one RSP breaks (M2) and removes the accumulation — a signature that §4 verifies
 167 empirically.

168 4 Empirical Validation

169 The mechanism of §3 runs from early *branching* through KV cache-growth narrowing to $\text{Pass}@N$
 170 accumulation under independent resampling. We verify this flow across three scenes. First, §4.2 tests
 171 whether RSPs preserve baseline accuracy and even recover it in some settings. Next, §4.3 examines
 172 whether the predicted early diversification and later stabilization appear in entropy and attention
 173 signals. Finally, §4.4 asks whether the $\text{Pass}@N$ gain depends on *independence* of the RSP draw.

174 4.1 Experimental Setup

175 We evaluate on three mathematical reasoning benchmarks, MATH-500 [Lightman et al., 2024],
 176 GSM8K [Cobbe et al., 2021], and AIME24, using LLaMA-3.1-8B-Instruct [Grattafiori et al., 2024]
 177 and the Qwen2.5-Math model family [Yang et al., 2024] across instruct models, base models, and
 178 base models with mismatched chat templates. All experiments use greedy decoding to isolate the
 179 effect of RSPs from sampling stochasticity. RSPs follow the definition of §3.1, using the default *suffix*
 180 position with a freshly drawn $\mathbf{H}$ per batch. Since Theorem 1 identifies the RSP length L as the design
 181 parameter capping the maximum perturbation influence, we ablate $L \in \{10, 15, 20\}$ per model to
 182 identify an appropriate perturbation strength. Tables 1 and 2 report results at the per-model selected
 183 length. The mechanism analyses in §4.3–§4.4 use a fixed setting ($L=10$, *suffix*), which insulates
 184 them from this selection effect. Full experimental details and per-position results are in Appendix A.

Table 1: Accuracy (%) across model configurations and benchmarks. Values in parentheses indicate the difference from the baseline. Full per-position breakdown (including *infix*) is in Appendix A.

Model	MATH-500			GSM8K			AIME24		
	Baseline	RSP(<i>prefix</i>)	RSP(<i>suffix</i>)	Baseline	RSP(<i>prefix</i>)	RSP(<i>suffix</i>)	Baseline	RSP(<i>prefix</i>)	RSP(<i>suffix</i>)
<i>Instruct Models</i>									
LLaMA-3.1-8B-Inst	52.20	45.00 (−7.2)	49.40 (−2.8)	85.75	76.35 (−9.4)	86.73 (+1.0)	6.67	3.33 (−3.3)	10.00 (+3.3)
Qwen2.5-Math-7B-Inst	83.20	81.40 (−1.8)	83.40 (+0.2)	95.45	94.92 (−0.5)	95.68 (+0.2)	13.33	13.33 (+0.0)	16.67 (+3.3)
Qwen2.5-Math-1.5B-Inst	73.00	74.80 (+1.8)	74.20 (+1.2)	85.29	85.29 (+0.0)	84.69 (−0.6)	10.00	13.33 (+3.3)	20.00 (+10.0)
<i>Base Models (Raw Text)</i>									
Qwen2.5-Math-7B	72.20	71.60 (−0.6)	55.60 (−16.6)	84.15	84.31 (+0.2)	54.13 (−30.0)	6.67	16.67 (+10.0)	13.33 (+6.7)
Qwen2.5-Math-1.5B	61.40	64.40 (+3.0)	54.60 (−6.8)	77.86	74.30 (−3.6)	58.61 (−19.3)	16.67	10.00 (−6.7)	3.33 (−13.3)
<i>Base Models + ChatML (Format Mismatch)</i>									
Qwen2.5-Math-7B	52.20	59.20 (+7.0)	70.40 (+18.2)	58.30	56.41 (−1.9)	75.97 (+17.7)	23.33	3.33 (−20.0)	23.33 (+0.0)
Qwen2.5-Math-1.5B	34.40	63.40 (+29.0)	51.00 (+16.6)	37.45	67.02 (+29.6)	50.64 (+13.2)	13.33	20.00 (+6.7)	13.33 (+0.0)

4.2 How Does Untrained RSP Compare to Baselines & Optimized Soft Prompts?

We first examine how injecting random embeddings affects reasoning performance. Table 1 reports accuracy at the default *suffix* position alongside *prefix* as a position ablation.

On instruct models RSP stays within a few percentage points of the baseline. The most striking pattern is format-mismatch recovery: when a ChatML template is applied to a base model not trained for that format, *Suffix* recovers +18.2/ +17.7 pp on Qwen2.5-Math-7B (MATH-500/GSM8K) and *prefix* recovers up to +29.0/ +29.6 pp on Qwen2.5-Math-1.5B — if simple noise were responsible, accuracy should have dropped further. In the base raw-text setting the picture inverts: *suffix* incurs sharp losses (−16.6/ −30.0 pp on Qwen2.5-Math-7B) while *prefix* changes accuracy less, suggesting base models treat EOS-adjacent signals as a derailment trigger that *prefix* avoids.

RSP’s *direct accuracy gains* concentrate in misaligned-input recovery and stay within $\pm$ a few percentage points on well-aligned instruct models. The effects this paper centers on are the token-, trajectory-, and outcome-level diversity in §4.4, with accuracy recovery as a secondary outcome.

Table 2: Comparison with prior soft prompt methods (TTSV, SoftCoT, LTPO) under unified evaluation. Baseline and RSP values are from Table 1. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Benchmark	Baseline	RSP	TTSV	SoftCoT	LTPO
Qwen2.5-Math-7B-Instruct	MATH-500	83.20	84.20	<u>83.80</u>	82.80	83.00
	GSM8K	95.45	95.68	95.30	95.00	94.84
	AIME24	13.33	<u>16.67</u>	23.30	<u>16.67</u>	13.33
Qwen2.5-Math-1.5B-Instruct	MATH-500	73.00	<u>75.20</u>	<u>75.20</u>	76.00	72.40
	GSM8K	85.29	85.75	<u>85.70</u>	84.46	84.53
	AIME24	<u>10.00</u>	20.00	6.70	6.67	<u>10.00</u>

We further compare RSP with TTSV, SoftCoT, and LTPO — prior soft prompt methods that optimize embeddings under different objectives. Table 2 reports a unified evaluation under the same fallback-based answer-extraction pipeline (Appendix A.2). RSP does not consistently dominate the trained methods — TTSV, for instance, beats RSP on AIME24 with Qwen2.5-Math-7B-Instruct (23.30% vs 16.67%). What we want to highlight is precisely the converse claim: even without training, RSP reaches a comparable range, and that fact is itself this paper’s main evidence that part of the soft-prompt effect comes from the *structure of injection* rather than the learned content.

4.3 How Does RSP Affect the Output Distribution?

We next examine how RSP injection reshapes the output distribution. The setting is Qwen2.5-Math-1.5B-Instruct on MATH-500, and we measure per-token entropy, top-1 probability, and varentropy across generation steps. Figure 3 compares these metrics for the first 5% of generation steps across RSP variants and prior soft prompt methods (LTPO, SoftCoT, TTSV); full curves are in Appendix A.4.




Figure 3: Mean entropy, top-1 probability, and varentropy during the first 5% of generation steps (Qwen2.5-Math-1.5B-Instruct, MATH-500). Shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: prior soft prompt methods (LTPO, SoftCoT, TTSV).

210 Latent prompt methods including RSP share the same signature: *a rise in early-generation entropy*
 211 *followed by convergence to the baseline later in generation*. A natural mechanistic explanation is
 212 the out-of-vocabulary (OOV) effect — when continuous embeddings outside the vocabulary enter the
 213 input, the model has to attend to a never-before-seen key position on top of its familiar token
 214 distribution, redistributing probability mass across multiple candidates within the early next-token
 215 distribution. All three RSP injection positions show early-stage entropy $\uparrow$, top-1 probability $\downarrow$, and
 216 varentropy $\uparrow$, jointly consistent with multimodal distributions [Ahmed et al., 2026]. The change is
 217 confined to the early stage and all three metrics return to baseline in the middle and final portions,
 218 consistent with the *early influence with automatic KV cache-growth decay* pattern predicted in §3.3.

219 Prior methods trained under different objectives (LTPO, SoftCoT, TTSV) share this signature. Even
 220 TTSV, whose reward explicitly *lowers* mean reasoning-trajectory entropy, shows early-stage entropy
 221 comparable to the baseline. The entropy value itself is not a quality score — higher entropy is not
 222 always better; rather, the signature reveals a *shared mechanism* in which OOV embedding injection
 223 produces the same fingerprint independently of the optimization objective, supporting the hypothesis
 224 that the key driver of latent reasoning is the *act of injection*, not the learned content.

225 4.4 Does RSP Induce Trajectory Diversity?

226 Section 4.3 empirically established that RSP reshapes the early-stage output distribution. We now
 227 examine whether this shift translates into reasoning diversity at three complementary levels of analysis:
 228 token rankings, semantic content of trajectories, and task-level outcomes. The three analyses converge
 229 on a directional picture: RSP’s first-token perturbation propagates into semantically more diverse
 230 reasoning trajectories, which in turn yields greater output diversity at the task level. All three analyses
 231 share the setting: Qwen2.5-Math-1.5B-Instruct, MATH-500, *suffix* injection, and $L = 10$.

232 **Token-level: distribution beyond temperature.**

233 We quantify how much RSP shifts the first-token distribution — the critical
 234 forking point for reasoning trajectories [Wang et al., 2025]. Using the baseline at
 235 $\tau=1.0$ as reference, we compare baselines at $\tau=2.0, 3.0$ (16 samples per problem)
 236 against RSP at $\tau=1.0$ (averaged over 16 seeds) on three metrics: Spearman rank correlation ρ to detect rank reorganization, the
 237 probability mass placed outside the baseline’s top-10 to measure support expansion,
 238 and Jensen–Shannon divergence to quantify the overall magnitude of distributional change (formal
 239 definitions in Appendix B). Table 3 shows the contrast: RSP’s distributional shift falls between $\tau=2.0$
 240 and $\tau=3.0$ in magnitude, but the mechanism differs — temperature is a monotone transform that
 241 preserves ranking ($\rho = 1$ at any τ), while RSP partially preserves but reranks ($\rho = 0.709$). The
 242 Pass@1 row exposes the cost: matching RSP’s magnitude with temperature toward $\tau=3.0$ collapses
 243 accuracy to 1.42%, whereas RSP preserves 73.30%. Thus, RSP produces a fundamentally different

Table 3: First-token distribution metrics measured against the baseline at $\tau=1.0$, comparing baselines at higher temperatures ($\tau=2.0, 3.0$) with RSP at $\tau=1.0$. RSP metrics are averaged over 16 random seeds; Acc reports mean $\pm$ std across 16 samples per problem.

Metric	$\tau=2.0$	$\tau=3.0$	RSP
Spearman ρ	1.000	1.000	0.709
Mass outside top-10	5.18%	29.11%	21.86%
JS divergence	0.060	0.218	0.131
Acc (%)	20.77 $\pm$ 1.47	1.42 $\pm$ 0.35	73.30 $\pm$ 1.07

250 perturbation than temperature scaling. Moreover, RSP’s accuracy variability across 16 seeds at $\tau=1.0$
 251 is ± 1.07 , comparable to the ± 1.47 that temperature sampling produces across 16 samples at $\tau=2.0$.

252 **Trajectory-level: semantic diversity.** Token-level reranking
 253 raises a second question: do first-token perturbations diverge
 254 into distinct trajectories, or converge to similar ones? We sample
 255 64 problems from MATH-500 and generate 300 independent
 256 trajectories per problem under Baseline and RSP at $\tau = 1.0$.
 257 Each trajectory is encoded by BGE-M3 [Chen et al., 2024] into a
 258 dense semantic vector, and we compute three metrics: *pairwise*
 259 *cosine distance* (overall semantic spread), *inter-cluster distance*
 260 between centroids of DBSCAN [Ester et al., 1996], a density-
 261 based clustering algorithm (separation between distinct trajectory groups), and *intra-cluster distance*
 262 (coherence within groups). Table 4 reveals three concurrent patterns: pairwise distance rises by
 263 $1.4\times$, inter-cluster distance jumps by $5.4\times$, and intra-cluster distance is essentially unchanged —
 264 the trajectory set becomes more diverse and splits into more separated groups while preserving
 265 within-group coherence. RSP also yields larger pairwise distance than baseline on 56 of 64 problems.

Table 4: Semantic diversity metrics (64 problems, 300 trajectories per condition).

Metric	Baseline	RSP
Pairwise cos. dist.	0.0612	0.0879
Inter-cluster dist.	0.0183	0.0982
Intra-cluster dist.	0.0526	0.0528

266 **Outcome-level: Pass@N scaling.** Does this token- and trajectory-level diversity translate into
 267 task-level outcome diversity, measured by Pass@N [Chen et al., 2021] (the probability that at least
 268 one of N sampled solutions is correct)? We generate 16 samples per problem on MATH-500 and
 269 AIME24 and compare four conditions: (i) *Baseline* at $\tau = 1.0$; (ii) *RSP (single seed)* at $\tau = 1.0$, with
 270 one RSP shared across all samples; (iii) *RSP (indep. seed)* with greedy decoding, a fresh RSP per
 271 sample; and (iv) *RSP (indep. seed)* at $\tau = 1.0$, a fresh RSP per sample.




Figure 4: Pass@N scaling on (a) MATH-500 and (b) AIME24 with Qwen2.5-Math-1.5B-Instruct, 16 samples per problem. *Baseline*: temperature sampling only. *RSP (single seed)*: single RSP shared across samples combined with temperature. *RSP (indep. seed)*: a different RSP per sample, with or without temperature.

272 Condition (iv) achieves the highest Pass@N on both benchmarks, reaching 92.80% on MATH-
 273 500 and 36.67% on AIME24 at $N = 16$, while (ii) shows no improvement over the baseline —
 274 the diversity gain depends on varying the embeddings across samples rather than fixing a single
 275 perturbation. Pass@1 stays comparable to the baseline on MATH-500, and on the harder AIME24
 276 (iv) even surpasses it, suggesting that independent RSPs combined with temperature sampling expand
 277 trajectory diversity without sacrificing single-sample accuracy.

278 5 Application: DAPO Training with RSP

279 Sections 3–4 established what RSP does (attention-mediated branch opening) and how it behaves
 280 at inference (Pass@N gains, rank-changing diversity unattainable by temperature). Can we use the
 281 same effect at training time? Rollout diversity drives learning in sampling-based RL, and DAPO [Yu
 282 et al., 2025] is a recent GRPO [Shao et al., 2024] variant explicitly designed to preserve diversity at
 283 the loss level. We use it as a stringent testbed: does input-side branch opening compose with loss-side
 284 preservation to yield further gains? Details and per-step results are in Appendix F.

285 DAPO preserves rollout diversity at the loss level by reshaping how rollouts are weighted and
 286 clipped during long chain-of-thought RL; RSP induces diversity *before* the loss is computed by
 287 perturbing the rollout distribution at the input level (§4.4). Writing each rollout’s token sequence as

288 $y_i = (y_{i,1}, \dots, y_{i,|y_i|})$ (using y to avoid collision with the head output $\mathbf{o}_i^{(\ell)}$ of §3.2) and pairing each
289 rollout with an independently sampled RSP $\mathbf{H}_i$ (reused at update time to keep the update on-policy),
290 importance ratio $r_{i,t}(\theta) = \pi_\theta(y_{i,t} | q, \mathbf{H}_i, y_{i,<t}) / \pi_{\theta_{\text{old}}}(y_{i,t} | q, \mathbf{H}_i, y_{i,<t})$, group-relative advantage
291 $\hat{A}_{i,t}$, and asymmetric clip range $(\varepsilon_{\text{low}}, \varepsilon_{\text{high}})$, the DAPO + RSP objective is

$$\mathcal{J}_{\text{DAPO+RSP}}(\theta) = \mathbb{E} \left[\frac{1}{\sum_i |y_i|} \sum_{i,t} \min \left(r_{i,t} \hat{A}_{i,t}, \text{clip}(r_{i,t}, 1-\varepsilon_{\text{low}}, 1+\varepsilon_{\text{high}}) \hat{A}_{i,t} \right) \right]. \quad (3)$$

Table 5: Per-benchmark accuracy (%) at each method’s peak step on Qwen2.5-Math-7B. **Bold** denotes the better of the two RL methods on each benchmark. Avg is the unweighted five-benchmark mean.

Method	GSM8K	MATH-500	College	Minerva	AIME 24	Avg
Base	57.2	52.6	18.3	13.6	8.6	30.06
DAPO	86.3	78.2	41.4	37.5	22.6	53.20
DAPO + RSP	87.7	78.6	42.2	39.7	23.6	54.36

292 We train Qwen2.5-Math-7B on a MATH Level-3–5 subset (~8.5K prompts) with SimpleRL-Zoo
293 [Zeng et al., 2025] under Eq. (3): each step uses $G = 8$ rollouts per prompt at $\tau=1.0$, batch 1,024,
294 4 PPO mini-batches of 256, for 100 steps. *DAPO* is compared against *DAPO + RSP*, where *suffix*
295 *RSP* ($L=20$) is applied during rollouts only and evaluation uses no RSP. We evaluate on five math
296 benchmarks (GSM8K [Cobbe et al., 2021], MATH-500 [Lightman et al., 2024], College Math,
297 Minerva Math [Lewkowycz et al., 2022], AIME 2024) and report the unweighted mean; AIME 2024
298 uses Avg@32 at $\tau=1.0$, the others greedy.

299 Figure 5 and Table 5 show two patterns: a late-
300 stage advantage and a delayed early reward growth.
301 DAPO + RSP reaches 54.36% at step 90 (vs DAPO’s
302 53.20% peak, +1.16 pp) and widens the gap to +3.10 pp
303 at step 100 with DAPO’s post-peak decline delayed; ear-
304 lier the curves cross around step 20. The patterns follow
305 from the methods acting at separate stages: DAPO pre-
306 serves diversity at the loss level over already-generated
307 rollouts while RSP perturbs the hidden states that pro-
308 duce them (§3.2), exposing the policy to broader learn-
309 ing signals. This is the exploration–exploitation dy-
310 namics of RL [Sutton and Barto, 2018] induced from
311 the input. These results use a single training seed and
312 should be read as a feasibility demonstration; multi-seed
313 evaluation is future work.




Figure 5: Five-benchmark average accuracy on Qwen2.5-Math-7B. DAPO + RSP reaches a higher peak (step 90) and stays stable through step 100.

314 6 Conclusion

315 We studied Random Soft Prompts (RSPs), training-free isotropic Gaussian vectors fitted to pretrained
316 embedding statistics. The RSP contribution per attention layer reduces to a single scalar $w_{r,i}^{(\ell)}$ whose
317 fixed-gap upper bound shrinks under KV cache growth, and isotropic resampling reaches top- K
318 entry regions that rank-preserving temperature cannot (§3). Latent-prompt methods share the same
319 entropy signature, the gain vanishes when one RSP is shared across samples (§4), and DAPO + RSP
320 yields a modest late-training advantage (§5). This signature emerges from purely random input-side
321 perturbation, suggesting that part of what trained soft prompts have been credited with is a *structural*
322 *by-product of injection*, and that RSP composes orthogonally with sampling-based pipelines as a
323 near-zero-cost diversity expander. Open questions remain on generalization beyond RoPE decoders
324 and math reasoning, deep-layer sharpening outside Theorem 1’s envelope, multi-seed DAPO variance,
325 and held-out hyperparameter selection (§4.1).

References

- Farhan Ahmed, Yuya Jeremy Ong, and Chad DeLuca. LogitScope: A Framework for Analyzing LLM Uncertainty through Information Metrics. In *ICLR Workshop*, 2026.
- Nora Belrose, Igor Ostrovsky, Lev McKinney, Zach Furman, Logan Smith, Danny Halawi, Stella Biderman, and Jacob Steinhardt. Eliciting Latent Predictions from Transformers with the Tuned Lens. *arXiv:2303.08112*, 2023.
- Jianlyu Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. M3-Embedding: Multi-Linguality, Multi-Functionality, Multi-Granularity Text Embeddings Through Self-Knowledge Distillation. In *Findings of ACL*, 2024.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgen Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr, Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating Large Language Models Trained on Code. *arXiv:2107.03374*, 2021.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training Verifiers to Solve Math Word Problems. *arXiv:2110.14168*, 2021.
- Jeremy M. Cohen, Elan Rosenfeld, and J. Zico Kolter. Certified Adversarial Robustness via Randomized Smoothing. In *ICML*, 2019.
- Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu. A Density-Based Algorithm for Discovering Clusters in Large Spatial Databases with Noise. In *KDD*, 1996.
- Meire Fortunato, Mohammad Gheshlaghi Azar, Bilal Piot, Jacob Menick, Ian Osband, Alex Graves, Vlad Mnih, Remi Munos, Demis Hassabis, Olivier Pietquin, Charles Blundell, and Shane Legg. Noisy Networks for Exploration. In *ICLR*, 2018.
- Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. Transformer Feed-Forward Layers Are Key-Value Memories. In *EMNLP*, 2021.
- Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh Nagarajan. Think before You Speak: Training Language Models with Pause Tokens. In *ICLR*, 2024.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurelien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Roziere, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song, Danielle Pintz, Danny Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina Lobanova, Emily Dinan, Eric Michael Smith, Filip Radenovic, Francisco Guzmán, Frank Zhang, Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, Graeme Nail, Gregoire Mialon, Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Iliyan Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang, Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelfer van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang, Jiawen Liu, Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua

377 Johnstun, Joshua Saxe, Junteng Jia, Kalyan Vasuden Alwala, Karthik Prasad, Kartikeya Upasani,
 378 Kate Plawiak, Ke Li, Kenneth Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz
 379 Malik, Kuenley Chiu, Kunal Bhalla, Kushal Lakhotia, Lauren Rantala-Yearly, Laurens van der
 380 Maaten, Lawrence Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo,
 381 Lukas Blecher, Lukas Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat
 382 Singh, Manohar Paluri, Marcin Kardas, Maria Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya
 383 Pavlova, Melanie Kambadur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman
 384 Goyal, Narjes Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri Chatterji, Ning Zhang,
 385 Olivier Duchenne, Onur Celebi, Patrick Alrassy, Pengchuan Zhang, Pengwei Li, Petar Vasic,
 386 Peter Weng, Prajjwal Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Puxin Xu,
 387 Qing He, Qingxiao Dong, Ragavan Srinivasan, Raj Ganapathy, Ramon Calderer, Ricardo Silveira
 388 Cabral, Robert Stojnic, Roberta Raileanu, Rohan Maheswari, Rohit Girdhar, Rohit Patel, Romain
 389 Sauvestre, Ronnie Polidoro, Roshan Sumbaly, Ross Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar
 390 Hosseini, Sahana Chennabasappa, Sanjay Singh, Sean Bell, Seohyun Sonia Kim, Sergey Edunov,
 391 Shao-liang Nie, Sharan Narang, Sharath Raparthy, Sheng Shen, Shengye Wan, Shruti Bhosale,
 392 Shun Zhang, Simon Vandenhende, Soumya Batra, Spencer Whitman, Sten Sootla, Stephane
 393 Collot, Suchin Gururangan, Sydney Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha,
 394 Thomas Georgiou, Thomas Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal
 395 Karn, Vedanuj Goswami, Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Gouget,
 396 Virginie Do, Vish Vogeti, Vitor Albiero, Vladan Petrovic, Weiwei Chu, Wenhan Xiong, Wenxin
 397 Fu, Whitney Meers, Xavier Martinet, Xiaodong Wang, Xiaofang Wang, Xiaoqing Ellen Tan, Xide
 398 Xia, Xinfeng Xie, Xuchao Jia, Xuwei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei,
 399 Yi Wen, Yiwen Song, Yuchen Zhang, Yue Li, Yuning Mao, Zacharie Delpierre Coudert, Zheng Yan,
 400 Zhengxing Chen, Zoe Papakipos, Aaditya Singh, Aayushi Srivastava, Abha Jain, Adam Kelsey,
 401 Adam Shajnfeld, Adithya Gangidi, Adolfo Victoria, Ahuva Goldstand, Ajay Menon, Ajay Sharma,
 402 Alex Boesenber, Alexei Baevski, Allie Feinstein, Amanda Kallet, Amit Sangani, Amos Teo,
 403 Anam Yunus, Andrei Lupu, Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew
 404 Poulton, Andrew Ryan, Ankit Ramchandani, Annie Dong, Annie Franco, Anuj Goyal, Aparajita
 405 Saraf, Arkabandhu Chowdhury, Ashley Gabriel, Ashwin Bharambe, Assaf Eisenman, Azadeh
 406 Yazdan, Beau James, Ben Maurer, Benjamin Leonhardi, Bernie Huang, Beth Loyd, Beto De Paola,
 407 Bhargavi Paranjape, Bing Liu, Bo Wu, Boyu Ni, Braden Hancock, Bram Wasti, Brandon Spence,
 408 Brani Stojkovic, Brian Gamido, Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu,
 409 Changan Wang, Changkyu Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris
 410 Tindal, Christoph Feichtenhofer, Cynthia Gao, Damon Civin, Dana Beaty, Daniel Kreymer, Daniel
 411 Li, David Adkins, David Xu, Davide Testuggine, Delia David, Devi Parikh, Diana Liskovich,
 412 Didem Foss, Dingkan Wang, Duc Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine
 413 Montgomery, Eleonora Presani, Emily Hahn, Emily Wood, Eric-Tuan Le, Erik Brinkman, Esteban
 414 Arcaute, Evan Dunbar, Evan Smothers, Fei Sun, Felix Kreuk, Feng Tian, Filippos Kokkinos, Firat
 415 Ozgenel, Francesco Caggioni, Frank Kanayet, Frank Seide, Gabriela Medina Florez, Gabriella
 416 Schwarz, Gada Badeer, Georgia Swee, Gil Halpern, Grant Herman, Grigory Sizov, Guangyi Zhang,
 417 Guna Lakshminarayanan, Hakan Inan, Hamid Shojanazeri, Han Zou, Hannah Wang, Hanwen Zha,
 418 Haroun Habeeb, Harrison Rudolph, Helen Suk, Henry Aspegren, Hunter Goldman, Hongyuan
 419 Zhan, Ibrahim Damlaj, Igor Molybog, Igor Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai
 420 Gat, Jake Weissman, James Geboski, James Kohli, Janice Lam, Japhet Asher, Jean-Baptiste Gaya,
 421 Jeff Marcus, Jeff Tang, Jennifer Chan, Jenny Zhen, Jeremy Reizenstein, Jeremy Teboul, Jessica
 422 Zhong, Jian Jin, Jingyi Yang, Joe Cummings, Jon Carvill, Jon Shepard, Jonathan McPhie, Jonathan
 423 Torres, Josh Ginsburg, Junjie Wang, Kai Wu, Kam Hou U, Karan Saxena, Kartikay Khandelwal,
 424 Katayoun Zand, Kathy Matosich, Kaushik Veeraraghavan, Kelly Michelena, Keqian Li, Kiran
 425 Jagadeesh, Kun Huang, Kunal Chawla, Kyle Huang, Lailin Chen, Lakshya Garg, Lavender A,
 426 Leandro Silva, Lee Bell, Lei Zhang, Liangpeng Guo, Licheng Yu, Liron Moshkovich, Luca
 427 Wehrstedt, Madian Khabsa, Manav Avalani, Manish Bhatt, Martynas Mankus, Matan Hasson,
 428 Matthew Lennie, Matthias Reso, Maxim Groshev, Maxim Naumov, Maya Lathi, Meghan Keneally,
 429 Miao Liu, Michael L. Seltzer, Michal Valko, Michelle Restrepo, Mihir Patel, Mik Vyatskov,
 430 Mikayel Samvelyan, Mike Clark, Mike Macey, Mike Wang, Miquel Jubert Hermoso, Mo Metanat,
 431 Mohammad Rastegari, Munish Bansal, Nandhini Santhanam, Natascha Parks, Natasha White,
 432 Navyata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier, Nikhil Mehta, Nikolay Pavlovich
 433 Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, Olivia Hart, Omkar Salpekar, Ozlem
 434 Kalinli, Parkin Kent, Parth Parekh, Paul Saab, Pavan Balaji, Pedro Rittner, Philip Bontrager,
 435 Pierre Roux, Piotr Dollar, Polina Zvyagina, Prashant Ratanchandani, Pritish Yuvraj, Qian Liang,

436 Rachad Alao, Rachel Rodriguez, Rafi Ayub, Raghotham Murthy, Raghu Nayani, Rahul Mitra,
 437 Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin Battey, Rocky Wang, Russ
 438 Howes, Ruty Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu, Samyak Datta, Sara Chugh,
 439 Sara Hunt, Sargun Dhillon, Sasha Sidorov, Satadru Pan, Saurabh Mahajan, Saurabh Verma, Seiji
 440 Yamamoto, Sharadh Ramaswamy, Shaun Lindsay, Sheng Feng, Shenghao Lin,
 441 Shengxin Cindy Zha, Shishir Patil, Shiva Shankar, Shuqiang Zhang, Shuqiang Zhang, Sinong Wang,
 442 Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala, Stephanie Max, Stephen Chen, Steve Kehoe,
 443 Steve Satterfield, Sudarshan Govindaprasad, Sumit Gupta, Summer Deng, Sungmin Cho, Sunny
 444 Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman, Tal Remez, Tamar Glaser, Tamara
 445 Best, Thilo Koehler, Thomas Robinson, Tianhe Li, Tianjun Zhang, Tim Matthews, Timothy Chou,
 446 Tzook Shaked, Varun Vontimitta, Victoria Ajayi, Victoria Montanez, Vijai Mohan, Vinay Satish
 447 Kumar, Vishal Mangla, Vlad Ionescu, Vlad Poenaru, Vlad Tiberiu Mihailescu, Vladimir Ivanov,
 448 Wei Li, Wenchen Wang, Wenwen Jiang, Wes Bouaziz, Will Constable, Xiaocheng Tang, Xiaojian
 449 Wu, Xiaolan Wang, Xilun Wu, Xinbo Gao, Yaniv Kleinman, Yanjun Chen, Ye Hu, Ye Jia, Ye Qi,
 450 Yenda Li, Yilin Zhang, Ying Zhang, Yossi Adi, Youngjin Nam, Yu Wang, Yu Zhao, Yuchen Hao,
 451 Yundi Qian, Yunlu Li, Yuzi He, Zach Rait, Zachary DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhenyu
 452 Yang, Zhiwei Zhao, and Zhiyu Ma. The Llama 3 Herd of Models. *arXiv:2407.21783*, 2024.

453 Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian.
 454 Training Large Language Models to Reason in a Continuous Latent Space. In *ICLR*, 2025.

455 Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Thai, Junhao Shen, Jinyi Hu, Xu Han,
 456 Yujie Huang, Yuxiang Zhang, Jie Liu, Lei Qi, Zhiyuan Liu, and Maosong Sun. OlympiadBench: A
 457 Challenging Benchmark for Promoting AGI with Olympiad-Level Bilingual Multimodal Scientific
 458 Problems. In *ACL*, 2024.

459 Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin de Laroussilhe, Andrea
 460 Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-Efficient Transfer Learning for NLP.
 461 In *ICML*, 2019.

462 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang,
 463 and Weizhu Chen. LoRA: Low-Rank Adaptation of Large Language Models. In *ICLR*, 2022.

464 Neel Jain, Ping-yeh Chiang, Yuxin Wen, John Kirchenbauer, Hong-Min Chu, Gowthami Somepalli,
 465 Brian R. Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Aniruddha Saha, Micah Goldblum,
 466 Jonas Geiping, and Tom Goldstein. NEFTune: Noisy Embeddings Improve Instruction Finetuning.
 467 In *ICLR*, 2024.

468 Xinyue Kang, Diwei Shi, and Li Chen. Model Whisper: Steering Vectors Unlock Large Language
 469 Models’ Potential in Test-time. In *AAAI*, 2026.

470 Brian Lester, Rami Al-Rfou, and Noah Constant. The Power of Scale for Parameter-Efficient Prompt
 471 Tuning. In *EMNLP*, 2021.

472 Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay
 473 Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, Yuhuai Wu, Behnam
 474 Neyshabur, Guy Gur-Ari, and Vedant Misra. Solving Quantitative Reasoning Problems with
 475 Language Models. In *NeurIPS*, 2022.

476 Xiang Lisa Li and Percy Liang. Prefix-Tuning: Optimizing Continuous Prompts for Generation. In
 477 *ACL-IJCNLP*, 2021.

478 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike,
 479 John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s Verify Step by Step. In *ICLR*, 2024.

480 Xiao Liu, Yanan Zheng, Zhengxiao Du, Ming Ding, Yujie Qian, Zhilin Yang, and Jie Tang. GPT
 481 Understands, Too. *AI Open*, 5:208–215, 2024.

482 Aman Madaan and Amir Yazdanbakhsh. Text and Patterns: For Effective Chain of Thought, It Takes
 483 Two to Tango. *arXiv:2209.07686*, 2022.

484 Alexander Robey, Eric Wong, Hamed Hassani, and George J. Pappas. SmoothLLM: Defending Large
 485 Language Models Against Jailbreaking Attacks. *Transactions on Machine Learning Research*,
 486 2025.

487 Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang,
488 Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the Limits of
489 Mathematical Reasoning in Open Language Models. *arXiv:2402.03300*, 2024.

490 Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng,
491 Haibin Lin, and Chuan Wu. HybridFlow: A Flexible and Efficient RLHF Framework. In *EuroSys*,
492 2025.

493 Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov.
494 Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *Journal of Machine*
495 *Learning Research*, 15:1929–1958, 2014.

496 Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2
497 edition, 2018.

498 Shenzhi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen,
499 Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen
500 Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 Rule: High-Entropy Minority Tokens Drive
501 Effective Reinforcement Learning for LLM Reasoning. In *NeurIPS*, 2025.

502 Yige Xu, Xu Guo, Zhiwei Zeng, and Chunyan Miao. SoftCoT: Soft Chain-of-Thought for Efficient
503 Reasoning with LLMs. In *ACL*, 2025.

504 An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu,
505 Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu,
506 Xingzhang Ren, and Zhenru Zhang. Qwen2.5-Math Technical Report: Toward Mathematical
507 Expert Model via Self-Improvement. *arXiv:2409.12122*, 2024.

508 Wengao Ye, Yan Liang, and Lianlei Shan. Thinking on the Fly: Test-Time Reasoning Enhancement
509 via Latent Thought Policy Optimization. In *ICLR*, 2026.

510 Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian
511 Fan, Gaohong Liu, Juncai Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming
512 Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Ru Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu,
513 Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Yuxuan Song, Xiangpeng Wei, Hao Zhou,
514 Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Yonghui Wu, and Mingxuan Wang. DAPO:
515 An Open-Source LLM Reinforcement Learning System at Scale. In *NeurIPS*, 2025.

516 Weihao Zeng, Yuzhen Huang, Qian Liu, Wei Liu, Keqing He, Zejun Ma, and Junxian He. SimpleRL-
517 Zoo: Investigating and Taming Zero Reinforcement Learning for Open Base Models in the Wild.
518 In *COLM*, 2025.

519 Guibin Zhang, Muxin Fu, and Shuicheng Yan. MemGen: Weaving Generative Latent Memory for
520 Self-Evolving Agents. In *ICLR*, 2026.

A Experimental Setup and Full Accuracy Breakdown

All experiments are conducted on a single NVIDIA A6000 40GB GPU with greedy decoding, a maximum generation length of 3,072 tokens for MATH-500 and GSM8K, and 4,096 tokens for AIME24. Batch size is fixed per model (16 for LLaMA-3.1-8B-Instruct and Qwen2.5-Math-7B variants, 32 for Qwen2.5-Math-1.5B variants). The RSP token count L is selected per (model, benchmark) pair from $\{10, 15, 20\}$ (Table 7); the per-position breakdown is in Table 6.

Table 6: Accuracy (%) across model configurations, injection positions, and benchmarks. Values in parentheses indicate the difference from the baseline. **Bold** denotes the best result and underline denotes the second best for each model–benchmark pair.

Model	Mode	MATH-500	GSM8K	AIME24
<i>Instruct Models</i>				
LLaMA-3.1-8B-Instruct	Baseline	52.20	85.75	<u>6.67</u>
	Prefix	45.00 (−7.2)	76.35 (−9.4)	3.33 (−3.3)
	Suffix	<u>49.40</u> (−2.8)	86.73 (+1.0)	10.00 (+3.3)
Qwen2.5-Math-7B-Instruct	Baseline	83.20	95.45	<u>13.33</u>
	Prefix	81.40 (−1.8)	94.92 (−0.5)	<u>13.33</u> (+0.0)
	Suffix	83.40 (+0.2)	95.68 (+0.2)	16.67 (+3.3)
Qwen2.5-Math-1.5B-Instruct	Baseline	73.00	<u>85.29</u>	10.00
	Prefix	74.80 (+1.8)	85.29 (+0.0)	<u>13.33</u> (+3.3)
	Suffix	74.20 (+1.2)	84.69 (−0.6)	20.00 (+10.0)
<i>Base Models (Raw Text)</i>				
Qwen2.5-Math-7B	Baseline	72.20	84.15	6.67
	Prefix	<u>71.60</u> (−0.6)	84.31 (+0.2)	16.67 (+10.0)
	Suffix	55.60 (−16.6)	54.13 (−30.0)	<u>13.33</u> (+6.7)
Qwen2.5-Math-1.5B	Baseline	<u>61.40</u>	77.86	16.67
	Prefix	64.40 (+3.0)	<u>74.30</u> (−3.6)	<u>10.00</u> (−6.7)
	Suffix	54.60 (−6.8)	58.61 (−19.3)	3.33 (−13.3)
<i>Base Models + ChatML (Format Mismatch)</i>				
Qwen2.5-Math-7B	Baseline	52.20	58.30	23.33
	Prefix	59.20 (+7.0)	56.41 (−1.9)	<u>3.33</u> (−20.0)
	Suffix	70.40 (+18.2)	75.97 (+17.7)	23.33 (+0.0)
Qwen2.5-Math-1.5B	Baseline	34.40	37.45	<u>13.33</u>
	Prefix	63.40 (+29.0)	67.02 (+29.6)	20.00 (+6.7)
	Suffix	<u>51.00</u> (+16.6)	<u>50.64</u> (+13.2)	<u>13.33</u> (+0.0)

Table 7: Number of RSP tokens (L) used for each model and benchmark.

Model	MATH-500	GSM8K	AIME24
LLaMA-3.1-8B-Instruct	15	20	15
Qwen2.5-Math-7B-Instruct	20	20	20
Qwen2.5-Math-1.5B-Instruct	20	10	20
Qwen2.5-Math-7B	10	10	20
Qwen2.5-Math-1.5B	10	10	20
Qwen2.5-Math-1.5B (ChatML)	20	20	10
Qwen2.5-Math-7B (ChatML)	20	20	20

A.1 RSP Injection Positions and Prompt Templates

Below we show the prompt structure for each injection position across all model types. **Blue text** indicates RSP embeddings, and **purple text** indicates special tokens.

A.1.1 Prefix

RSP embeddings are concatenated before the prompt embeddings. No text modification is applied.

ChatML (Qwen Instruct / Base + ChatML).

```
[Random Embeddings ( $L$  tokens)]
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{<|im_end|>
<|im_start|>user
{question}<|im_end|>
<|im_start|>assistant
```

LLaMA Chat Template.

```
[Random Embeddings ( $L$  tokens)]
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{<|eot_id|>
<|start_header_id|>user<|end_header_id|>
{question}<|eot_id|>
<|start_header_id|>assistant<|end_header_id|>
```

Raw Text (Qwen Base).

```
[Random Embeddings ( $L$  tokens)]
{question}
Please reason step by step, and put your final answer within \boxed{<|eot_id|>
```

A.1.2 Infix

A description text and special tokens are inserted into the prompt. The special token positions are then replaced with RSP embeddings at the embedding level.

ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{<|im_end|>
<|im_start|>user
{question}

There are { $L$ } special tokens that contain compressed latent reasoning information that
might be useful for your reasoning. If these tokens are useful for your case, you can
use them as reference. If these tokens are not useful for your case, you can ignore
them and focus back to solving the problem.

Here are the { $L$ } special tokens: <special_token_1>...<special_token_ $L$ >
<|im_end|>
<|im_start|>assistant
```

LLaMA Chat Template.

```
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{<|eot_id|>
<|start_header_id|>user<|end_header_id|>
{question}

There are { $L$ } special tokens that contain compressed latent reasoning information that
might be useful for your reasoning. If these tokens are useful for your case, you can
use them as reference. If these tokens are not useful for your case, you can ignore
them and focus back to solving the problem.

Here are the { $L$ } special tokens:
<reserved_special_token_0>...
<reserved_special_token_ $L$ >
<|eot_id|>
<|start_header_id|>assistant<|end_header_id|>
```

545 Raw Text (Qwen Base).

```
{question}
There are {L} special tokens that contain compressed latent reasoning information that
might be useful for your reasoning. If these tokens are useful for your case, you can
use them as reference. If these tokens are not useful for your case, you can ignore
them and focus back to solving the problem.

Here are the {L} special tokens: <special_token_1>...<special_token_L>

Please reason step by step, and put your final answer within \boxed{}
```

547 A.1.3 Suffix

548 For chat-templated models, RSP embeddings are inserted immediately before the end-of-turn token.
 549 For raw text models, RSP embeddings are concatenated at the end of the prompt.

550 ChatML (Qwen Instruct / Base + ChatML).

```
<|im_start|>system
Please reason step by step, and put your final answer within \boxed{}.<|im_end|>
<|im_start|>user
{question}[Random Embeddings (L tokens)]<|im_end|>
<|im_start|>assistant
```

552 LLaMA Chat Template.

```
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
Cutting Knowledge Date: December 2023
Today Date: 26 Jul 2024
Please reason step by step, and put your final answer within \boxed{}.<|eot_id|>
<|start_header_id|>user<|end_header_id|>
{question}[Random Embeddings (L tokens)]<|eot_id|>
<|start_header_id|>assistant<|end_header_id|>
```

554 Raw Text (Qwen Base).

```
{question}
Please reason step by step, and put your final answer within \boxed{}. [Random
Embeddings (L tokens)]
```

556 A.2 Answer Extraction and Grading

557 The final answer is extracted from the model output through a fallback chain. Each step is attempted
 558 in order; if extraction fails, the next step is tried.

Answer Extraction Fallback Chain

1. "final answer is \$...\$. I hope" pattern (Minerva format)
2. \boxed{...}: last non-empty box is used; empty \boxed{} are skipped
3. Text following "the answer is"
4. Text following "final answer is"
5. Last number in the output (regex fallback)

560 Extracted answers are normalized by removing \$, \left, \right, and unit strings. Grading
 561 is performed via symbolic equivalence checking: exact string match, numerical comparison
 562 (`math.isclose, rel_tol=1e-4`), and SymPy symbolic simplification.

563 A.3 Reproduced Prior Work Hyperparameters

564 All prior methods are reproduced on Qwen2.5-Math-1.5B-Instruct and Qwen2.5-Math-7B-Instruct,
 565 evaluated with greedy decoding (temperature = 0) and our unified answer extraction pipeline.

TTSV. *Prefix* length 20, trained for 20 epochs with AdamW optimizer, batch size 2, gradient accumulation 8 steps, and linear learning rate schedule. Learning rate is 1e-3 for Qwen-1.5B and 1e-5 for Qwen-7B.

LTPO. Table 8 reports the per-benchmark hyperparameters.

Table 8: LTPO hyperparameters.

Parameter	GSM8K	MATH-500	AIME24
tokens	8	8	8
steps	20	20	20
top_k	10	10	10
sigma	20	20	5
sigma_decay	0.95	0.95	0.95
lr	1e-2	5e-3	5e-2

SoftCoT. Projection module trained for 10 epochs. Thought tokens: 32 during training, 4 during evaluation. Batch size 4 for GSM8K, 1 for MATH with gradient accumulation 4. The small (assistant) model is Qwen2.5-Math-1.5B-Instruct in both configurations, while the large model is Qwen2.5-Math-1.5B-Instruct (learning rate 2e-5) or Qwen2.5-Math-7B-Instruct (learning rate 5e-6). For MATH-500 and AIME24 evaluation, we use the projection module trained on GSM8K, as it yields higher accuracy than the one trained on MATH.

A.4 Full Entropy Curves

Figure 6 shows the full normalized (0–100%) generation trajectories for entropy, top-1 probability, and varentropy, and Table 9 reports the corresponding scalar statistics including overall means, first-5% means, and the average number of generated tokens per problem. All methods converge to similar levels after the initial generation stage, confirming that the distributional changes induced by RSP are localized to the early reasoning phase.

Table 9: Per-step statistics for Qwen2.5-Math-1.5B-Instruct on MATH-500 across the full generation and the first 5% of generation steps, together with the average number of generated tokens per problem. Top-1 Prob (overall) is reported as a weighted average over the first-10%/middle/last-10% segment means with weights 0.1/0.8/0.1. Extension of Figure 3 (main text) and Figure 6.

Metric	Baseline	<i>Prefix</i>	<i>Infix</i>	<i>Suffix</i>	LTPO	SoftCoT	TTSV
Tokens (mean)	575.7	558.4	549.9	573.3	592.3	594.4	577.4
Entropy (first 5%)	0.1332	0.1366	0.1402	0.1941	0.1662	0.4218	0.1359
Entropy (overall)	0.0871	0.0881	0.0899	0.1049	0.0909	0.1059	0.0864
Top-1 Prob (first 5%)	0.9535	0.9523	0.9525	0.9373	0.9437	0.9156	0.9534
Top-1 Prob (overall)	0.9684	0.9681	0.9678	0.9647	0.9683	0.9651	0.9685
Varentropy (first 5%)	0.2181	0.2207	0.2463	0.4010	0.2953	2.2234	0.2174
Varentropy (overall)	0.1353	0.1375	0.1422	0.2038	0.1452	0.2404	0.1361




Figure 6: Mean entropy (top), top-1 probability (middle), and varentropy (bottom) over the full normalized generation trajectory (0–100%). Averaged over MATH-500, shading indicates ± 1 SEM. Solid lines: RSP variants and the baseline. Dashed lines: prior soft prompt methods. All methods converge after the initial stage.

582 B Token-Level Distribution Metrics

583 This appendix provides formal definitions for the metrics used in Section 4.4. Numerical results
 584 are reported in Table 3 of the main text. We extract first-token logits via a single forward pass (no
 585 autoregressive generation) and store the top-100 token ids and logit values per problem. Temperature
 586 conditions reuse the baseline logits scaled by $1/\tau$, requiring no separate inference. RSP draws
 587 16 independent seeds per problem (MATH-500), and reported metrics are the mean across the 16
 588 seeds. Pass@1 accuracy is computed over 16 samples per problem, with the baseline at $\tau=1.0$
 589 reaching $73.92 \pm 0.78\%$. Let P_{base} denote the baseline next-token distribution at $\tau=1.0$ and P_{target} the
 590 candidate distribution being compared (e.g., baselines at $\tau=2.0, 3.0$ or RSP at $\tau=1.0$). All metrics
 591 are computed analytically from saved logits at the first generation step.

592 **Spearman rank correlation.** Spearman’s ρ is the Pearson correlation coefficient between the token
 593 ranks of two distributions:

$$\rho = \text{Pearson}(\text{rank}(P_{\text{target}}), \text{rank}(P_{\text{base}})). \quad (4)$$

594 Because temperature scaling $P^{(\tau)} \propto P^{1/\tau}$ is a strictly monotone transform, the rank order of tokens
 595 is preserved exactly, so $\rho = 1$ for any $\tau > 0$ as a mathematical identity. Any value $\rho < 1$ therefore
 596 indicates rank reorganization beyond what temperature scaling can produce.

597 **Mass outside top- K .** Let $\text{TopK}(P_{\text{base}})$ be the set of K tokens with the highest probability under
 598 P_{base} . The probability mass that the candidate distribution places *outside* this set is

$$\text{MassOutside}_K = 1 - \sum_{t \in \text{TopK}(P_{\text{base}})} P_{\text{target}}(t). \quad (5)$$

599 This directly measures support expansion: how much probability the candidate assigns to tokens that
 600 are not in the baseline’s preferred set. Reported values use $K = 10$.

601 **Jensen–Shannon divergence.** The symmetrized, bounded version of Kullback–Leibler divergence:

$$\text{JS}(P, Q) = \frac{1}{2} \text{KL}(P \parallel M) + \frac{1}{2} \text{KL}(Q \parallel M), \quad M = \frac{1}{2}(P + Q). \quad (6)$$

602 We use base-2 logarithms so that $\text{JS} \in [0, 1]$. Tokens absent from the stored top-100 are assigned a
 603 sentinel logit (-10^9) before softmax, which is negligible for our setting.

C Prompt-Law Properties Used by the Theory

This appendix isolates the only properties of the RSP law used in the main text: centeredness, direction-agnostic covariance under a variance budget, and positive probability on every nonempty open set. We avoid stronger semantic claims because Section 3.2 needs only these geometric and measure-theoretic facts.

C.1 Centered Perturbations Do Not Impose Systematic First-Order Drift

Let z_i^{Base} denote the no-RSP output logit at vocab token i and $z_i^{\text{RSP}}(\bar{h})$ the corresponding logit when the centered RSP perturbation takes value $\bar{h}$. The (true) vocab-logit gap under RSP is $\Delta_{ij}(\bar{h}) := z_i^{\text{RSP}}(\bar{h}) - z_j^{\text{RSP}}(\bar{h})$, distinct from the no-RSP baseline gap $\Delta_{ij}^{\text{Base}} := z_i^{\text{Base}} - z_j^{\text{Base}}$. Note that $\bar{h} = 0$ corresponds to the centered RSP at its entrywise mean rather than to the no-RSP forward pass, so in general $\Delta_{ij}(0) \neq \Delta_{ij}^{\text{Base}}$. The transformer’s logit map is non-linear in $\bar{h}$, but a first-order Taylor expansion around $\bar{h} = 0$ yields the local surrogate

$$\Delta_{ij}^{\text{lin}}(\bar{h}) := \Delta_{ij} + b_{ij}^\top \bar{h},$$

with $\Delta_{ij} := \Delta_{ij}(0)$ and $b_{ij} := \nabla_{\bar{h}}(z_i^{\text{RSP}} - z_j^{\text{RSP}})|_{\bar{h}=0}$. The proposition below is stated for the surrogate Δ_{ij}^{lin} . For the true non-linear gap $\Delta_{ij}(\bar{h})$, the first-order term has zero mean under centered perturbations, but the Taylor remainder $r_{ij}(\bar{h}) = O(\|\bar{h}\|^2)$ may contribute a second-order mean shift; we make no claim about its sign or magnitude.

Proposition 2 (No systematic first-order drift in the surrogate). *If $\mathbb{E}[\bar{h}] = 0$, then*

$$\mathbb{E}[\Delta_{ij}^{\text{lin}}(\bar{h})] = \Delta_{ij}$$

for every token pair (i, j) . If a deterministic centered prompt $\bar{h}_0$ is reused across runs, then

$$\Delta_{ij}^{\text{lin}}(\bar{h}_0) - \Delta_{ij} = b_{ij}^\top \bar{h}_0$$

is a fixed directional bias.

Proof. The centered case is immediate from linearity of expectation:

$$\mathbb{E}[\Delta_{ij}^{\text{lin}}(\bar{h})] = \Delta_{ij} + b_{ij}^\top \mathbb{E}[\bar{h}] = \Delta_{ij}.$$

The deterministic case follows by substitution. $\square$

This proposition rules out a run-averaged first-order bias. It does not say that individual prompt draws leave the logits unchanged.

C.2 Proof of Proposition 1

For a unit vector b , the first-order perturbation energy available along b is $\text{Var}_{h \sim D}(b^\top h) = b^\top \Sigma_D b$. The worst-case directional energy is therefore the minimum Rayleigh quotient of Σ_D .

Proof. For any symmetric covariance matrix Σ_D ,

$$\min_{\|b\|=1} b^\top \Sigma_D b = \lambda_{\min}(\Sigma_D).$$

Let the eigenvalues of Σ_D be $\lambda_1, \dots, \lambda_d$. Then

$$\lambda_{\min}(\Sigma_D) \leq \frac{1}{d} \sum_{m=1}^d \lambda_m = \frac{\text{tr}(\Sigma_D)}{d} \leq \frac{\rho^2}{d}.$$

Equality holds if and only if all eigenvalues are equal, i.e., $\Sigma_D = (\rho^2/d)\mathbf{I}_d$. $\square$

The proof uses only covariance. We choose a Gaussian law for RSP because it adds the open-set reachability property proved next.

635 **C.3 Open-Set Reachability of Isotropic Gaussian Prompts**

636 **Proposition 3** (Open-set reachability). *Let $\bar{h} \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_d)$ with $\sigma > 0$. Then every nonempty open*
 637 *set $U \subseteq \mathbb{R}^d$ satisfies*

$$\Pr(\bar{h} \in U) > 0.$$

638 *Proof.* The Gaussian density

$$(2\pi\sigma^2)^{-d/2} \exp\left(-\frac{\|\bar{h}\|^2}{2\sigma^2}\right)$$

639 is strictly positive at every point of $\mathbb{R}^d$. Every nonempty open set contains a ball of positive Lebesgue
 640 measure, and integrating a strictly positive density over that ball gives positive probability. $\square$

641 This is the only support property used in the informal branching argument of §3.2 and in Ap-
 642 pendix D.3.3.

643 D Proofs for the Transformer-Level Mechanism

644 This appendix follows the same order as the main theory section: exploration, annealing, and
645 performance gain. The prompt-law facts used before these proofs are collected in Appendix C.

646 D.1 Exploration: attention decomposition and perturbation size

647 **Proposition 4** (Attention decomposition). *For one attention head at query position i in layer ℓ , let*
648 *$\alpha_{ij}^{(\ell)}$ be the attention weights over $n \geq 1$ unmasked real tokens and $L \geq 1$ random tokens, with finite*
649 *attention logits. Define*

$$w_{r,i}^{(\ell)} := \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)}$$

650 and, for $j \in [n]$,

$$\tilde{\alpha}_{ij}^{(\ell)} := \frac{\alpha_{ij}^{(\ell)}}{1 - w_{r,i}^{(\ell)}}.$$

651 Then

$$\mathbf{o}_i^{(\ell)} = (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{o}}_i^{(\ell)} + \boldsymbol{\eta}_i^{(\ell)},$$

652 where

$$\tilde{\mathbf{o}}_i^{(\ell)} := \sum_{j=1}^n \tilde{\alpha}_{ij}^{(\ell)} \mathbf{v}_j^{(\ell)} \quad \text{and} \quad \boldsymbol{\eta}_i^{(\ell)} := \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)}.$$

653 *Proof.* The head output is

$$\mathbf{o}_i^{(\ell)} = \sum_{j=1}^n \alpha_{ij}^{(\ell)} \mathbf{v}_j^{(\ell)} + \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)}.$$

654 By the softmax positivity and the existence of at least one unmasked real token, $\sum_{j=1}^n \alpha_{ij}^{(\ell)} =$
655 $1 - w_{r,i}^{(\ell)} > 0$, so

$$\sum_{j=1}^n \alpha_{ij}^{(\ell)} \mathbf{v}_j^{(\ell)} = (1 - w_{r,i}^{(\ell)}) \sum_{j=1}^n \frac{\alpha_{ij}^{(\ell)}}{1 - w_{r,i}^{(\ell)}} \mathbf{v}_j^{(\ell)} = (1 - w_{r,i}^{(\ell)}) \tilde{\mathbf{o}}_i^{(\ell)}.$$

656 Substituting this identity gives the decomposition. □

657 **Corollary 1** (Magnitude scales with random-token attention). *For every realization of the random*
658 *values,*

$$\|\boldsymbol{\eta}_i^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_{j \in [L]} \|\mathbf{v}_{r_j}^{(\ell)}\|.$$

659 *Proof.* By the triangle inequality,

$$\|\boldsymbol{\eta}_i^{(\ell)}\| = \left\| \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \mathbf{v}_{r_j}^{(\ell)} \right\| \leq \sum_{j=1}^L \alpha_{i,n+j}^{(\ell)} \|\mathbf{v}_{r_j}^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_{j \in [L]} \|\mathbf{v}_{r_j}^{(\ell)}\|.$$

660 □

661 Corollary 1 is the only quantitative fact needed in the main text. No independence assumption
662 between attention weights and random keys is required. Once $w_{r,i}^{(\ell)}$ is small, the random contribution
663 must be small as well.

664 D.2 Corollaries of Theorem 1 on KV cache growth

665 The fixed-gap decay bound yields the corollaries below.

666 **Corollary 2** (Sequence-wise annealing of the fixed-gap envelope). *For fixed L and $\Delta_i^{(\ell)}$, the upper*
 667 *bound*

$$f(n) = \frac{L}{L + n \exp(\Delta_i^{(\ell)} / \sqrt{d_{\text{head}}})}$$

668 *is strictly decreasing in n and tends to 0 as $n \rightarrow \infty$. KV cache growth alone therefore narrows the*
 669 *largest RSP attention mass compatible with that gap.*

670 During autoregressive decoding $\Delta_i^{(\ell)}$ also moves because queries, keys, and hidden states co-evolve.
 671 The accurate reading is not “the realized mass decreases at every step” but “the envelope compatible
 672 with a fixed gap shrinks as the KV cache grows.”

673 *Proof.* Differentiate $f(n)$ with respect to n :

$$f'(n) = -\frac{L \exp(\Delta_i^{(\ell)} / \sqrt{d_{\text{head}}})}{\left(L + n \exp(\Delta_i^{(\ell)} / \sqrt{d_{\text{head}}})\right)^2} < 0.$$

674 Thus $f(n)$ is strictly decreasing, and its denominator diverges linearly in n while its numerator is
 675 fixed, so $f(n) \rightarrow 0$. $\square$

676 **Corollary 3** (Vanishing random contribution under annealing). *If $n \rightarrow \infty$ while L and $\Delta_i^{(\ell)}$ remain*
 677 *fixed, $w_{r,i}^{(\ell)} \rightarrow 0$, and for every realization of the random values*

$$\|\boldsymbol{\eta}_i^{(\ell)}\| \leq w_{r,i}^{(\ell)} \max_{j \in [L]} \|\mathbf{v}_{r_j}^{(\ell)}\| \rightarrow 0.$$

678 *Proof.* $w_{r,i}^{(\ell)} \rightarrow 0$ by Corollary 2. Apply Corollary 1 to bound $\|\boldsymbol{\eta}_i^{(\ell)}\|$. The maximum norm is finite
 679 for any realization of finitely many sampled values. $\square$

680 D.3 Performance gain: from local admission to Pass@N

681 From here we work at the vocab-logit and task levels, dropping the per-layer index ℓ . The main text
 682 uses a first-order surrogate to state two claims: the §3.2 claim that a baseline-excluded vocab token
 683 can enter the local top- K set, and the §3.3 closing that independent reruns turn such local openings
 684 into Pass@ N gains. The proofs below use only one support condition on the centered prompt law D :

$$(S1) \quad D(U) > 0 \quad \text{for every nonempty open set } U \subseteq \mathbb{R}^d.$$

685 Proposition 3 shows that the isotropic Gaussian law of §3.2 satisfies (S1).

686 D.3.1 Autoregressive Formulation

687 In autoregressive decoding, the joint distribution over a sequence $y_{1:T}$ factorizes as

$$\Pr(y_{1:T} \mid x) = \prod_{t=1}^T p(y_t \mid x, y_{<t}),$$

688 where each step is governed by a prefix-conditioned distribution. Any perturbation introduced by
 689 RSP affects generation through a sequence of local changes to $p(y_t \mid x, y_{<t})$, rather than a single
 690 global distribution.

691 D.3.2 Local Linearization of Logits under RSP

692 Throughout this subsection, $\bar{h} \in \mathbb{R}^d$ denotes *one* centered prompt vector from the RSP definition
 693 (Eq. (1) in §3.1), treated as a per-token surrogate. The actual implementation uses a length- L prompt
 694 $\mathbf{H} = (h_1, \dots, h_L)$ with L independent draws; the arguments below isolate how one local perturbation
 695 can shift the logits. We model the perturbed logits as a function $z_i(\bar{h})$, assuming local smoothness in
 696 a neighborhood of $\bar{h} = 0$:

$$z_i(\bar{h}) = z_i(0) + \nabla_{\bar{h}} z_i(0)^\top \bar{h} + r_i(\bar{h}),$$

697 where $r_i(\bar{h}) = O(\|\bar{h}\|^2)$. Defining $a_i := \nabla_{\bar{h}} z_i(0)$ and $z_i := z_i(0)$, we obtain the local affine
 698 surrogate

$$z_i^{\text{lin}}(\bar{h}) := z_i + a_i^\top \bar{h}.$$

699 This approximation is used as a first-order surrogate for branch opening, not as an exact global model
 700 of the transformer logits.

701 D.3.3 Top- K entry region (Remark)

702 This complements the informal branching argument in §3.2. Define the entry region

$$\mathcal{G}_{i,K} := \{\bar{h} \in \mathbb{R}^d : \#\{j \neq i : z_j^{\text{lin}}(\bar{h}) > z_i^{\text{lin}}(\bar{h})\} \leq K - 1\}.$$

703 If $\mathcal{G}_{i,K}$ contains a nonempty open set U , then by (S1) $D(U) > 0$, and $U \subseteq \mathcal{G}_{i,K}$ gives $\Pr_{\bar{h}}(\bar{h} \in$
 704 $\mathcal{G}_{i,K}) \geq D(U) > 0$. A useful sufficient condition is the existence of some $\bar{h}_0$ at which token i strictly
 705 outranks all but at most $K - 1$ other tokens in the affine surrogate; by continuity, a neighborhood of
 706 $\bar{h}_0$ is then contained in $\mathcal{G}_{i,K}$. This is a rank-changing event unreachable by the monotone temperature
 707 transform, but it does not by itself imply (M1).

708 D.3.4 Pass@ N ensemble bound (Remark)

709 The bound $\Pr(\text{Pass@}N \text{ on } x) \geq 1 - (1 - p_{\min}(x))^N$ in the main text is a standard independent-
 710 Bernoulli union argument. Let A_n be the event that the n -th run reaches a correct trajectory on x ;
 711 under (M1) $\Pr(A_n) \geq p_{\min}(x)$, and under (M2)

$$\Pr\left(\bigcap_{n=1}^N A_n^c\right) \leq (1 - p_{\min}(x))^N, \quad \Pr\left(\bigcup_{n=1}^N A_n\right) \geq 1 - (1 - p_{\min}(x))^N \geq 1 - e^{-N p_{\min}(x)}.$$

712 This is a generic fact, not specific to RSP; RSP's role is restricted to supplying mechanisms that
 713 satisfy (M1) and (M2).

E Attention Mass Analysis

This appendix formalizes the per-token attention mass reported in Figure 2 and the exclusion criteria applied to the reasoning span and the layer axis. For each of the 500 MATH-500 problems, a fresh RSP $\mathbf{H} \in \mathbb{R}^{L \times d}$ with $L = 10$ is sampled independently, so the reported heatmaps average over both problem variability and RSP-vector variability.

E.1 Per-Token Attention Mass

For each problem, we measure attention on the concatenated sequence [prompt; $\mathbf{H}$; generation], where the generation is the trajectory produced by the same model under matching RSP injection. Let $\alpha_{i,j}^{(\ell,a)}$ denote the post-softmax attention weight at layer ℓ , head index a (using a to avoid collision with the RSP matrix $\mathbf{H}$), from query position i to key position j , satisfying $\sum_j \alpha_{i,j}^{(\ell,a)} = 1$ under the causal mask. Let Q and R denote the index sets of question and RSP tokens with sizes $|Q|$ and $|R| = L$, and let N_{head} be the number of heads. The per-token attention mass that query i directs to each group at layer ℓ is

$$\text{PTM}_Q[\ell, i] = \frac{1}{|Q| N_{\text{head}}} \sum_{a=1}^{N_{\text{head}}} \sum_{j \in Q} \alpha_{i,j}^{(\ell,a)}, \quad \text{PTM}_R[\ell, i] = \frac{1}{|R| N_{\text{head}}} \sum_{a=1}^{N_{\text{head}}} \sum_{j \in R} \alpha_{i,j}^{(\ell,a)}. \quad (7)$$

Normalization by $|Q|$ and $|R|$ controls for group size, so both quantities express how much attention a single question (resp. RSP) token receives on average under the same softmax denominator. Question tokens thereby serve as a size-matched baseline that absorbs sequence-length effects common to both groups.

E.2 Heatmap Aggregation

We partition the layer indices and the reasoning position indices each into five equal-size bins $\{B_L^{(b)}\}_{b=1}^5$ and $\{B_P^{(b)}\}_{b=1}^5$. For sample s and target group $\bullet \in \{Q, R\}$, the value in cell (b_L, b_P) is

$$M_{\bullet}^{(s)}[b_L, b_P] = \frac{1}{|B_L^{(b_L)}| \cdot |B_P^{(b_P)}|} \sum_{\ell \in B_L^{(b_L)}} \sum_{i \in B_P^{(b_P)}} \text{PTM}_{\bullet}[\ell, i]. \quad (8)$$

Figure 2 reports the sample average of Eq. (8) over the 500 valid samples.

E.3 Excluding the `\boxed{\dots}` Span

The reasoning position range terminates strictly before the final `\boxed{\dots}` span. Chain-of-thought outputs decompose into a *text* (content) component and a *pattern* (template) component that contribute independently to downstream performance [Madaan and Yazdanbakhsh, 2022]. The `\boxed{\dots}` wrapper is a canonical pattern: a short, stereotyped span whose role is to complete the answer format rather than to extend reasoning. Including it would therefore inject a template-driven attention signature unrelated to reasoning dynamics.

E.4 Excluding the Final Two Layers

The layer axis further excludes the last two transformer layers. This choice is motivated by a standard late-layer effect rather than by any model-specific tuning decision.

Output-projection specialization. Late-layer hidden states lie in an approximately linear relationship with vocabulary logits and therefore function as a progressive linear readout rather than as representation-transforming blocks [Belrose et al., 2023]. Consistently, late-layer feed-forward modules have been characterized as output-distribution shaping mechanisms [Geva et al., 2021]. Attention in these layers therefore reflects output formation rather than reasoning computation.

Excluding the final two layers keeps the heatmap focused on reasoning-phase dynamics instead of readout-phase dynamics. This exclusion is not load-bearing for the main claim: the sequence-direction attenuation emphasized in the main text is also visible without it, and Theorem 1 concerns sequence-direction attenuation rather than layer-wise monotonicity.

754 E.5 Additional Suffix Heatmaps

755 For the late-layer reasons discussed above, the pattern in Figure 7 does not generalize uniformly
 756 across every cell; nevertheless, the overall trend is consistent across both models: question-token
 757 attention varies broadly across layers, whereas RSP attention shows sequence-wise decay consistent
 758 with Theorem 1 alongside an empirical layer-wise attenuation that the theorem itself does not predict.




Figure 7: Per-token RSP attention mass under *suffix* injection for the remaining two models (500 MATH-500 samples each). Axes and preprocessing match Figure 2: x-axis is reasoning position binned into five quantiles, y-axis is layer depth binned into five quantiles (top = shallow), color encodes the 500-sample mean per-token attention assigned to RSP tokens, and tokens inside `\boxed{}` spans and the last two layers are excluded.

759 F DAPO Implementation

760 This appendix complements Section 5 with the DAPO loss modifications relative to GRPO, the
 761 DAPO + RSP implementation, and full per-step results.

762 F.1 From GRPO to DAPO

763 The vanilla GRPO objective [Shao et al., 2024] for a group of G rollouts $\{y_i\}_{i=1}^G$ (each $y_i =$
 764 $(y_{i,1}, \dots, y_{i,|y_i|})$ a token sequence; we use y rather than o throughout this appendix to avoid collision
 765 with the head output $\mathbf{o}_i^{(\ell)}$ of §3.2) with rewards $\{R_i\}_{i=1}^G$ is

$$\mathcal{J}_{\text{GRPO}}(\theta) = \mathbb{E} \left[\frac{1}{G} \sum_{i=1}^G \frac{1}{|y_i|} \sum_{t=1}^{|y_i|} \left(\min(r_{i,t} \hat{A}_{i,t}, \text{clip}(r_{i,t}, 1-\varepsilon, 1+\varepsilon) \hat{A}_{i,t}) - \beta D_{\text{KL}}(\pi_\theta \| \pi_{\text{ref}}) \right) \right], \quad (9)$$

766 with importance ratio $r_{i,t} = \pi_\theta(y_{i,t} | q, y_{i,<t}) / \pi_{\theta_{\text{old}}}(y_{i,t} | q, y_{i,<t})$, group-relative advantage $\hat{A}_{i,t} =$
 767 $(R_i - \text{mean}(\{R_i\})) / \text{std}(\{R_i\})$, clipping range ε , and KL penalty βD_{KL} against the reference π_{ref} .

768 We build on the SimpleRL-Zoo [Zeng et al., 2025] training pipeline and implement DAPO on
 769 top of VeRL [Sheng et al., 2025]. DAPO modifies Eq. (9) in four ways: (i) token-level loss
 770 aggregation $1 / \sum_i |y_i|$ instead of $1 / |y_i|$, (ii) asymmetric clipping with $(\varepsilon_{\text{low}}, \varepsilon_{\text{high}}) = (0.2, 0.28)$, (iii)
 771 an implementation-side negative-advantage dual-clipping safeguard $\max(L_{\text{clipped}}, c\hat{A})$ with $c = 10$
 772 (active only when $\hat{A} < 0$; omitted from Eq. (10) below for readability), and (iv) KL terms removed.
 773 The resulting DAPO objective is

$$\mathcal{J}_{\text{DAPO}}(\theta) = \mathbb{E} \left[\frac{1}{\sum_i |y_i|} \sum_{i,t} \min \left(r_{i,t} \hat{A}_{i,t}, \text{clip}(r_{i,t}, 1-\varepsilon_{\text{low}}, 1+\varepsilon_{\text{high}}) \hat{A}_{i,t} \right) \right]. \quad (10)$$

774 F.2 DAPO + RSP Implementation

775 For DAPO + RSP, we additionally inject a per-rollout RSP $\mathbf{H}_i \in \mathbb{R}^{L \times d}$ ($L = 20$) into each rollout.
 776 Each $\mathbf{H}_i$ is generated deterministically from a per-step rollout seed and injected via a forward hook;
 777 it is not a trainable parameter, so the learning signal is the policy adapting to arbitrary $\mathbf{H}_i$ rather than
 778 a learned prompt. Including $\mathbf{H}_i$ in the log-probabilities at both rollout and update time keeps the
 779 update on-policy with respect to the RSP-conditioned distribution (avoiding off-policy mismatch).
 780 With $G = 8$ rollouts per prompt, each rollout receives an independently drawn $\mathbf{H}_i$.

781 F.3 Data, Batch, and Hyperparameters

Table 10: DAPO / DAPO + RSP training setup on Qwen2.5-Math-7B.

Field	Value
Train data	MATH Level 3–5 subset (<code>simplelr_math_35</code> , ~8,500 prompts)
Prompt template	<code>qwen-boxed</code>
Train batch size	1024
PPO mini-batch size	256 (4 PPO updates per rollout)
Rollouts per prompt G	8
Max prompt / response length	2,028 / 2,048
Rollout sampling	temperature 1.0, top- p 1.0, top- k 50
$(\epsilon_{\text{low}}, \epsilon_{\text{high}})$	(0.2, 0.28)
Dual clip c	10.0
Loss aggregation	token-mean ($1 / \sum_i y_i $)
KL terms	disabled
Entropy coefficient	0
Optimizer	AdamW, learning rate 5×10^{-7} (constant), no warmup
Total training	~10 epochs, ≥ 80 optimization steps
Save / eval frequency	every 10 steps
RSP (DAPO + RSP only)	<i>suffix</i> injection, $L = 20$, freshly drawn per rollout
Hardware	4× NVIDIA B200 GPUs
Wall-clock training time	~12 hours per run

782 F.4 Evaluation Protocol

783 Each saved checkpoint is converted to HuggingFace format and evaluated using the SimpleRL-
784 Zoo `eval_math_nodes.sh` pipeline, which selects sampling parameters per benchmark to reduce
785 variance on the smaller competition sets:

Table 11: Per-benchmark evaluation protocol used by SimpleRL-Zoo `sh/eval.sh`.

Benchmark	# Problems	Temperature	n_{sampling}	Metric
AIME 2024	30	1.0	32	Avg@32
AMC 2023	40	1.0	32	Avg@32
MATH-500	500	0.0	1	accuracy (mean@1)
GSM8K	1,319	0.0	1	accuracy
OlympiadBench	675	0.0	1	accuracy
Minerva Math	272	0.0	1	accuracy
GaoKao 2023-EN	385	0.0	1	accuracy

786 The maximum number of generated tokens is 16,000 across all benchmarks. The reported average is
787 the unweighted mean over the five benchmarks (GSM8K, MATH-500, College Math, Minerva Math,
788 AIME 2024).

Table 12: Five-benchmark average accuracy (%) at every checkpoint, computed over GSM8K, MATH-500, College Math, Minerva Math, and AIME 2024. **Bold** marks each method’s peak.

Step	DAPO	DAPO + RSP
0	30.06	30.06
10	46.40	46.22
20	49.02	49.54
30	49.58	50.72
40	50.62	52.06
50	50.88	52.40
60	51.84	52.96
70	52.52	53.86
80	53.20	54.04
90	52.52	54.36
100	50.54	53.64

790 G Broader Impacts

791 This work provides empirical and theoretical analysis of how random embedding injection affects
792 the reasoning behavior of large language models. The contributions are primarily foundational,
793 with potential downstream implications for reinforcement learning-based post-training of reasoning
794 models such as GRPO.

795 **Positive impacts.** RSP-induced trajectory diversity could improve the sample efficiency of RL-
796 based LLM post-training by providing richer learning signals within group-relative advantage estima-
797 tion. This has the potential to reduce compute costs associated with alignment and reasoning-oriented
798 fine-tuning, making such methods more accessible to research groups with limited resources.

799 **Potential concerns.** Methods that expand the effective output support of LLMs may increase the
800 reachability of low-probability tokens, which includes both desirable alternative reasoning paths
801 and potentially unsafe or off-distribution outputs. Practitioners applying RSP in production systems
802 should combine it with existing safety filtering and alignment mechanisms rather than treating it as a
803 standalone diversity source.

804 **Limitations on impact scope.** Since this work focuses on analysis rather than deploying a new
805 model or dataset, the direct societal impact is limited. The broader implications described above are
806 contingent on future work integrating RSP into applied training pipelines.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and §1 state four scoped claims, each tied to a section: (i) RSP’s contribution at each attention layer is captured by a single scalar with a fixed-gap upper-envelope cache-growth bound (Theorem 1, §3.3); (ii) under a local affine surrogate, isotropic re-sampling reaches top- K entry regions unattainable by rank-preserving temperature sampling (§3.2); (iii) RSP raises early-token entropy and stabilizes later, with Pass@ N gains under independent re-sampling (§4); and (iv) the diversity transfers to RL training when combined with DAPO (§5). We do not claim a theorem that directly explains the format-mismatch recovery (Table 1); the wrong-attractor hypothesis there is explicitly informal.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Limitations are noted in several places rather than in a dedicated section: (a) the multi-path Pass@ N argument requires the task-side assumption $p_{\min}(x) > 0$, flagged in §3.2 and §6; (b) the +18–29 pp recovery on the format-mismatch rows of Table 1 is empirically observed but not formally derived from any theorem in §3, with the wrong-attractor account stated as a hypothesis to be quantified in future work (§4.2); (c) Table 1 reports the best-position accuracy across {Prefix, Infix, Suffix}, and per-position variance is large with several cases below baseline (Appendix A), so position is treated as a hyperparameter rather than something we justify principally; (d) DAPO + RSP results (§5) use a single training seed under a reduced DAPO configuration that omits Dynamic Sampling and overlong reward shaping; (e) main accuracy comparisons are single-seed without bootstrap or significance tests (item 7). Future work in §6 addresses training-time comparisons, domain expansion, and principled position selection.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: Theorem 1 (cache-growth attenuation) carries an inline proof in §3.3. Proposition 1 is stated in §3.2 with its proof in Appendix C; the supporting top- K entry and Pass@ N ensemble arguments around it are presented as informal discussion in §3.2 with formal write-ups in the appendix. Each statement explicitly lists its assumptions: finite logits and $n, L \geq 1$ for Theorem 1; finite second-moment budget for Proposition 1; and the locally affine surrogate plus $p_{\min}(x) > 0$ task-side condition together with mutual independence of the N runs for the Pass@ N argument.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: §4.1 and Appendix A specify benchmarks (MATH-500, GSM8K, AIME 2024 in the main table; further sets used for the DAPO study), models (LLaMA-3.1-8B-Instruct and Qwen2.5-Math 1.5B/7B variants), the RSP injection grid (Prefix/Infix/Suffix; $L \in \{10, 15, 20\}$), greedy decoding, and the answer extraction pipeline. Appendix F provides the full DAPO + RSP training recipe (batch of 1,024 prompts, $G = 8$ rollouts per prompt at $\tau = 1.0$, four PPO mini-batch updates of 256 prompts each, 100 training steps, suffix RSP $L = 20$ during rollouts only).

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: All datasets and base models used in the paper are publicly available and cited (MATH-500, GSM8K, AIME 2024, AMC 2023, OlympiadBench, Minerva Math, GaoKao 2023-EN, College Math; Qwen2.5-Math 1.5B/7B and LLaMA-3.1-8B-Instruct). RSP itself is described to a level of detail sufficient for re-implementation (Eq. (1) and §4.1). However, our RSP injection harness and DAPO + RSP training pipeline are not yet released; we will release an anonymized implementation alongside the camera-ready submission.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: §4.1 and Appendix A list all benchmarks, models, RSP injection settings (positions, lengths, per-row L choice in Appendix Table 7), maximum generation length per benchmark, and decoding strategy. Appendix F provides the full DAPO + RSP training settings, including optimizer, batch composition, mini-batch updates, asymmetric clipping range $(\epsilon_{\text{low}}, \epsilon_{\text{high}}) = (0.2, 0.28)$, dual clipping safeguard, and RSP injection protocol during rollouts.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: Pass@ N experiments use $N = 16$ samples per problem (§4.4), Table 3 reports mean $\pm$ standard deviation over 16 RSP seeds, and AIME 2024 is scored by Avg@32 to reduce variance. However, Table 1 (Table 1) reports single-run accuracies without bootstrap confidence intervals or paired tests, and the DAPO + RSP training in §5 uses a single training seed. Multi-seed evaluation, particularly on small competition sets (AIME 2024, AMC 2023) where variance is highest, is left for future work.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Inference experiments (Tables 1, 2; Figure 3) run on a single NVIDIA A6000 40GB GPU with greedy decoding (Appendix A). DAPO + RSP training (§5) runs on $4 \times$ NVIDIA B200 GPUs for approximately 12 hours per run (Appendix F, hardware row of the configuration table).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work uses publicly available math reasoning benchmarks and pretrained language models, involves no human subjects or personal data, and does not introduce a release that bypasses standard model-safety considerations. We have reviewed the NeurIPS Code of Ethics and identify no conflicts.

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Appendix G discusses both directions: positive impact via more sample-efficient RL post-training of reasoning LLMs, and a potential downside that broadening the effective output support also increases reachability of undesirable trajectories, which we argue should keep RSP coupled with existing safety filtering rather than replacing it.

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release new pretrained models, generative systems, or scraped datasets. RSP is a generation-time perturbation applied to existing pretrained models and introduces no separate artifact requiring safeguards.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: All datasets (MATH-500 [Lightman et al., 2024], GSM8K [Cobbe et al., 2021], AIME 2024, AMC 2023, OlympiadBench [He et al., 2024], Minerva Math [Lewkowycz et al., 2022], GaoKao 2023-EN, College Math) and pretrained models (Qwen2.5-Math [Yang et al., 2024], LLaMA-3.1 [Grattafiori et al., 2024]) are publicly available and cited at first use. The training pipeline builds on SimpleRL-Zoo [Zeng et al., 2025] and VeRL [Sheng et al., 2025], also cited. Each asset is used under its original publicly stated terms.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: The paper does not release new datasets or pretrained models. The implementation code planned for camera-ready release (item 5) will be accompanied by configuration files and a README; no derived dataset or model checkpoint is part of the submission.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper involves no crowdsourcing and no research with human subjects. All evaluations use static, publicly available math reasoning benchmarks.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: As noted in item 14, the paper involves no human subjects, so IRB review is not applicable.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

962 Answer: [N/A]
963 Justification: Pretrained LLMs (Qwen2.5-Math, LLaMA-3.1) appear in the paper as the
964 experimental subject of analysis, not as a non-standard component of the methods themselves.
965 The methods (RSP definition, attention decomposition, DAPO + RSP training) do not rely
966 on auxiliary LLMs in any non-standard or originality-affecting role.